



大数据导论

Introduction to Big Data



大数据深度学习基础

叶允明

计算机科学与技术学院

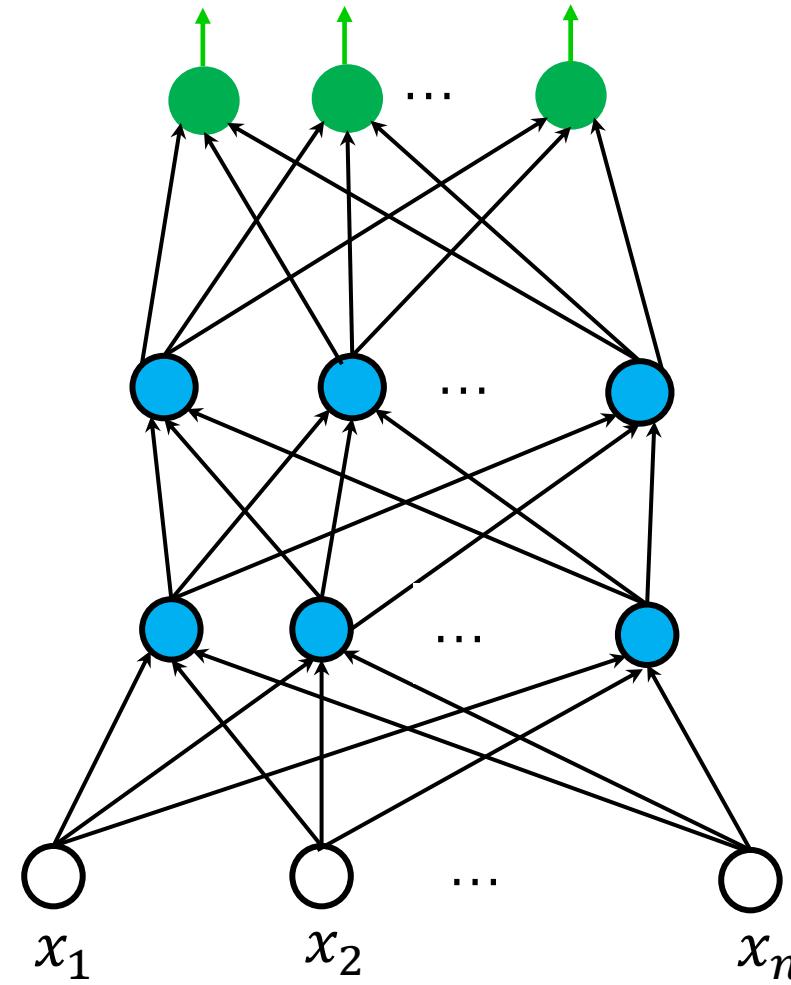
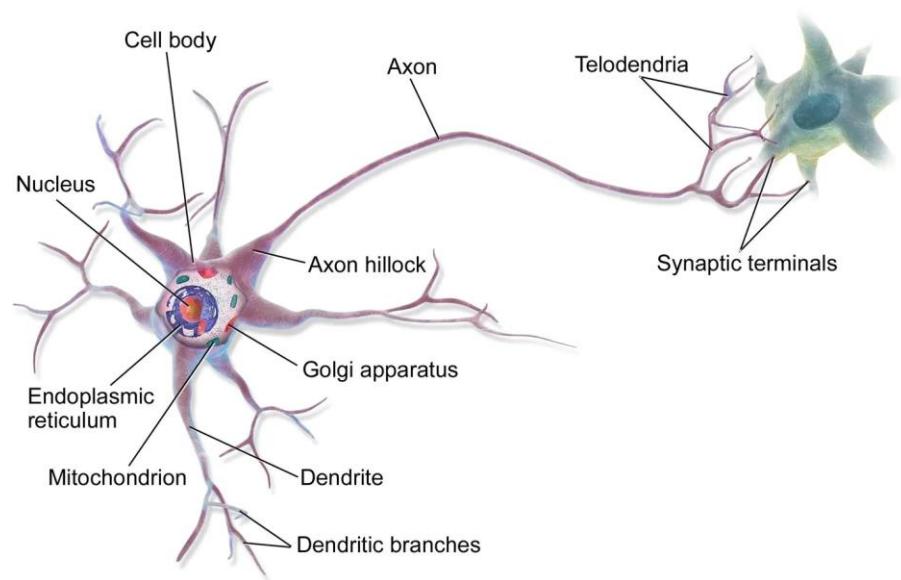
哈尔滨工业大学（深圳）

目录

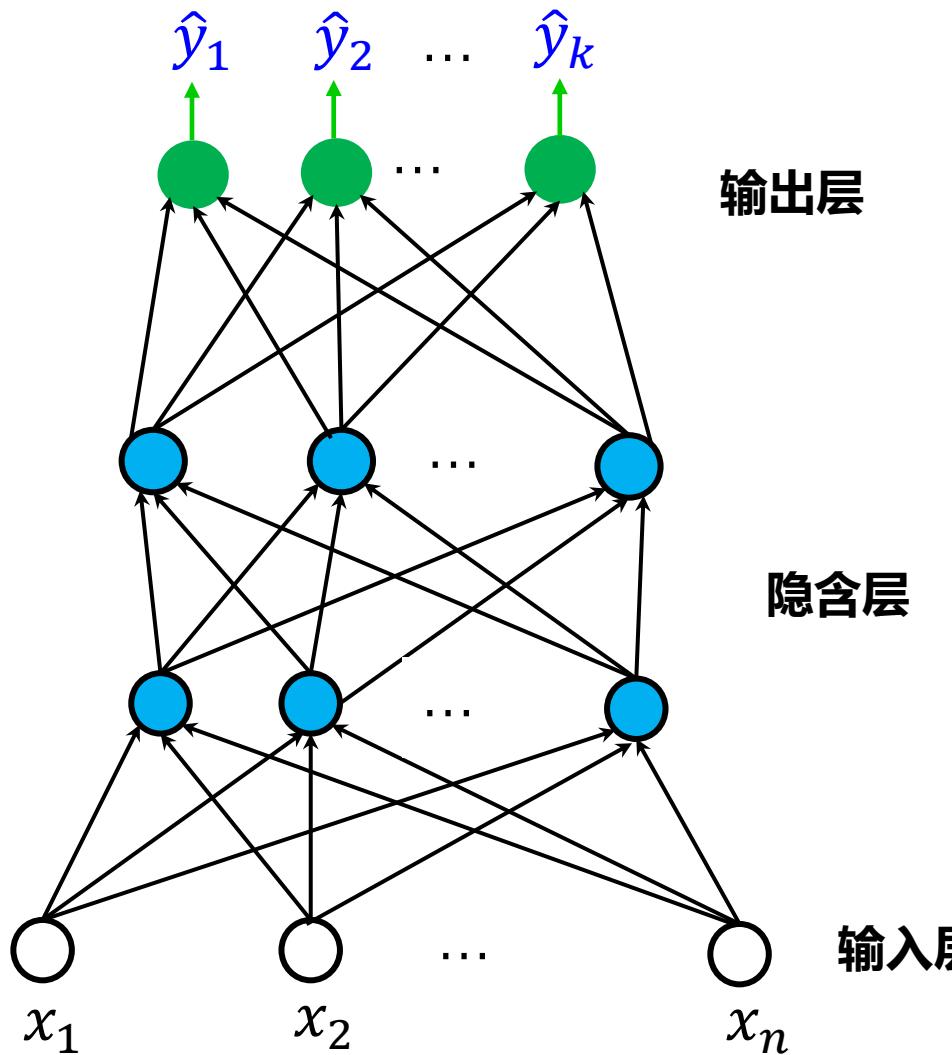
- 神经网络与深度学习基础
- 如何学习出更好的神经网络

神经网络与深度学习基础

人工神经网络(Artificial Neural Network)



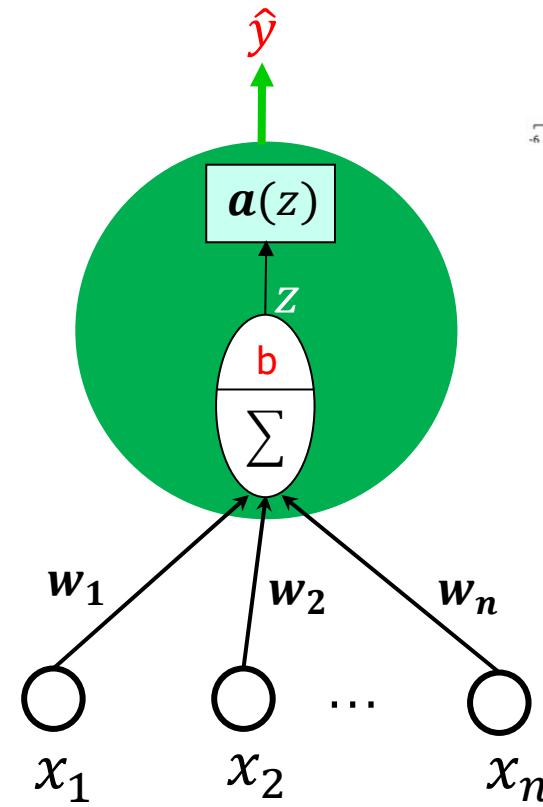
神经网络的常用结构



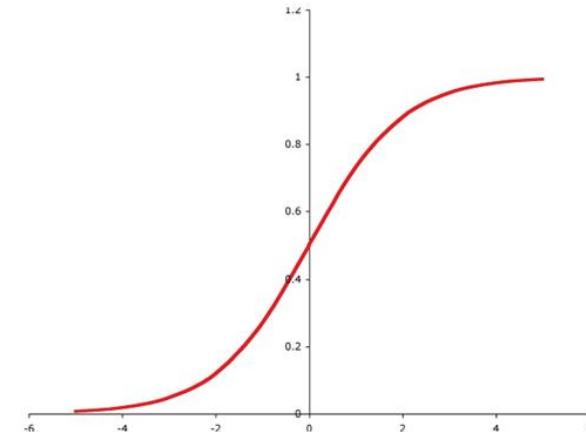
输出层

隐含层

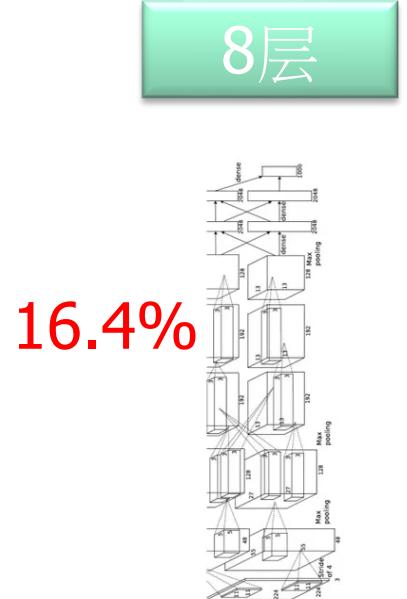
输入层



$$a(z) = \frac{1}{1 + e^{-z}}$$

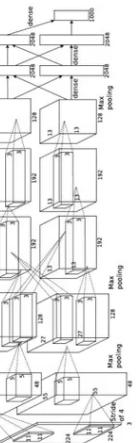


深度学习：深度神经网络

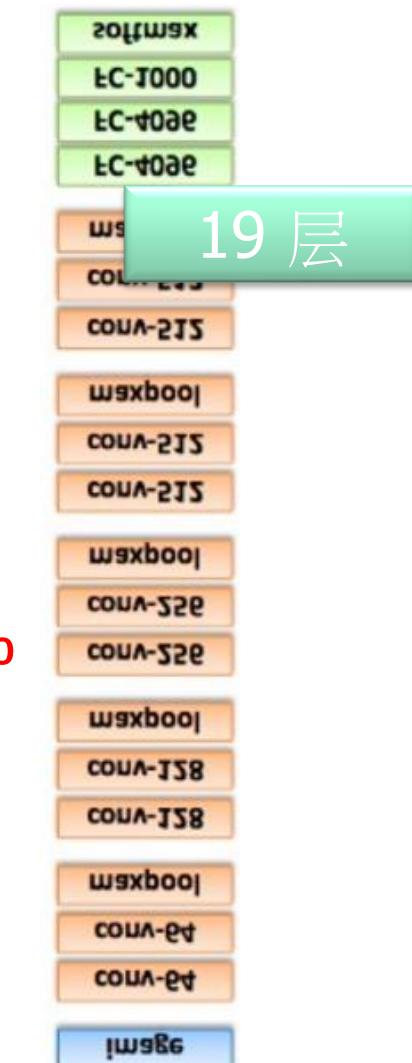


16.4%

7.3%



VGG (2014)



19 层

22 层

6.7%

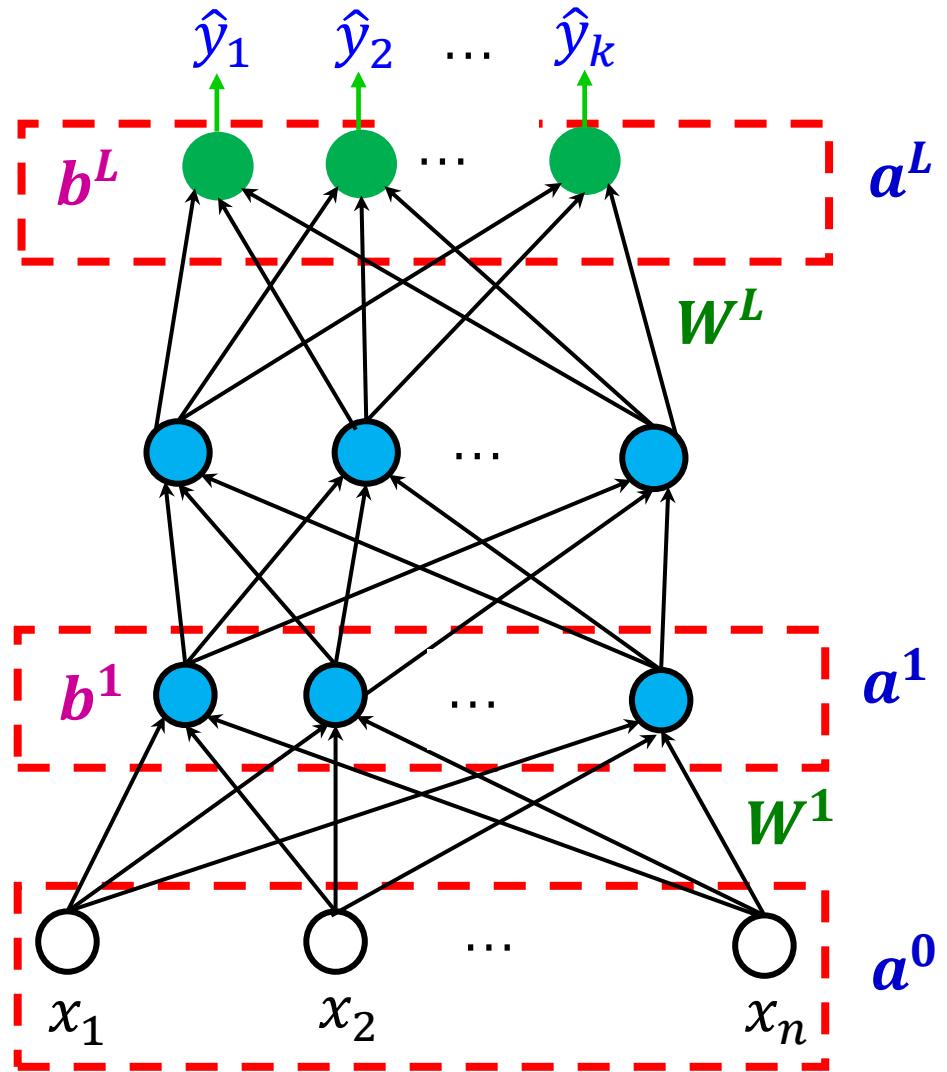
.....

据推测：GPT-4
在120层中总共
包含了1.8万亿
参数！

AlexNet (2012)

GoogleNet (2014)

全连接前馈神经网络的计算过程



$$a^L = \sigma(W^L a^{L-1} + b^L)$$

d^L is the total number of neurons of L layer

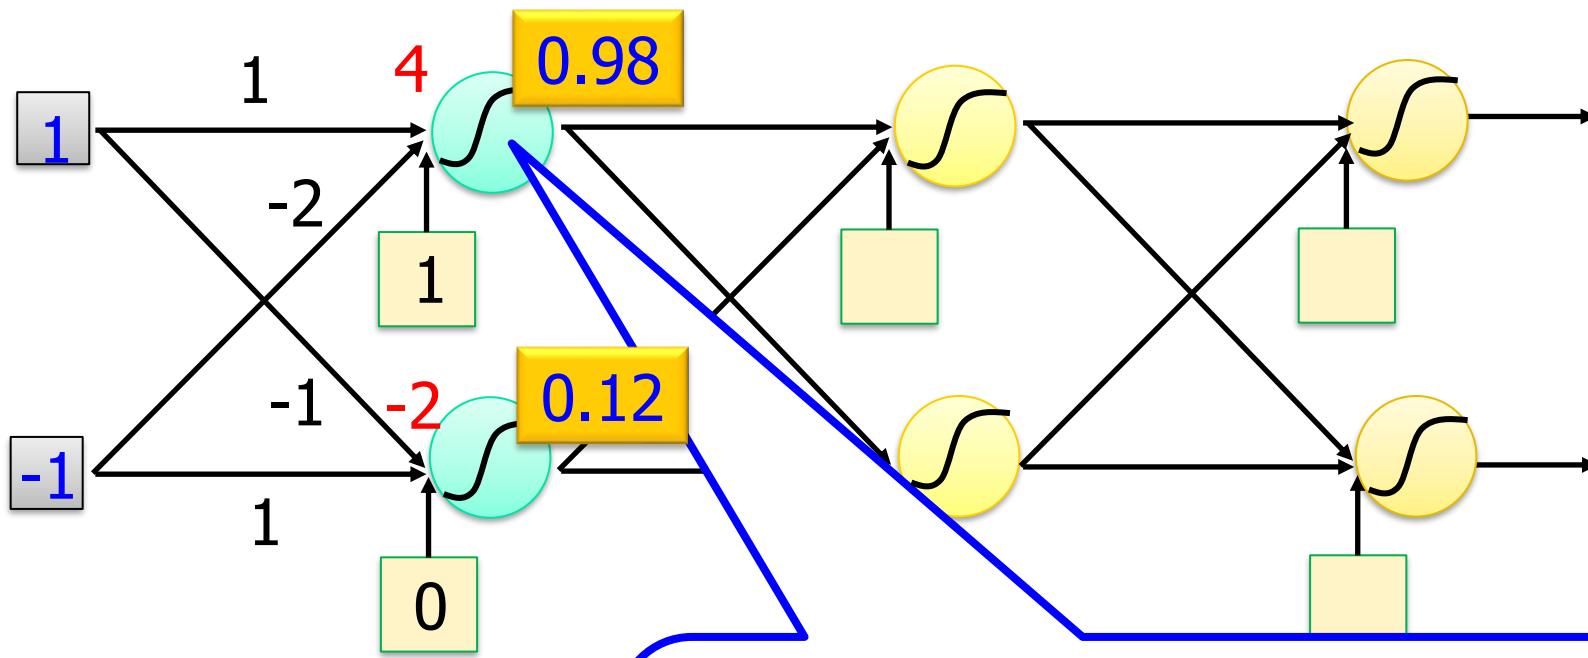
$$a^1 = \sigma(W^1 a^0 + b^1)$$

$$W^1 = \begin{bmatrix} (\mathbf{w}_1^1)^T \\ (\mathbf{w}_2^1)^T \\ \vdots \\ (\mathbf{w}_i^1)^T \\ \vdots \\ (\mathbf{w}_{d^1}^1)^T \end{bmatrix} = \begin{bmatrix} w_{11}^1 & \cdots & w_{1d^0}^1 \\ \cdots & w_{ij}^1 & \cdots \\ w_{d^1d^0}^1 & \cdots & w_{d^1d^0}^1 \end{bmatrix}$$

$$a(z) = \sigma(z) = \frac{1}{1 + e^{-z}}$$

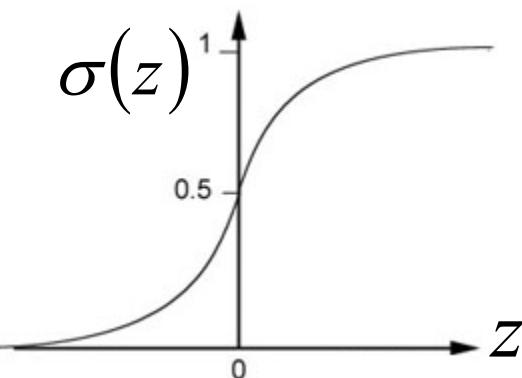
$$\mathbf{b}^1 = \begin{bmatrix} b_1^1 \\ b_2^1 \\ \vdots \\ b_i^1 \\ \vdots \\ b_{d^1}^1 \end{bmatrix}$$

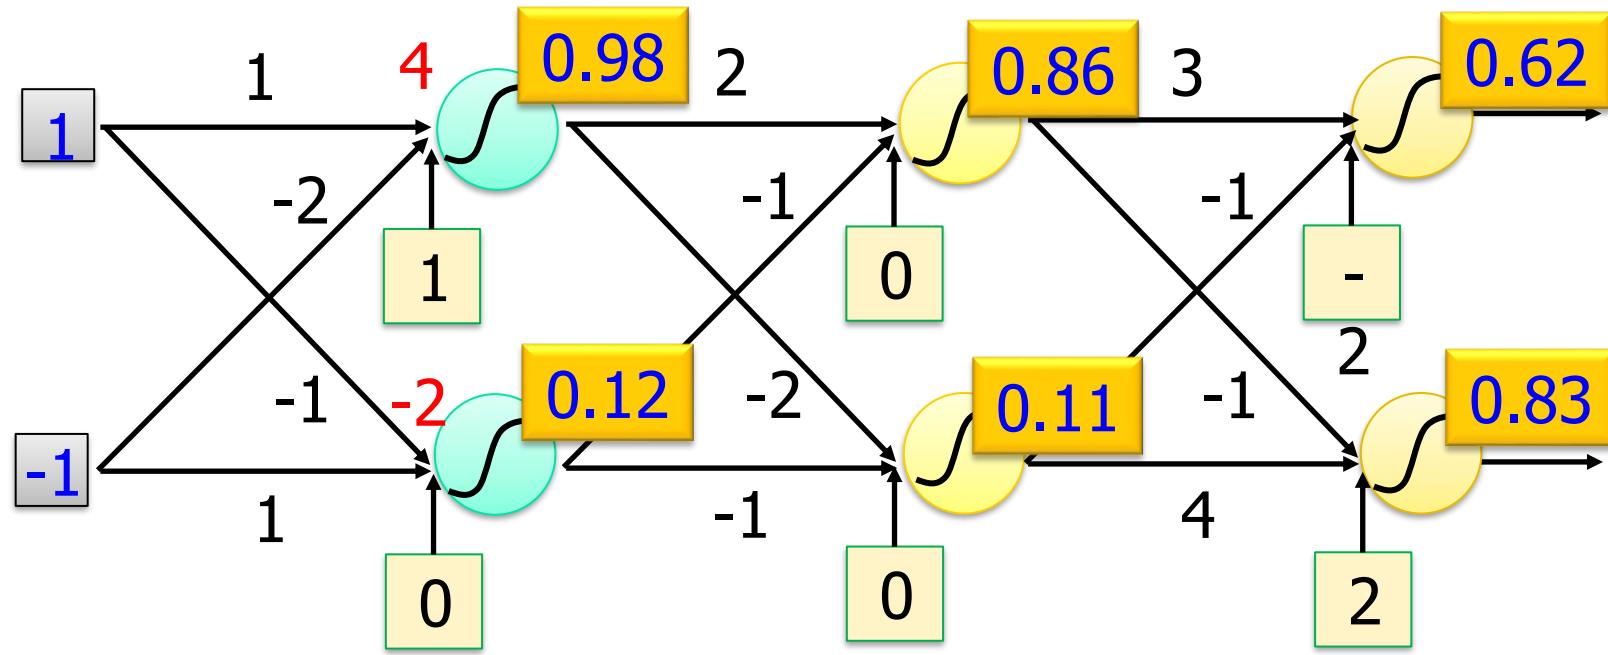
$$\mathbf{w}_i^1 = \begin{bmatrix} w_{i1}^1 \\ w_{i2}^1 \\ \vdots \\ w_{ij}^1 \\ \vdots \\ w_{id^0}^1 \end{bmatrix}$$



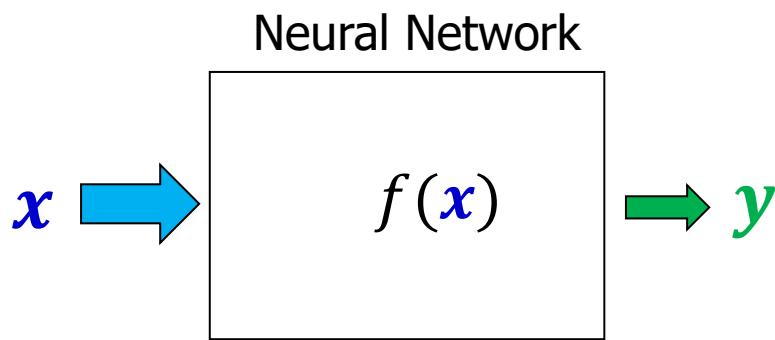
Sigmoid Function

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

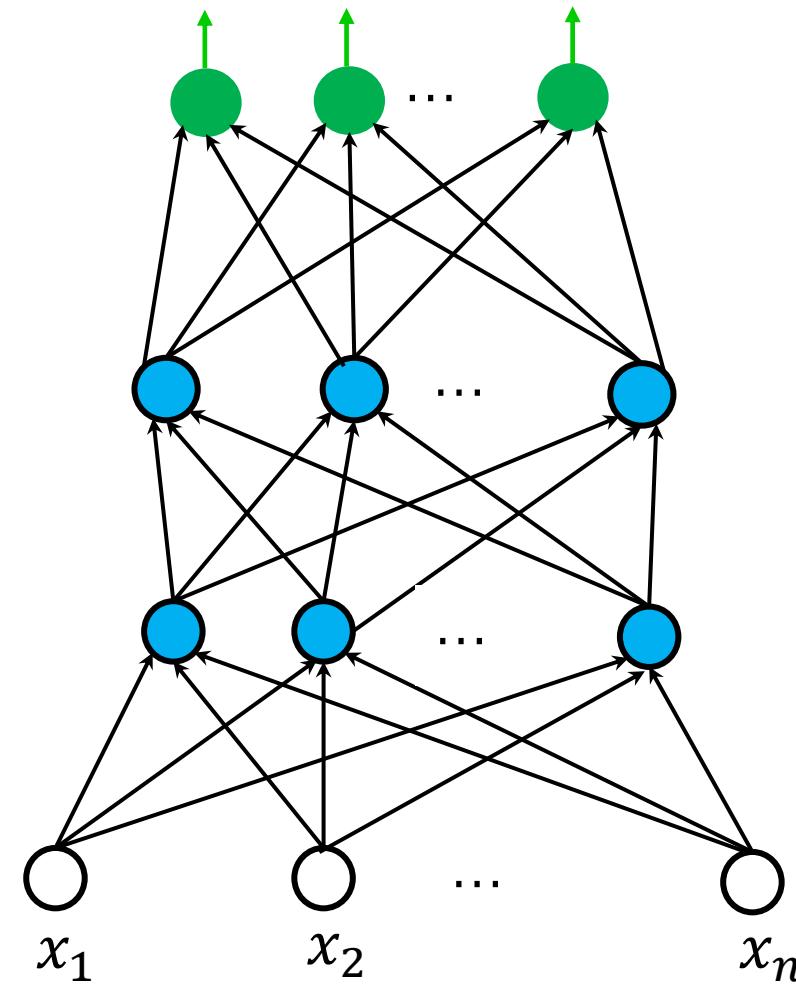




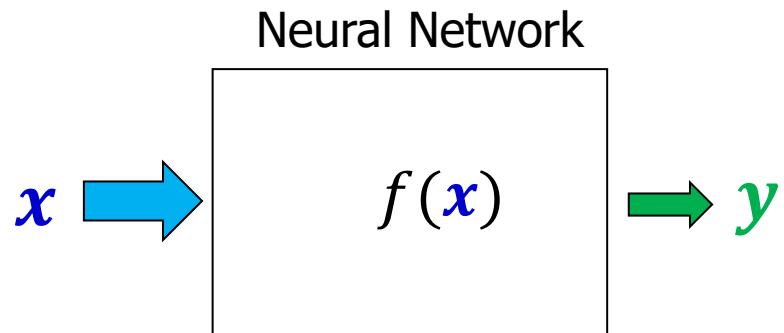
神经网络的本质



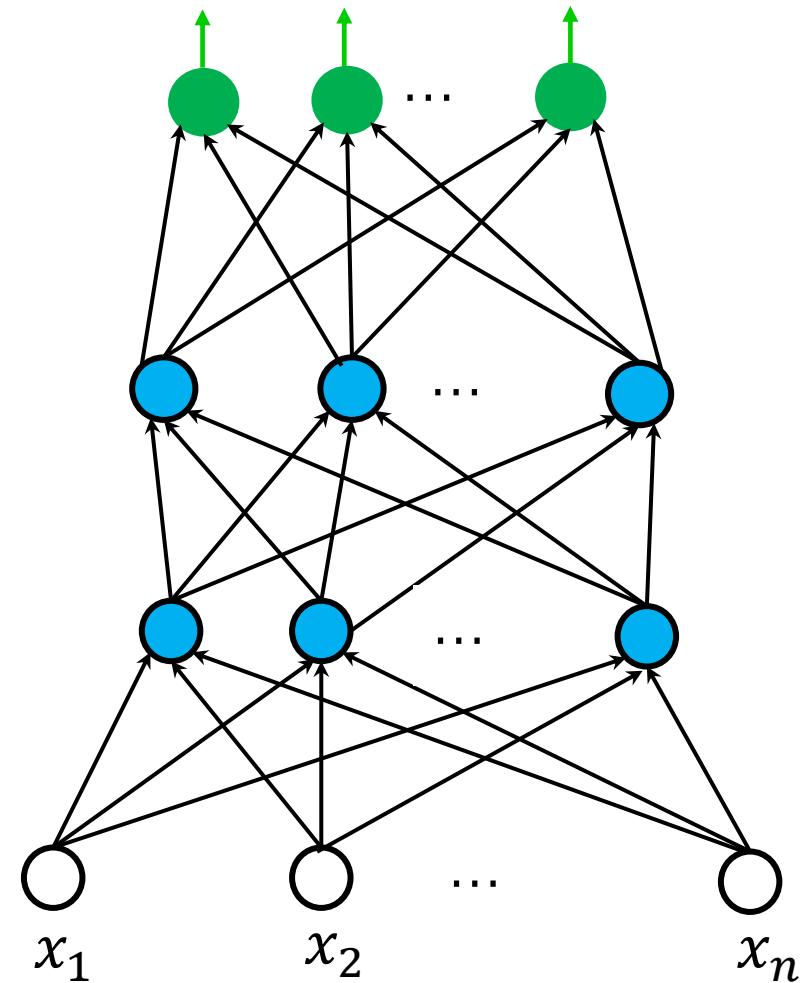
$$f(\mathbf{x}) = f_n(f_{n-1}(\dots f_2(f_1(\mathbf{x})) \dots))$$



神经网络能做什么？



- 表征学习 (Representation learning)
- 执行决策 (Decision making)



为什么神经网络具有普适的应用价值

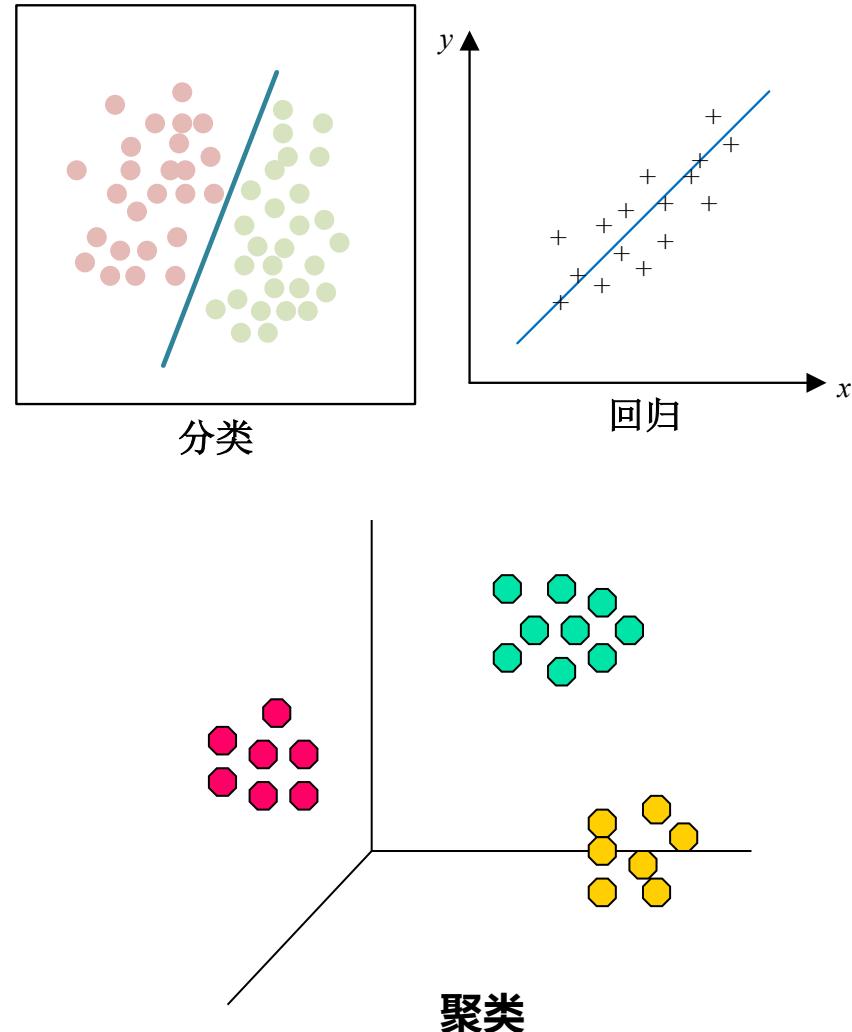
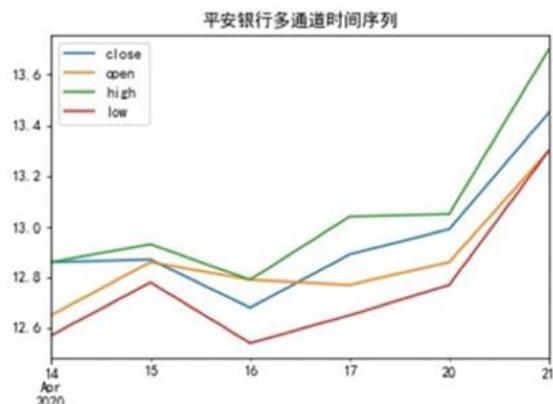
- 事物的向量化表示具有通用性
- 各种应用任务都可转化为函数 $f(x)$ 拟合问题

Tid	Refund	Marital Status	Taxable Income	Cheat
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes



据中央气象台消息，今年第11号台风“轩岚诺”（超强台风级）的中心今天（8月31日）早晨5点钟位于日本冲绳县那霸市偏东方向约340公里的西北太平洋洋面上，就是北纬26.1度、东经131.1度，中心附近最大风力17级以上（62米/秒），中心最低气压为915百帕，七级风圈半径220~230公里，十级风圈半径70公里，十二级风圈半径40公里。

预计，“轩岚诺”将以每小时30公里左右的速度向西偏南转西南方向移动，9月1~2日在琉球群岛以东洋面停滞或回旋，而后转向北偏西方向移动，3日夜间移入东海东南部海面。未来4~5天“轩岚诺”将维持超强台风级的强度，最大强度可达17级以上（62~70米/秒）。



聚类

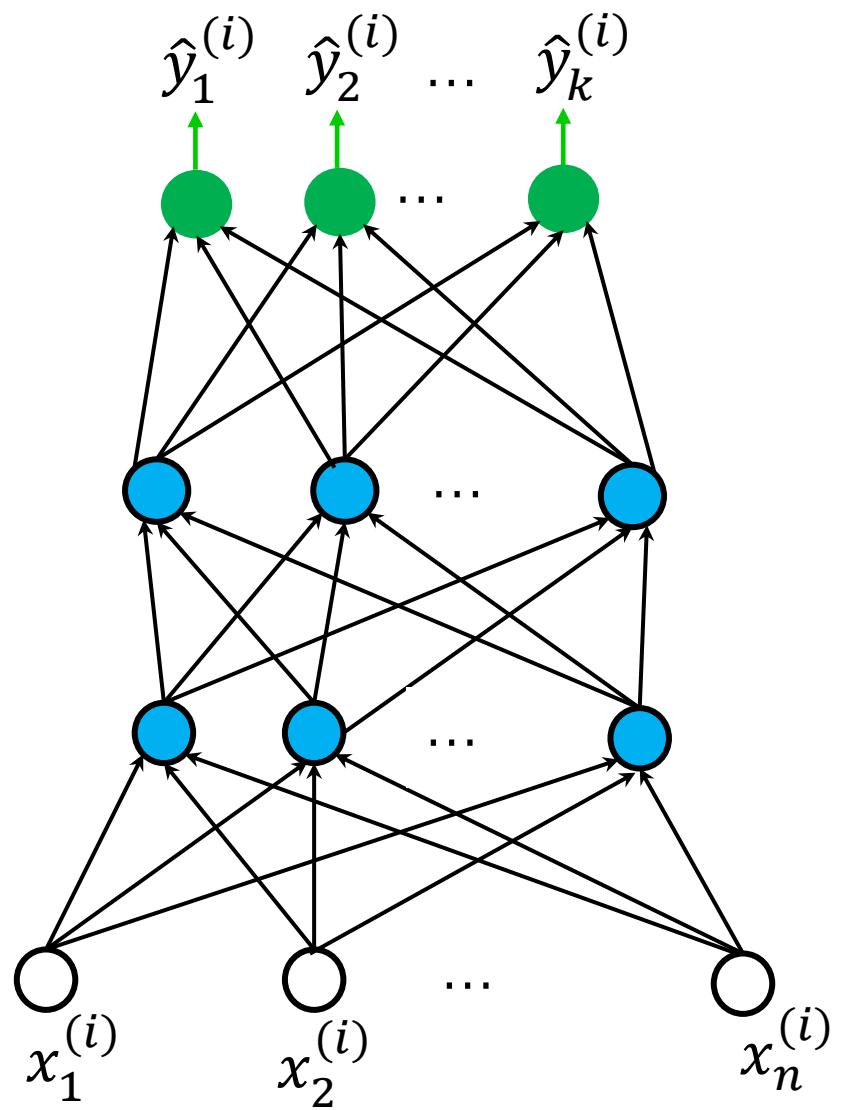
如何学习出神经网络模型： loss function

- Given $\mathbb{X} = \{\langle x^{(1)}, y^{(1)} \rangle, \langle x^{(2)}, y^{(2)} \rangle \dots, \langle x^{(m)}, y^{(m)} \rangle\}$
- Translate $y^{(i)}$ to be a k-dim one-hot vector $\hat{y}^{(i)}$

$$\begin{bmatrix} \mathbf{y}^{(1)} \\ \mathbf{y}^{(2)} \\ \dots \\ \mathbf{y}^{(i)} \\ \dots \\ \mathbf{y}^{(m)} \end{bmatrix} \quad \longleftrightarrow \quad \begin{bmatrix} \hat{\mathbf{y}}^{(1)} \\ \hat{\mathbf{y}}^{(2)} \\ \dots \\ \hat{\mathbf{y}}^{(i)} \\ \dots \\ \hat{\mathbf{y}}^{(m)} \end{bmatrix}$$

Minimize loss:

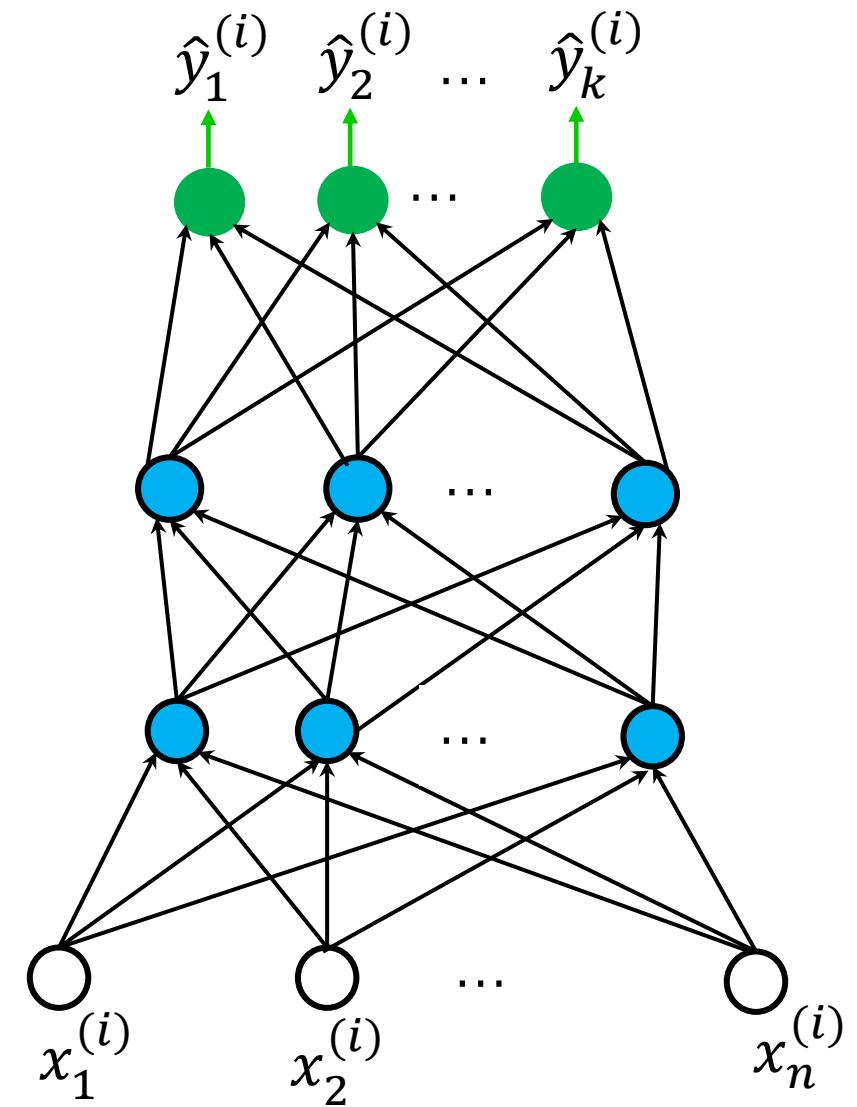
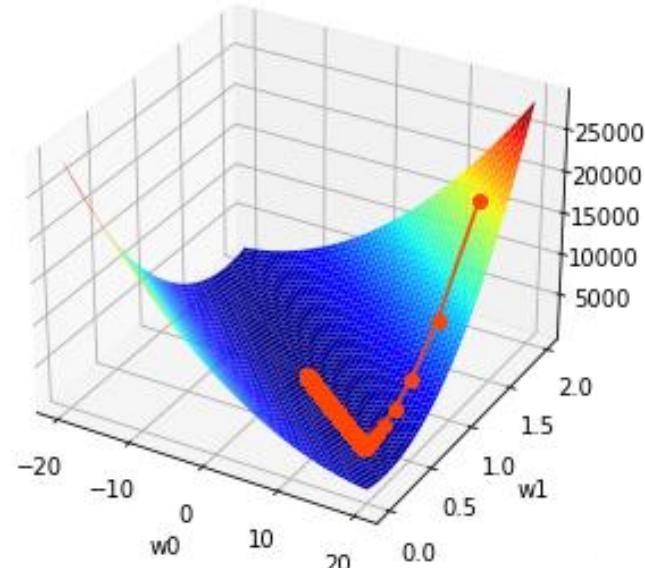
$$L(W, b) = \frac{1}{2} \sum_{i=1}^m \| \mathbf{y}^{(i)} - \hat{\mathbf{y}}^{(i)} \|^2$$



如何学习出神经网络模型：优化问题

$$L(W, b) = \frac{1}{2} \sum_{i=1}^m \| \mathbf{y}^{(i)} - \hat{\mathbf{y}}^{(i)} \|^2$$

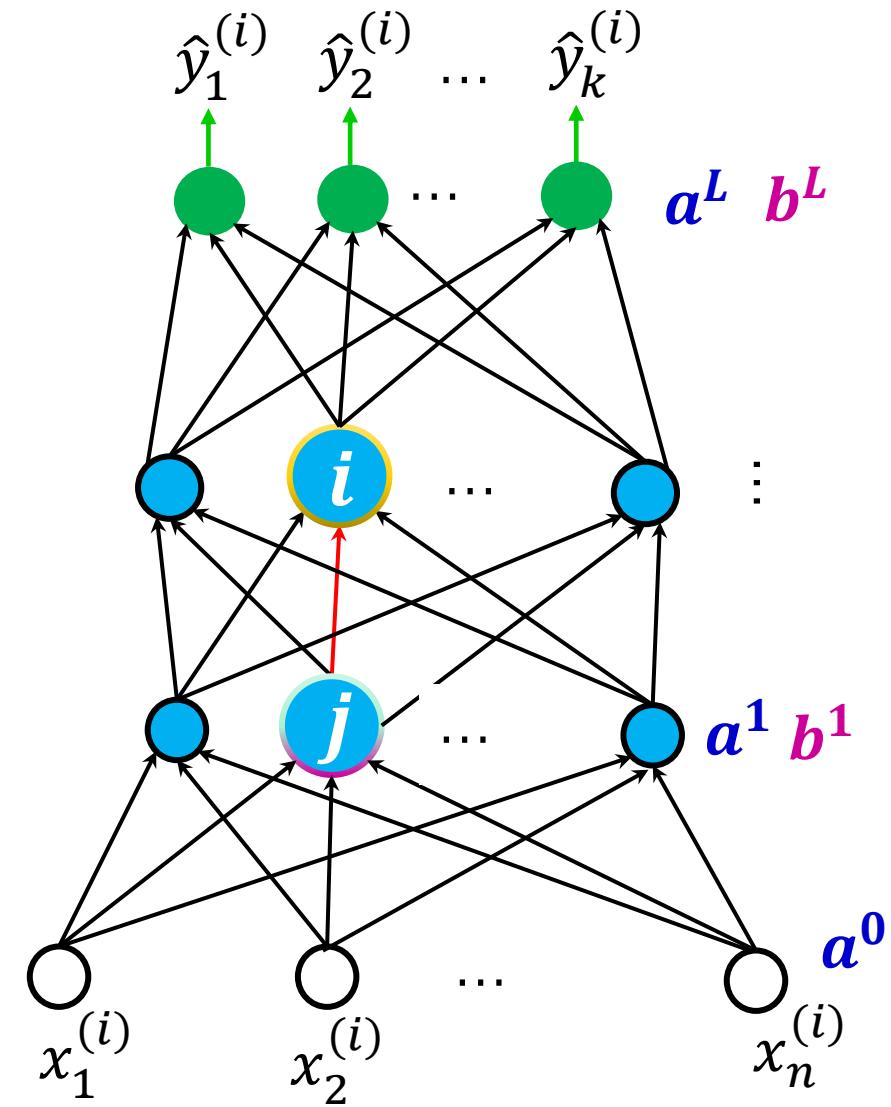
$$\begin{aligned}\hat{\mathbf{y}}^{(i)} &= f(\mathbf{x}^{(i)}) \\ &= \sigma(W^L \sigma(W^{L-1} \dots \sigma(W^2 \sigma(W^1 a^0 + b^1) + b^2) \dots + b^L))\end{aligned}$$



如何学习出神经网络模型：梯度下降法

$$L(W, b) = \frac{1}{2} \sum_{i=1}^m \| \mathbf{y}^{(i)} - \hat{\mathbf{y}}^{(i)} \|^2$$

$$\left[\begin{array}{cccc} w_{11}^1 & \dots & w_{1d^0}^1 \\ \dots & & \dots \\ w_{ij}^1 & \dots & w_{d^1d^0}^1 \\ \dots & & \dots \\ w_{d^1d^0}^1 & \dots & w_{d^1d^0}^1 \end{array} \right] \dots \quad \nabla L = \left[\begin{array}{c} \frac{\partial L}{\partial w_{11}^1} \\ \frac{\partial L}{\partial w_{12}^1} \\ \vdots \\ \frac{\partial L}{\partial w_{ij}^r} \\ \vdots \end{array} \right]$$
$$w_{ij}^{r(new)} = w_{ij}^{r(old)} - \eta \frac{\partial L}{\partial w_{ij}^r}$$



$$L(\mathbf{W}, \mathbf{b}) = \frac{1}{2} \sum_{i=1}^m \| \mathbf{y}^{(i)} - \hat{\mathbf{y}}^{(i)} \|^2$$

$$\frac{\partial L}{\partial w_{ij}^r} = ?$$

$$\hat{\mathbf{y}} = f(\mathbf{x})$$

$$a(\mathbf{W}^L a(\mathbf{W}^{L-1} \dots a(\mathbf{W}^r \dots a(\mathbf{W}^1 \mathbf{x} + \mathbf{b}^1) \dots + \mathbf{b}^r) \dots + \mathbf{b}^L)$$

$$w_{ij}^r \cdot a_j^{r-1} \rightarrow z_i^r \rightarrow a_i^r \dots a^L = \hat{\mathbf{y}} \rightarrow L(.)$$

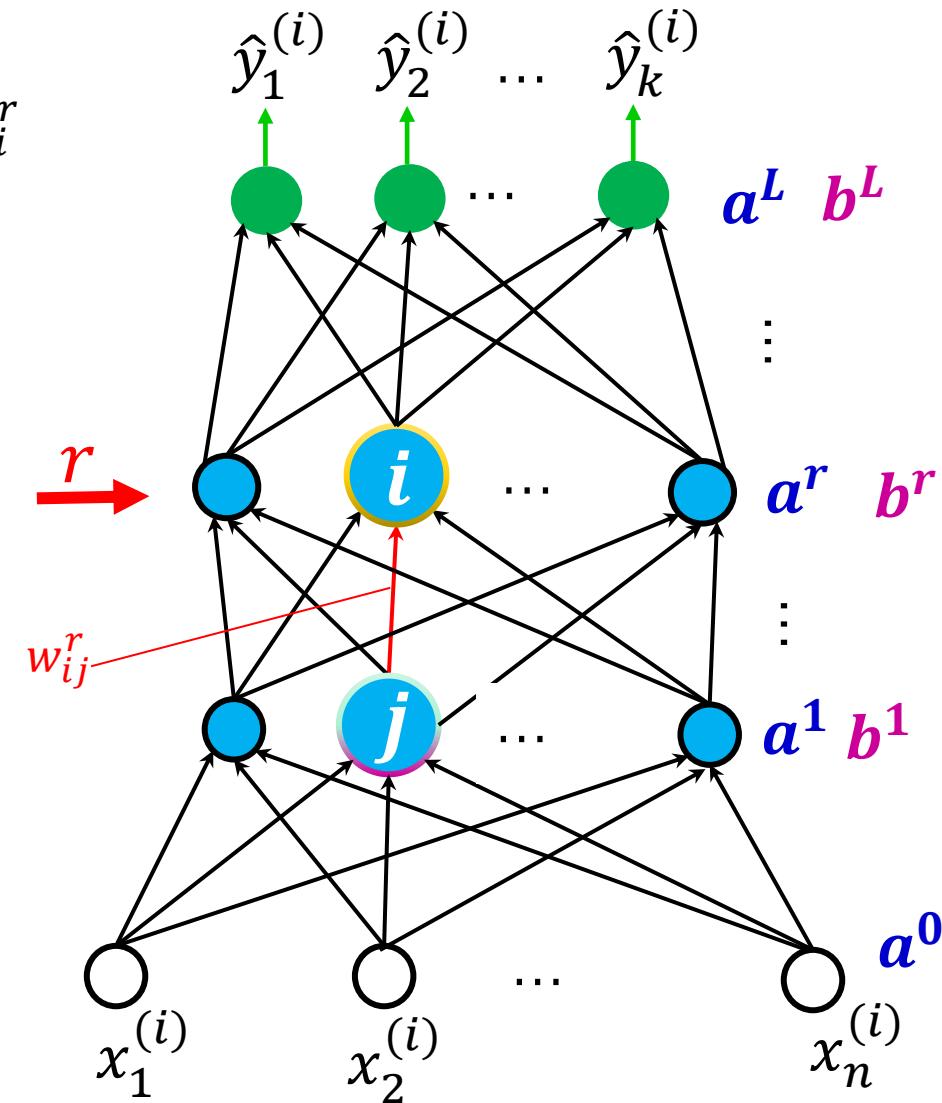
Chain rule for derivation:

$$y = f(u), \quad u = g(x)$$

$$\frac{dy}{dx} = \frac{dy}{du} \frac{du}{dx}$$

$$a_i^r = a(z_i^r):$$

$$z_i^r = \sum_{j=1}^{d^{r-1}} w_{ij}^r \cdot a_j^{r-1} + b_i^r$$



BP (Back Propagation) 算法: basic idea

$$L(\mathbf{W}, \mathbf{b}) = \frac{1}{2} \sum_{i=1}^m \| \mathbf{y}^{(i)} - \hat{\mathbf{y}}^{(i)} \|^2$$

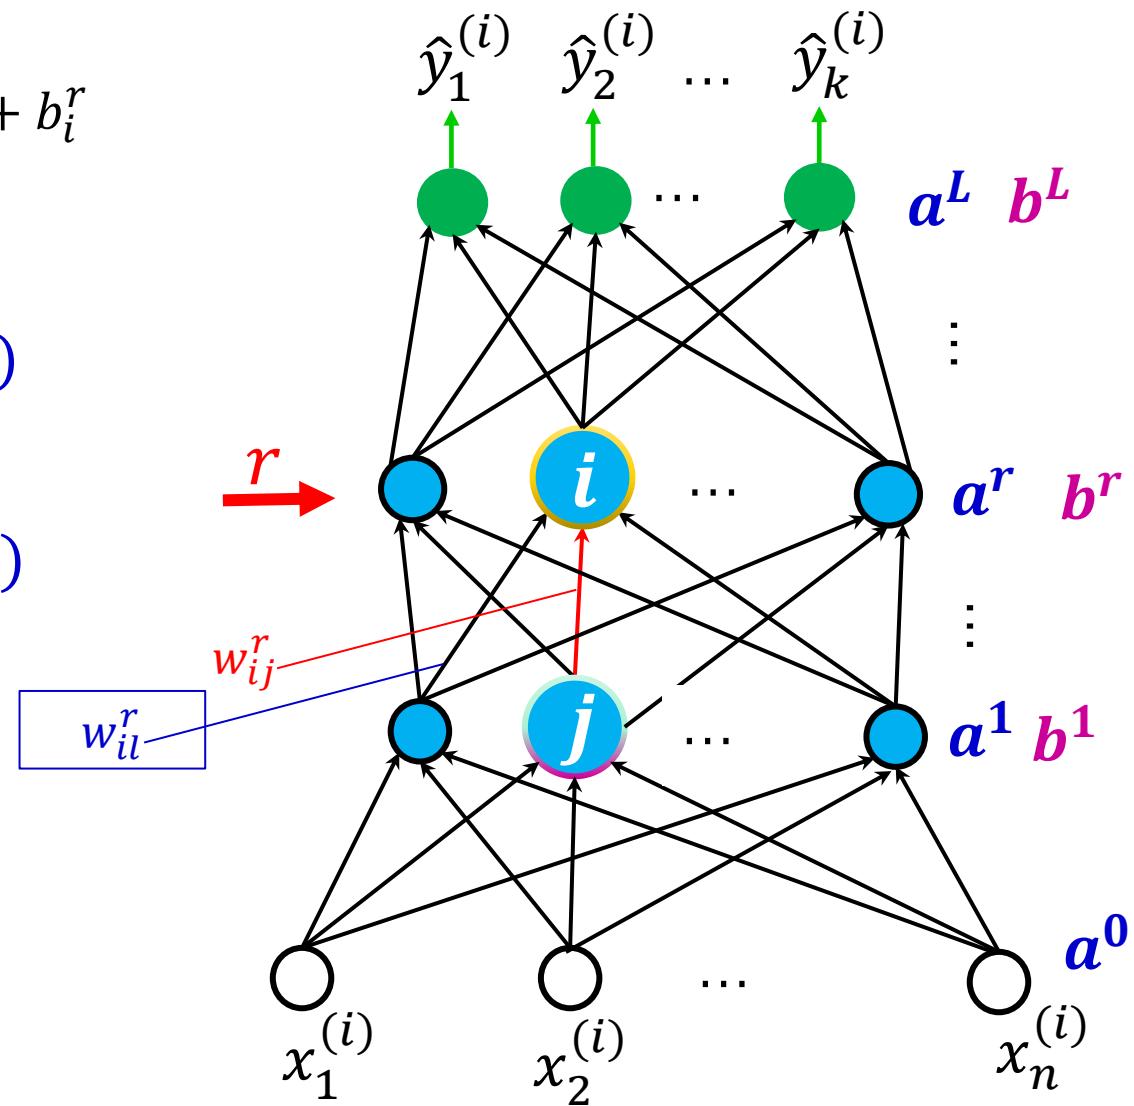
$$z_i^r = \sum_{j=1}^{d^{r-1}} w_{ij}^r \cdot a_j^{r-1} + b_i^r$$

$$w_{ij}^r \cdot a_j^{r-1} \rightarrow z_i^r \rightarrow a_i^r \dots a^L = \hat{\mathbf{y}} \rightarrow L(.)$$

$$w_{il}^r \cdot a_l^{r-1} \rightarrow z_i^r \rightarrow a_i^r \dots a^L = \hat{\mathbf{y}} \rightarrow L(.)$$

$$\frac{\partial L}{\partial w_{ij}^r} = \frac{\partial L}{\partial a^L} \dots \frac{\partial a^{r+1}}{\partial a_i^r} \frac{\partial a_i^r}{\partial z_i^r} \frac{\partial z_i^r}{\partial w_{ij}^r}$$

$$\frac{\partial L}{\partial b_i^r} = \frac{\partial L}{\partial a^L} \dots \frac{\partial a^{r+1}}{\partial a_i^r} \frac{\partial a_i^r}{\partial z_i^r} \frac{\partial z_i^r}{\partial b_i^r}$$



$$L(W, b) = \frac{1}{2} \sum_{i=1}^m \| \mathbf{y}^{(i)} - \hat{\mathbf{y}}^{(i)} \|^2$$

$$z_i^r = \sum_{j=1}^{d^{r-1}} w_{ij}^r \cdot a_j^{r-1} + b_i^r$$

$$\frac{\partial L}{\partial w_{ij}^r} = \underbrace{\frac{\partial L}{\partial \mathbf{a}^L} \cdots \frac{\partial \mathbf{a}^{r+1}}{\partial a_i^r}}_{\text{red bracket}} \frac{\partial a_i^r}{\partial z_i^r} \frac{\partial z_i^r}{\partial w_{ij}^r}$$

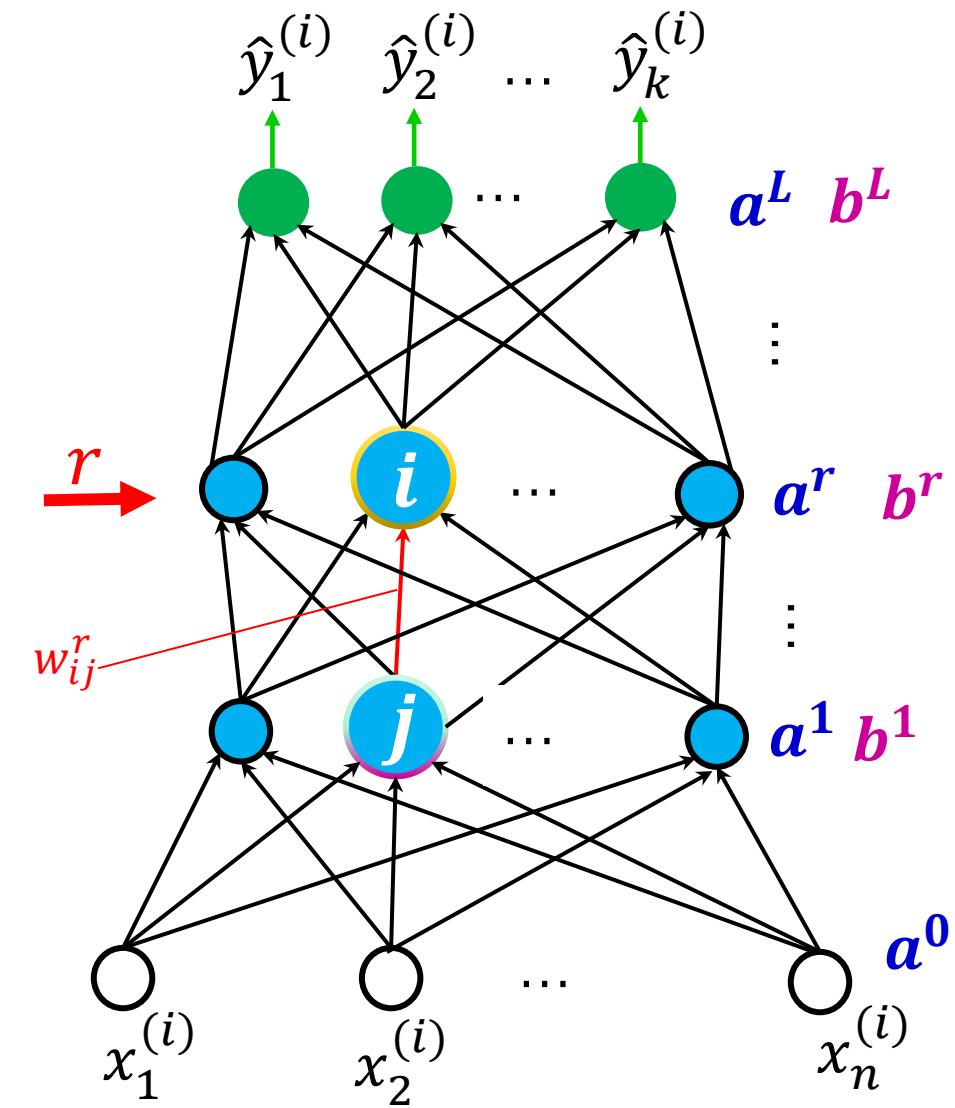
$$\delta_i^r = \frac{\partial L}{\partial z_i^r}$$

$$\frac{\partial L}{\partial w_{ij}^r} = \delta_i^r \times a_j^{r-1}$$

$$\frac{\partial z_i^r}{\partial w_{ij}^r} = a_j^{r-1}$$

$$\frac{\partial a_i^r}{\partial z_i^r} : \quad \downarrow \quad a: \quad \sigma = \frac{1}{1 + e^{-z}}$$

$$\frac{\partial a_i^r}{\partial z_i^r} = a_i^r (1 - a_i^r)$$



$$L(\mathbf{W}, \mathbf{b}) = \frac{1}{2} \sum_{i=1}^m \| \mathbf{y}^{(i)} - \hat{\mathbf{y}}^{(i)} \|^2$$

$$z_i^r = \sum_{j=1}^{d^{r-1}} w_{ij}^r \cdot a_j^{r-1} + b_i^r$$

$$\frac{\partial L}{\partial w_{ij}^r} = \underbrace{\frac{\partial L}{\partial \mathbf{a}^L} \cdots \frac{\partial \mathbf{a}^{r+1}}{\partial a_i^r}}_{\delta_i^r} \frac{\partial a_i^r}{\partial z_i^r} \frac{\partial z_i^r}{\partial w_{ij}^r}$$

$$\delta_i^r = \frac{\partial L}{\partial z_i^r}$$

$$\frac{\partial L}{\partial w_{ij}^r} = \delta_i^r \times a_j^{r-1}$$

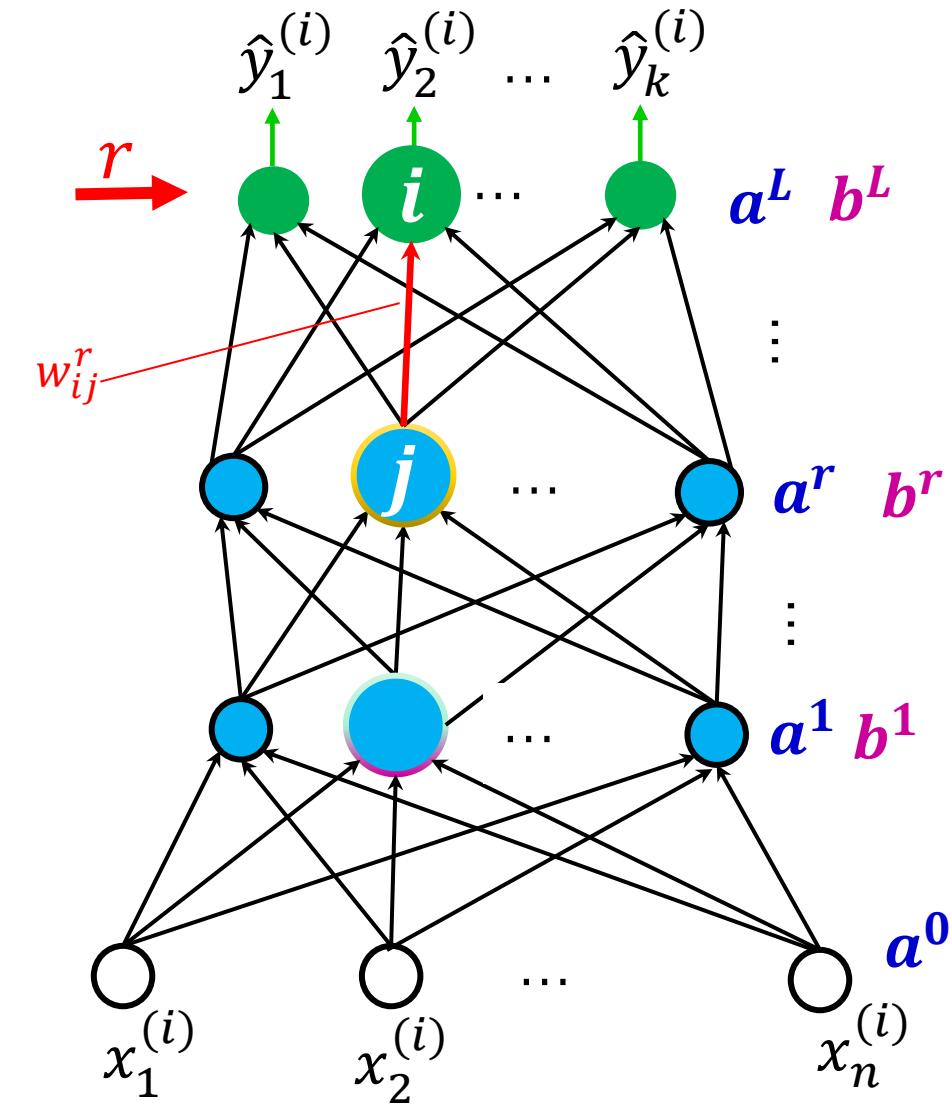
Output Layer($r = L$):

$$\delta_i^r = -(y_i - a_i^L) a_i^L (1 - a_i^L)$$

$$\frac{\partial z_i^r}{\partial w_{ij}^r} = a_j^{r-1}$$

$$\frac{\partial a_i^r}{\partial z_i^r} : \quad a: \sigma = \frac{1}{1 + e^{-z}}$$

$$\frac{\partial a_i^r}{\partial z_i^r} = a_i^r (1 - a_i^r)$$



$$L(\mathbf{W}, \mathbf{b}) = \frac{1}{2} \sum_{i=1}^m \| \mathbf{y}^{(i)} - \hat{\mathbf{y}}^{(i)} \|^2$$

$$z_i^r = \sum_{j=1}^{d^{r-1}} w_{ij}^r \cdot a_j^{r-1} + b_i^r$$

$$\frac{\partial L}{\partial w_{ij}^r} = \underbrace{\frac{\partial L}{\partial \mathbf{a}^L} \cdots \frac{\partial \mathbf{a}^{r+1}}{\partial a_i^r}}_{\delta_i^r} \frac{\partial a_i^r}{\partial z_i^r} \frac{\partial z_i^r}{\partial w_{ij}^r}$$

$$\delta_i^r = \frac{\partial L}{\partial z_i^r}$$

$$\frac{\partial L}{\partial w_{ij}^r} = \delta_i^r \times a_j^{r-1}$$

Output Layer ($r = L$):

$$\delta_i^r = -(y_i - a_i^L) a_i^L (1 - a_i^L)$$

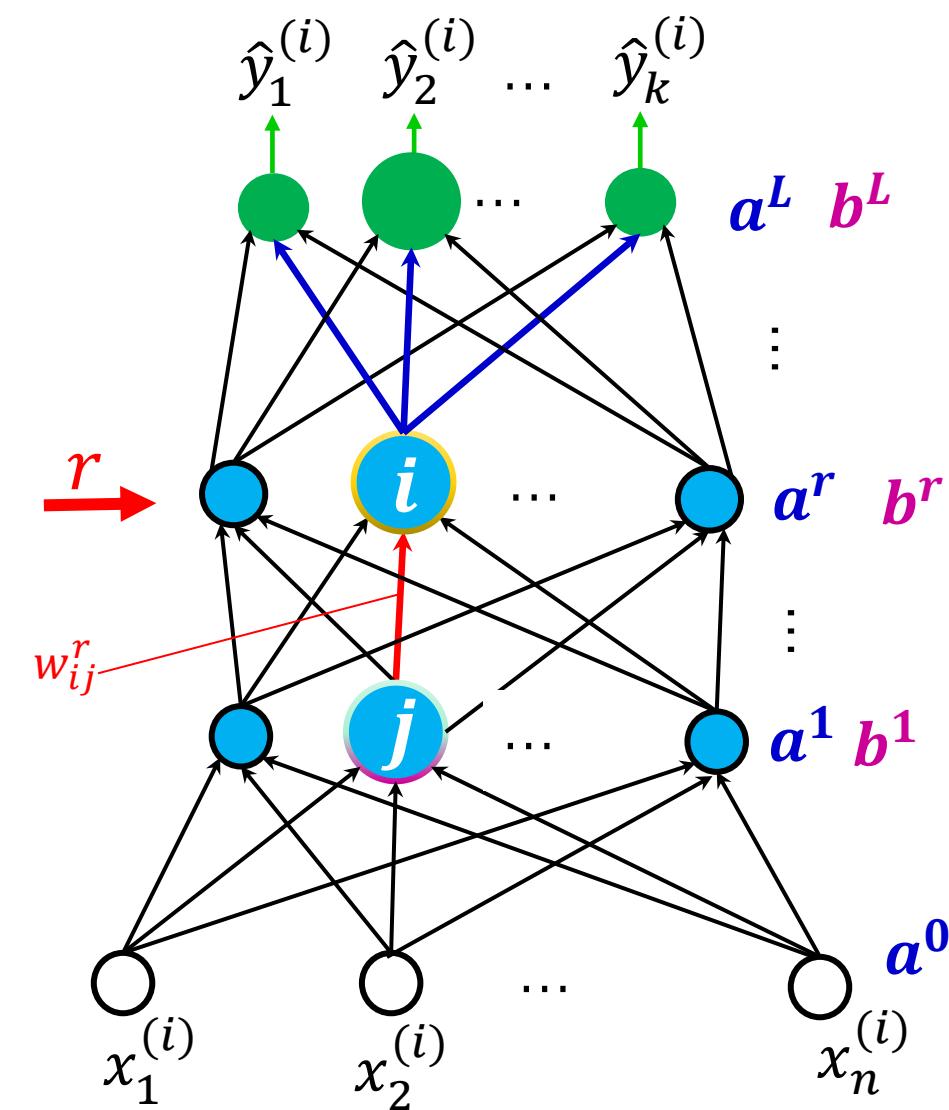
Hidden Layer ($r < L$):

$$\delta_i^r = a_i^r (1 - a_i^r) \sum_{s \in \text{Next}(i)} \delta_s^{r+1} w_{si}^{r+1}$$

$$\frac{\partial z_i^r}{\partial w_{ij}^r} = a_j^{r-1}$$

$$\frac{\partial a_i^r}{\partial z_i^r} : \quad a: \sigma = \frac{1}{1 + e^{-z}}$$

$$\frac{\partial a_i^r}{\partial z_i^r} = a_i^r (1 - a_i^r)$$



批处理反向传播 (BP with Batch Training)

- Given $\mathbb{X} = \{\langle \mathbf{x}^{(1)}, y^{(1)} \rangle, \langle \mathbf{x}^{(2)}, y^{(2)} \rangle \dots, \langle \mathbf{x}^{(m)}, y^{(m)} \rangle\}$, and hyper parameters

➤ Translate $y^{(i)}$ to be a k-dim one-hot vector $\mathbf{y}^{(i)}$ for classification task

1. Initialize network weights with small random values
2. **Do**

① Initialize Loss: $L(\mathbf{W}, \mathbf{b}) = 0$

$$L(\mathbf{W}, \mathbf{b}) = \frac{1}{2} \sum_{i=1}^m \| \mathbf{y}^{(i)} - \hat{\mathbf{y}}^{(i)} \|^2$$

② For each sample $\langle \mathbf{x}^{(N)}, y^{(N)} \rangle$

a. Compute the output of $\hat{\mathbf{y}}^{(N)} = f(\mathbf{x}^{(N)})$

$$\frac{\partial L}{\partial w_{ij}^r} = \delta_i^r \times a_j^{r-1}$$

b. Update the loss: $L(\mathbf{W}, \mathbf{b}) = L(\mathbf{W}, \mathbf{b}) + \frac{1}{2} (y^{(N)} - \hat{y}^{(N)})^2$

$$\delta_i^r = -(y_i - a_i^L) a_i^L (1 - a_i^L)$$

③ Calculate the gradient with $L(\mathbf{W}, \mathbf{b})$ and current weights $w_{ij}^{r(0)}$

$$\delta_i^r = a_i^r (1 - a_i^r) \sum_{s \in Next(i)} \delta_s^{r+1} w_{si}^{r+1}$$

④ Update the weights: $w_{ij}^{r(new)} = w_{ij}^{r(old)} - \eta \frac{\partial L}{\partial w_{ij}^r}$

3. Until stopping criteria satisfied

随机梯度下降算法 (BP with SGD)

- Given $\mathbb{X} = \{\langle \mathbf{x}^{(1)}, y^{(1)} \rangle, \langle \mathbf{x}^{(2)}, y^{(2)} \rangle \dots, \langle \mathbf{x}^{(m)}, y^{(m)} \rangle\}$, and hyper parameters

➤ Translate $y^{(i)}$ to be a k-dim one-hot vector $\mathbf{y}^{(i)}$ for classification task

1. Initialize network weights with small random values
2. **Do**

① **For each sample** $\langle \mathbf{x}^{(N)}, y^{(N)} \rangle$

a. **Compute the output of** $\hat{\mathbf{y}}^{(N)} = f(\mathbf{x}^{(N)})$

b. **Compute the loss of** $\mathbf{x}^{(N)}$: $l(W, b) = \frac{1}{2} (y^{(N)} - \hat{\mathbf{y}}^{(N)})^2$

c. **Calculate the gradient with** $l(W, b)$ **and current weights** $w_{ij}^{r(\text{old})}$

d. **Update the weights:** $w_{ij}^{r(\text{new})} = w_{ij}^{r(\text{old})} - \eta \frac{\partial L}{\partial w_{ij}^r}$

3. **Until stopping criteria satisfied**

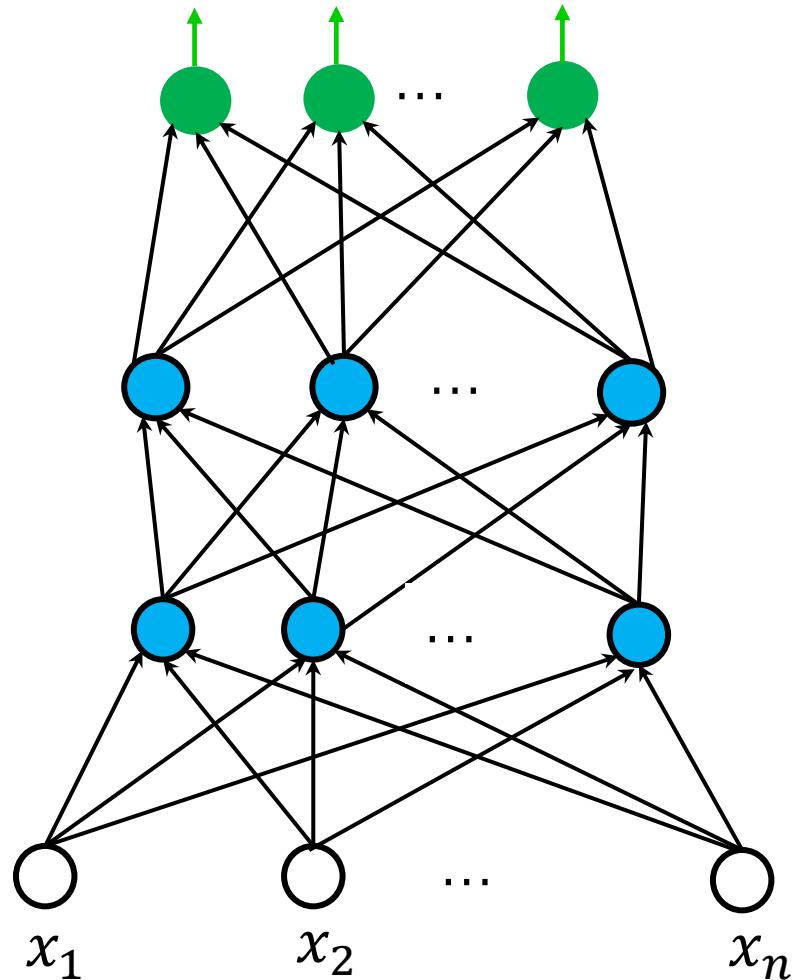
$$L(W, b) = \frac{1}{2} \sum_{i=1}^m \| \mathbf{y}^{(i)} - \hat{\mathbf{y}}^{(i)} \|^2$$

$$\frac{\partial L}{\partial w_{ij}^r} = \delta_i^r \times a_j^{r-1}$$

$$\delta_i^r = -(y_i - a_i^L) a_i^L (1 - a_i^L)$$

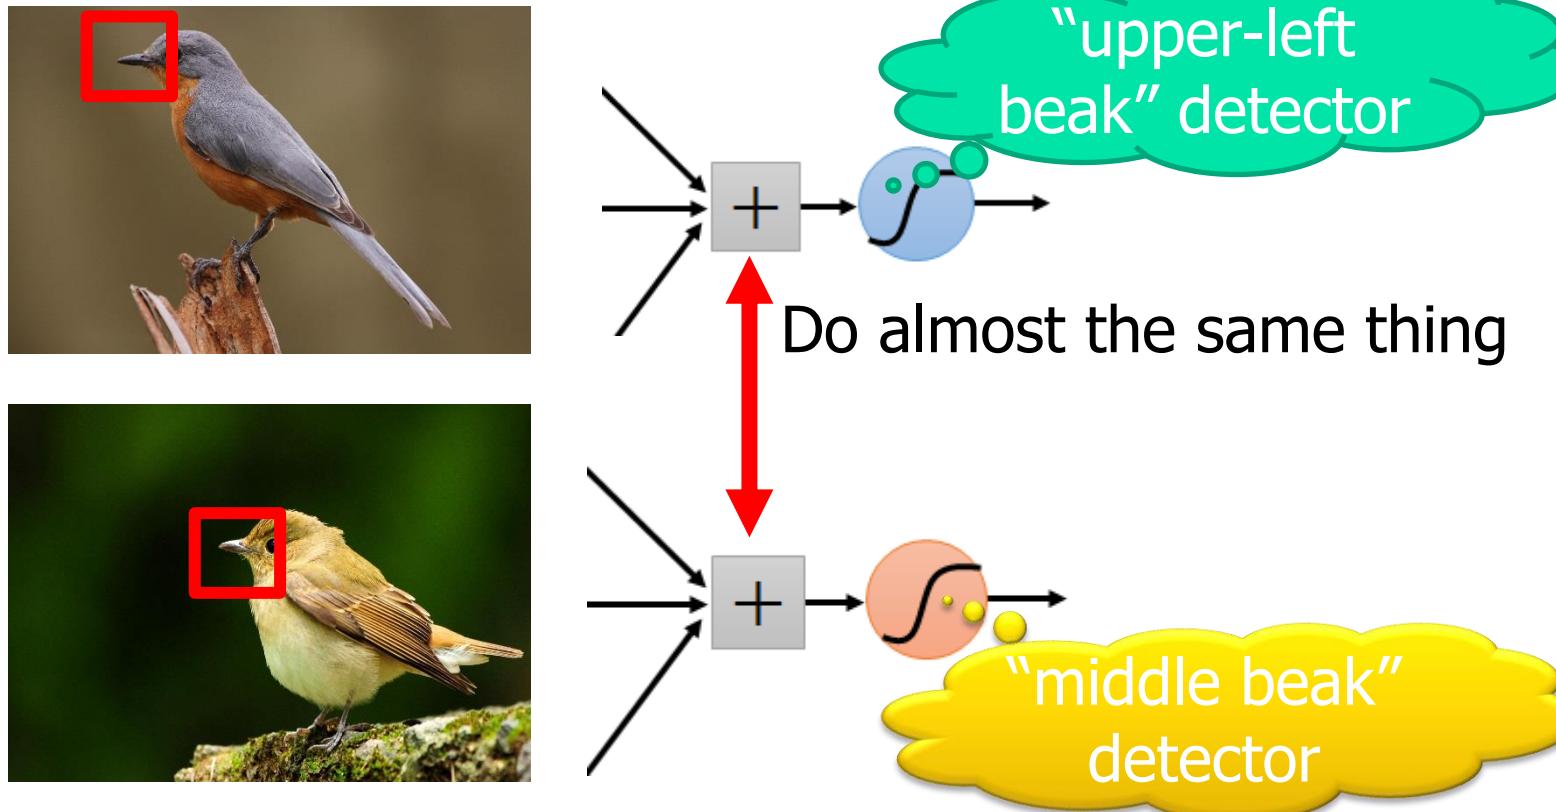
$$\delta_i^r = a_i^r (1 - a_i^r) \sum_{s \in \text{Next}(i)} \delta_s^{r+1} w_{si}^{r+1}$$

全连接网络的问题



卷积神经网络 (CNN) : 可挖掘局部模式

- The same patterns appear in different regions.



卷积神经网络 (CNN) : 可挖掘局部模式

- Subsampling the pixels will not change the object

bird



subsampling

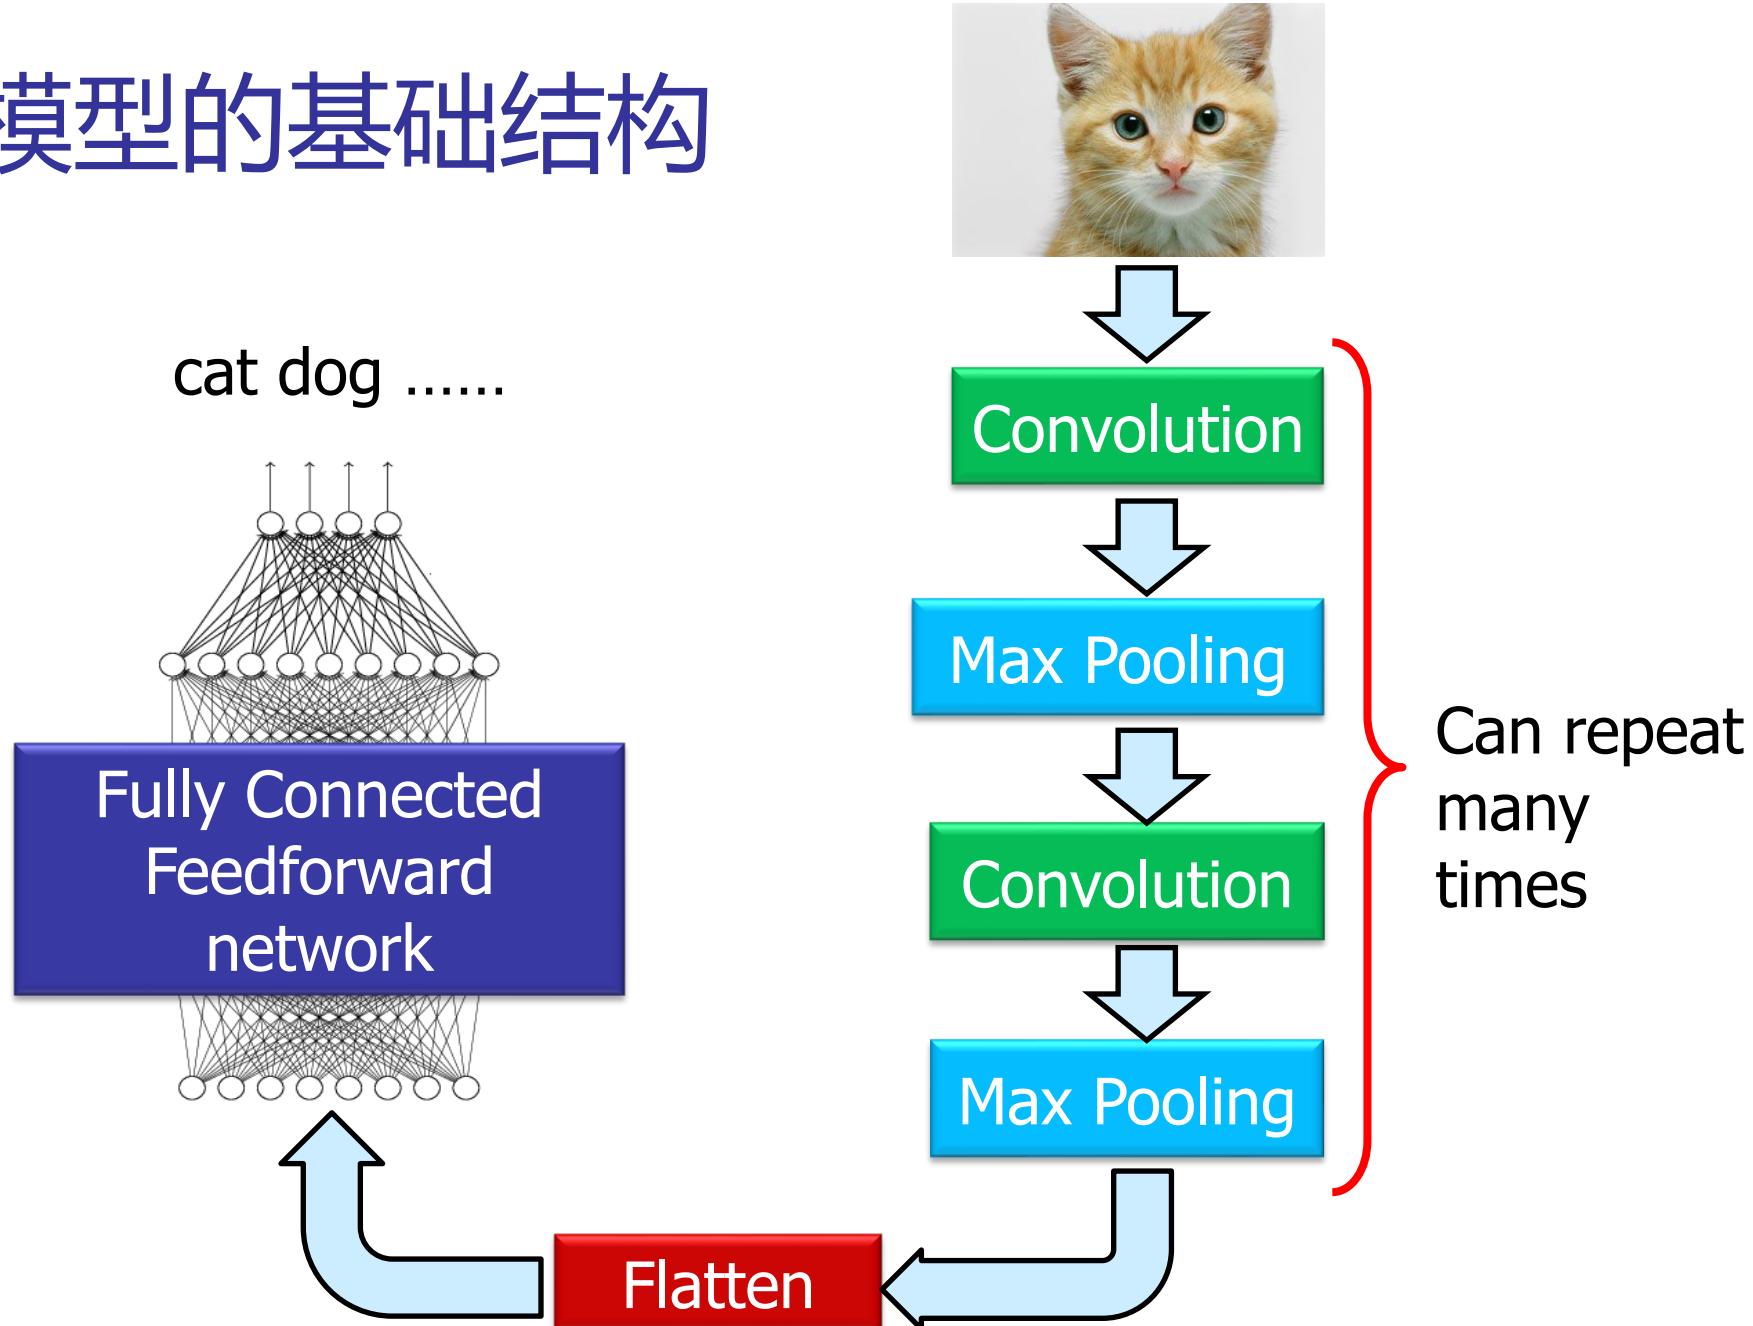
bird



We can subsample the pixels to make image smaller

Less parameters for the network to process the image

CNN模型的基础结构



CNN – 卷积 (Convolution)

Those are the network parameters to be learned.

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1
Matrix

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2
Matrix

⋮ ⋮

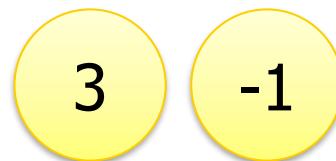
Each filter detects a small pattern (3 x 3).

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0



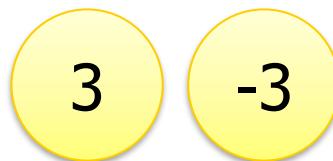
6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

If stride=2

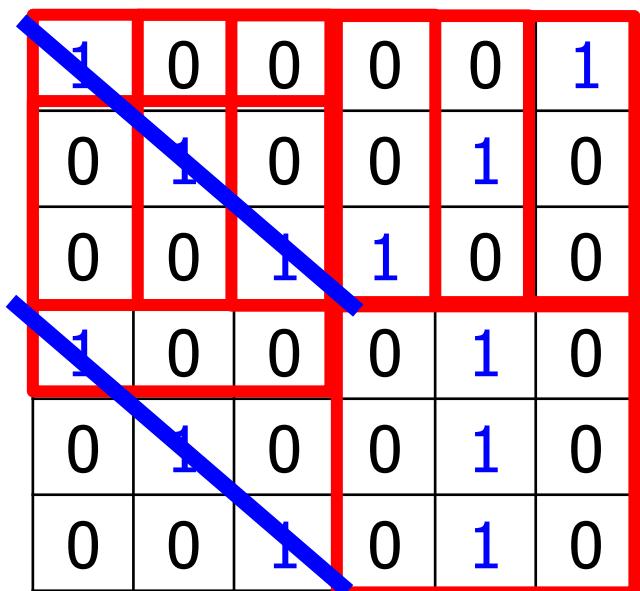
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0



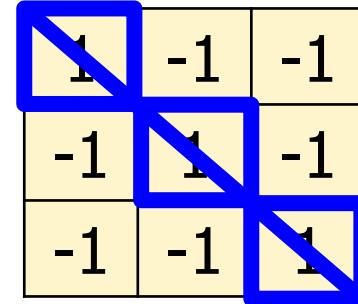
We set stride=1 below

6 x 6 image

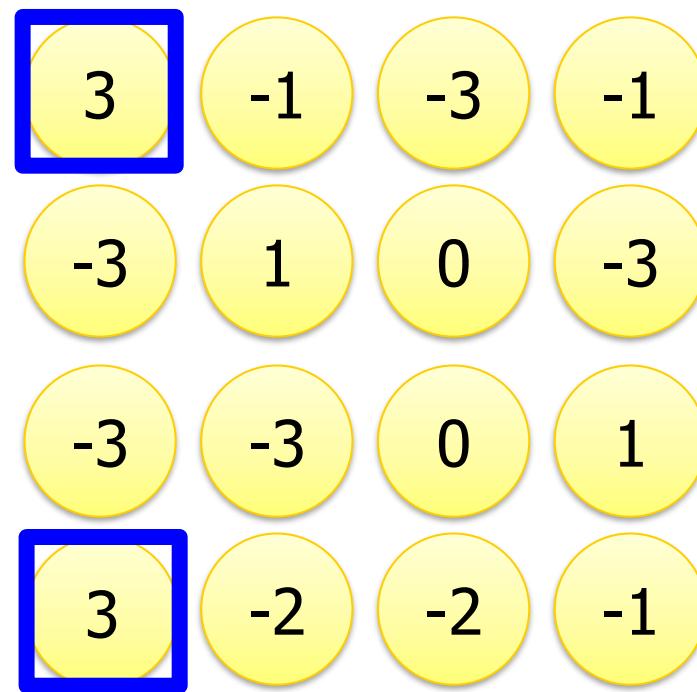
stride=1



6 x 6 image



Filter 1



stride=1

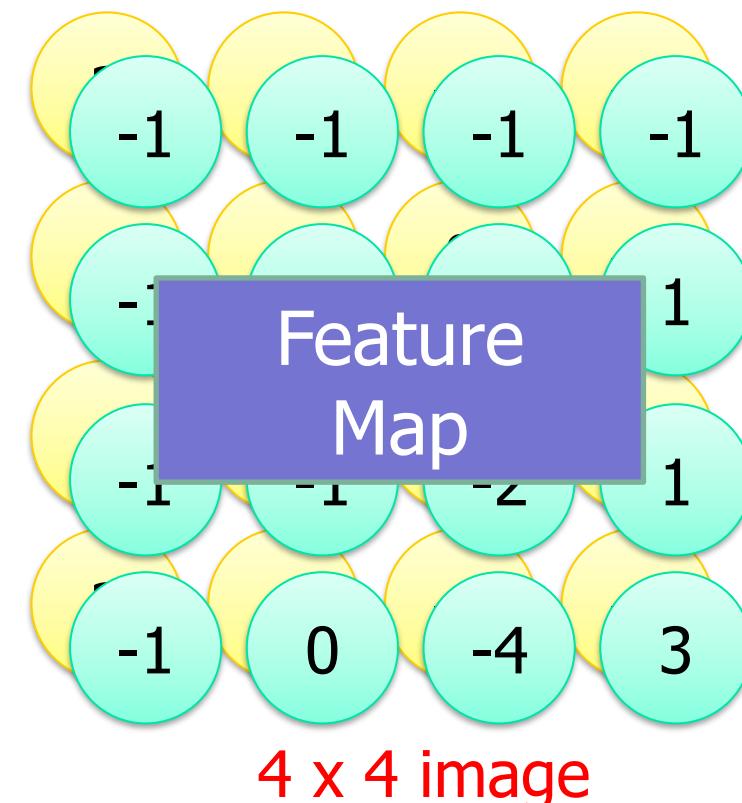
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

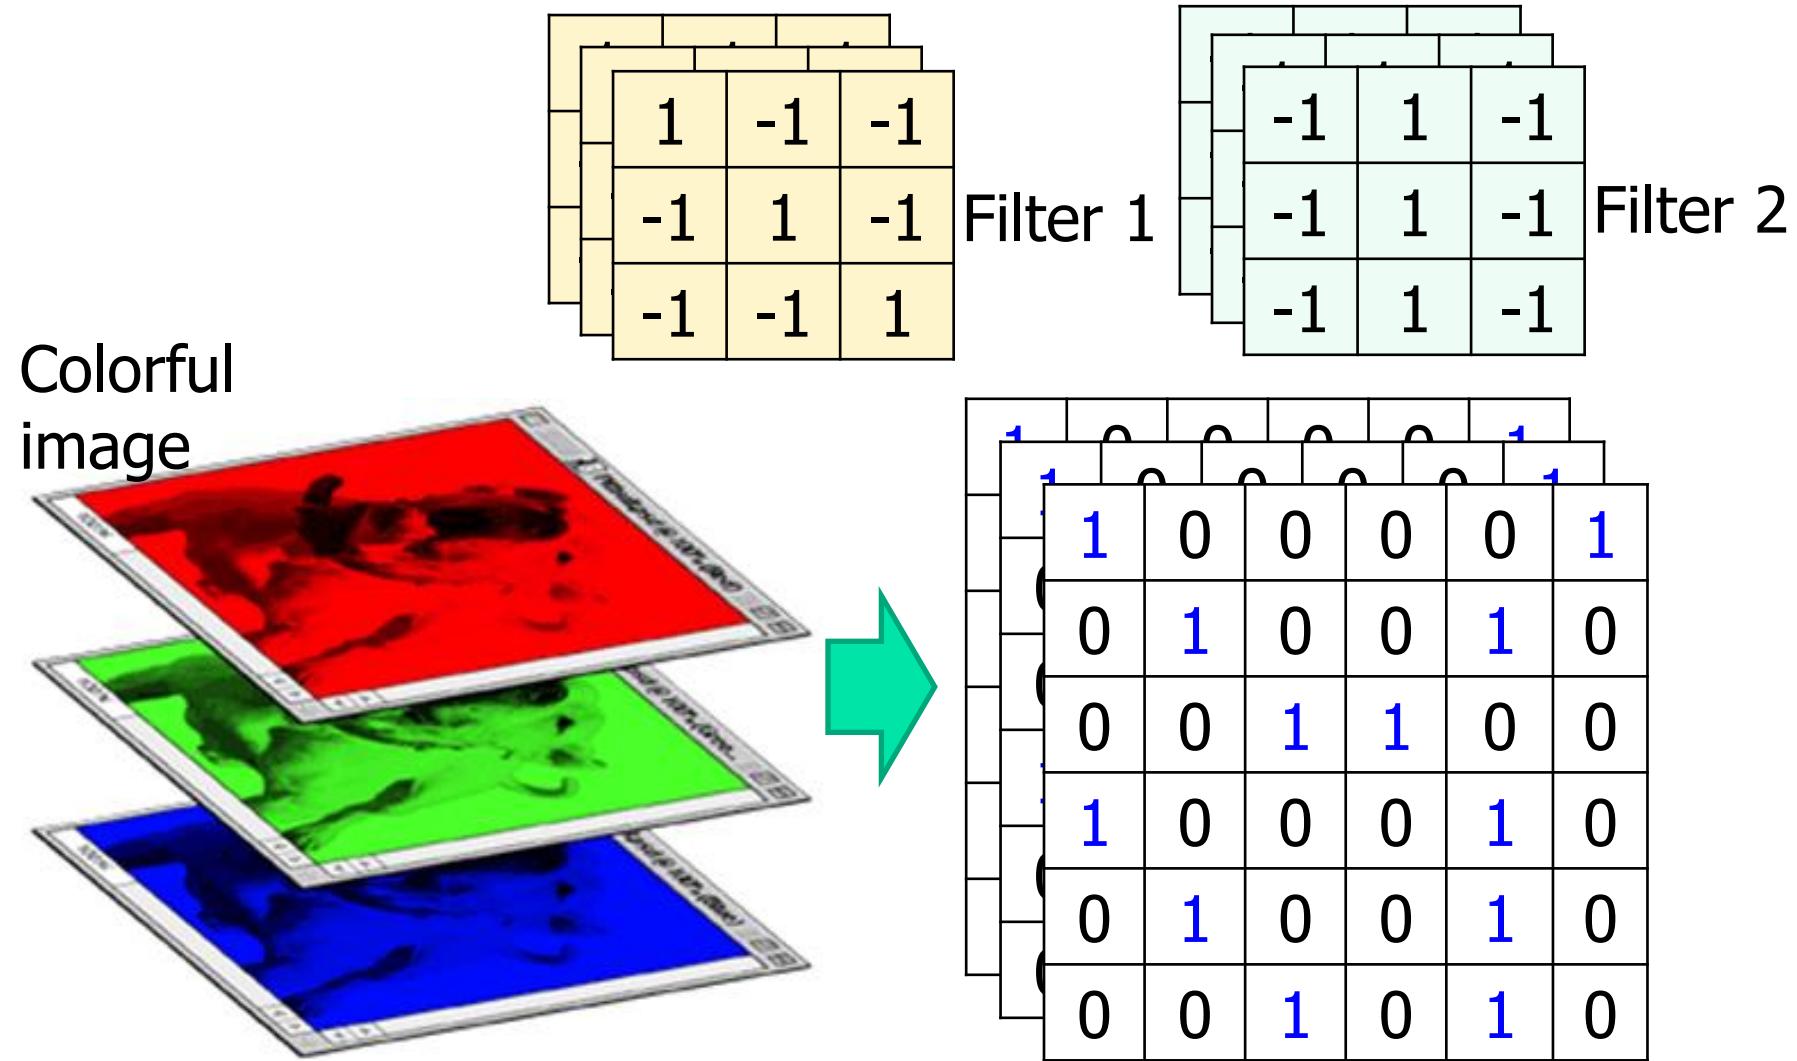
-1	1	-1
-1	1	-1
-1	1	-1

Filter 2

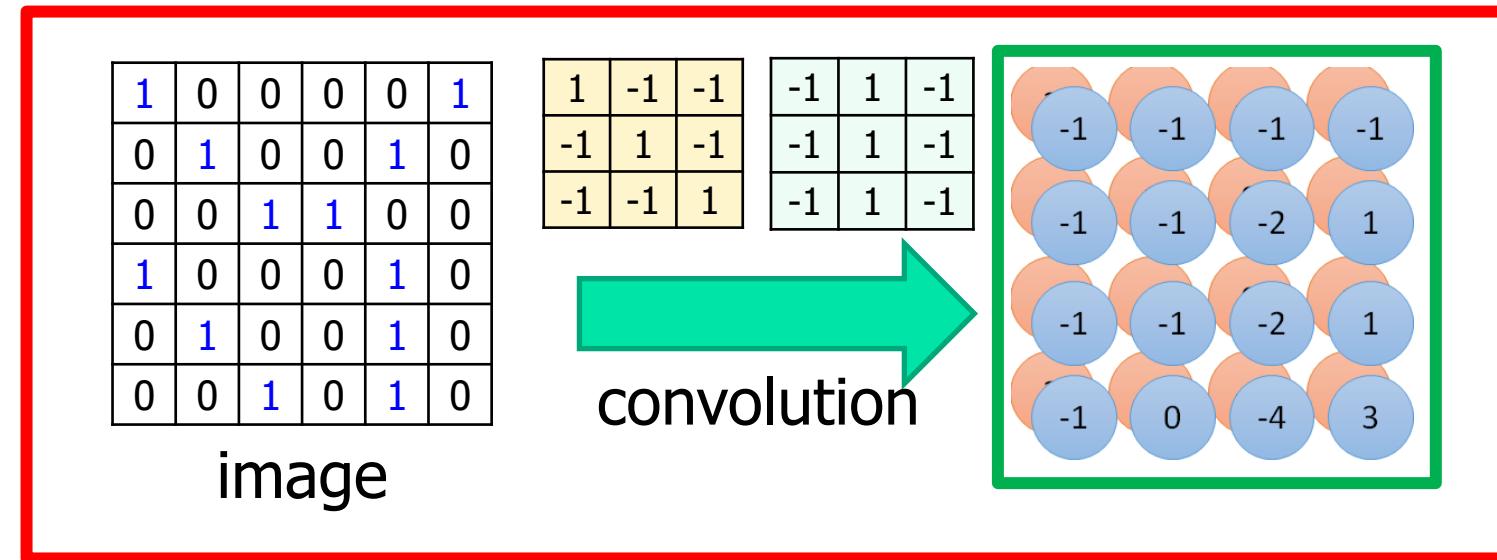
Do the same process
for every filter



CNN – Colorful image

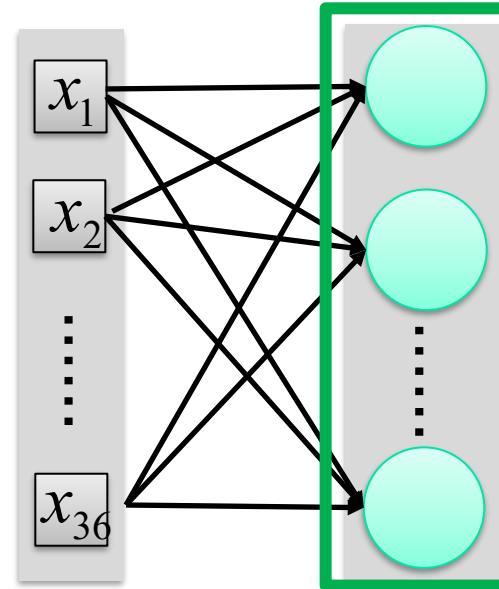


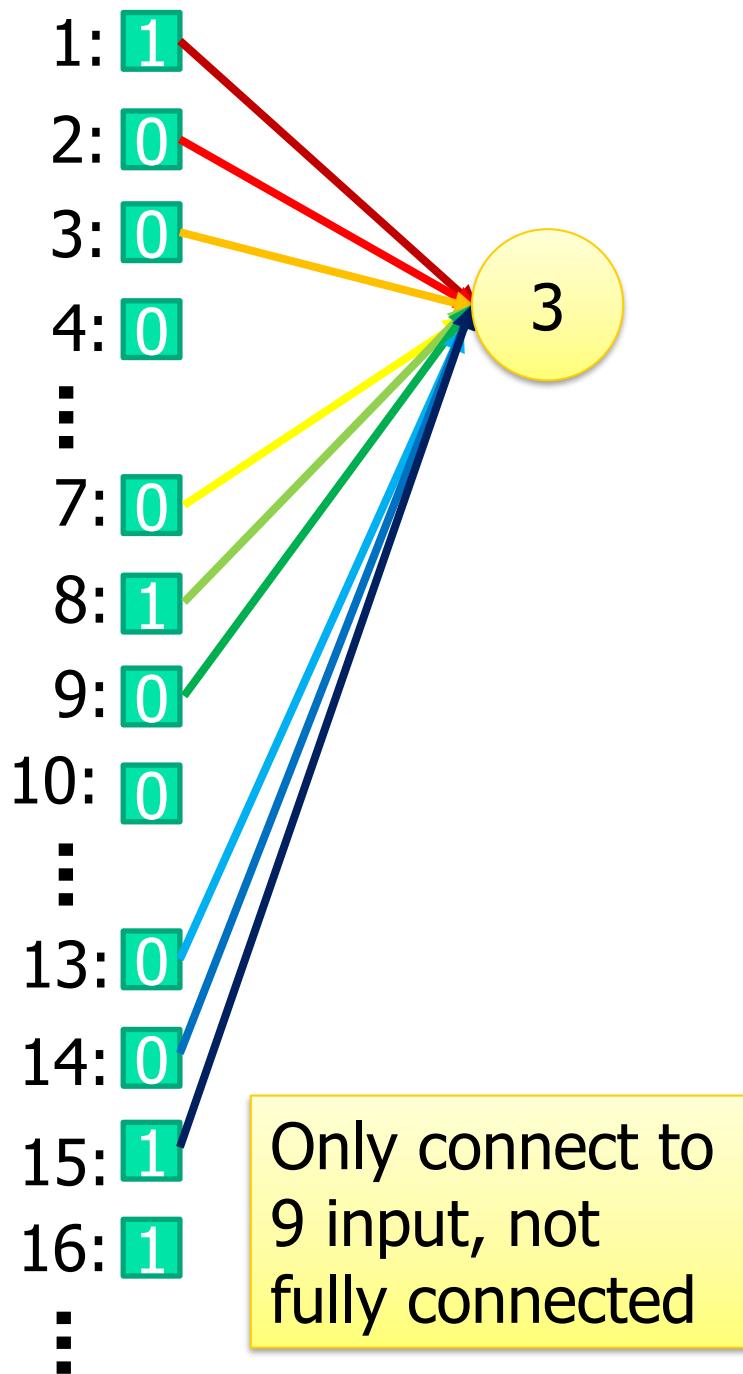
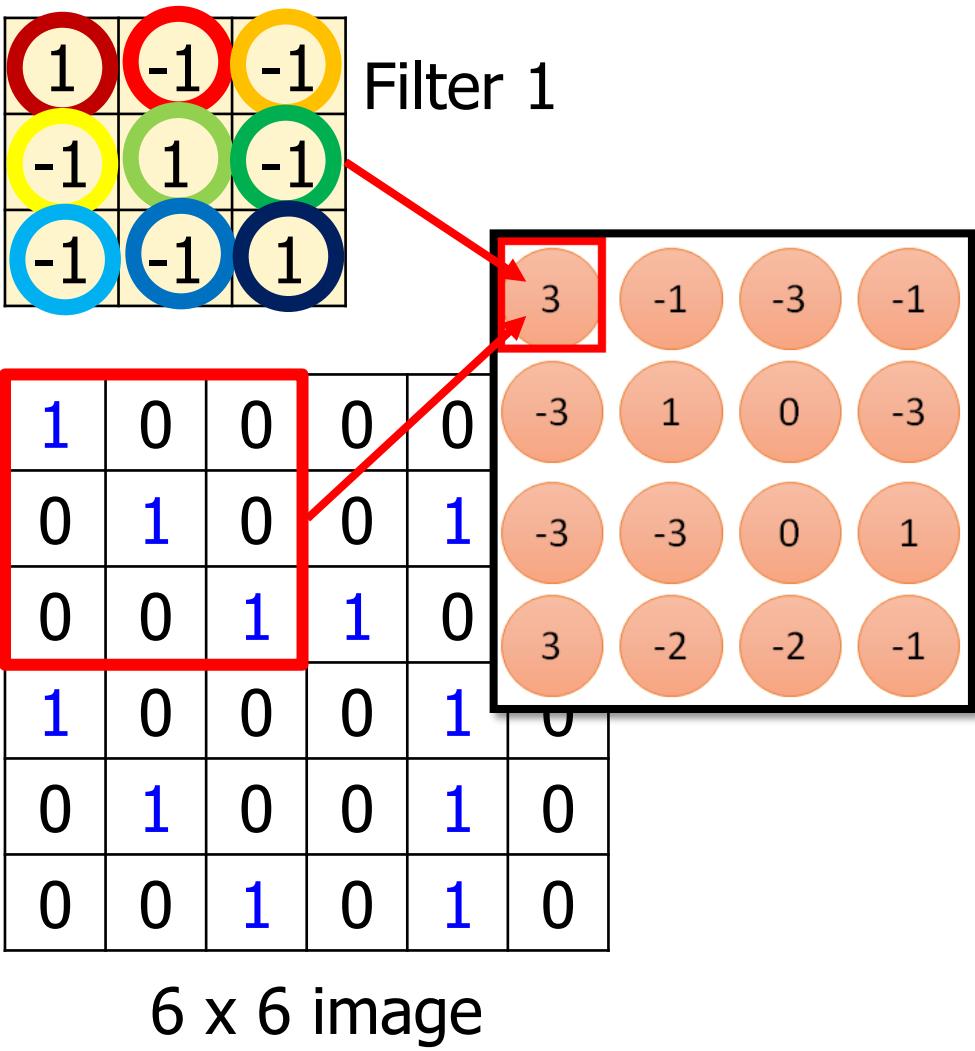
Convolution v.s. Fully Connected

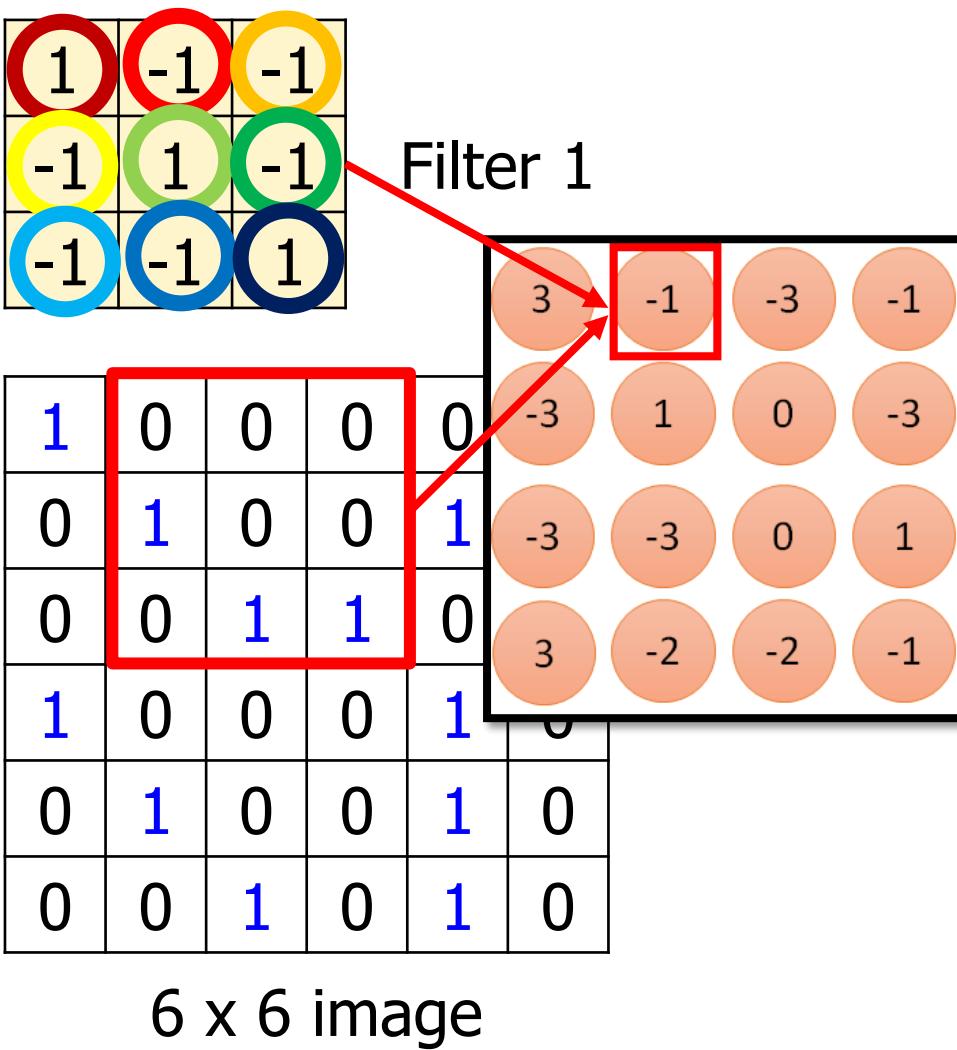


Fully-connected

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

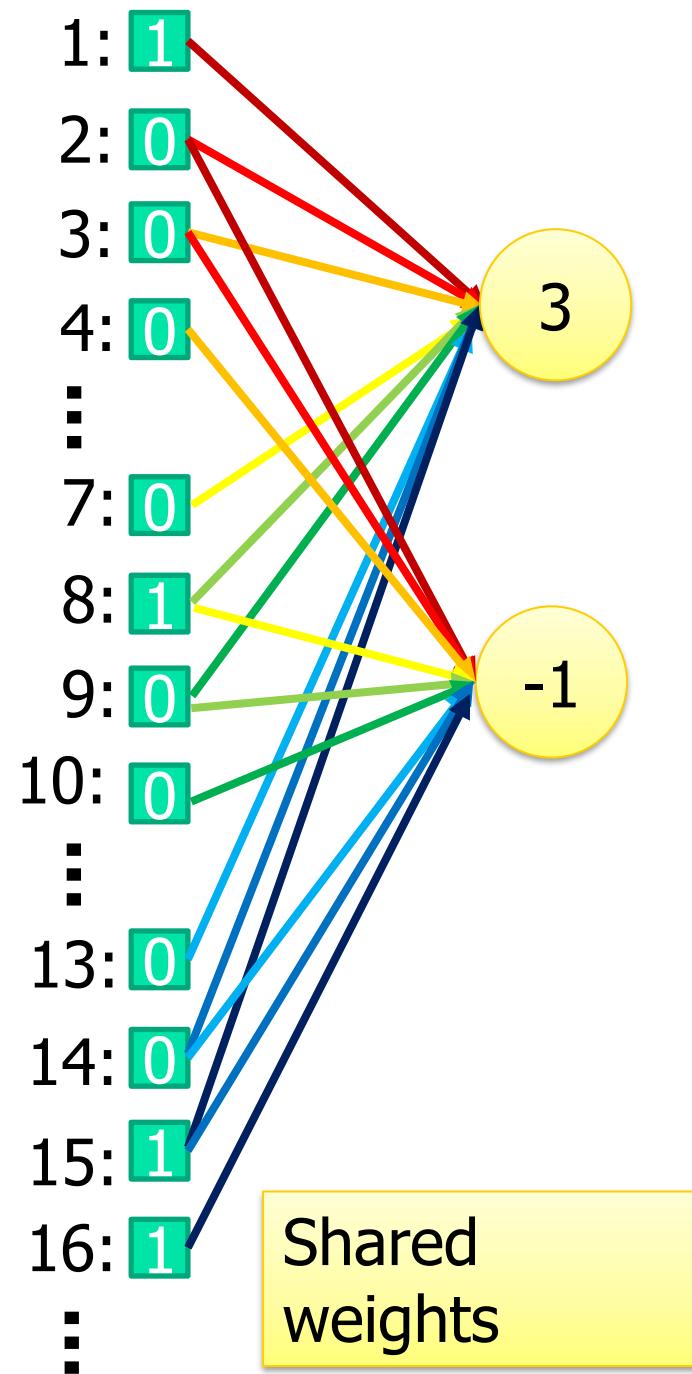






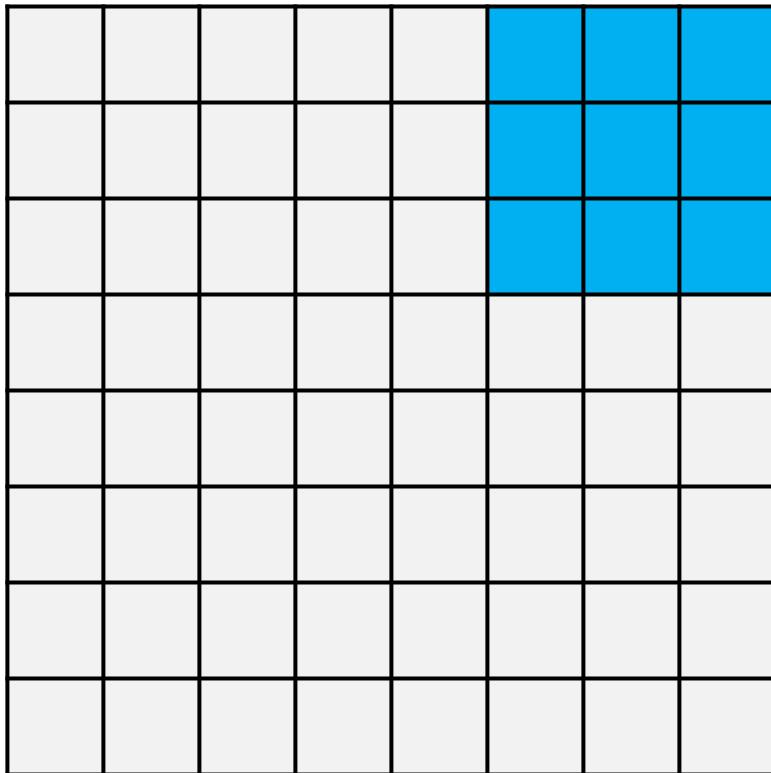
Less parameters!

Even less parameters!

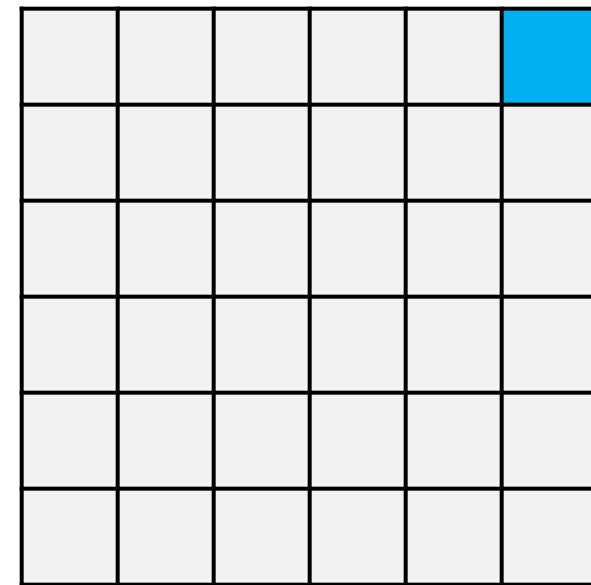


Convolution: 填充 (padding)

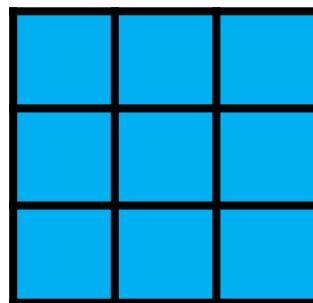
input: $H \times W = 8 \times 8$



output: $H \times W = 6 \times 6$



filter



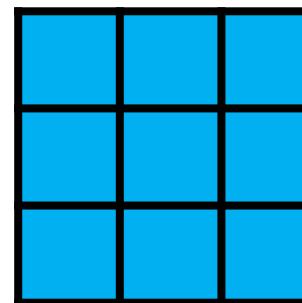
$$H_{\text{out}} = H_{\text{in}} - K_h + 1$$

Convolution: 填充 (padding)

input: 8×8 , + pad

0	0	0	0	0	0	0	0	0	0
0									0
0									0
0									0
0									0
0									0
0									0
0									0
0	0	0	0	0	0	0	0	0	0

filter



output: $H \times W = 8 \times 8$

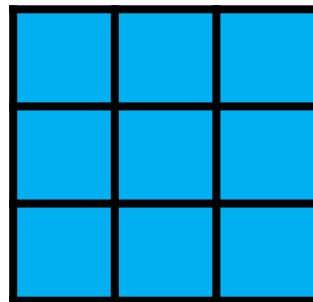
$$H_{\text{out}} = H_{\text{in}} + 2\text{pad}_h - K_h + 1$$

Convolution: 步长 (stride)

input

stride = 2

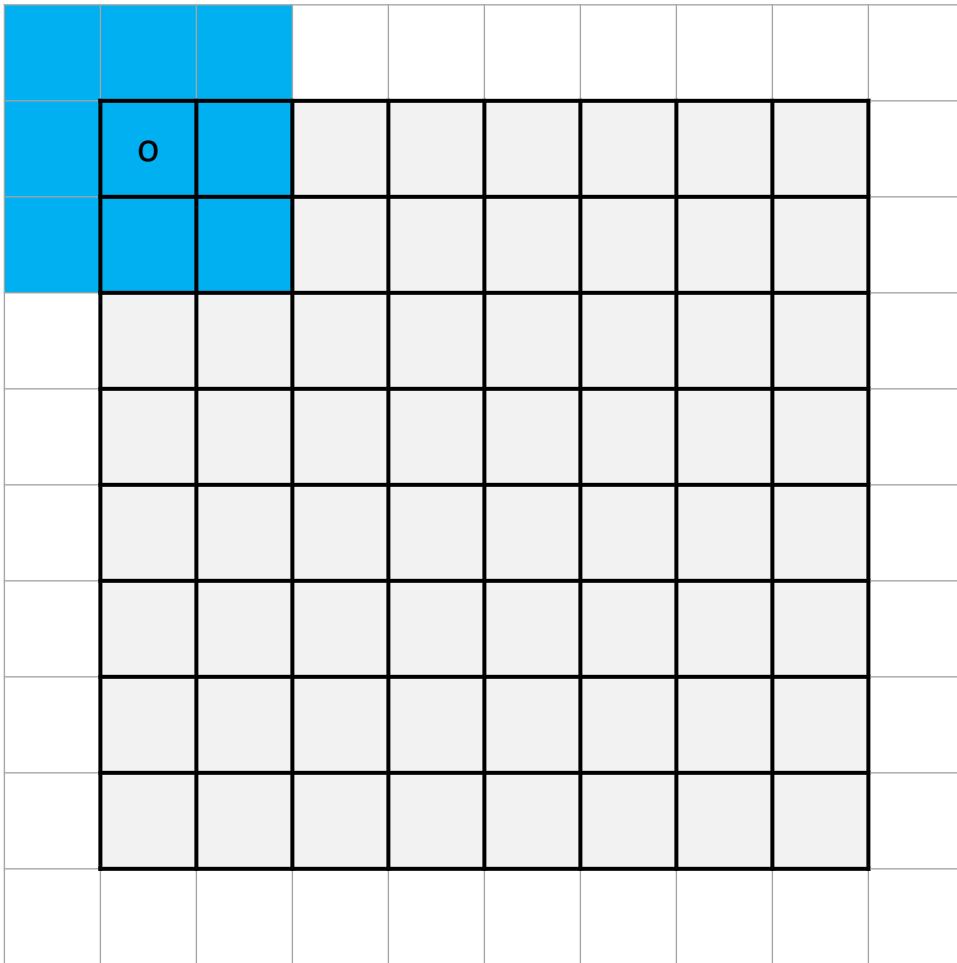
filter



output

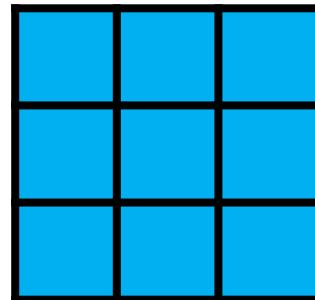
Convolution: 步长 (stride)

input

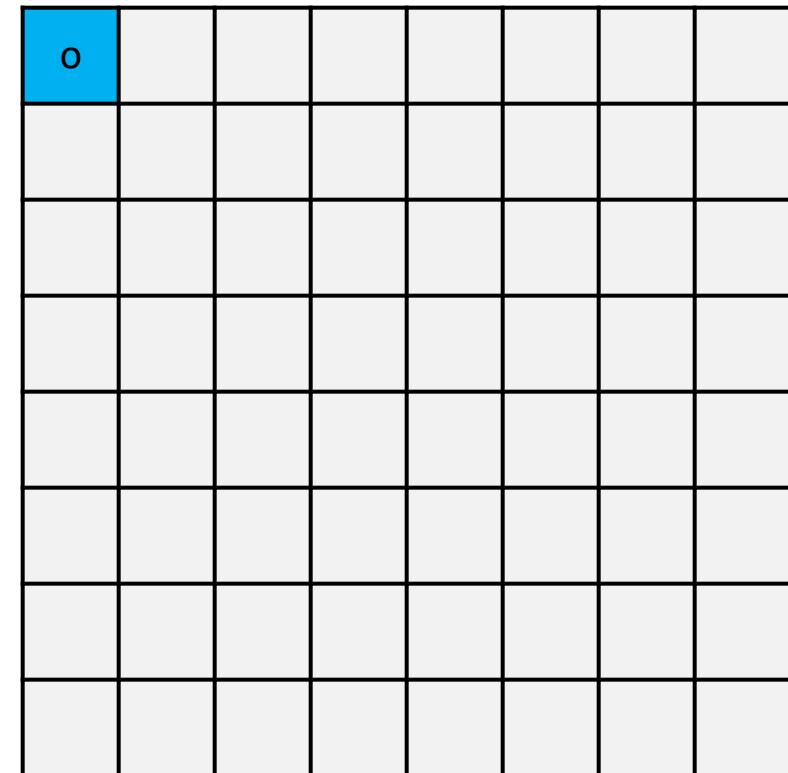


stride = 2

filter

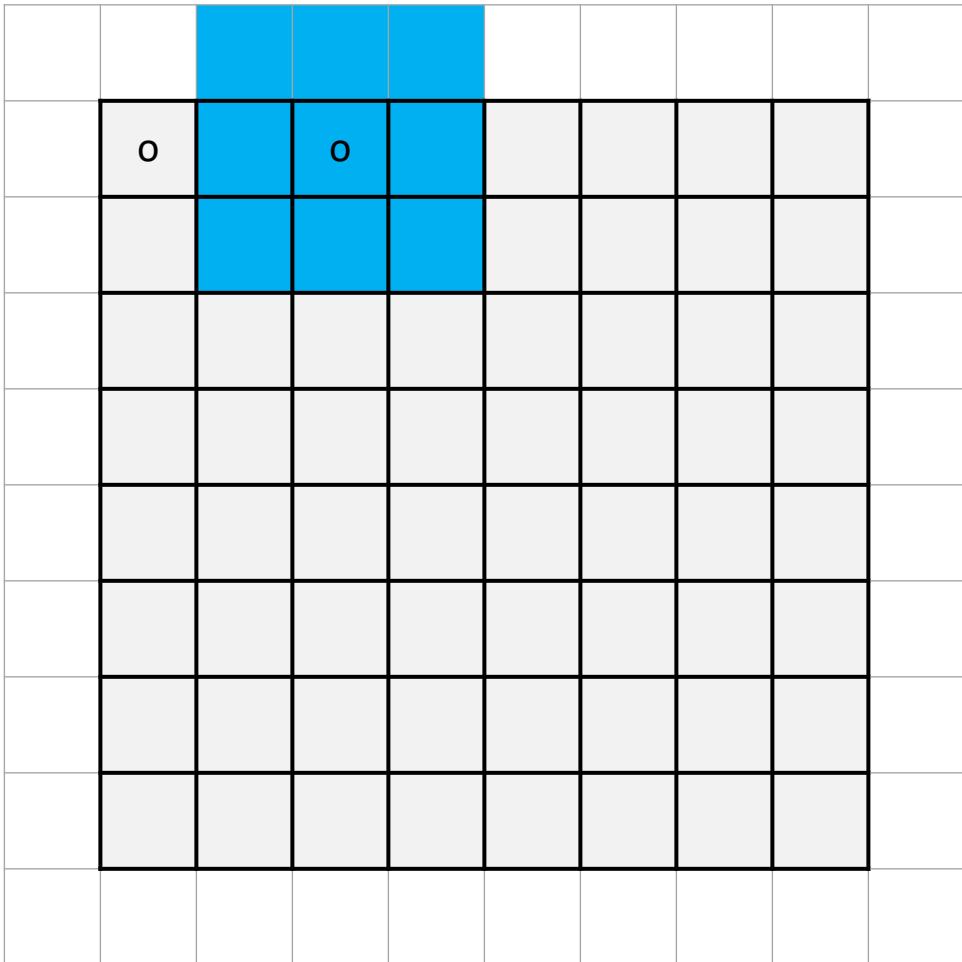


output



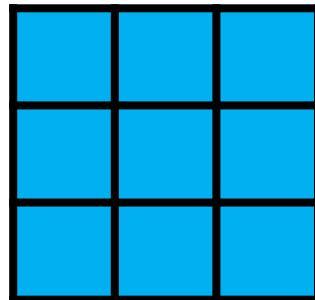
Convolution: 步长 (stride)

input

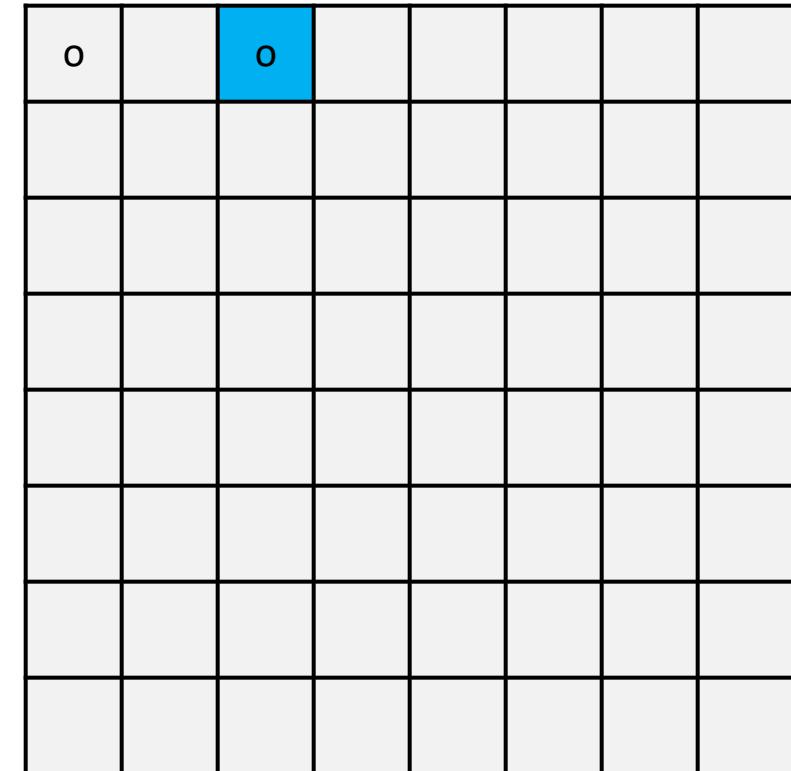


stride = 2

filter

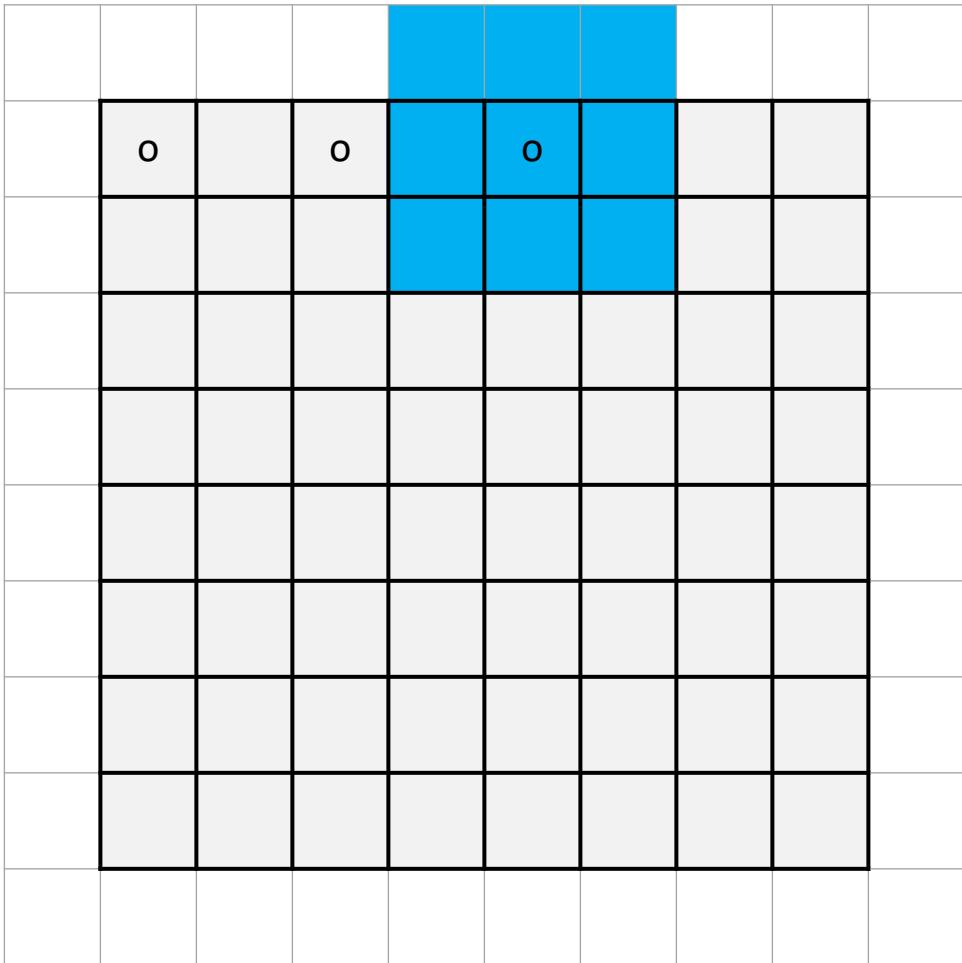


output



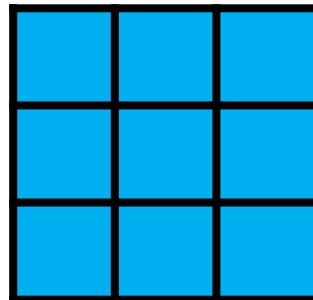
Convolution: 步长 (stride)

input

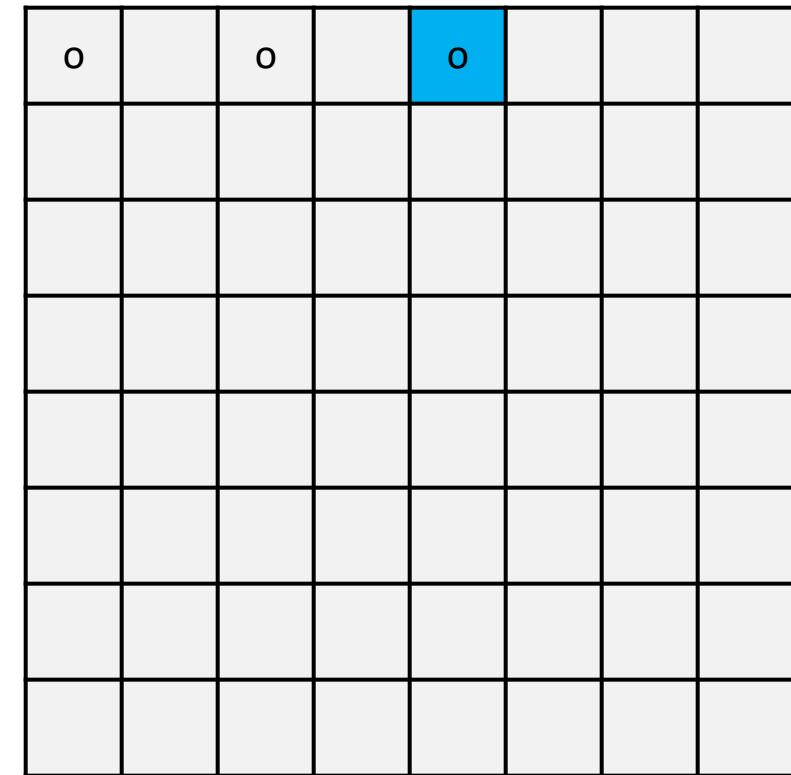


stride = 2

filter



output

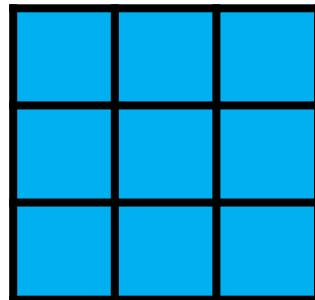


Convolution: 步长 (stride)

input

stride = 2

filter



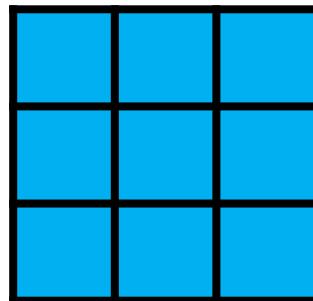
output

Convolution: 步长 (stride)

input

stride = 2

filter



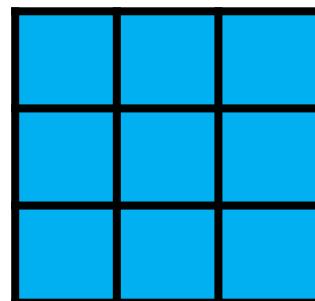
output

Convolution: 步长 (stride)

input

o		o		o		o		
o		o		o		o		
o		o		o	o	o	o	
o		o		o	o	o	o	

filter



output

o		o		o		o	
o		o		o		o	
o		o		o		o	
o		o		o	o	o	o

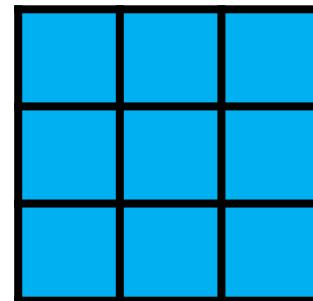
Convolution: 步长 (stride)

input

o		o		o		o	
o		o		o		o	
o		o		o		o	
o		o		o		o	

stride = 2

filter



output: $H \times W = 4 \times 4$

o		o		o	o
o		o		o	o
o		o		o	o
o		o		o	o

Convolution: 步长 (stride)

input

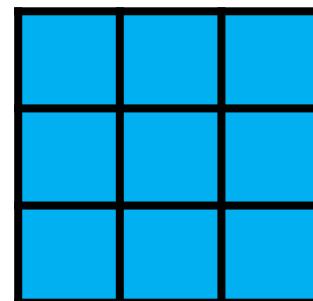
o		o		o		o		
o		o		o		o		
o		o		o		o		
o		o		o		o		

stride = 2

- reduces feature map size
- compress and abstract

output: $H \times W = 4 \times 4$

filter



o	o	o	o
o	o	o	o
o	o	o	o
o	o	o	o

$$H_{out} = \lfloor (H_{in} + 2\text{pad}_h - K_h) / \text{str} \rfloor + 1$$

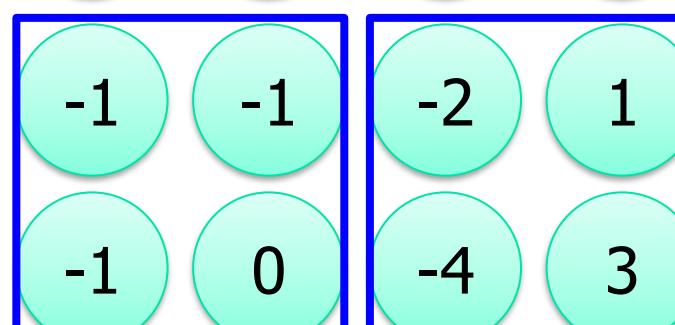
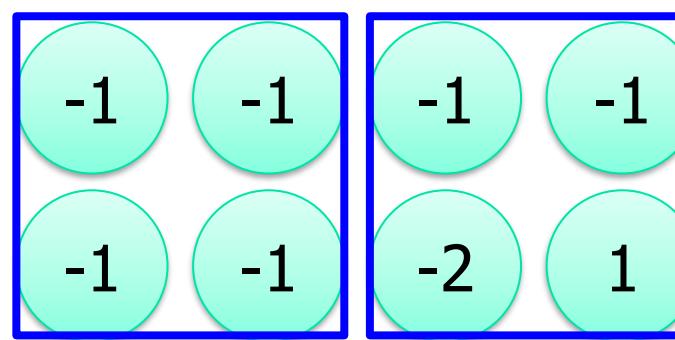
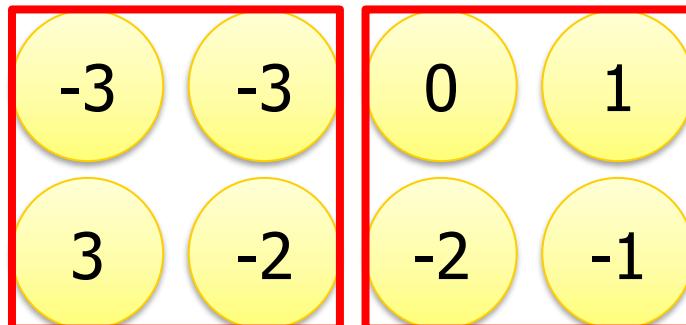
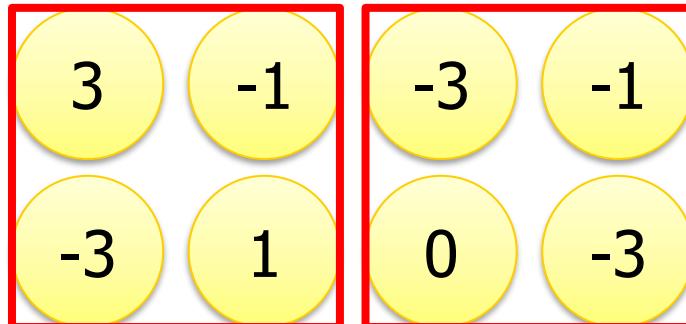
CNN – Max Pooling

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

-1	1	-1
-1	1	-1
-1	1	-1

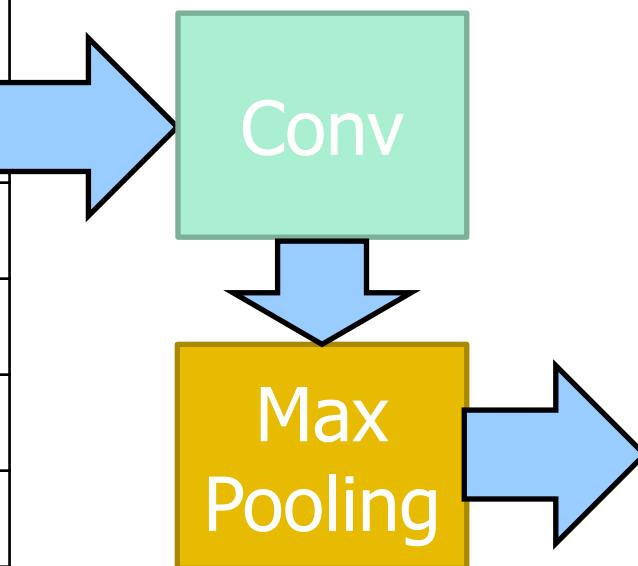
Filter 2



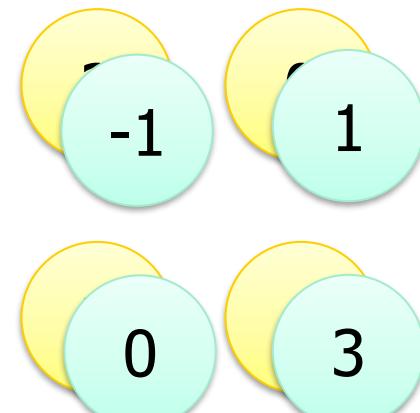
CNN – Max Pooling

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image



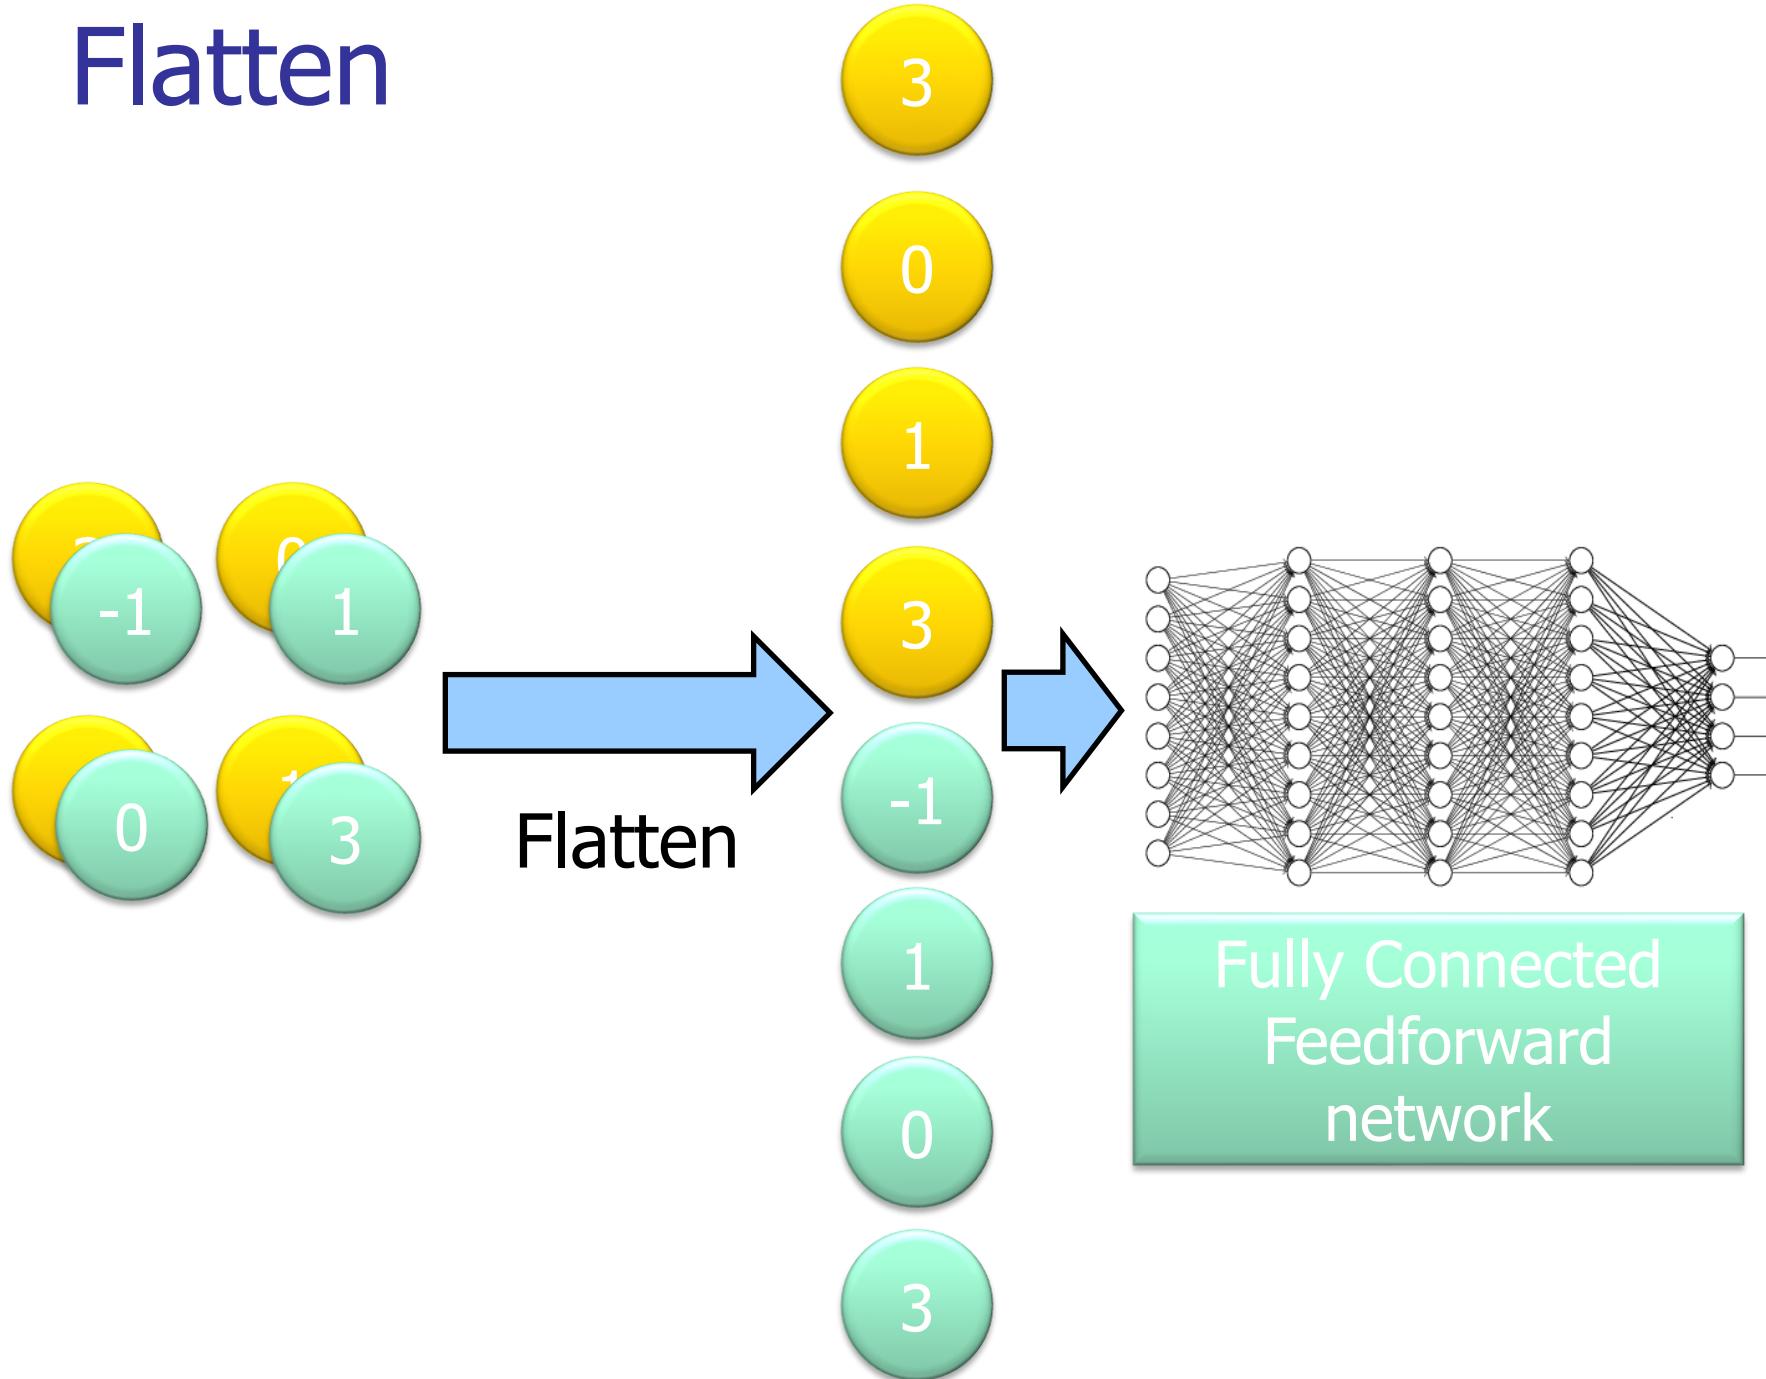
New image
but smaller



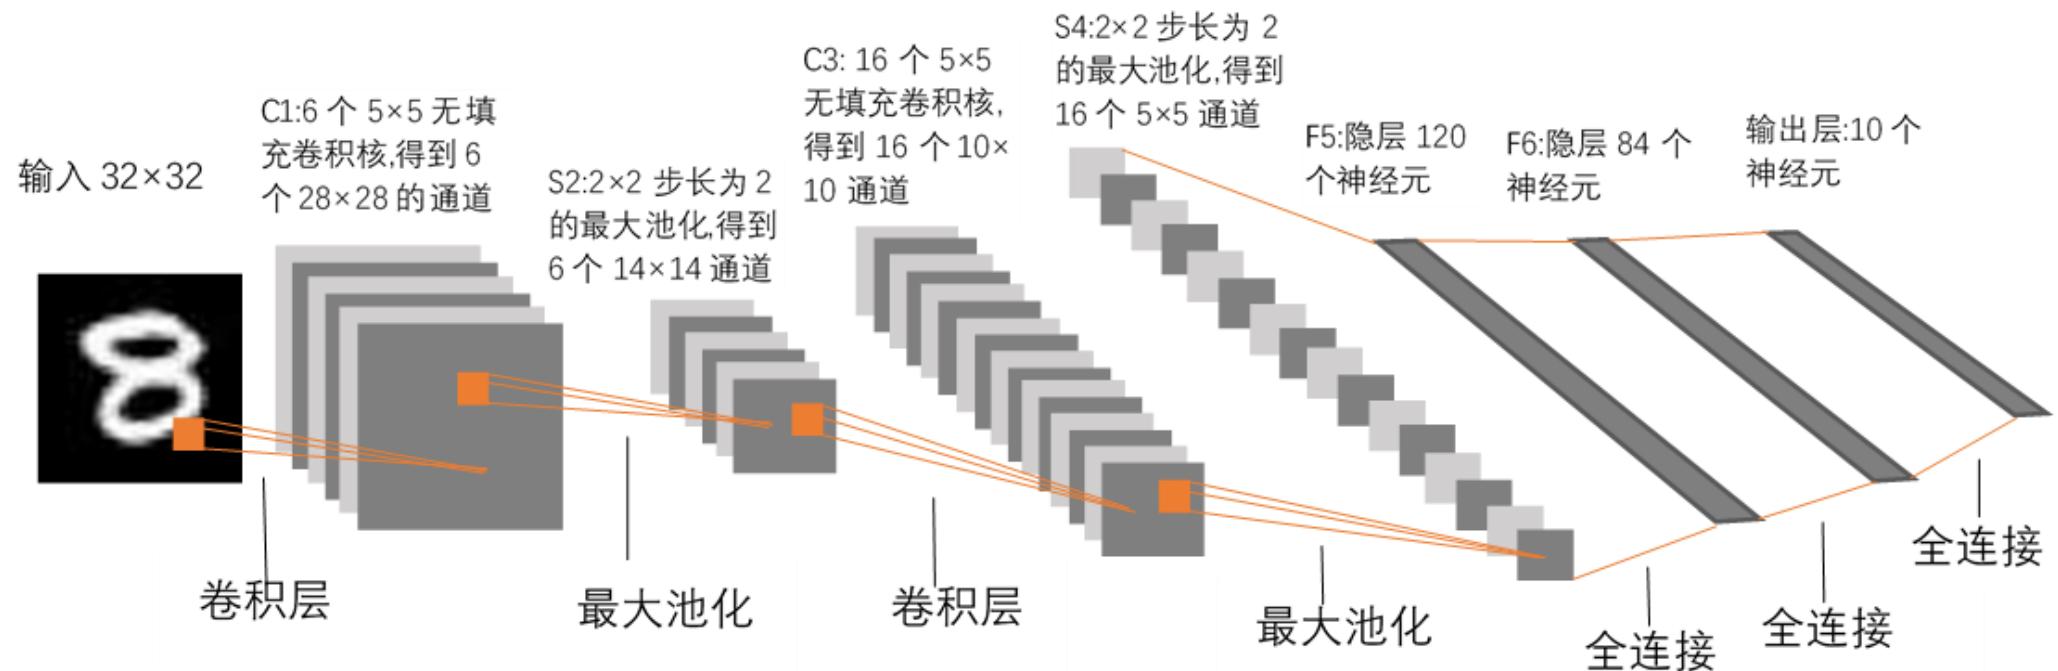
2 x 2 image

Each filter
is a channel

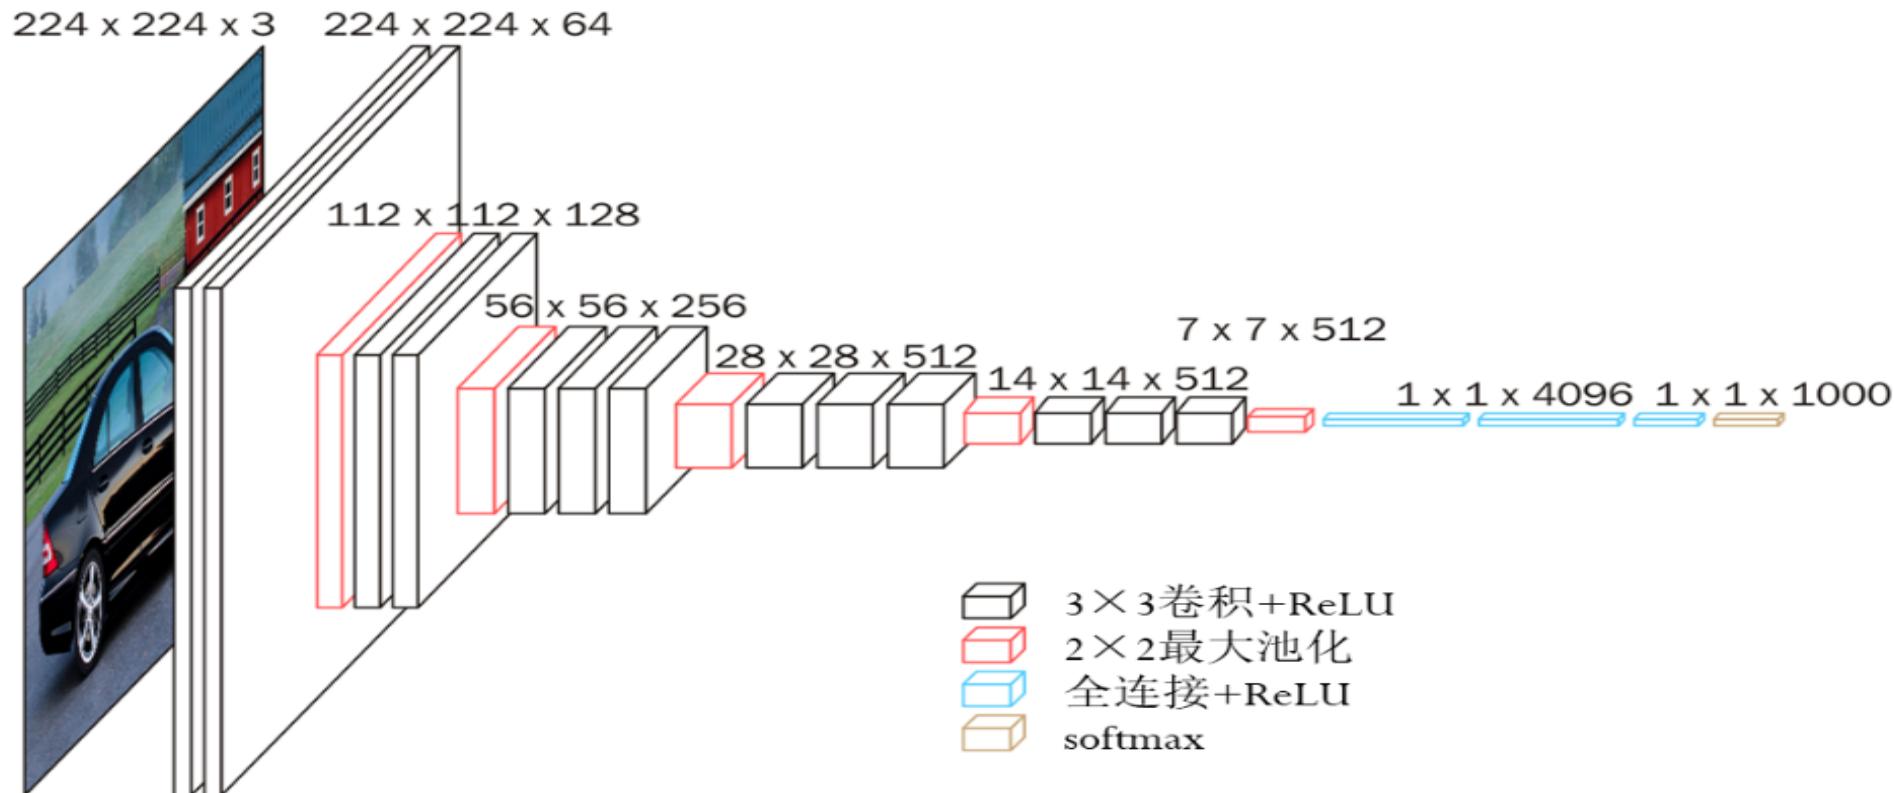
Flatten



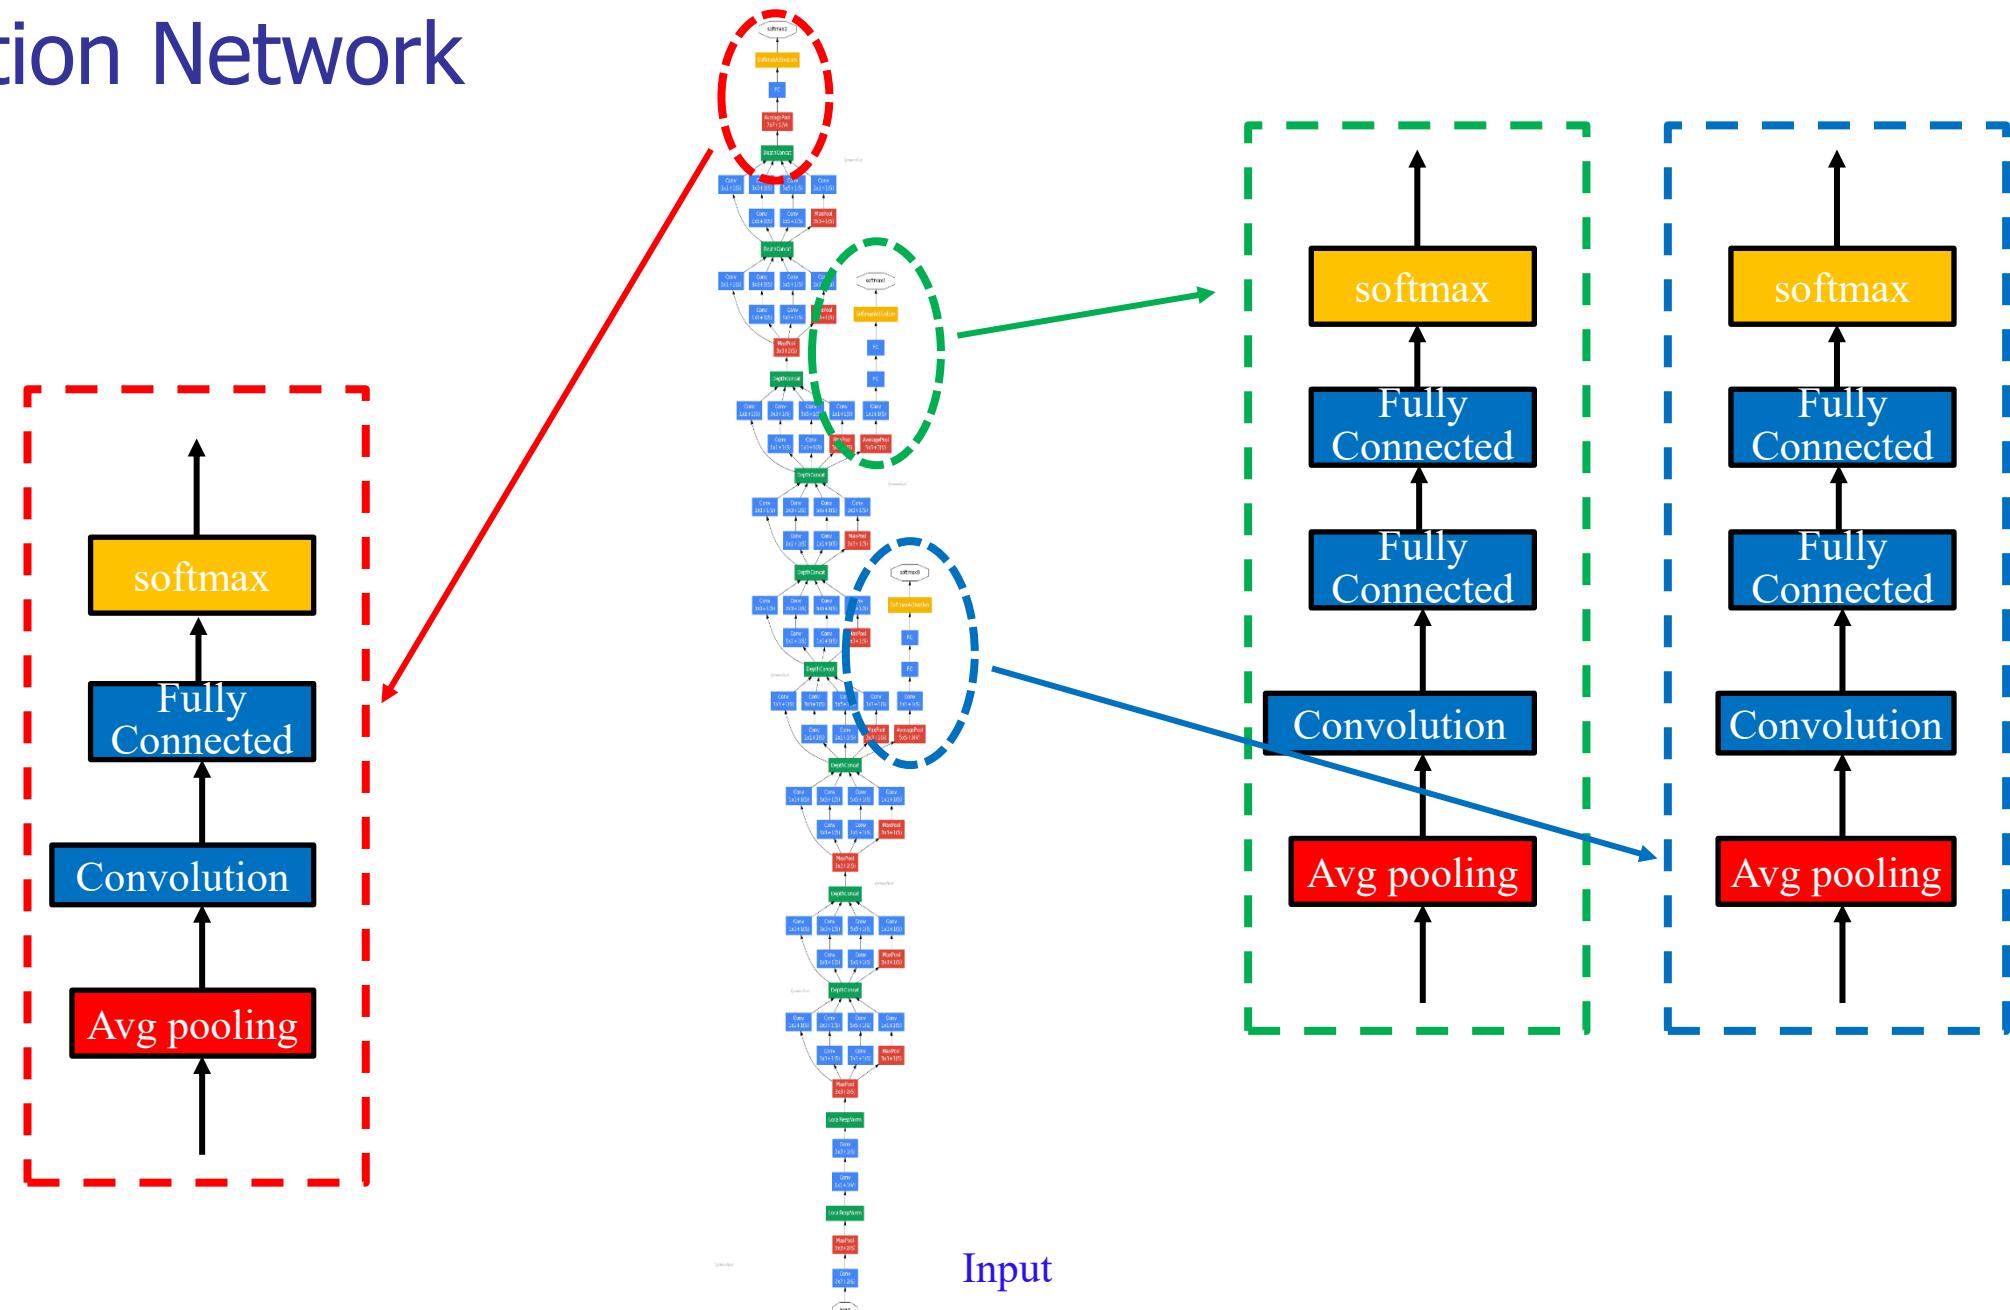
LeNet

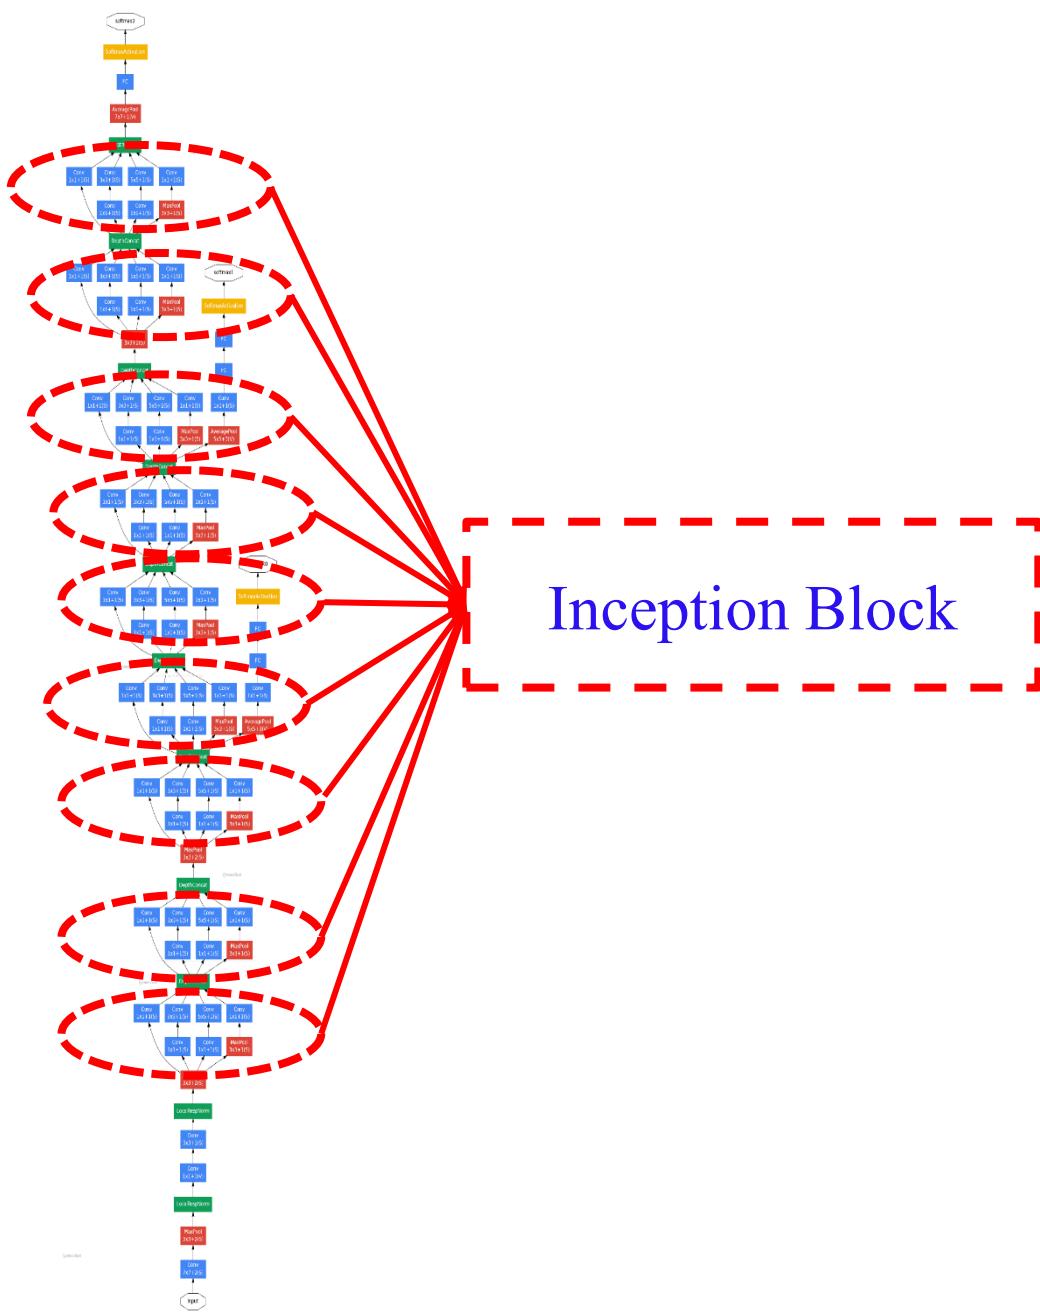


VGG-16



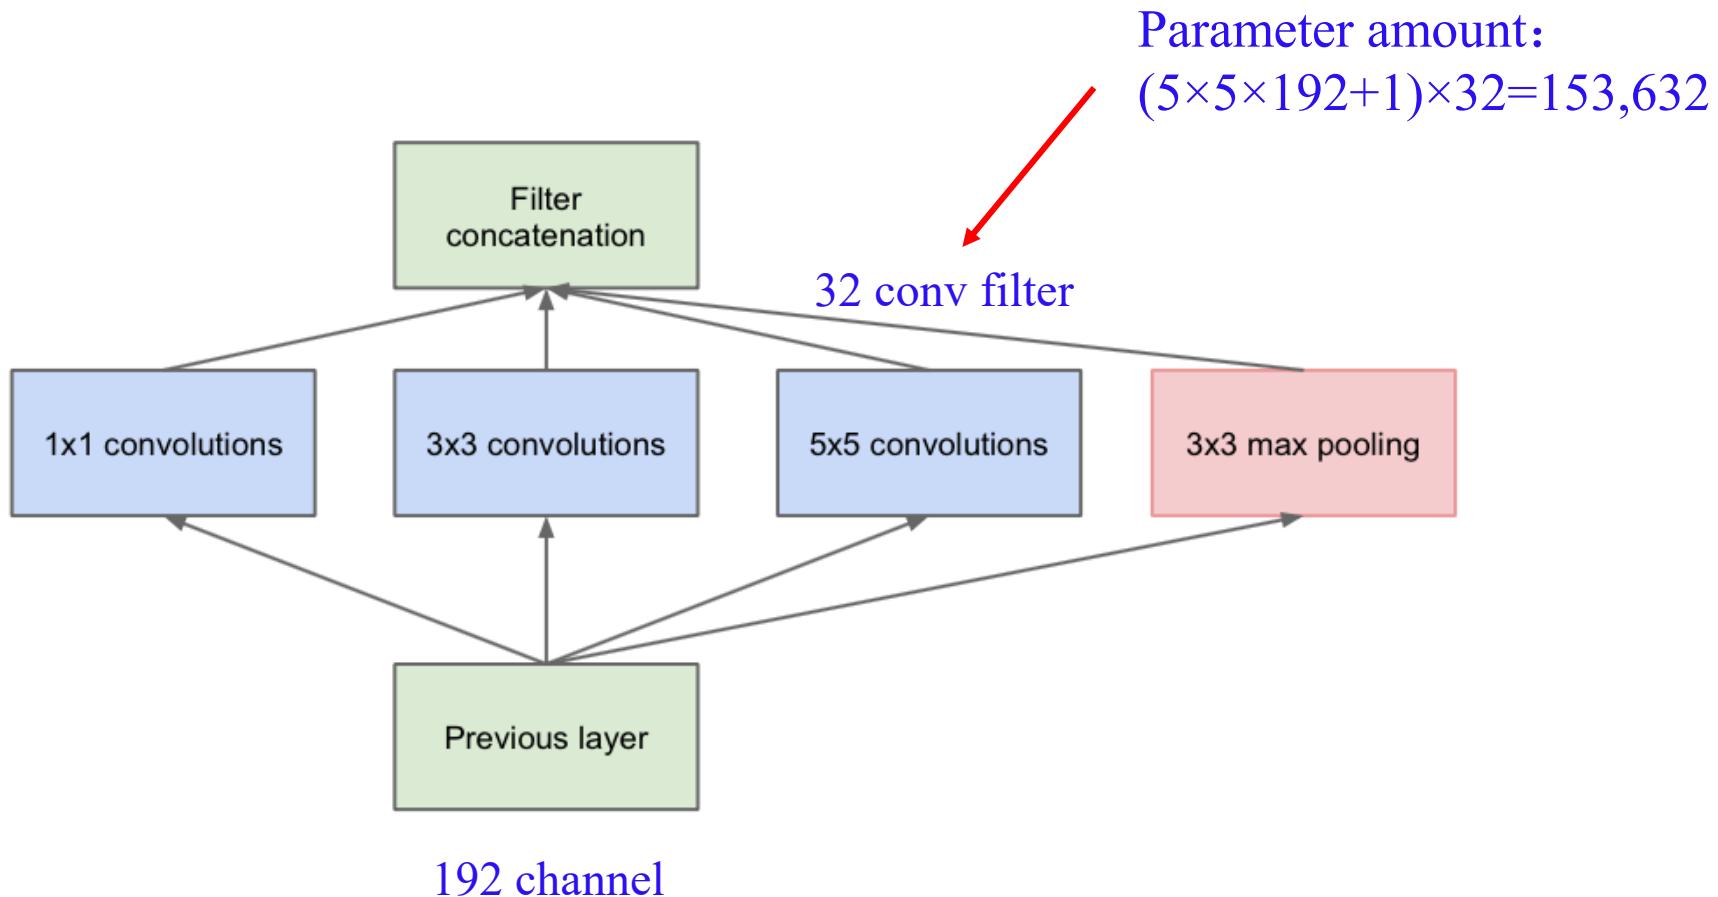
Inception Network



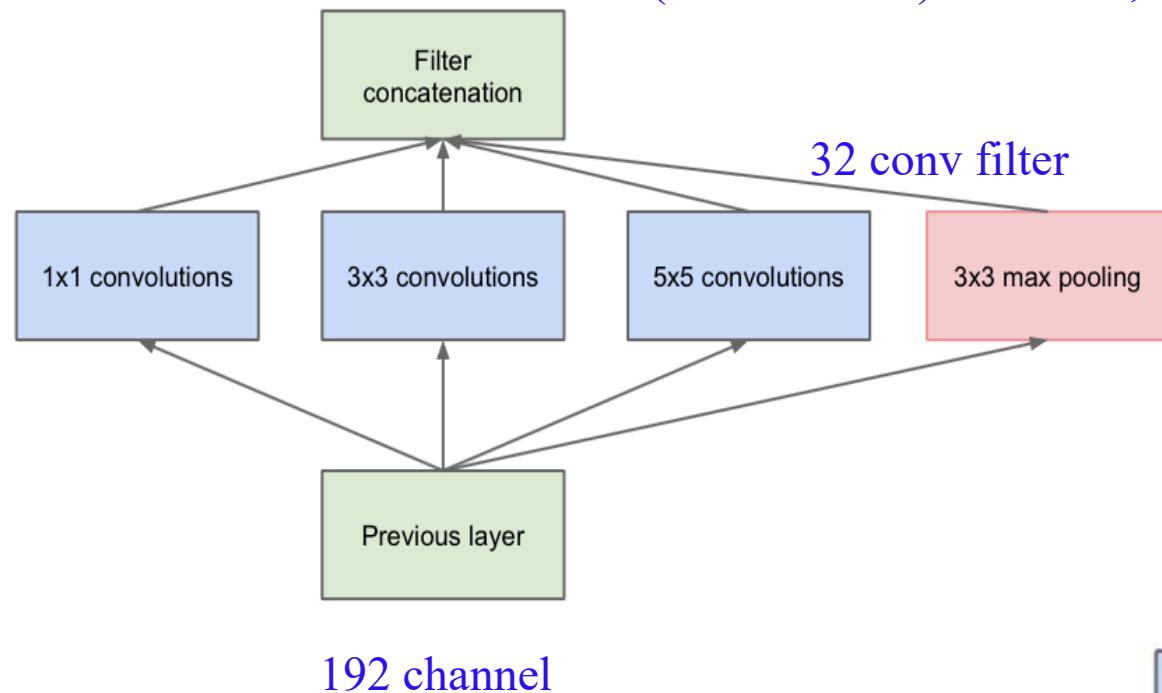


Inception Block

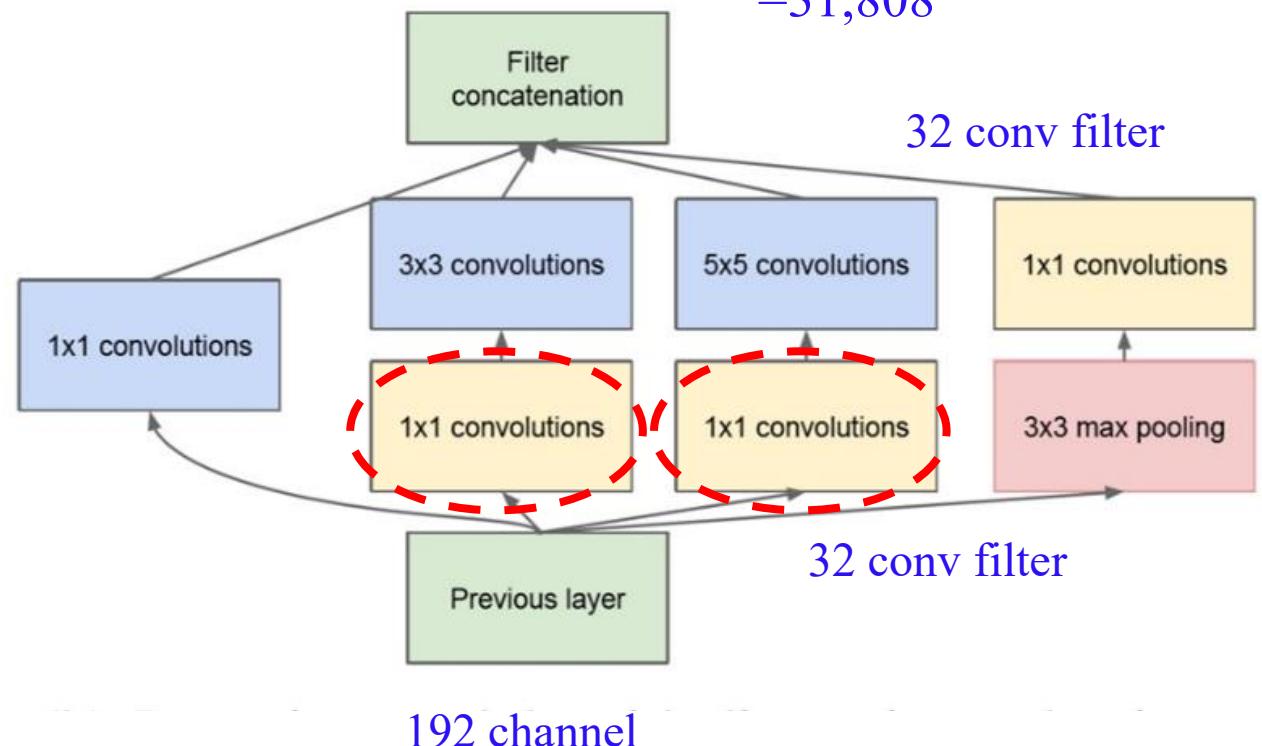
Inception Block

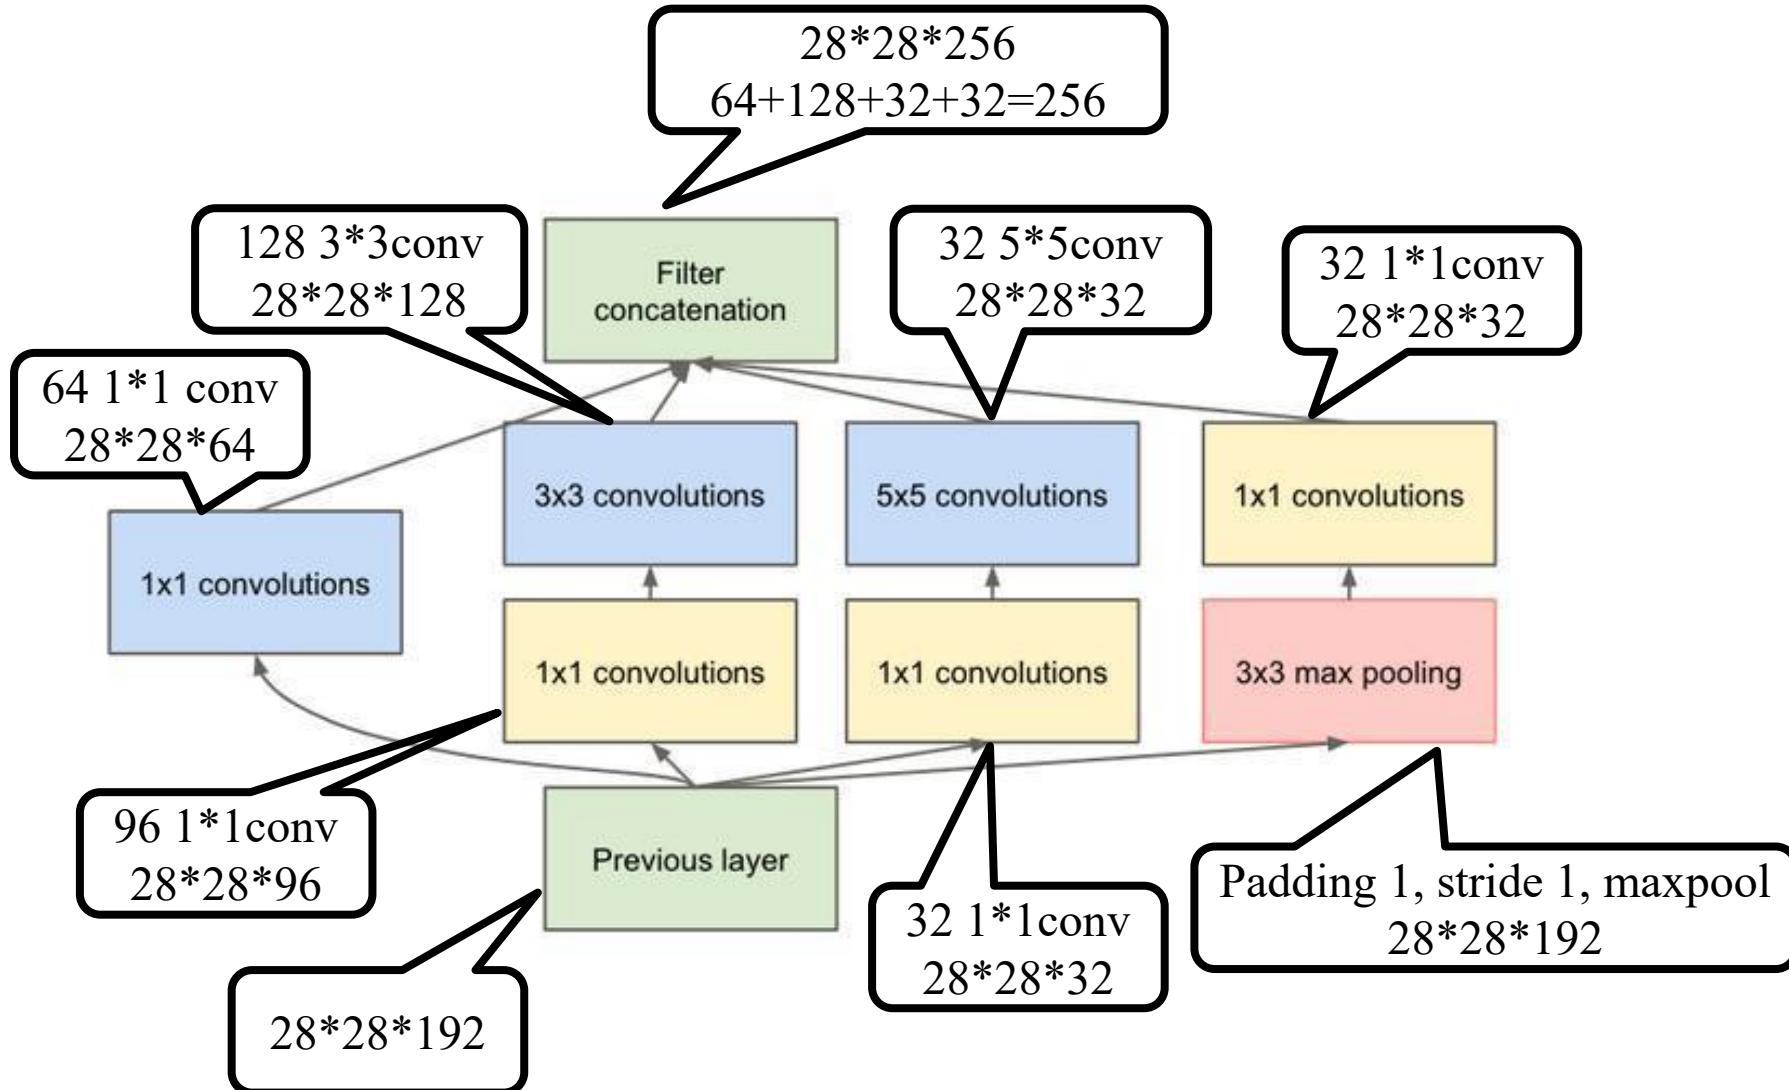


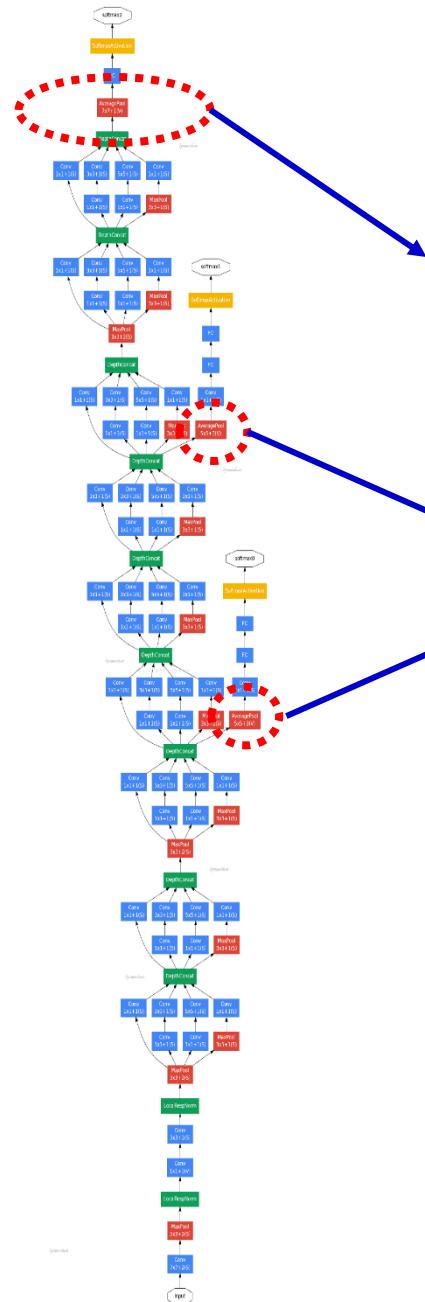
Parameters:
 $(5 \times 5 \times 192 + 1) \times 32 = 153,632$



Parameters:
 $(1 \times 1 \times 192 + 1) \times 32 +$
 $(5 \times 5 \times 32 + 1) \times 32$
 $= 31,808$



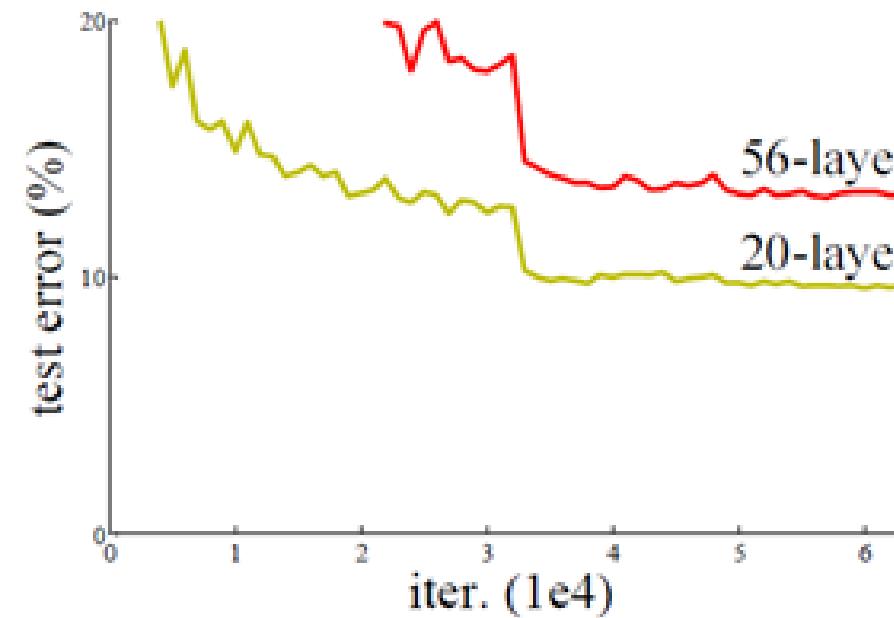
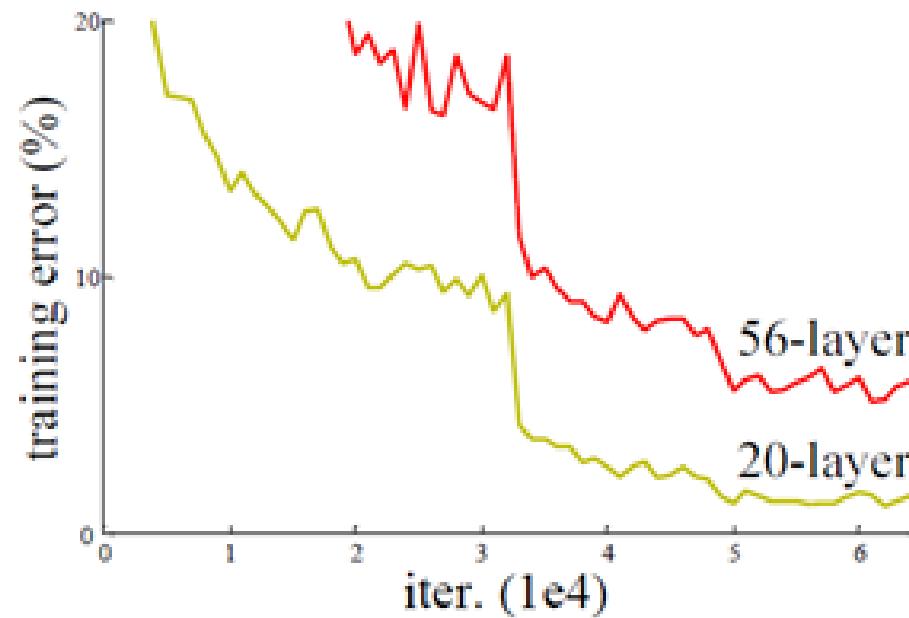




7×7 average pooling
Stride 1, no padding

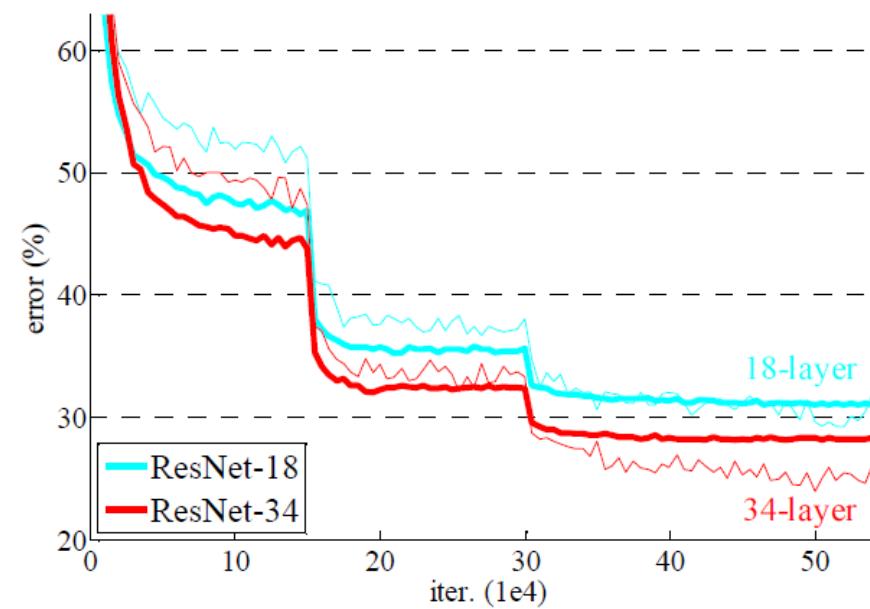
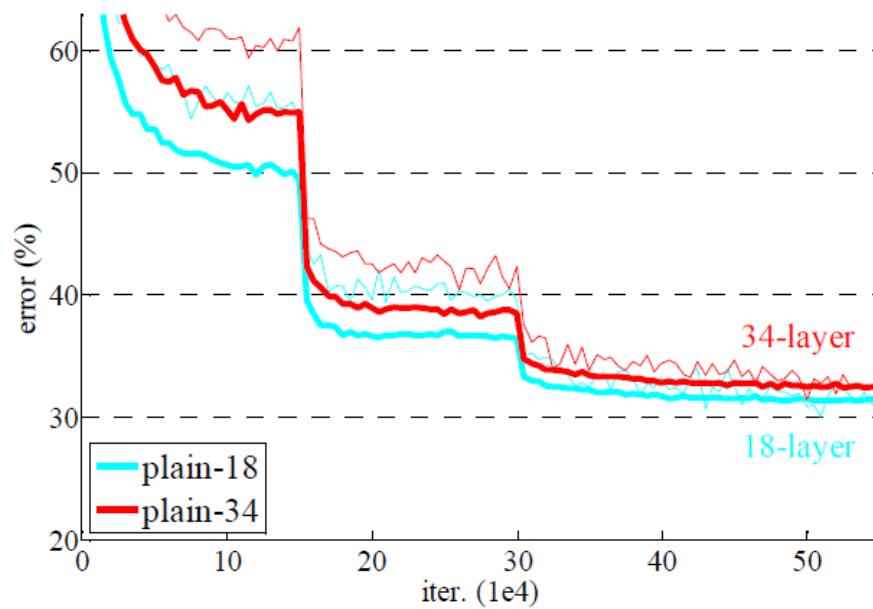
5×5 average pooling
Stride 3, no padding

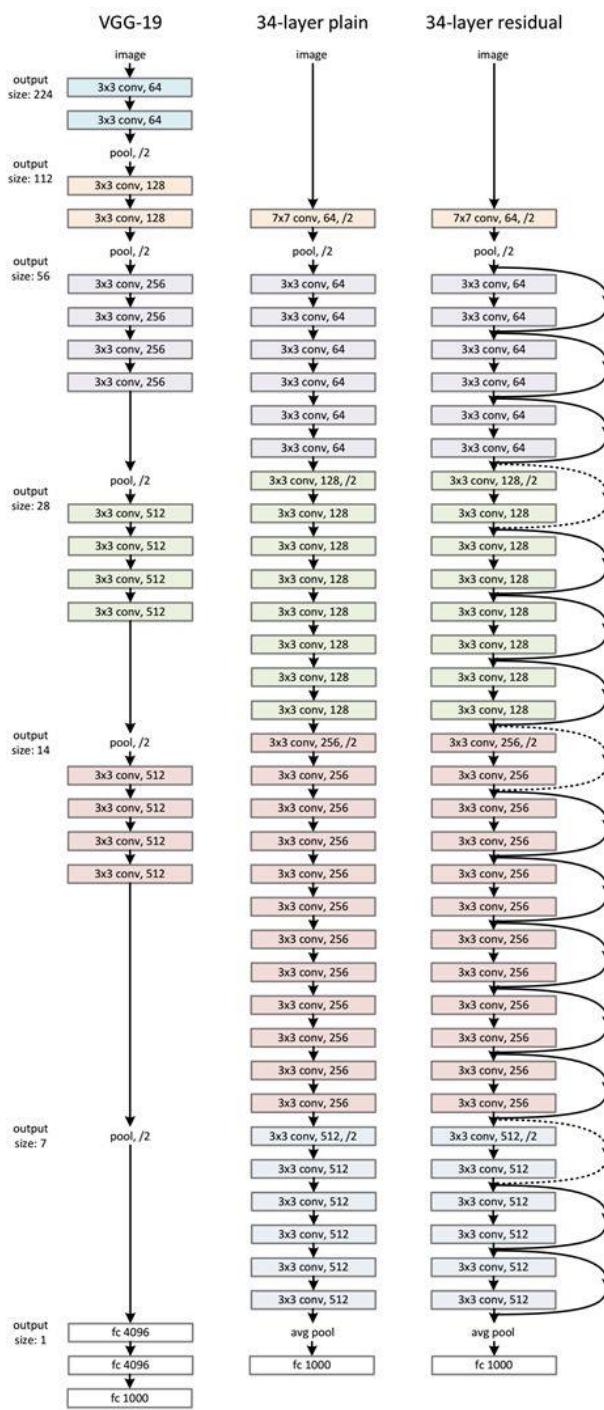
Degradation Problem in Deep Networks

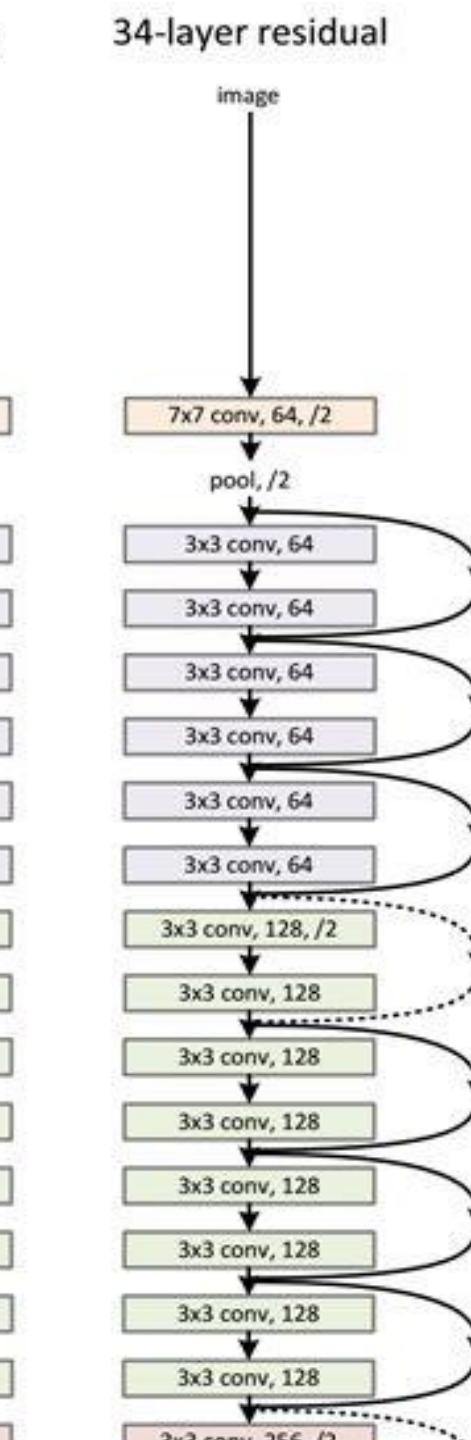
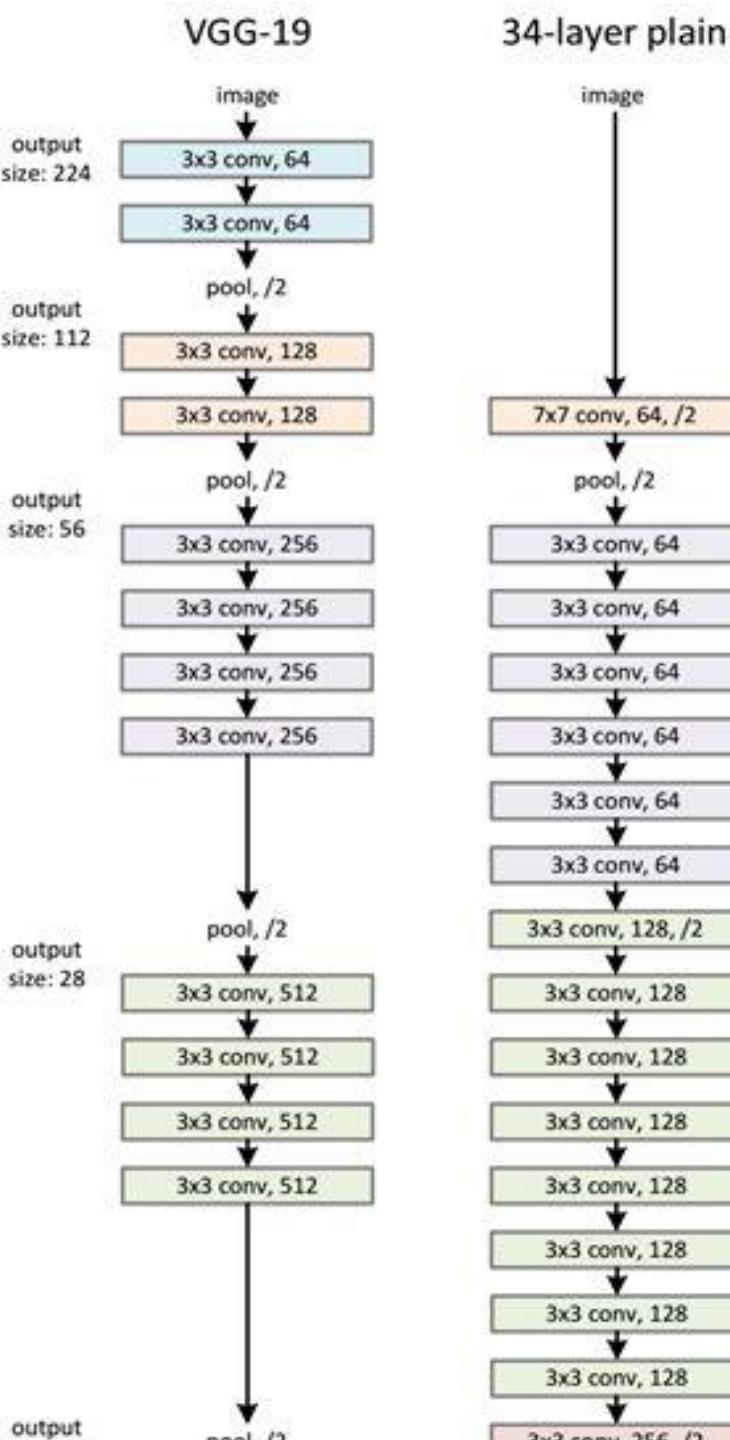
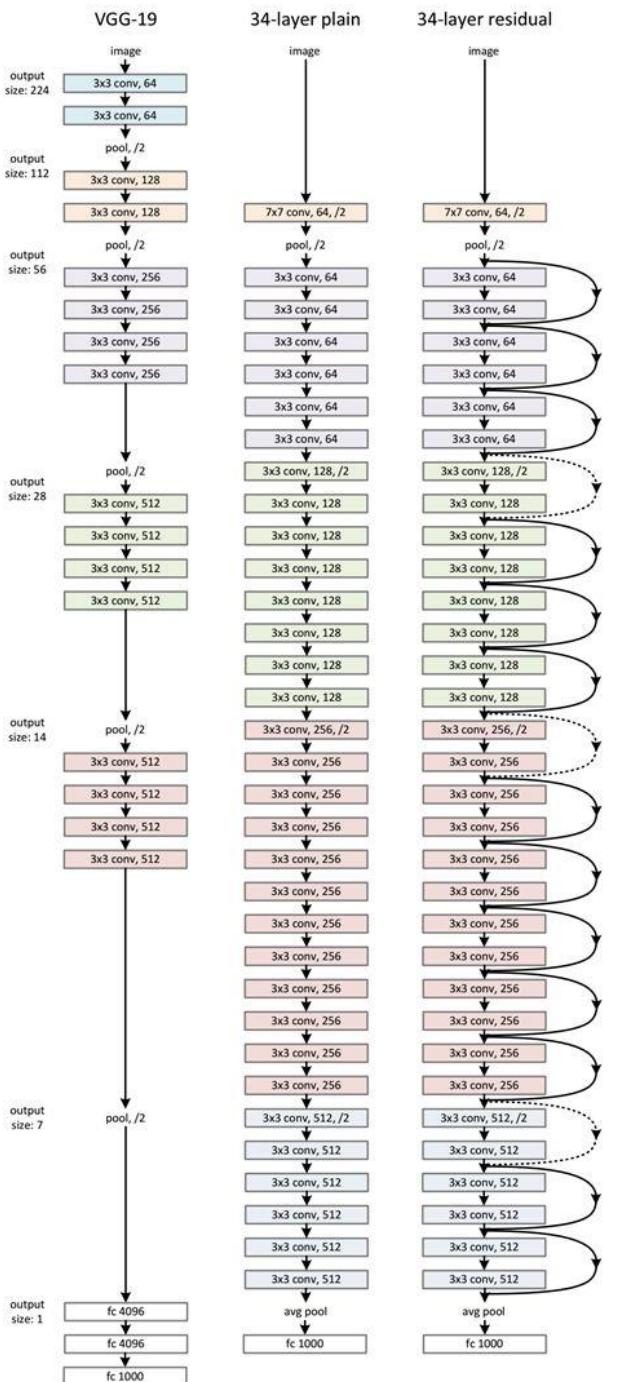


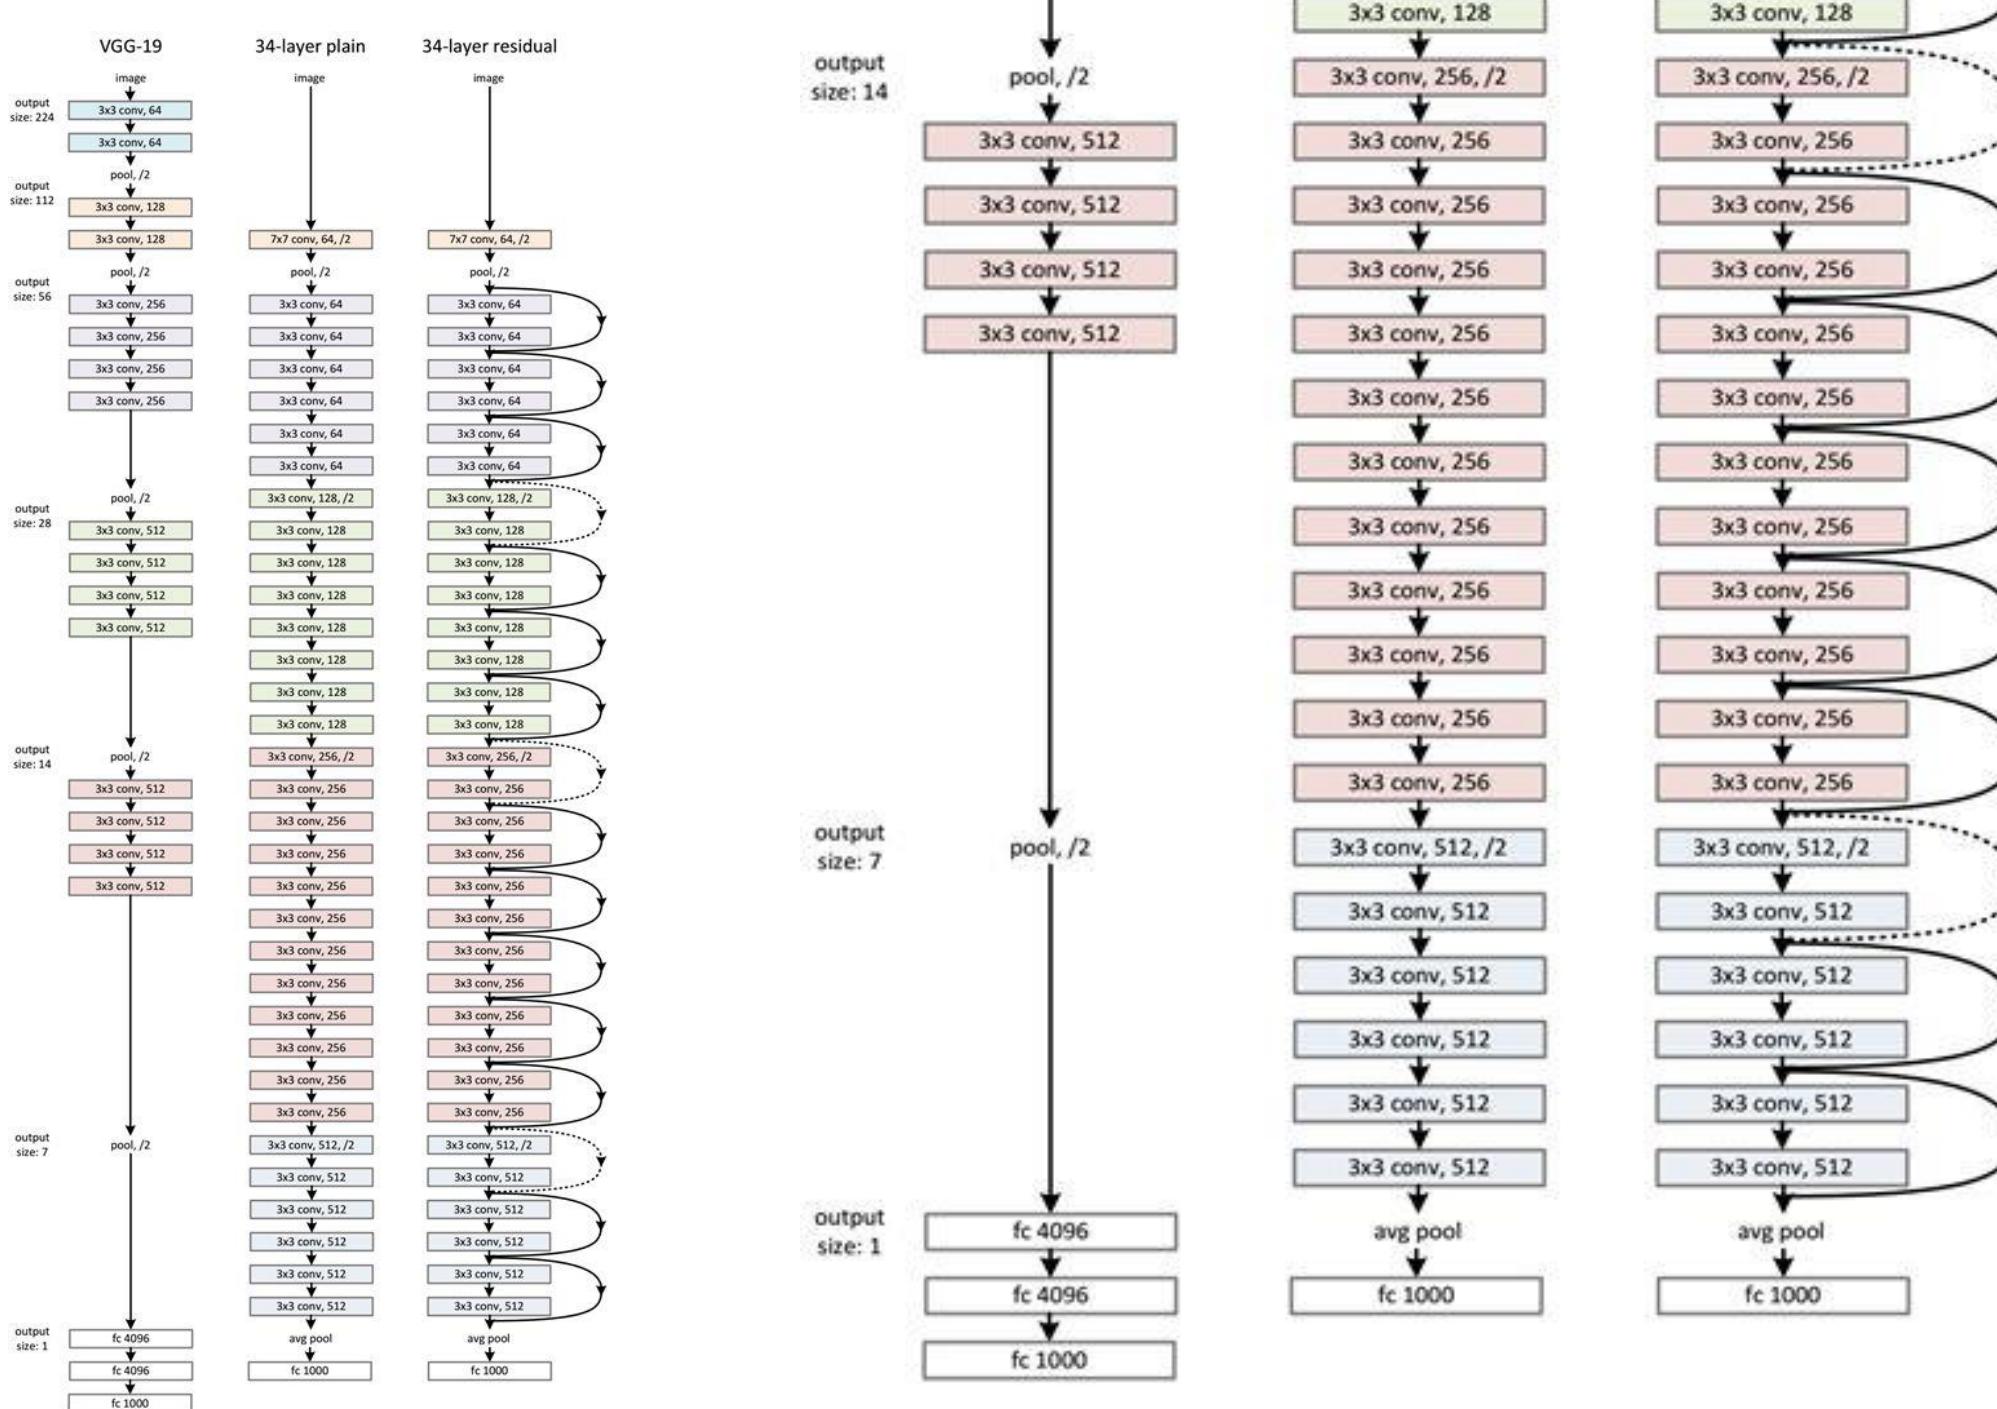
Deep Residual Learning

- Kaiming He, Xiangyu Zhang, Shaoqing Ren, Jian Sun. Deep Residual Learning for Image Recognition. CVPR 2016.

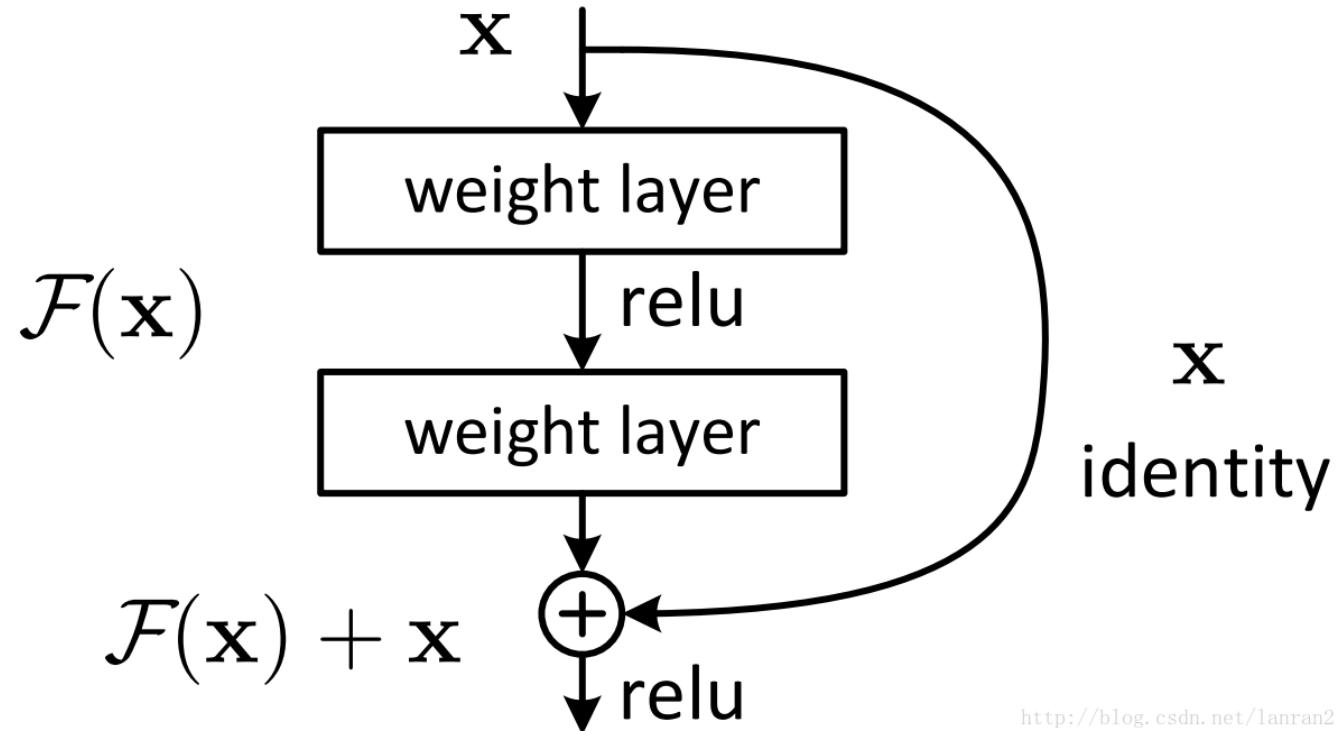




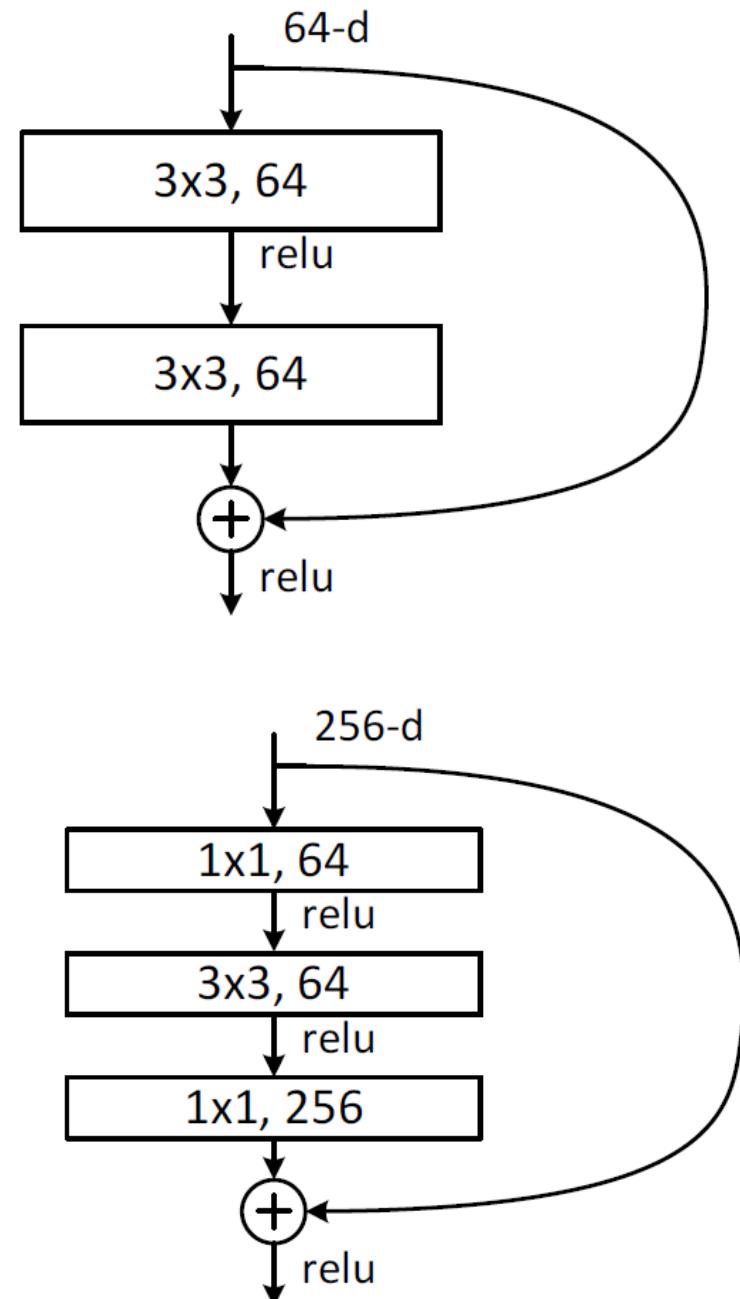




Identity Mapping by Shortcuts



<http://blog.csdn.net/lanran2>



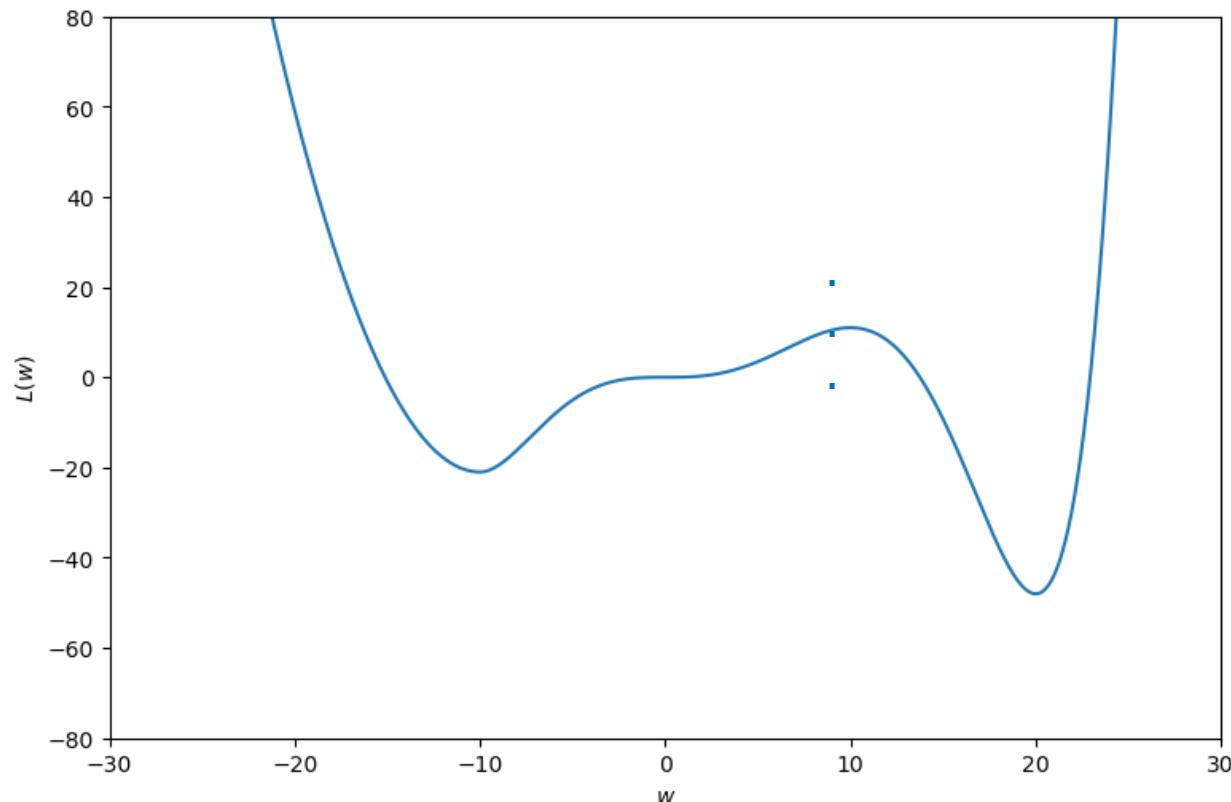
Some Techniques for Learning Better Deep Networks

Techniques for Learning Better NNs

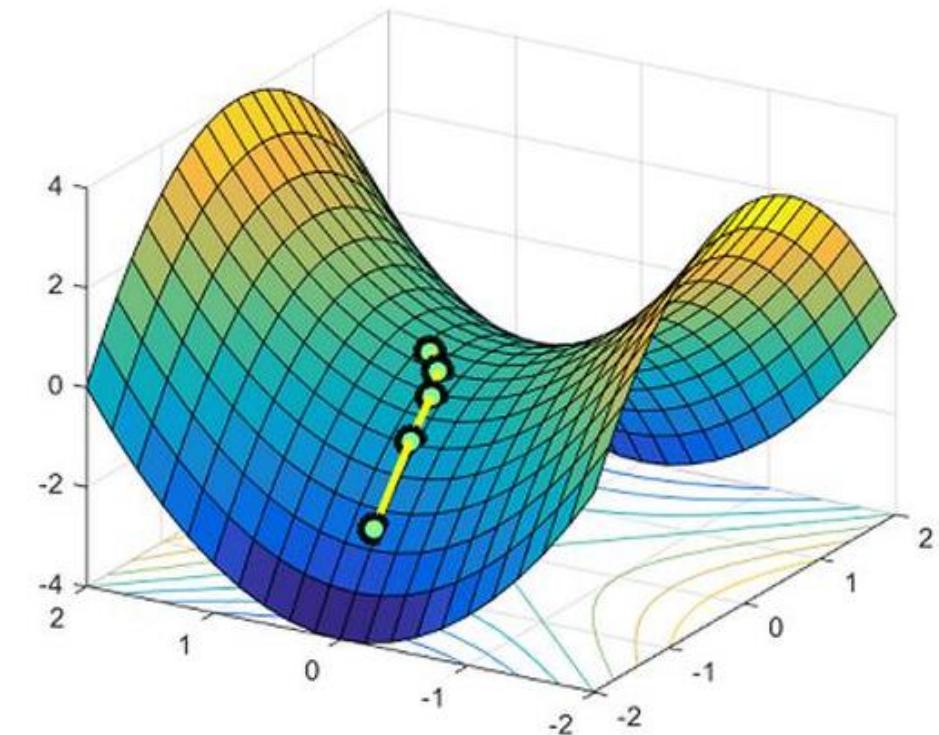
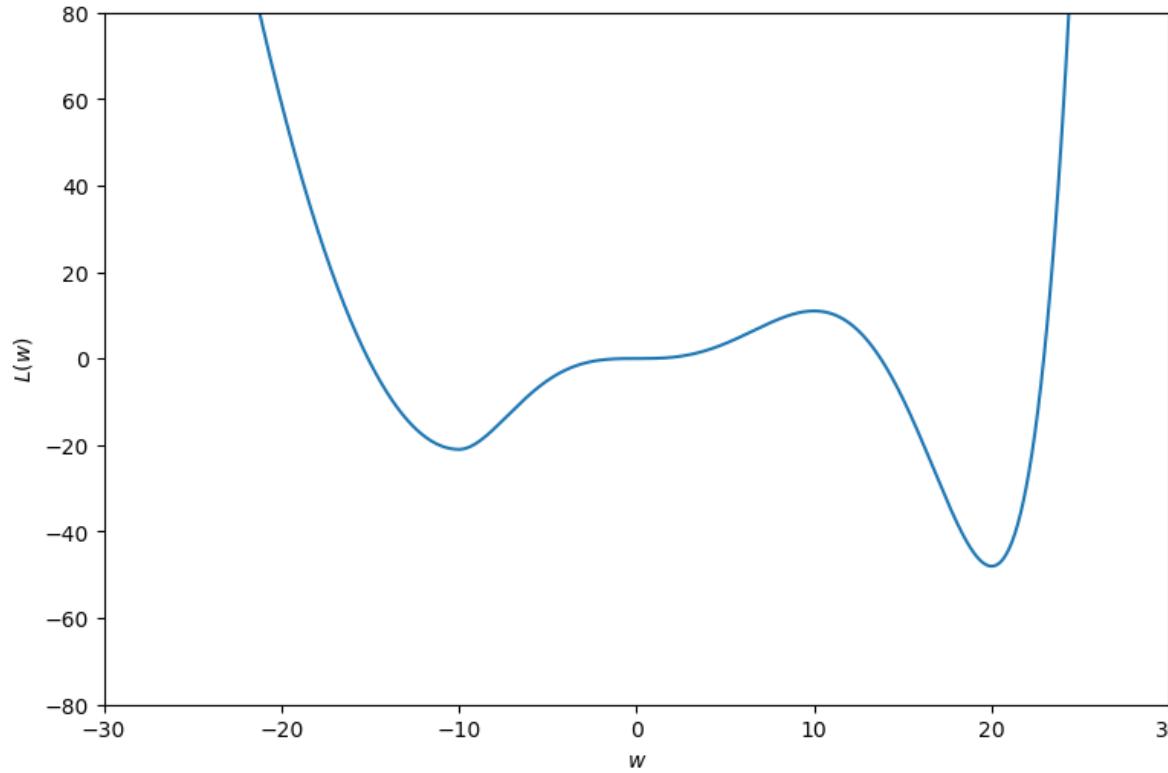
- Design good network structures
- Design good loss?
- Optimization techniques
- Others
 - Data preprocessing
 - Data augmentation

Remember Two Facts

- A NN is a complicated function
 - We can study it from function properties
- Optimizing a NN is usually a non-convex problem



Stationary Point and Saddle Points in High-dimensional Space



Flat Minima in NN

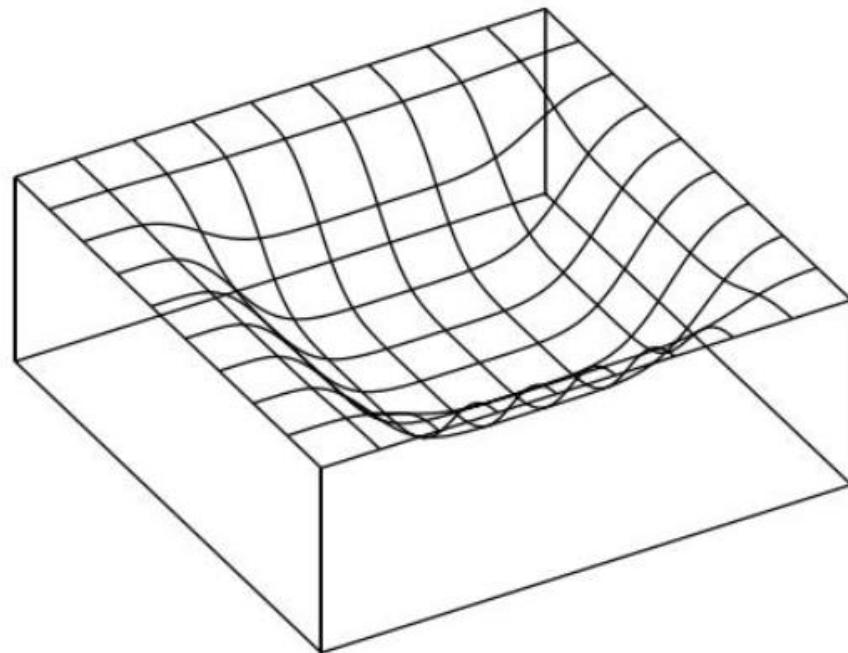


Figure 1: Example of a “flat” minimum.

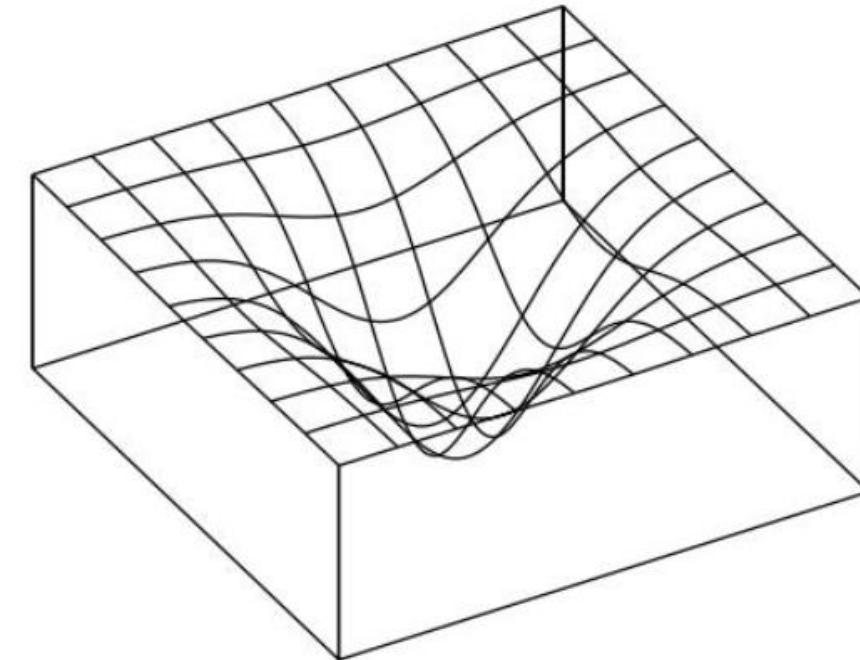


Figure 2: Example of a “sharp” minimum.

Loss Landscape of NN

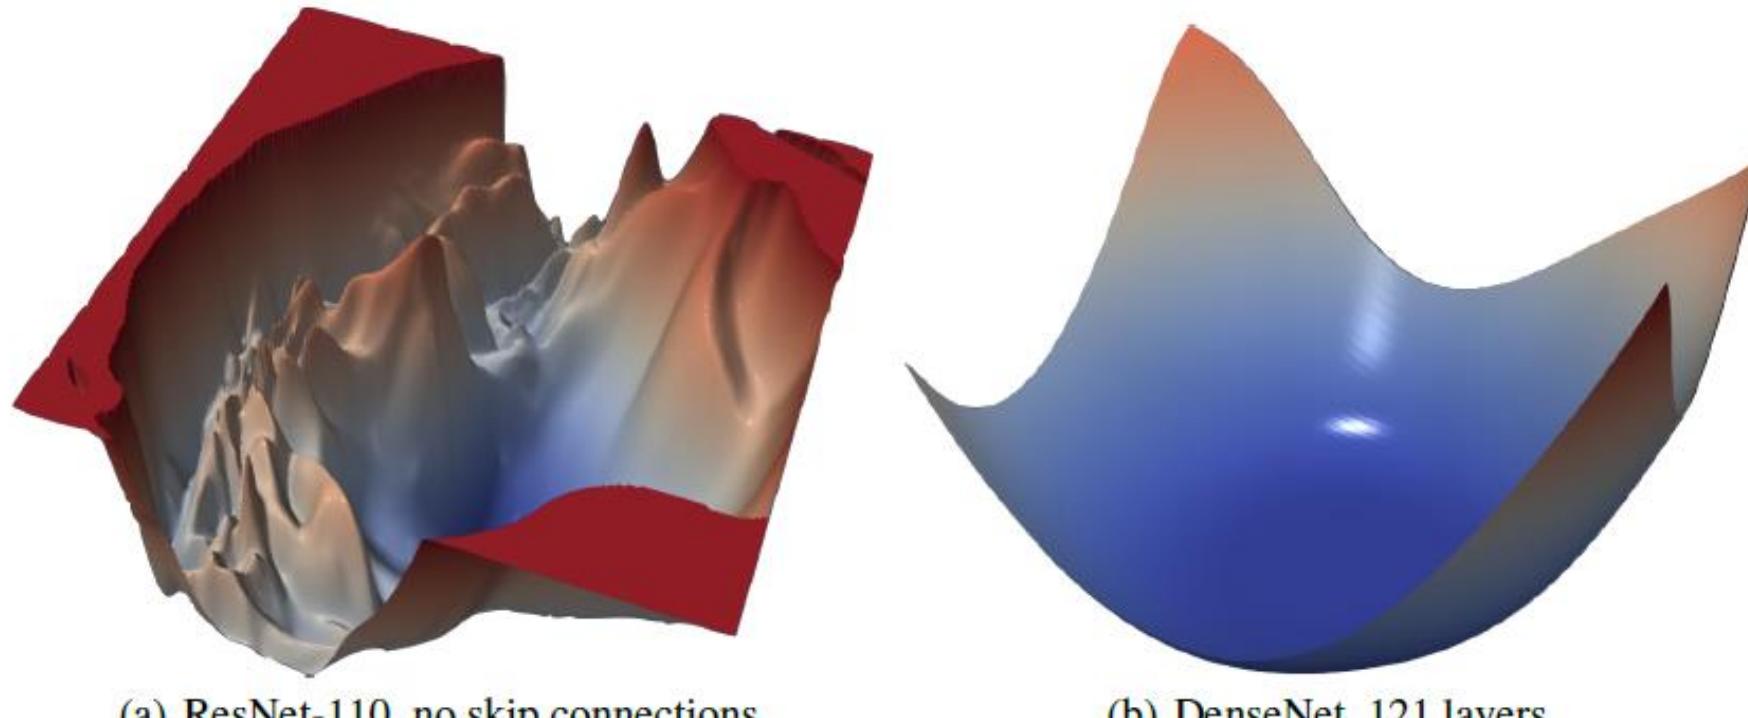


Figure 4: The loss surfaces of ResNet-110-noshort and DenseNet for CIFAR-10.

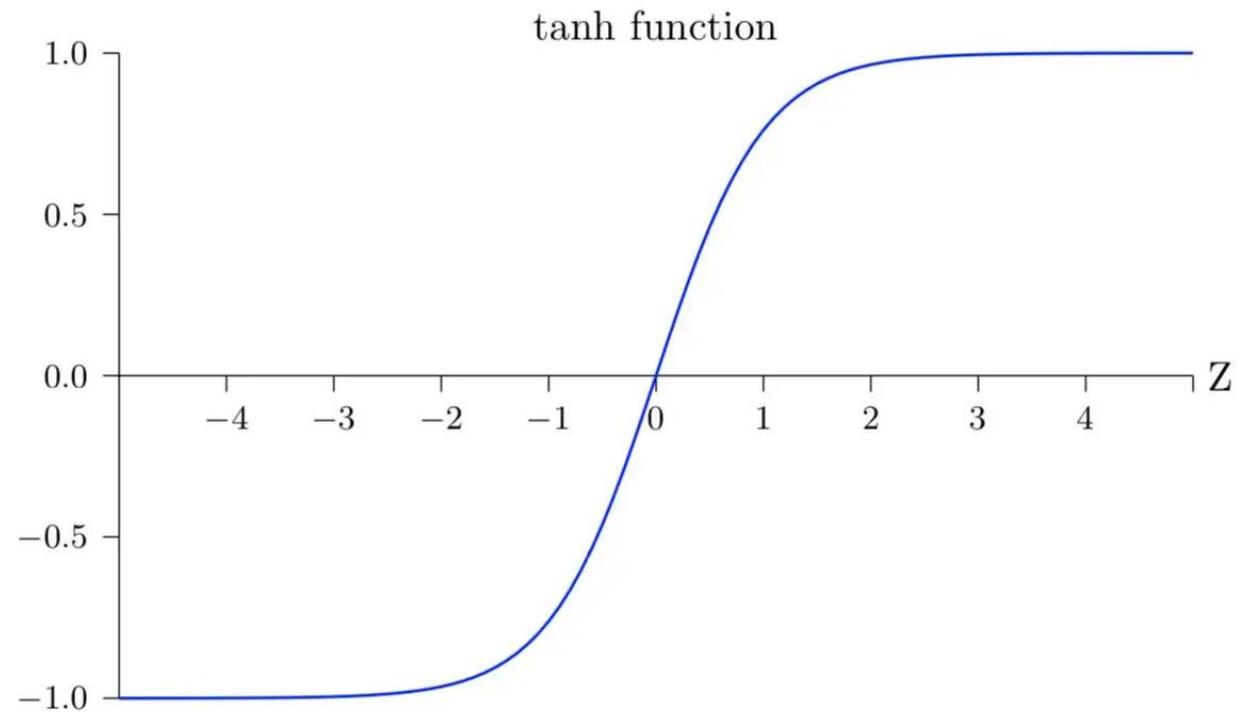
Network Design Issues

- How many layers?
- How many neurons in a layer?
- Connection structure:
 - other than fully-connected
- Activation Functions

Activation Function: Hyperbolic Tangent

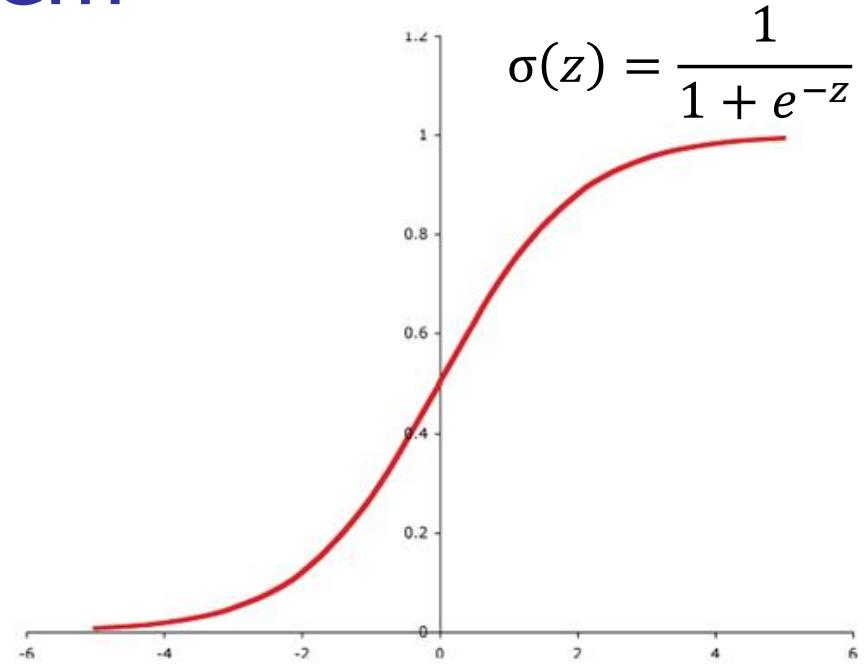
$$a(z) = \tanh(z)$$

$$= \frac{e^z - e^{-z}}{e^z + e^{-z}}$$

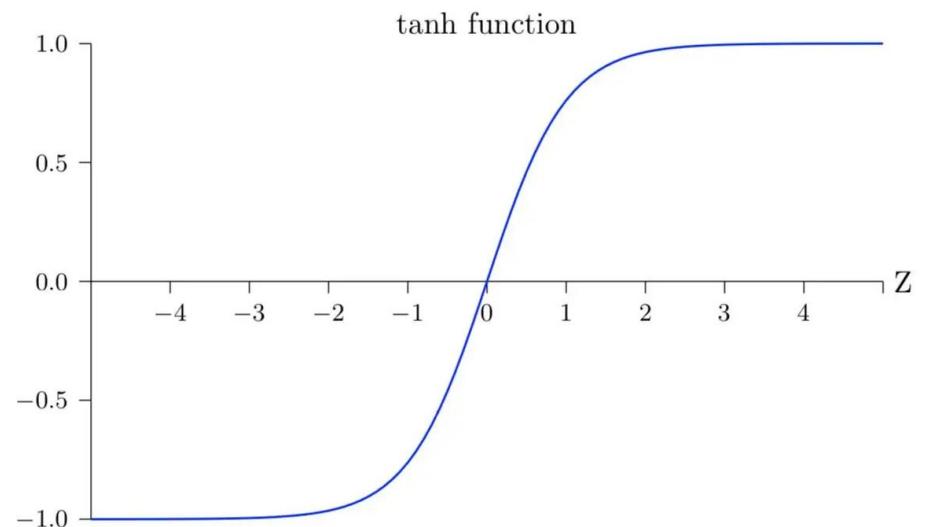


Gradient Vanishing Problem

$$\frac{\partial \mathbf{L}}{\partial w_{ij}^r} = \frac{\partial \mathbf{L}}{\partial \mathbf{a}^L} \cdots \cdots \frac{\partial \mathbf{a}^{r+1}}{\partial a_i^r} \frac{\partial a_i^r}{\partial z_i^r} \frac{\partial z_i^r}{\partial w_{ij}^r}$$

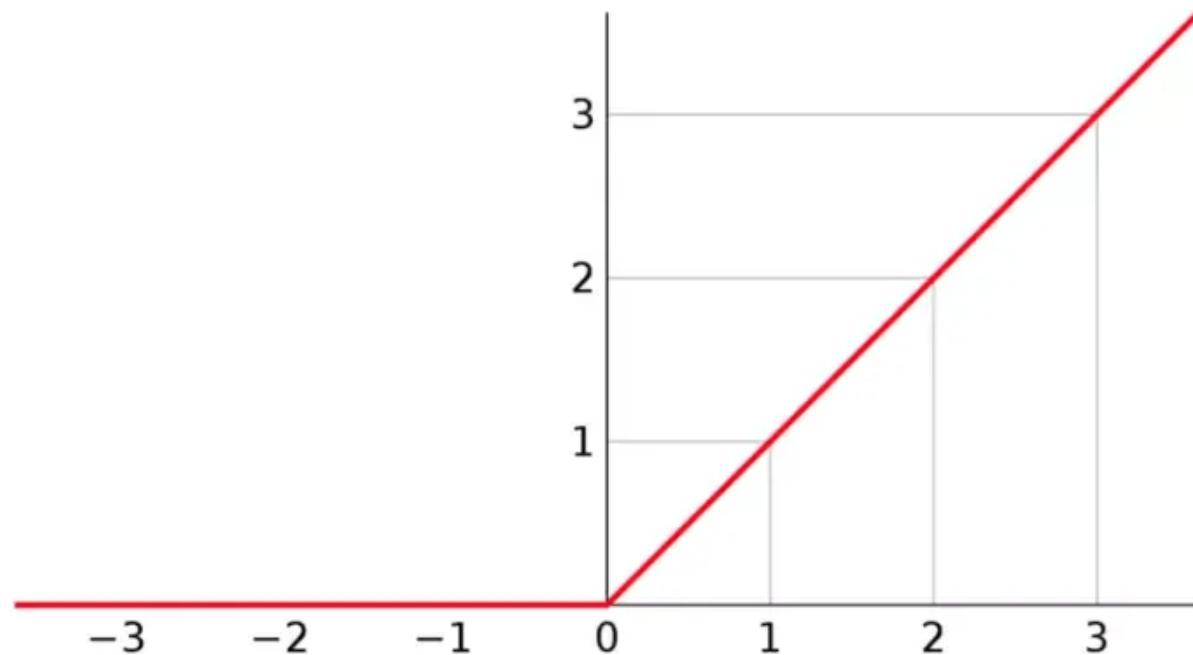


Sigmoid function are now often used as gate unit



Activation Function: ReLU

$$\begin{aligned}a(z) &= \text{ReLU}(z) \\&= \max(0, z)\end{aligned}$$



Design Good Loss Functions

- Different tasks should have specific loss functions

- MSE
- Cross entropy
- Focal loss
-

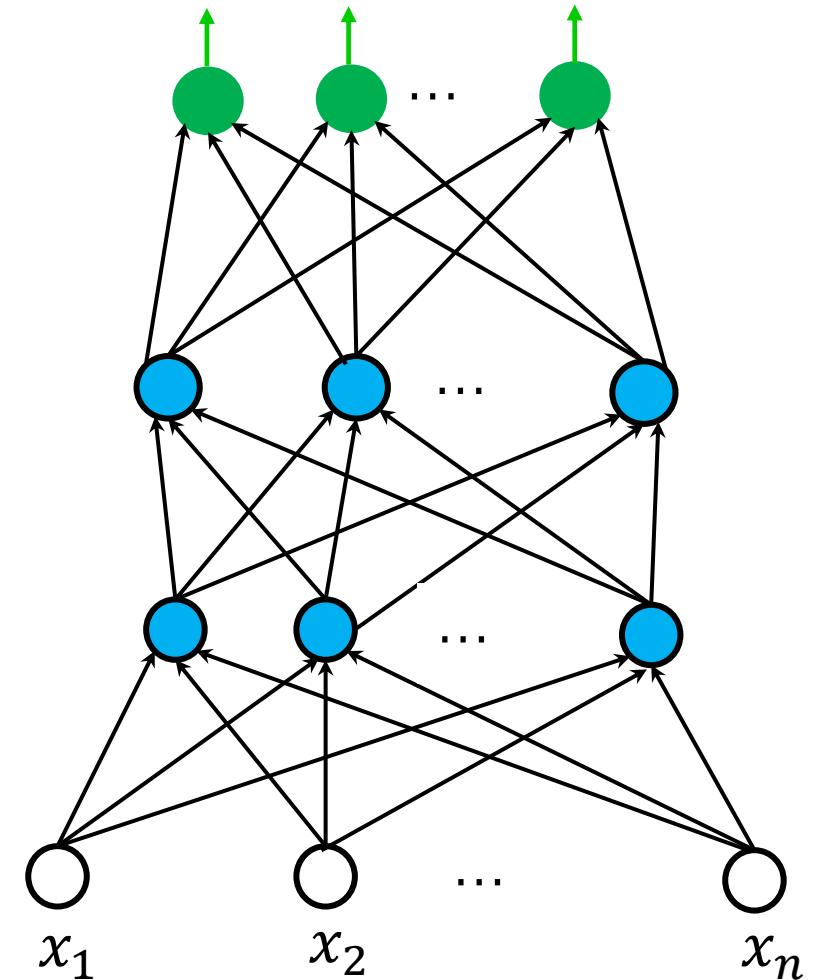
$$L(W) = \frac{1}{2} \sum_{i=1}^m \| \mathbf{y}^{(i)} - \hat{\mathbf{y}}^{(i)} \|^2$$

$$L(W) = -\frac{1}{m} \sum_{i=1}^m \mathbf{y}^{(i)} \ln(\hat{\mathbf{y}}^{(i)})$$

- Adding regularization terms: L1, L2,

Parameter Initialization Techniques

- Initialize weights with Small random numbers
 - Set bias to zero initially
- Layer-wise pretraining
- Pretrained model+finetune



Employ Mini-batch instead of SGD

- Bias and Variance
- Small mini-batch usually don't change the bias expectation of SGD
- Mini-batch can fully exploit the GPU power
- Big mini-batch can have lower variance without big loss of bias expectation
 - Can have bigger learning rate
 - Correspondingly, with small mini-batch, learning rate should also be small
- Can set the learning rate linearly with mini-batch size

Gradient Updating Optimization

- The problem of gradient estimation error

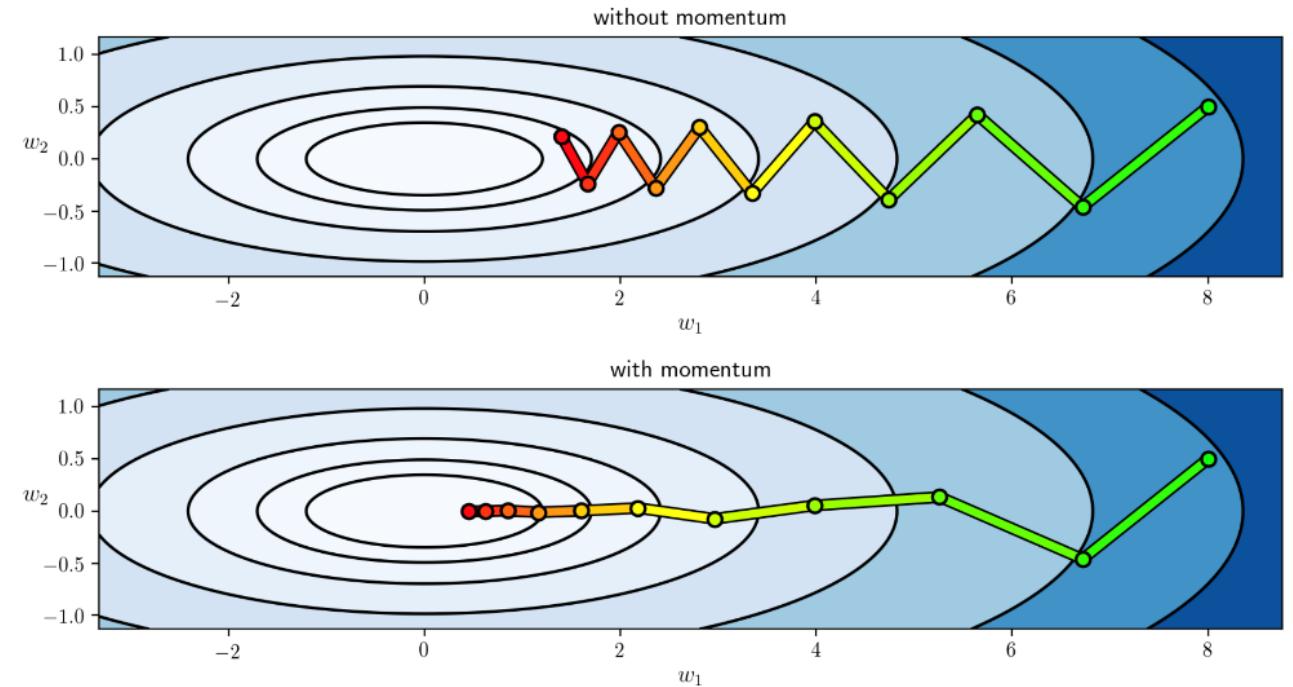
- Momentum+

Adaptive Learning Rate

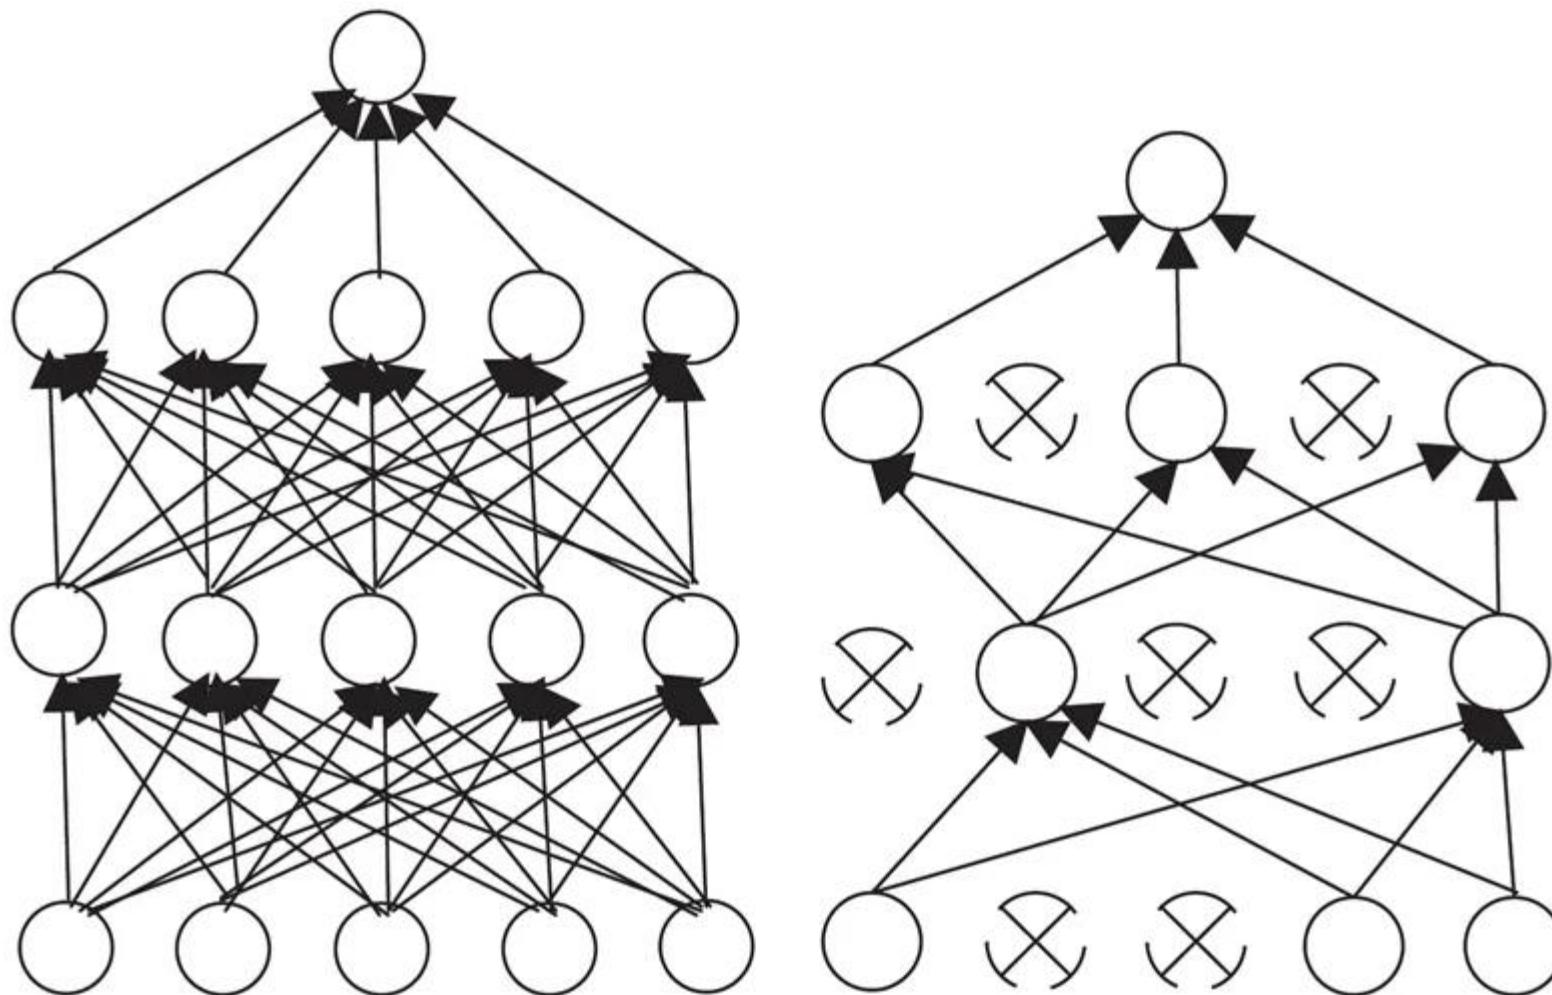
- Adam

$$\Delta\theta_t = \rho\Delta\theta_{t-1} - \alpha\mathbf{g}_t$$

- Gradient clipping



Dropout



Other Classical Techniques

- Batch normalization
- Layer normalization
- Weight decay
- Early stopping
- Hyperparameter optimization
 - Grid search
 - Bayesian optimization

Acknowledgements

- Some text, figures and formulations are from WWW.
- This lecture is distributed for nonprofit purpose.

Thank You for Your Attention

Contact me at: yym@hit.edu.cn

Tel: 26033008, 13760196623

Address: Rm.1402, H# Building

