



哈爾濱工業大學(深圳)

HARBIN INSTITUTE OF TECHNOLOGY, SHENZHEN

操作系统

Operating Systems

刘川意 教授

哈尔滨工业大学 (深圳)

2025年9月

Module 3: 调度

1. Mechanism: Limited Direct Execution

1.1 Direct Execution的局限性 – 基础

- 问题1: Restricted Operation – 基础
- 问题2: Switching Between Process – 基础

1.2 Limited Direct Execution – 基础

2. Policy: Scheduling Algorithms

2.1 Introduction – 基础

2.2 The Multi-Level Feedback Queue – 基础

2.3 Proportional Share – 进阶

OS如何可控的实现CPU虚拟化

- OS内核实现物理CPU在进程间共享的基本思路: time sharing.
- 要解决2个核心问题
 - **Performance:** How can we implement virtualization without adding excessive overhead to the system? 即引入的额外性能尽可能小
 - **Control:** How can we run processes efficiently while retaining control over the CPU? 即不要“跑飞”了, 不要被“劫持”了

用户program直接运行在物理CPU上

- Just run the program directly on the CPU.

OS	Program
<ol style="list-style-type: none">1. Create entry for process list2. Allocate memory for program3. Load program into memory4. Set up stack with argc / argv5. Clear registers6. Execute call main()9. Free memory of process10. Remove from process list	<ol style="list-style-type: none">7. Run main()8. Execute return from main()

Without *limits* on running programs,
the OS wouldn't be in control of anything and
thus would be "**just a library**"

让program无限制的运行在物理CPU上，则OS沦为library

Problem 1: Restricted Operation

- What if a process wishes to perform some kind of restricted operation such as ... 如果一个进程希望执行某种受限操作, 例如
 - Issuing an I/O request to a disk 向磁盘发出 I/O 请求
 - Gaining access to more system resources such as CPU or memory 获得对更多系统资源 (例如 CPU 或内存) 的访问权限
- **Solution:** Using 保护模式 (protected control transfer)
 - **User mode:** Applications do not have full access to hardware resources. 应用程序没有对硬件资源的完全访问权限
 - **Kernel mode:** The OS has access to the full resources of the machine 操作系统可以访问机器的全部资源

System Call

- Allow the kernel to **carefully expose** certain key pieces of functionality to user program, such as ...
允许内核谨慎地向用户程序公开某些关键功能，例如
 - Accessing the file system 访问文件系统
 - Creating and destroying processes 创建和销毁进程
 - Communicating with other processes 与其他进程通信
 - Allocating more memory 分配更多内存
- **Trap** instruction
 - Jump into the kernel 进入内核态
 - Raise the privilege level to kernel mode 将权限级别提升到内核模式
- **Return-from-trap** instruction
 - Return into the calling user program 返回到调用用户程序
 - Reduce the privilege level back to user mode 将权限级别降低回用户模式

Limited Direction Execution Protocol

OS @ boot (kernel mode)	Hardware	
initialize trap table 初始化trap表	remember address of syscall handler 记住系统调用处理程 序的地址	
OS @ run (kernel mode)	Hardware	Program (user mode)
<ul style="list-style-type: none"> ➤ Create entry for process list 为进程列表创建条目 ➤ Allocate memory for program 为程序分配内存 ➤ Load program into memory 将程序加载到内存中 ➤ Setup user stack with argv 使用 argv 设置用户堆栈 ➤ Fill kernel stack with reg/PC 使用 reg/PC 填充内核堆栈 ➤ return-from-trap <ul style="list-style-type: none"> ➤ restore regs from kernel stack 从内核堆栈恢复寄存器 ➤ move to user mode 移动到用户模式 ➤ jump to main 跳转到主程序 		<ul style="list-style-type: none"> ➤ Run main() ➤ ... ➤ Call system ➤ trap into OS

Limited Direction Execution Protocol (Cont.)

OS @ run (kernel mode)	Hardware	Program (user mode)
	<ul style="list-style-type: none">➤ save regs to kernel stack 将寄存器保存到内核堆栈➤ move to kernel mode 移动到内核态➤ jump to trap handler 跳转到陷阱处理程序	
<ul style="list-style-type: none">➤ Handle trap➤ Do work of syscall➤ return-from-trap	<ul style="list-style-type: none">➤ restore regs from kernel stack 从内核堆栈恢复寄存器➤ move to user mode 移动到用户态➤ jump to PC after trap 陷阱后跳转到PC	<ul style="list-style-type: none">➤ ...➤ return from main➤ trap (via exit())
<ul style="list-style-type: none">➤ Free memory of process➤ Remove from process list		

Problem 2: Switching Between Processes

- How can the OS **regain control** of the CPU so that it can switch between processes? 操作系统如何重新获得对 CPU 的控制权，以便在进程之间切换
 - A cooperative Approach: **Wait for system calls**
一种合作的方法：等待系统调用
 - A Non-Cooperative Approach: **The OS takes control**
一种非合作方法：操作系统控制

A cooperative Approach: Wait for system calls

- Processes periodically give up the CPU by making **system calls** such as yield. **进程周期性地通过系统调用（例如yield）放弃CPU**
 - The OS decides to run some other task. **操作系统决定运行一些其他任务。**
 - Application also transfer control to the OS when they do something illegal. **当应用程序做一些非法的事情时，它们也会将控制权转移给操作系统。**
 - Divide by zero **除0**
 - Try to access memory that it shouldn't be able to access **非法访问**
- **Example:** Early versions of the Macintosh OS, The old Xerox Alto system
示例：Macintosh OS 的早期版本，旧的 Xerox Alto 系统

A process gets stuck in an infinite loop.
→ **Reboot the machine**

A Non-Cooperative Approach: OS Takes Control

□ A timer interrupt

- During the boot sequence, the OS start the timer.
在启动序列中，操作系统启动计时器
- The timer raise an interrupt every so many milliseconds.
计时器每隔这么多毫秒引发一次中断
- When the interrupt is raised :
 - ▶ The currently running process is halted. 当前运行的进程被暂停
 - ▶ Save enough of the state of the program. 保存足够的程序状态
 - ▶ A pre-configured interrupt handler in the OS runs.
操作系统中预配置的中断处理程序运行

A **timer interrupt** gives OS the ability to
run again on a CPU.

Saving and Restoring Context

- Scheduler makes a decision: 调度器做出决定
 - Whether to continue running the **current process**, or switch to a **different one**. 是继续运行当前进程，还是切换到其他进程
 - If the decision is made to switch, the OS executes context switch. 如果决定切换，则操作系统执行上下文切换。

Context Switch

- A low-level piece of assembly code 一段低级的汇编代码
 - **Save a few register values** for the current process onto its kernel stack 将当前进程的一些寄存器值保存到其内核堆栈中
 - General purpose registers 通用寄存器
 - PC
 - kernel stack pointer 内核堆栈指针
 - **Restore a few** for the soon-to-be-executing process from its kernel stack 从内核堆栈中为即将执行的进程恢复一些
 - **Switch to the kernel stack** for the soon-to-be-executing process 切换到即将执行的进程的内核堆栈

Module 3: 调度

1. Mechanism: Limited Direct Execution

1.1 Direct Execution的局限性 – 基础

- 问题1: Restricted Operation – 基础
- 问题2: Switching Between Process – 基础

1.2 Limited Direct Execution – 基础

2. Policy: Scheduling Algorithms

2.1 Introduction – 基础

2.2 The Multi-Level Feedback Queue – 基础

2.3 Proportional Share – 进阶

Limited Direction Execution Protocol (Timer interrupt)

OS @ boot
(kernel mode)

Hardware

initialize trap table

remember address of ...
syscall handler
timer handler

start interrupt timer

start timer
interrupt CPU in X ms

OS @ run
(kernel mode)

Hardware

Program
(user mode)

Process A

...

timer interrupt

save regs(A) to k-stack(A)
move to kernel mode
jump to trap handler

Limited Direction Execution Protocol (Timer interrupt) (Cont.)

OS @ run (kernel mode)	Hardware	Program (user mode)
---------------------------	----------	------------------------

Handle the trap

Call switch() routine

save regs(A) to proc-struct(A)

restore regs(B) from proc-struct(B)

switch to k-stack(B)

return-from-trap (into B)

Hardware

Program
(user mode)

restore regs(B) from k-stack(B)

move to user mode

jump to B's PC

Process B

...

The xv6 Context Switch Code

```

1 # void swtch(struct context **old, struct context *new);
2 #
3 # Save current register context in old
4 # and then load register context from new.
5 .globl swtch
6 swtch:
7     # Save old registers
8     movl 4(%esp), %eax          # put old ptr into eax
9     popl 0(%eax)              # save the old IP
10    movl %esp, 4(%eax)         # and stack
11    movl %ebx, 8(%eax)         # and other registers
12    movl %ecx, 12(%eax)
13    movl %edx, 16(%eax)
14    movl %esi, 20(%eax)
15    movl %edi, 24(%eax)
16    movl %ebp, 28(%eax)
17
18     # Load new registers
19    movl 4(%esp), %eax          # put new ptr into eax
20    movl 28(%eax), %ebp         # restore other registers
21    movl 24(%eax), %edi
22    movl 20(%eax), %esi
23    movl 16(%eax), %edx
24    movl 12(%eax), %ecx
25    movl 8(%eax), %ebx
26    movl 4(%eax), %esp          # stack is switched here
27    pushl 0(%eax)              # return addr put in place
28    ret                         # finally return into new ctxt

```

Worried About Concurrency?

- What happens if, during interrupt or trap handling, another interrupt occurs?
如果在中断或陷阱处理期间发生另一个中断，会发生什么
- OS handles these situations: **操作系统处理这些情况**
 - **Disable interrupts** during interrupt processing **在中断处理期间禁用中断**
 - Use a number of sophisticated **locking** schemes to protect concurrent access to internal data structures.
使用许多复杂的锁定方案来保护对内部数据结构的并发访问

Module 3: 调度

1. Mechanism: Limited Direct Execution

2. Policy: Scheduling Algorithms

2.1 Introduction – 基础

- 概述 – 基础
- 经典调度算法 – 基础
 - ▶ 先进先出 FIFO – 基础
 - ▶ 最短任务优先 SJF – 基础
 - ▶ 轮转调度 RR – 基础

2.2 The Multi-Level Feedback Queue – 基础

2.3 Proportional Share – 进阶

Scheduling: Introduction

- Workload assumptions (工作负载初始化假设, 后续算法逐步放宽这些假设):
 1. Each job runs for the **same amount of time**. 每个作业运行相同的时间
 2. All jobs **arrive** at the same time. 所有工作同时到达
 3. All jobs only use the **CPU** (i.e., they perform no I/O). 所有作业仅使用 CPU (例如, 它们不执行 I/O)
 4. The **run-time** of each job is known. 每个作业的运行时间是已知的

Scheduling Metrics 调度指标

- Performance metric: Turnaround time (周转时间)
 - The time at which **the job completes** minus the time at which **the job arrived** in the system. (任务完成时间减去任务到达系统的时间)

$$T_{turnaround} = T_{completion} - T_{arrival}$$


- Another metric is fairness (公平).
 - Performance and fairness are often at odds in scheduling.
性能和公平性在调度中经常不一致
- The time from **when the job arrives** to the **first time it is scheduled**.
响应时间是指:从任务到达系统到首次运行(首次被调度)的时间

$$T_{response} = T_{firstrun} - T_{arrival}$$

- STCF and related disciplines are not particularly good for response time. STCF最短完成时间优先策略, 用响应时间来评估不是好方法

First In, First Out 先进先出 (FIFO)


- First Come, First Served (FCFS)
 - Very simple and easy to implement
- Example:
 - A arrived just before B which arrived just before C.
 - Each job runs for 10 seconds.



$$\text{Average turnaround time} = \frac{10 + 20 + 30}{3} = 20 \text{ sec}$$

Why FIFO is not that great? – Convoy effect (护航效应)


- Let's relax assumption 1: Each job **no longer** runs for the same amount of time. 让我们放宽假设 1: 每个作业不再运行相同的时间
- Example:
 - A arrived just before B which arrived just before C.
 - A runs for 100 seconds, B and C run for 10 each.



$$\text{Average turnaround time} = \frac{100 + 110 + 120}{3} = 110 \text{ sec}$$

Shortest Job First (最短任务优先 SJF)


- Run the shortest job first, then the next shortest, and so on
首先运行最短的作业，然后运行下一个最短的作业，依此类推
 - Non-preemptive (非抢占式) scheduler
- Example:
 - A arrived just before B which arrived just before C.
 - A runs for 100 seconds, B and C run for 10 each.



$$\text{Average turnaround time} = \frac{10 + 20 + 120}{3} = 50 \text{ sec}$$

SJF with Late Arrivals from B and C

- Let's relax assumption 2: Jobs can arrive at any time. 让我们放宽假设
2: 工作可以随时到达
- Example:
 - A arrives at t=0 and needs to run for 100 seconds.
A到达t=0, 需要运行100秒
 - B and C arrive at t=10 and each need to run for 10 seconds
[B,C arrive]



$$\text{Average turnaround time} = \frac{100 + (110 - 10) + (120 - 10)}{3} = 103.33 \text{ sec}$$

Shortest Time-to-Completion First

—— 最短完成时间优先 STCF


- Add **preemption** (抢占) to SJF
 - Also known as Preemptive Shortest Job First (PSJF)
也称为抢占式最短任务优先
- A new job enters the system:
 - Determine of the remaining jobs and new job **确定剩余工作和新工作**
 - Schedule the job which has the least time left **安排剩余时间最少的作业**

Shortest Time-to-Completion First

—— 最短完成时间优先 STCF

□ Example:

- A arrives at t=0 and needs to run for 100 seconds.
- B and C arrive at t=10 and each need to run for 10 seconds



$$\text{Average turnaround time} = \frac{(120 - 0) + (20 - 10) + (30 - 10)}{3} = 50 \text{ sec}$$

Shortest Job First (最短任务优先 SJF)

□ 【例题】现在有三个同时到达的作业J1、J2和J3，它们的执行时间分别是T1、T2和T3，而且 $T_1 < T_2 < T_3$ 。系统按单道方式运行且采用短作业优先调度算法，则平均周转时间是（ ）。

- A. $T_1+T_2+T_3$
- B. $(3T_1+2T_2+T_3)/3$
- C. $(T_1+T_2+T_3)/3$
- D. $(T_1+2T_2+3T_3)/3$

答案：B

系统采用短作业优先调度算法，作业的执行顺序为J1、J2和J3，它们的周转时间分别为 T_1 ， T_1+T_2 和 $T_1+T_2+T_3$ 。

所以平均周转时间为 $(3T_1+2T_2+T_3)/3$ 。

Shortest Job First (最短任务优先 SJF)

- 【例题】假设4个任务到达系统的时刻和运行时间见下表
系统在t=2时开始作业调度。若分别采用先来先服务和最
短任务优先调度算法，则选中的任务分别是（ ）

A、J₂、J₃

B、J₁、J₄

C、J₂、J₄

D、J₁、J₃

作业	到达时间t	运行时间
J ₁	0	3
J ₂	1	3
J ₃	1	2
J ₄	3	1

答案： D

解析：先来先服务调度算法时作业来得越早，优先级越高，因此会选择J₁。短作业优先调度算法是作业运行时间越短，优先级越高，因此会选择J₃。

New scheduling metric: Response time

——新的调度指标：响应时间

- The time from **when the job arrives** to the **first time it is scheduled**.

响应时间是指：从任务到达系统到首次运行(首次被调度)的时间

$$T_{response} = T_{firstrun} - T_{arrival}$$

- STCF and related disciplines are not particularly good for response time.
STCF 和相关策略对响应时间不是特别好

How can we build a scheduler that is
sensitive to response time?


Round Robin (RR) Scheduling

——轮转调度


- Time slicing Scheduling
 - Run a job for a **time slice**(时间片) and then switch to the next job in the **run queue** until the jobs are finished.
 - ▶ Time slice is sometimes called a scheduling quantum(调度量子).
 - It repeatedly does so until the jobs are finished.
 - The length of a time slice must be a *multiple of* the timer-interrupt period (时钟中断周期). **时间片的长度必须是定时器中断周期的倍数**

RR Scheduling Example

- A, B and C arrive at the same time.
- They each wish to run for 5 seconds.



$$T_{average\ response} = \frac{0 + 5 + 10}{3} = 5\ sec$$



$$T_{average\ response} = \frac{0 + 1 + 2}{3} = 1\ sec$$

Round Robin (RR) Scheduling

——轮转调度

□ Time slicing Scheduling

- Run a job for a time slice(时间片) and then switch to the next job in the run queue until the jobs are finished.

将一个作业运行一个时间片（时间片），然后切换到运行队列中的下一个作业，直到作业完成

- ▶ Time slice is sometimes called a scheduling quantum(调度量子).

- It repeatedly does so until the jobs are finished.

它反复这样做，直到作业完成

- The length of a time slice must be a *multiple of* the timer-interrupt period (时钟中断周期). 时间片的长度必须是定时器中断周期的倍数

RR is fair, but performs poorly on metrics
such as turnaround time

RR 是公平的，但在指标上表现不佳
比如周转时间

Round Robin (RR) Scheduling

——轮转调度

【例题】下列有关时间片的进程调度的描述中，错误的是（ ）

- A. 时间片越短，进程切换的次数越多，系统开销也越大
- B. 当前进程的时间片用完后，该进程状态由执行态变为阻塞态
- C. 时钟中断发生后，系统会修改当前的进程在时间片内的剩余时间
- D. 影响时间片大小的主要因素包括响应时间、系统开销和进程数量等

答案：B

解析：进程切换带来系统开销，切换次数越多，系统开销越大，选项A正确。当前进程的时间片用完后，其状态由执行态变为就绪态，选项B错误。时钟中断是系统中特定的周期性时钟节拍，操作系统通过它来确定时间间隔，实现时间的延时跟任务的超时，选项C正确。现代操作系统为了保证性能最优，通常根据响应时间、系统开销、进程数目、进程运行时间、进程切换开销等因素确定时间片大小，选项D正确。

Round Robin (RR) Scheduling

——轮转调度

【例题】假设某系统使用时间片轮转调度算法进行CPU调度，时间片大小为5ms，系统共有10个进程，初始时均处于就绪队列，执行结束前仅处于执行态或就绪态。若队尾的进程P所需CPU时间为25ms，在不考虑系统开销的情况下，则进程P的周转时间为

- A. 200ms B. 205ms
- C. 250ms D. 295ms

答案: C

解析: 进程P所需CPU时间为25ms，因为时间片为5ms，进程P需要5轮时间片才能完成执行。

每轮中，10个进程均需执行一个时间片，总时间为：

$$10 \times 5\text{ms} = 50\text{ms}$$

5轮共250ms

The length of the time slice is critical.

- The shorter time slice
 - Better response time
 - The cost of context switching will dominate overall performance.
上下文切换的成本将主导整体性能


- The longer time slice
 - Amortize the cost of context switching 分摊上下文切换的成本
 - Worse response time

Deciding on the length of the time slice presents
a **trade-off** to a system designer


Incorporating I/O

- Let's relax assumption 3: **All programs can perform I/O**
- Example:
 - A and B need 50ms of CPU time each.
 - A runs for 10ms and then issues an I/O request
 - ▶ I/Os each take 10ms
 - B simply uses the CPU for 50ms and performs no I/O
 - 为显示I/O对调度的影响, 假设使用最简化的调度算法: The scheduler runs A first, then B after

Incorporating I/O (Cont.)



Poor Use of Resources



Maximize the
CPU utilization

Overlap Allows Better Use of Resources

Incorporating I/O (Cont.)

- When a job initiates an I/O request.
 - The job is blocked waiting for I/O completion. 作业被阻塞等待 I/O 完成
 - The scheduler should schedule another job on the CPU.
调度程序应该在 CPU 上调度另一个作业

- When the I/O completes
 - An interrupt is raised.
 - The OS moves the process from blocked back to the ready state.
操作系统将进程从阻塞状态移回就绪状态

Incorporating I/O (Cont.)

□ 【例题】一个多道批处理系统中仅有P1和P2两个作业，P2比P1晚5s到达，它们的计算和I/O操作顺序如下：

P1：计算60ms，I/O：80ms，计算20ms


P2：计算120ms，I/O：40ms，计算40ms

若不考虑调度和切换时间，则完成两个作业需要的时间最少是（ ）。

- A. 240ms B. 260ms C. 340ms D. 360ms

答案：B

由于P2比P1晚5ms到达，P1先占用CPU，作业运行的甘特图如下。



Module 3: 调度

1. Mechanism: Limited Direct Execution

2. Policy: Scheduling Algorithms

2.1 Introduction – 基础

□ 概述 – 基础

□ 经典调度算法 – 基础

2.2 The Multi-Level Feedback Queue – 基础

2.3 Proportional Share – 进阶

Multi-Level Feedback Queue (MLFQ)

——多级反馈队列

- A Scheduler that learns from the past to predict the future.
- Objective (目标):
 - Optimize **turnaround time** → Run shorter jobs first
 - Minimize **response time** without a 先验知识 (*priori knowledge*) of 任务运行时间 (*job length*)

MLFQ: Basic Rules


- MLFQ has a number of distinct **queues**.
- Each queue is assigned a different priority level.
每一个队列赋予不同的优先级
- A job that is ready to run is on a single queue.
 - A job **on a higher queue** is chosen to run. 不同队列, 高级队列先调度
 - Use round-robin scheduling among jobs in the same queue 同队列, 采用R-R调度

Rule 1: If $\text{Priority}(A) > \text{Priority}(B)$, A runs (B doesn't).
Rule 2: If $\text{Priority}(A) = \text{Priority}(B)$, A & B run in RR.

MLFQ: Basic Rules (Cont.)

- MLFQ varies the priority of a job based on its observed behavior. 任务的优先级动态变化
- Example:
 - A job repeatedly relinquishes (放弃) the CPU while waiting IOs → Keep its priority high 作业在等待 IO 时反复放弃 (放弃) CPU → 保持其优先级高
 - A job uses the CPU intensively for long periods of time → Reduce its priority. 作业长时间密集使用 CPU → 降低其优先级


MLFQ Example



Q7


Q6

Q5



Q3

Q2




Attempt 1: How to Change Priority

- MLFQ priority adjustment algorithm:
 - **Rule 3:** When a job enters the system, it is placed at the highest priority
当作业进入系统时，它被置于最高优先级
 - **Rule 4a:** If a job uses up an entire time slice while running, its priority is reduced (i.e., it moves down on queue).
计算密集型任务，优先级逐步调低
 - **Rule 4b:** If a job gives up the CPU before the time slice is up, it stays at the same priority level
交互型任务，优先级保留

In this manner, MLFQ approximates SJF

Example 1: A Single Long-Running Job

- A three-queue scheduler with time slice 10ms



Example 2: Along Came a Short Job

实例2：来了一个短任务

□ Assumption:


- **Job A:** A long-running CPU-intensive job
- **Job B:** A short-running interactive job (20ms runtime)
- A has been running for some time, and then B arrives at time $T=100$.



Example 3: What About I/O?

- Assumption:

- **Job A:** A long-running CPU-intensive job
- **Job B:** An interactive job that need the CPU only for 1ms before performing an I/O



The MLFQ approach keeps an interactive job at the highest priority

Problems with the Basic MLFQ


基础MLFQ的一些问题

- Starvation 首先, 会有饥饿问题
 - If there are “too many” interactive jobs in the system.
 - Long-running jobs will never receive any CPU time.
- Game the scheduler 其次, 聪明的用户可能重写程序来愚弄调度程序
 - After running 99% of a time slice, issue an I/O operation.
 - The job gains a higher percentage of CPU time.
- A program may change its behavior over time. 最后, 一个程序可能在不同时间段表现不同
 - CPU bound process → I/O bound process

Attempt 2: The Priority Boost


尝试2：提升优先级

- **Rule 5:** After some time period S, move all the jobs in the system to the topmost queue. 经过一段时间 S 后，将系统中的所有作业移动到最顶层队列
 - Example:
 - ▶ A long-running job(A) with two short-running interactive job(B, C)



Attempt 3: Better Accounting


- How to prevent gaming of our scheduler?
 - Rules 4a and 4b : which let a job retain its priority by relinquishing the CPU before the time slice expires 通过在时间片到期之前放弃 CPU 让作业保留其优先级
- Solution:
 - **Rule 4** (Rewrite Rules 4a and 4b): Once a job **uses up its time allotment** (时间配额) at a given level (regardless of how many times it has given up the CPU), **its priority is reduced**(i.e., it moves down on queue).
一旦作业在给定级别用完其时间分配 (时间损耗) (无论它放弃了多少次 CPU) , 它的优先级就会降低 (即, 它在队列中向下移动)



Tuning MLFQ And Other Issues

Lower Priority, Longer Quanta

- The high-priority queues → Short time slices
 - ▶ E.g., 10 or fewer milliseconds
- The Low-priority queue → Longer time slices
 - ▶ E.g., 100 milliseconds



Example) 10ms for the highest queue, 20ms for the middle,
40ms for the lowest

The Solaris MLFQ implementation

- For the Time-Sharing scheduling class (TS)
 - 60 Queues
 - Slowly increasing time-slice length
 - ▶ The highest priority: 20msec
 - ▶ The lowest priority: A few hundred milliseconds
 - Priorities boosted around every 1 second or so.

MLFQ: Summary

- The refined set of MLFQ rules:
 - **Rule 1:** If Priority(A) > Priority(B), A runs (B doesn't).
 - **Rule 2:** If Priority(A) = Priority(B), A & B run in RR.
 - **Rule 3:** When a job enters the system, it is placed at the highest priority.
 - **Rule 4:** Once a job uses up its time allotment at a given level (regardless of how many times it has given up the CPU), its priority is reduced(i.e., it moves down on queue).
 - **Rule 5:** After some time period S, move all the jobs in the system to the topmost queue.

MLFQ

【例题】系统采用二级反馈队列调度算法进行进程调度。就绪队列Q1采用时间片轮转调度算法，时间片为10ms；就绪队列Q2采用短进程优先调度算法；系统优先调度Q1队列中的进程，当Q1为空时系统才会调度Q2中的进程；新创建的进程首先进入Q1；Q1中的进程执行一个时间片后，若未结束，则转入Q2。若当前Q1、Q2为空，系统依次创建进程P1、P2后即开始进程调度，P1、P2需要的CPU时间为30ms和20ms，则进程P1、P2在系统中的平均等待时间为（）。


- A. 25ms B. 20ms C. 15ms D. 10ms

答案：C

解析：进程P1、P2依次进入Q1后，将依次被分配10ms的CPU时间，两个进程执行完一个时间片后转入Q2队列。此时P1还需要20ms，P2还需要10ms，根据Q2的SJF算法，P2会被优先调度进入Q1，10ms后P2执行结束，之后P1再被调度执行，20ms后执行结束。

进程P1、P2的等待时间为图中的虚横线部分。

$$\begin{aligned} \text{平均等待时间} &= (\text{P1等待时间} + \text{P2等待时间}) / 2 \\ &= (20 + 10) / 2 = 15 \end{aligned}$$



Module 3: 调度

1. Mechanism: Limited Direct Execution

2. Policy: Scheduling Algorithms

2.1 Introduction – 基础

□ 概述 – 基础

□ 经典调度算法 – 基础

2.2 The Multi-Level Feedback Queue – 基础

2.3 Proportional Share – 进阶

Proportional Share Scheduler

比例份额调度

- Fair-share scheduler 公平份额调度程序
 - Guarantee that each job obtain a *certain percentage* of CPU time.
保证每个作业获得一定比例的CPU时间
 - Not optimized for turnaround or response time
未针对周转时间或响应时间进行优化

Basic Concept

- Tickets (彩票数)
 - Represent the share of a resource that a process should receive
表示进程应接收的资源份额
 - The percent of tickets represents its share of the system resource in question.
彩票数的百分比表示其在相关系统资源中的份额
- Example
 - There are two processes, A and B.
 - ▶ Process A has 75 tickets → receive 75% of the CPU
 - ▶ Process B has 25 tickets → receive 25% of the CPU

Lottery scheduling 彩票调度

- The scheduler picks a winning ticket.
 - Load the state of that *winning process* and runs it.
加载该获胜进程的状态并运行它
- Example
 - There are 100 tickets
 - ▶ Process A has 75 tickets: 0 ~ 74
 - ▶ Process B has 25 tickets: 75 ~ 99

Scheduler's winning tickets:	63 85 70 39 76 17 29 41 36 39 10 99 68 83 63
Resulting scheduler:	A B A A B A A A A A B A B A

The longer these two jobs compete,
The more likely they are to achieve the desired percentages.

Ticket Mechanisms 彩票机制

□ Ticket currency 彩票货币

- A user allocates tickets among their own jobs in whatever currency they would like. 拥有彩票的用户将彩票分配给他们的不同工作。
- The system converts the currency into the correct global value.
之后，系统会自动地将货币兑换为正确的全局变量。
- Example
 - ▶ There are 200 tickets (Global currency)
 - ▶ User A has 100 tickets
 - ▶ User B has 100 tickets

User A → 500 (A's currency) to A1 → 50 (global currency)
→ 500 (A's currency) to A2 → 50 (global currency)

User B → 10 (B's currency) to B1 → 100 (global currency)

Ticket Mechanisms 彩票机制 (Cont.)


- Ticket transfer 彩票转让
 - A process can temporarily hand off its tickets to another process.
一个进程可以暂时将它的彩票交给另一个进程。

- Ticket inflation 彩票通胀
 - A process can temporarily raise or lower the number of tickets it owns.
进程可以临时增加或减少拥有的彩票数量。
 - If any one process needs *more CPU time*, it can boost its tickets.
如果任何一个进程需要更多的 CPU 时间，它可以增加它的彩票。

Implementation 实现

- Example: There are three processes, A, B, and C.

- Keep the processes in a list:



```


1      // counter: used to track if we've found the winner yet
2      int counter = 0;
3
4      // winner: use some call to a random number generator to
5      // get a value, between 0 and the total # of tickets
6      int winner = getrandom(0, totaltickets);
7
8      // current: use this to walk through the list of jobs
9      node_t *current = head;
10
11     // loop until the sum of ticket values is > the winner
12     while (current) {
13         counter = counter + current->tickets;
14         if (counter > winner)
15             break; // found the winner
16         current = current->next;
17     }
18     // 'current' is the winner: schedule it...
  
```

Example: Lottery Scheduling

- U : unfairness metric 不公平指标
 - The time the first job completes divided by the time that the second job completes. 两个任务完成时刻相除得到U的值
- Example:
 - There are two jobs, each jobs has runtime 10.
 - ▶ First job finishes at time 10
 - ▶ Second job finishes at time 20
 - $U = \frac{10}{20} = 0.5$
 - U will be close to 1 when both jobs finish at nearly the same time.
当两个工作几乎同时完成时, U 将接近 1

Lottery Fairness Study 彩票公平性研究

- There are two jobs.
- Each job has the same number of tickets (100).



When the job length is not very long,
average unfairness can be quite severe.

Stride Scheduling 步长调度

- Stride of each process 每个进程的步长(与彩票的数量成反比):
 - $(\text{A large number}) / (\text{the number of tickets of the process})$
 - Example: A large number = 10,000
 - ▶ Process A has 100 tickets → stride of A is 100
 - ▶ Process B has 50 tickets → stride of B is 200
 - ▶ Process C has 250 tickets → stride of C is 40
- A process runs, increment a counter(=pass value) for it by its stride.
进程运行后, 计数器(称为行程 pass)值每次增加它的步长大小
 - Pick the process to run that has the lowest pass value

```
current = remove_min(queue);           // pick client with minimum pass
schedule(current);                   // use resource for quantum
current->pass += current->stride;    // compute next pass using stride
insert(queue, current);             // put back into the queue
```

A pseudo code implementation

Stride Scheduling Example

Pass(A) (stride=100)	Pass(B) (stride=200)	Pass(C) (stride=40)	Who Runs?
0	0	0	A
100	0	0	B
100	200	0	C
100	200	40	C
100	200	80	C
100	200	120	A
200	200	120	C
200	200	160	C
200	200	200	...

If new job enters with pass value 0,
It will **monopolize** the CPU!