



哈爾濱工業大學(深圳)

HARBIN INSTITUTE OF TECHNOLOGY, SHENZHEN

操作系统

Operating Systems

刘川意 教授

liuchuanyi@hit.edu.cn

哈尔滨工业大学 (深圳)

2025年9月

并发与同步

1. **Concurrency(并发) Introduction**
2. **Locks**
3. **基于Lock的并发数据结构**
4. **Condition Variables 条件变量**
5. **Semaphore 信号量**
6. **常见并发问题**
7. **基于事件的并发**

Concurrency为什么放到OS中讲?

□ History !

- OS Kernel是第一个并发程序，如：[write\(\)设计](#)，中断对shared structures的影响 (page tables, process lists, file system structures, and virtually every kernel data structure has to be carefully accessed)

- 很多并发处理技术是在OS中发明和实现的
- multi-threads进程中，应用程序也需要考虑并发

并发相关的重要术语

- **Critical Section(临界区)**, a piece of code that accesses a *shared* resource, usually a variable or data structure. 临界区是一段访问共享资源（通常是变量或数据结构）的代码。
- **Race Condition(条件竞争)** arises if multiple threads of execution enter the critical section at roughly the same time; both attempt to update the shared data structure, leading to a surprising (and perhaps undesirable) outcome. 如果多个正在执行的线程同时进入临界区，则会出现条件竞争；这些线程都试图更新共享的数据结构，会造成意料之外的结果。
- **An Indeterminate(不确定的)** program consists of one or more race conditions; the output of the program varies from run to run, depending on which threads ran when. The outcome is thus not **deterministic**, something we usually expect from computer systems. 不确定的进程由一个或多个条件竞争组成；程序每一次运行的输出都有可能不同，具体取决于每个线程运行的时间。
- **Mutual Exclusion(互斥)** primitives(原语), guarantee that only a single thread ever enters a critical section, thus avoiding races, and resulting in deterministic program outputs. 互斥原语保证只有单个线程进入临界区，从而避免竞争，并导致确定性的进程输出。

Review: Thread


- 轻量化执行环境, new abstraction for a single running process
- Multi-threaded 程序的特点:
 - A multi-threaded program has more than one point of execution.
 - Multiple PCs (Program Counter)
 - They **share** the same **address space**.

Context switch between threads

- Each thread has its own program counter and set of registers. 每个线程拥有独立的PC和寄存器。
 - One or more **thread control blocks(TCBs)** are needed to store the state of each thread. 需要一个或多个线程控制块 (TCB) 来存储每个线程的状态。
- When switching from running one (T1) to running the other (T2),
 - The register state of T1 be saved.
 - The register state of T2 restored.
 - The address space remains the same. 地址空间保持不变。

The stack of the relevant thread

- There will be one stack per thread.




A Single-Threaded Address Space

The code segment:
where instructions live

The heap segment:
contains malloc'd data (**malloc**
动态申请的内存数据)
dynamic data structures (**动态
的数据结构**)
(it grows downward)

(it grows upward)
The stack segment:
contains local variables (**局部
变量**), arguments to
routines (**函数参数**), return
values (**返回值**), etc.



Two threaded Address Space

badcnt.c: Improper Synchronization

```

7. /* Global shared variable */
8. volatile long cnt = 0; /* Counter */

17. int main(int argc, char **argv)
18. {
19.     long niters;
20.     pthread_t tid1, tid2;

21.     niters = atoi(argv[1]);
22.     pthread_create(&tid1, NULL,
23.                     thread, &niters);
24.     pthread_create(&tid2, NULL,
25.                     thread, &niters);
26.     pthread_join(tid1, NULL);
27.     pthread_join(tid2, NULL);

28.     /* Check result */
29.     if (cnt != (2 * niters))
30.         printf("BOOM! cnt=%ld\n", cnt);
31.     else
32.         printf("OK cnt=%ld\n", cnt);
33.     exit(0);
34. }
```

badcnt.c

```

38. /* Thread routine */
39. void *thread(void *vargp)
40. {
41.     long i, niters =
42.         *((long *)vargp);
43.
44.     for (i = 0; i < niters; i++)
45.         cnt++;
46.
47.     return NULL;
48. }
```

```
[zs_cao@localhost conc] $ ./badcnt 10000
OK cnt=20000
[zs_cao@localhost conc] $ ./badcnt 10000
BOOM! cnt=17302
[zs_cao@localhost conc] $ ./badcnt 10000
OK cnt=20000
```

线程并发执行的问题

Assembly Code for Counter Loop

编译:

- gcc -s badcnt.c -o badcnt.s
- vim badcnt.s

```
for (i = 0; i < niters; i++)
    cnt++;
```

```

94     movq  %rdi, -24(%rbp)
95     movq  -24(%rbp), %rax
96     movq  (%rax), %rax
97     movq  %rax, -8(%rbp)
98     movq  $0, -16(%rbp)
99     jmp   .L6
100 .L7:
101     movq  cnt(%rip), %rax
102     addq  $1, %rax
103     movq  %rax, cnt(%rip)
104     addq  $1, -16(%rbp)
105 .L6:
106     movq  -16(%rbp), %rax
107     cmpq  -8(%rbp), %rax
108     jl    .L7
109     movl  $0, %eax
110     popq  %rbp
```

$\} H_i : \text{Head}$

$\} L_i : \text{Load cnt}$
 $\} U_i : \text{Update cnt}$
 $\} S_i : \text{Store cnt}$

$\} T_i : \text{Tail}$

Assembly Code for Counter Loop

汇编:

- gcc -c badcnt.s -o badcnt.o
- objdump -dx badcnt.o

00000000000000ed <thread>:

ed: 55	push %rbp	
ee: 48 89 e5	mov %rsp,%rbp	
f1: 48 89 7d e8	mov %rdi,-0x18(%rbp)	
f5: 48 8b 45 e8	mov -0x18(%rbp),%rax	
f9: 48 8b 00	mov (%rax),%rax	
fc: 48 89 45 f8	mov %rax,-0x8(%rbp)	
100: 48 c7 45 f0 00 00 00	movq \$0x0,-0x10(%rbp)	
107: 00		
108: eb 17	jmp 121 <thread+0x34>	
10a: 48 8b 05 00 00 00 00	mov 0x0(%rip),%rax # 111	L_i
<thread+0x24>		
	10d: R_X86_64_PC32 cnt-0x4	
111: 48 83 c0 01	add \$0x1,%rax	
115: 48 89 05 00 00 00 00	mov %rax,0x0(%rip) # 11c	S_i
<thread+0x2f>		
	118: R_X86_64_PC32 cnt-0x4	
11c: 48 83 45 f0 01	addq \$0x1,-0x10(%rbp)	
121: 48 8b 45 f0	mov -0x10(%rbp),%rax	
125: 48 3b 45 f8	cmp -0x8(%rbp),%rax	
129: 7c df	jl 10a <thread+0x1d>	
12b: b8 00 00 00 00	mov \$0x0,%eax	
130: 5d	pop %rbp	
131: c3	retq	

```
for (i = 0; i < niters; i++)
    cnt++;
```

Assembly Code for Counter Loop

□ 链接:

- gcc -o badcnt.c -o badcnt -lpthread
- objdump -d badcnt

0000000000000957 <thread>:

957: 55	push	%rbp	怎么计算?
958: 48 89 e5	mov	%rsp,%rbp	
95b: 48 89 7d e8	mov	%rdi,-0x18(%rbp)	
95f: 48 8b 45 e8	mov	-0x18(%rbp),%rax	
963: 48 8b 00	mov	(%rax),%rax	
966: 48 89 45 f8	mov	%rax,-0x8(%rbp)	
96a: 48 c7 45 f0 00 00 00	movq	\$0x0,-0x10(%rbp)	
971: 00			
972: eb 17	jmp	98b <thread+0x34>	
974: 48 8b 05 b5 06 20 00	mov	0x2006b5(%rip),%rax	L_i
# 201030 <cnt>			
97b: 48 83 c0 01	add	\$0x1,%rax	
97f: 48 89 05 aa 06 20 00	mov	%rax,0x2006aa(%rip)	
# 201030 <cnt>			U_i
986: 48 83 45 f0 01	addq	\$0x1,-0x10(%rbp)	
98b: 48 8b 45 f0	mov	-0x10(%rbp),%rax	
98f: 48 3b 45 f8	cmp	-0x8(%rbp),%rax	
993: 7c df	j1	974 <thread+0x1d>	T_i
995: b8 00 00 00 00	mov	\$0x0,%eax	
99a: 5d	pop	%rbp	
99b: c3	retq		

Race condition

- 把上述示例简化一下：
- counter = counter + 1 (default is 50)
- We expect the result is 52. However,

OS	Thread1	Thread2	(after instruction)		
			PC	%eax	counter
	before critical section		100	0	50
	mov 0x8049a1c, %eax		105	50	50
	add \$0x1, %eax		108	51	50
interrupt					
	save T1's state		100	0	50
	restore T2's state		105	50	50
		mov 0x8049a1c, %eax	108	51	50
		add \$0x1, %eax	113	51	51
interrupt					
	save T2's state		108	51	50
	restore T1's state		113	51	51
	mov %eax, 0x8049a1c				51

Critical section

- A piece of code that **accesses a shared variable** and must not be concurrently executed by more than one thread. 一段访问共享变量的代码不能并发地被超过一个线程执行。
- Multiple threads executing critical section can result in a race condition. 多线程执行临界区代码会引起条件竞争。
- Need to support **atomicity** for critical sections (**mutual exclusion**). 需要支持访问临界区的原子性（互斥）。

Critical section

【例题】下列对临界区的论述中，正确的是（ ）。

- A. 临界区是指进程中用于实现进程互斥的那段代码。
- B. 临界区是指进程中用于实现进程同步的那段代码。
- C. 临界区是指进程中用于实现进程通信的那段代码。
- D. 临界区是指进程中用于访问临界资源的那段代码。

答案：D

解析：【PPT第4页】**Critical Section(临界区)**, a piece of code that accesses a *shared* resource, usually a variable or data structure.

临界区是一段访问共享资源（通常是变量或数据结构）的代码。

Critical section

【例题】下列准则中，实现临界区互斥机制必须遵循的是（ ）。

- I. 两个进程不能同时进入临界区
- II. 允许进程访问空闲的临界资源
- III. 进程等待进入临界区的时间是有限的
- IV. 不能进入临界区的执行态进程立即放弃CPU

- A. 仅I、IV
- B. 仅II、III
- C. 仅I、II、III
- D. 仅I、III、IV

答案：C

IV选项，不一定必须满足，因为某些机制（如自旋锁）允许忙等待。

Critical section

【例题】两个旅行社甲和乙为旅客到某航空公司订飞机票，形成互斥资源的是（）。

- A. 旅行社
- B. 航空公司
- C. 飞机票
- D. 旅行社与航空公司

答案：C

解析：一张飞机票不能售给不同的旅客，因此飞机票是互斥资源，其他因素只是为完成飞机票订票的中间过程，与互斥资源无关。

Locks

- Ensure that any such critical section executes as if it were a single atomic instruction (**execute a series of instructions atomically**). 保证任意临界区代码像原子操作一样执行 (即原子地执行一系列指令) 。


```
1  lock_t mutex;
2  . . .
3  lock(&mutex);
4  balance = balance + 1; → Critical section
5  unlock(&mutex);
```


Concurrent Execution(并发执行)

- Key idea: In general, any sequentially consistent interleaving is possible, but some give an unexpected result! 通常，任何顺序的交错执行都是有可能的，但有些会给出意想不到的结果！

- I_i denotes that thread i executes instruction I . I_i 表示线程*i*执行指令*I*。
- $\%rdx_i$ is the content of $\%rdx$ in thread i 's context. $\%rdx_i$ 是线程*i*中 $\%rdx$ 的值。

i (thread)	$instr_i$	$\%rdx_1$	$\%rdx_2$	cnt
1	H_1	-	-	0
1	L_1	0	-	0
1	U_1	1	-	0
1	S_1	1	-	1
2	H_2	-	-	1
2	L_2	-	1	1
2	U_2	-	2	1
2	S_2	-	2	2
2	T_2	-	2	2
1	T_1	1	-	2

 Thread 1
 critical section

 Thread 2
 critical section

OK

Concurrent Execution (cont)

- Incorrect ordering: two threads increment the counter, but the result is 1 instead of 2. 两个线程同时增加counter的值，但结果是1而不是2。


i (thread)	$instr_i$	$\%rdx_1$	$\%rdx_2$	cnt
1	H_1	-	-	0
1	L_1	0	-	0
1	U_1	1	-	0
2	H_2			0
2	L_2	-	0	0
1	S_1	1	-	1
1	T_1	1	-	1
2	U_2	-	1	1
2	S_2	-	1	1
2	T_2	-	1	1

S1应该在L2之前执行

Oops!

Progress Graphs (进度图)

Thread 2



A *progress graph* depicts the discrete *execution state space* of concurrent threads. 进度图描述了并发线程的离散执行状态空间。


Each axis corresponds to the sequential order of instructions in a thread. 每个轴表示一个进程的指令执行顺序。

Each point corresponds to a possible *execution state* ($Inst_1, Inst_2$). 每个点表示一个可能的执行状态。

E.g., (L_1, S_2) denotes state where thread 1 has completed L_1 and thread 2 has completed S_2 .

Trajectories in Progress Graphs

Thread 2




A *trajectory* (轨道) is a sequence of legal state transitions (转换) that describes one possible concurrent execution of the threads. 轨道是一系列合法的状态转换，用于描述线程的一个可能的并发执行序列。

Example:

$H_1, L_1, U_1, H_2, L_2, S_1, T_1, U_2, S_2, T_2$

Critical Sections and Unsafe Regions



L , U , and S form a *critical section* (临界区) with respect to the shared variable cnt


Instructions in critical sections (wrt some shared variable) should not be interleaved

临界区中的指令（写入共享变量）不能被交错执行

Sets of states where such interleaving occurs form *unsafe regions*

可能出现交错的状态集合形成了不安全区域

Critical Sections and Unsafe Regions



Def: A trajectory(轨道) is *safe* iff it does not enter any unsafe region
一个轨道是安全的当且仅当它不进入任何不安全区域。

Claim: A trajectory is correct (wrt cnt) iff it is safe

Enforcing Mutual Exclusion

- Question: How can we guarantee a safe trajectory? 如何确保得到一个安全的轨道?
- Answer: We must **synchronize** (同步) the execution of the threads so that they can never have an unsafe trajectory. 我们需要对进程的执行进行同步, 保证进程不存在不安全的轨道。
 - i.e., need to guarantee **mutually exclusive access** (互斥地访问) for each critical section. 即需要保证对每个临界区的互斥访问。
- Classic solution:
 - Semaphores (信号量) (Edsger Dijkstra)
- Other approaches (out of our scope)
 - Mutex and condition variables (Pthreads)
 - Monitors (Java)

1. Concurrency Introduction
2. Locks
3. 基于Lock的并发数据结构
4. Condition Variables 条件变量
5. Semaphore 信号量
6. 常见并发问题
7. 基于事件的并发

Locks: The Basic Idea

- Ensure that any **critical section** executes as if it were **a single atomic instruction**. “全部或都不”
 - Eg. update of a shared variable

```
balance = balance + 1;
```

- Add some code around the critical section

```
1  lock_t mutex; // some globally-allocated lock 'mutex'  
2  ...  
3  lock (&mutex);  
4  balance = balance + 1;  
5  unlock (&mutex) ;
```

Lock变量

- Lock variable holds the state of the lock.
 - **available** (or **unlocked** or **free**)
 - ▶ No thread holds the lock.
 - **acquired** (or **locked** or **held**)
 - ▶ Exactly one thread holds the lock and presumably is in a critical section. 有且只有一个进程拥有锁，且这个进程很可能处在临界区中。

lock()原语的语义 (semantics)

- lock()
 - Try to acquire the lock.
 - If no other thread holds the lock, the thread will **acquire** the lock.
 - **Enter** the *critical section*.
 - ▶ This thread is said to be the owner of the lock.
 - Other threads are *prevented from* entering the critical section while the first thread that holds the lock is in there. 当一个拥有锁的进程进入一个临界区时，其他线程不能进入这个临界区。

Pthread Locks - mutex

- The name that the POSIX library uses for a lock.
 - Used to provide **mutual exclusion**(互斥) between threads.

```
1 pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;  
2  
3 Pthread_mutex_lock(&lock); // wrapper for pthread_mutex_lock()  
4 balance = balance + 1;  
5 Pthread_mutex_unlock(&lock);
```

- We may be using *different locks* to protect *different variables* → Increase **concurrency** (a more **fine-grained** approach). 我们可以使用不同的锁来保护不同的变量 → 提升并发性 (更细粒度的方法)。

Lock如何实现?

- Efficient locks provided mutual exclusion at **low cost**. 高效的锁提供最低开销的互斥。
- Building a lock need some help from the **hardware** and the **OS**. 构造一个锁需要硬件和操作系统的协助。

如何评价lock原语？

□ Mutual exclusion 正确性

- Does the lock work, preventing multiple threads from entering a *critical section*? 这个锁是否能保证多线程不能同时进入临界区？

□ Fairness 公平性

- Does each thread contending for the lock get a fair shot at acquiring it once it is free? (**Starvation**) 每个争夺锁的线程在锁被释放后，是否都能获得公平的获取锁机会？（**饥饿**）

□ Performance 性能

- The time overheads added by using the lock 使用锁带来的时间开销

Controlling Interrupts 基于中断控制的锁实现

□ Disable Interrupts for critical sections 在临界区禁用中断

- One of the earliest solutions used to provide mutual exclusion
- Invented for single-processor systems.

```
1 void lock() {  
2     DisableInterrupts();  
3 }  
4 void unlock() {  
5     EnableInterrupts();  
6 }
```

□ Problem:

- Require too much *trust* in applications
 - Greedy (or malicious) program could monopolize the processor. 贪婪的 (或恶意的) 程序会独占处理器。
- Do not work on multiprocessors 多处理器体系结构这种方式不work
- Code that masks or unmasks interrupts be executed *slowly* by modern CPUs 对中断mask或unmask的指令在现代CPU上执行速度很慢

Why hardware support needed?

- First attempt: Using a *flag* denoting whether the lock is held or not.
 - The code below has problems.

```
1  typedef struct __lock_t { int flag; } lock_t;  
2  
3  void init(lock_t *mutex) {  
4      // 0 → lock is available, 1 → held  
5      mutex->flag = 0;  
6  }  
7  
8  void lock(lock_t *mutex) {  
9      while (mutex->flag == 1) // TEST the flag  
10          ; // spin-wait (do nothing)  
11      mutex->flag = 1; // now SET it !  
12  }  
13  
14 void unlock(lock_t *mutex) {  
15     mutex->flag = 0;  
16 }
```

Why hardware support needed? (Cont.)

- Problem 1: No Mutual Exclusion (assume `flag=0` to begin)

Thread1

```
call lock()
while (flag == 1)
interrupt: switch to Thread 2
```

Thread2

```
call lock()
while (flag == 1)
flag = 1;
interrupt: switch to Thread 1
```

- Problem 2: Spin-waiting wastes time waiting for another thread.

```
flag = 1; // set flag to 1 (too!)
```

- So, we need an atomic instruction supported by **Hardware!**

- **test-and-set instruction**, also known as **atomic exchange**

基于Test-and-set硬件指令实现

- An instruction to support the creation of simple locks

```
1 int TestAndSet(int *ptr, int new) {  
2     int old = *ptr;           // fetch old value at ptr  
3     *ptr = new;             // store 'new' into ptr  
4     return old;             // return the old value  
5 }
```

- **return**(testing) old value pointed to by the `ptr`. 返回ptr指向的旧的值。
- **Simultaneously update**(setting) said value to `new`. 同步地将ptr指向的值设置为 `new`。
- This sequence of operations is performed atomically. 这一系列操作是原子地执行的。

A Simple Spin Lock(自旋锁) using test-and-set

```
1  typedef struct __lock_t {
2      int flag;
3  } lock_t;
4
5  void init(lock_t *lock) {
6      // 0 indicates that lock is available,
7      // 1 that it is held
8      lock->flag = 0;
9  }
10
11 void lock(lock_t *lock) {
12     while (TestAndSet(&lock->flag, 1) == 1)
13         ;           // spin-wait
14 }
15
16 void unlock(lock_t *lock) {
17     lock->flag = 0;
18 }
```

- **Note:** To work correctly on a *single processor*, it requires a preemptive scheduler.
- 在单处理器体系结构中，需要OS kernel实现抢占式调度策略来支持

Evaluating Spin Locks

□ Correctness(正确性): yes

- The spin lock only allows a single thread to entry the critical section. 自旋锁只允许一个进程进入临界区。

□ Fairness(公平性): no

- Spin locks don't provide any fairness guarantees. Spin locks不能为公平性提供任何保证。
- Indeed, a thread spinning may spin *forever*. 实际上，一个自旋锁可能永远自旋（即原地“打转”）。

□ Performance(性能):

- In the single CPU, performance overheads can be quite *painful*. 在单核CPU中，性能开销非常大。
- If the number of threads roughly equals the number of CPUs, spin locks work *reasonably well*. 如果线程数约等于CPU数，则自旋锁效果很好。

基于Compare-And-Swap硬件指令实现

- Test whether the value at the address(`ptr`) is equal to `expected`.
 - If so, update the memory location pointed to by `ptr` with the new value.
 - In either case, return the actual value at that memory location.

```
1 int CompareAndSwap(int *ptr, int expected, int new) {  
2     int actual = *ptr;  
3     if (actual == expected)  
4         *ptr = new;  
5     return actual;  
6 }
```

Compare-and-Swap hardware atomic instruction (C-style)

```
1 void lock(lock_t *lock) {  
2     while (CompareAndSwap(&lock->flag, 0, 1) == 1)  
3         ; // spin  
4 }
```

Spin lock with compare-and-swap

Compare-And-Swap (Cont.)

□ C-callable x86-version of compare-and-swap

```
1  char CompareAndSwap(int *ptr, int old, int new) {  
2      unsigned char ret;  
3  
4      // Note that sete sets a 'byte' not the word  
5      __asm__ __volatile__ (  
6          " lock\n"  
7          " cmpxchgl %2,%1\n"  
8          " sete %0\n"  
9          : "=q" (ret), "=m" (*ptr)  
10         : "r" (new), "m" (*ptr), "a" (old)  
11         : "memory");  
12      return ret;  
13  }
```

Load-Linked and Store-Conditional

```
1 int LoadLinked(int *ptr) {
2     return *ptr;
3 }
4
5 int StoreConditional(int *ptr, int value) {
6     if (no one has updated *ptr since the LoadLinked to this address) {
7         *ptr = value;
8         return 1; // success!
9     } else {
10        return 0; // failed to update
11    }
12 }
```

Load-linked And Store-conditional

- The store-conditional *only succeeds* if **no intermittent store** to the address has taken place. 仅当ptr不被修改时，store-conditional会成功（即返回1）。
 - ▶ **success**: return 1 and update the value at ptr to value.
 - ▶ **fail**: the value at ptr is not updates and 0 is returned.

Load-Linked and Store-Conditional(Cont.)

```
1 void lock(lock_t *lock) {
2     while (1) {
3         while (LoadLinked(&lock->flag) == 1)
4             ; // spin until it's zero
5         if (StoreConditional(&lock->flag, 1) == 1)
6             return; // if set-it-to-1 was a success: all done
7             otherwise: try it all over again
8     }
9 }
10
11 void unlock(lock_t *lock) {
12     lock->flag = 0;
13 }
```

Using LL/SC To Build A Lock

```
1 void lock(lock_t *lock) {
2     while (LoadLinked(&lock->flag) || !StoreConditional(&lock->flag, 1))
3         ; // spin
4 }
```

A more concise form of the lock() using LL/SC

Fetch-And-Add

- Atomically increment a value while returning the old value at a particular address.

```
1 int FetchAndAdd(int *ptr) {  
2     int old = *ptr;  
3     *ptr = old + 1;  
4     return old;  
5 }
```

Fetch-And-Add Hardware atomic instruction (C-style)

Ticket Lock

- Ticket lock can be built with fetch-and add.
 - Ensure progress for all threads. → fairness

```
1  typedef struct __lock_t {  
2      int ticket;  
3      int turn;  
4  } lock_t;  
5  
6  void lock_init(lock_t *lock) {  
7      lock->ticket = 0;  
8      lock->turn = 0;  
9  }  
10  
11 void lock(lock_t *lock) {  
12     int myturn = FetchAndAdd(&lock->ticket);  
13     while (lock->turn != myturn)  
14         ; // spin  
15 }  
16 void unlock(lock_t *lock) {  
17     FetchAndAdd(&lock->turn);  
18 }
```

So Much Spinning

- Hardware-based spin locks are **simple** and they work.
- In some cases, these solutions can be quite **inefficient**.
 - Any time a thread gets caught *spinning*, it **wastes an entire time slice** doing nothing but checking a value. 每当线程自旋时，它会浪费整个时间片，除了检查锁的值以外什么也不做。

How To Avoid *Spinning*?
We' ll need OS Support too!

如何解决“自旋空转”，办法1:Just Yield

- When you are going to spin, give up the CPU to another thread. **当线程将要自旋时，直接让出CPU。**
 - OS system call moves the caller from the *running state* to the *ready state*. **OS通过系统调用将线程从运行态变为就绪态。**
 - The cost of a **context switch** can be substantial and the **starvation** problem still exists. **上下文切换的开销很大，且饥饿问题依然存在。**

```
1 void init() {  
2     flag = 0;  
3 }  
4  
5 void lock() {  
6     while (TestAndSet(&flag, 1) == 1)  
7         yield(); // give up the CPU  
8 }  
9  
10 void unlock() {  
11     flag = 0;  
12 }
```

Lock with Test-and-set and Yield

办法2: Using Queues: Sleeping, not Spinning

- Queue to keep track of which threads are waiting to enter the lock. 用队列来追踪正在等待锁的线程。
- park()
 - Put a calling thread to sleep 将调用线程置于睡眠状态
- unpark(threadID)
 - Wake a particular thread as designated by threadID. 唤醒特定threadID的进程

Using Queues: Sleeping Instead of Spinning

```
1  typedef struct __lock_t { int flag; int guard; queue_t *q; } lock_t;
2
3  void lock_init(lock_t *m) {
4      m->flag = 0;
5      m->guard = 0;
6      queue_init(m->q);
7  }
8
9  void lock(lock_t *m) {
10     while (TestAndSet(&m->guard, 1) == 1)
11         ; // acquire guard lock by spinning
12     if (m->flag == 0) {
13         m->flag = 1; // lock is acquired
14         m->guard = 0;
15     } else {
16         queue_add(m->q, gettid());
17         m->guard = 0;
18         park();
19     }
20 }
21 ...
```

Lock With Queues, Test-and-set, Yield, And Wakeup

Using Queues: Sleeping Instead of Spinning

```
22 void unlock(lock_t *m) {  
23     while (TestAndSet(&m->guard, 1) == 1)  
24         ; // acquire guard lock by spinning  
25     if (queue_empty(m->q))  
26         m->flag = 0; // let go of lock; no one wants it  
27     else  
28         unpark(queue_remove(m->q)); // hold lock (for next thread!)  
29     m->guard = 0;  
30 }
```

Lock With Queues, Test-and-set, Yield, And Wakeup (Cont.)

Two-Phase Locks

- A two-phase lock realizes that spinning can be useful if the lock *is about to* be released. 二阶段锁意识到如果锁即将释放，则自旋可能是有用的。

- **First phase**

- The lock spins for a while, *hoping that* it can acquire the lock. 自旋一段时间，希望可以获取到锁定资源
- If the lock is not acquired during the first spin phase, a second phase is entered. 如果在第一阶段自旋期间没有获取锁定资源，则进入第二阶段

- **Second phase**

- The caller is put to sleep.
- The caller is only woken up when the lock becomes free later.

1. Concurrency Introduction
2. Locks
3. 基于Lock的并发数据结构
4. Condition Variables 条件变量
5. Semaphore 信号量
6. 常见并发问题
7. 基于事件的并发

Lock-based Concurrent Data structure

- Adding locks to a data structure makes the structure **thread safe**.
 - How locks are added determine both the **correctness** and **performance** of the data structure.

Example: Concurrent Counter without Lock

- Simple but not scalable

```
1      typedef struct __counter_t {
2          int value;
3      } counter_t;
4
5      void init(counter_t *c) {
6          c->value = 0;
7      }
8
9      void increment(counter_t *c) {
10         c->value++;
11     }
12
13     void decrement(counter_t *c) {
14         c->value--;
15     }
16
17     int get(counter_t *c) {
18         return c->value;
19     }
```


Add a single lock

- acquired when calling a routine manipulating the data structure.

```
1  typedef struct __counter_t {
2      int value;
3      pthread_lock_t lock;
4  } counter_t;
5
6  void init(counter_t *c) {
7      c->value = 0;
8      Pthread_mutex_init(&c->lock, NULL);
9  }
10
11 void increment(counter_t *c) {
12     Pthread_mutex_lock(&c->lock);
13     c->value++;
14     Pthread_mutex_unlock(&c->lock);
15 }
16
17 void decrement(counter_t *c) {
18     Pthread_mutex_lock(&c->lock);
19     c->value--;
20     Pthread_mutex_unlock(&c->lock);
21 }
22
23 int get(counter_t *c) {
24     Pthread_mutex_lock(&c->lock);
25     int rc = c->value;
26     Pthread_mutex_unlock(&c->lock);
27     return rc;
28 }
```

The performance cost of the simple approach

- Each thread updates a single shared counter.
 - Each thread updates the counter one million times.
 - iMac with four Intel 2.7GHz i5 CPUs.



Synchronized counter **scales poorly**.

同步计数器的扩展性差

Perfect Scaling

- Even though more work is done, it is **done in parallel**.
- The time taken to complete the task is *not increased*.

Sloppy counter

- The sloppy counter works by representing ...
 - A single **logical counter** via numerous local physical counters, on per CPU core 在每个CPU核心上通过多个本地物理计数器来表示单个逻辑计数器
 - A single **global counter**
 - There are **locks**:
 - ▶ One for each local counter and one for the global counter
- Example: on a machine with four CPUs
 - Four local counters
 - One global counter

The basic idea of sloppy counting

- When a thread running on a core wishes to increment the counter.
 - It increments its local counter.
 - Each CPU has its own local counter:
 - ▶ Threads across CPUs can update local counters *without contention*.
跨CPU的线程可以更新本地计数器而不会发生争用
 - ▶ Thus counter updates are **scalable**.
- The local values are periodically transferred to the global counter.**局部值会定期传输到全局计数器**
 - ▶ Acquire the global lock
 - ▶ Increment it by the local counter's value
 - ▶ The local counter is then reset to zero.

The basic idea of sloppy counting(Cont.)

- How often the local-to-global transfer occurs is determined by a threshold, S (sloppiness). 本地到全局传输发生的频率由阈值S决定。
 - The smaller S :
 - ▶ The more the counter behaves like the *non-scalable counter*. S 越小计数器的行为越像不可扩展的计数器
 - The bigger S :
 - ▶ The more scalable the counter.
 - ▶ The further off the global value might be from the *actual count*. 全局值可能离计数器的实际值偏差越大


Sloppy counter example

- Tracing the Sloppy Counters
 - The threshold S is set to 5.
 - There are threads on each of 4 CPUs
 - Each thread updates their local counters $L_1 \dots L_4$.

Time	L_1	L_2	L_3	L_4	\mathbf{G}
0	0	0	0	0	0
1	0	0	1	1	0
2	1	0	2	1	0
3	2	0	3	1	0
4	3	0	3	2	0
5	4	1	3	3	0
6	$5 \rightarrow 0$	1	3	4	5 (from L_1)
7	0	2	4	$5 \rightarrow 0$	10 (from L_4)

Importance of the threshold value's

- Each four threads increments a counter 1 million times on four CPUs.
 - Low S → Performance is **poor**, The global count is always quite **accurate**.
 - High S → Performance is **excellent**, The global count **lags**.
全局计数滞后



Scaling Sloppy Counters

Sloppy Counter Implementation

```
1  typedef struct __counter_t {
2      int global;                      // global count
3      pthread_mutex_t glock;          // global lock
4      int local[NUMCPUS];            // local count (per cpu)
5      pthread_mutex_t llock[NUMCPUS]; // ... and locks
6      int threshold;                 // update frequency
7  } counter_t;
8
9  // init: record threshold, init locks, init values
10 //       of all local counts and global count
11 void init(counter_t *c, int threshold) {
12     c->threshold = threshold;
13
14     c->global = 0;
15     pthread_mutex_init(&c->glock, NULL);
16
17     int i;
18     for (i = 0; i < NUMCPUS; i++) {
19         c->local[i] = 0;
20         pthread_mutex_init(&c->llock[i], NULL);
21     }
22 }
23
```

Sloppy Counter Implementation (Cont.)



清华大学
深圳研究生院
TSINGHUA UNIVERSITY
SHENZHEN GRADUATE SCHOOL

(Cont.)

```
24     // update: usually, just grab local lock and update local
amount
25     //           once local count has risen by 'threshold', grab
global
26     //           lock and transfer local values to it
27     void update(counter_t *c, int threadID, int amt) {
28         pthread_mutex_lock(&c->llock[threadID]);
29         c->local[threadID] += amt;           // assumes amt > 0
30         if (c->local[threadID] >= c->threshold) { // transfer
to global
31             pthread_mutex_lock(&c->glock);
32             c->global += c->local[threadID];
33             pthread_mutex_unlock(&c->glock);
34             c->local[threadID] = 0;
35         }
36         pthread_mutex_unlock(&c->llock[threadID]);
37     }
38
39     // get: just return global amount (which may not be perfect)
40     int get(counter_t *c) {
41         pthread_mutex_lock(&c->glock);
42         int val = c->global;
43         pthread_mutex_unlock(&c->glock);
44         return val;           // only approximate!
45     }
```

Concurrent Linked Lists

```
1 // basic node structure
2 typedef struct __node_t {
3     int key;
4     struct __node_t *next;
5 } node_t;
6
7 // basic list structure (one used per list)
8 typedef struct __list_t {
9     node_t *head;
10    pthread_mutex_t lock;
11 } list_t;
12
13 void List_Init(list_t *L) {
14     L->head = NULL;
15     pthread_mutex_init(&L->lock, NULL);
16 }
17
```

(Cont.)

Concurrent Linked Lists

(Cont.)

```
18     int List_Insert(list_t *L, int key) {
19         pthread_mutex_lock(&L->lock);
20         node_t *new = malloc(sizeof(node_t));
21         if (new == NULL) {
22             perror("malloc");
23             pthread_mutex_unlock(&L->lock);
24             return -1; // fail
26         new->key = key;
27         new->next = L->head;
28         L->head = new;
29         pthread_mutex_unlock(&L->lock);
30         return 0; // success
31     }
32     int List_Lookup(list_t *L, int key) {
33         pthread_mutex_lock(&L->lock);
34         node_t *curr = L->head;
35         while (curr) {
36             if (curr->key == key) {
37                 pthread_mutex_unlock(&L->lock);
38                 return 0; // success
39             }
40             curr = curr->next;
41         }
42         pthread_mutex_unlock(&L->lock);
43         return -1; // failure
44     }
```

Concurrent Linked Lists (Cont.)

并发链表

- The code **acquires** a lock in the insert routine upon entry. 在进入时获取插入操作例程中的锁
- The code **releases** the lock upon exit. 在退出时释放锁
 - If `malloc()` happens to *fail*, the code must also release the lock before failing the insert. 如果 `malloc` 失败，必须在插入操作失败前释放锁
 - This kind of exceptional control flow has been shown to be **quite error prone**. 这种异常的控制流已被证明非常容易出错
 - **Solution:** The lock and release *only surround* the actual critical section in the insert code 锁定和释放仅围绕插入操作代码的关键部分

Concurrent Linked List: Rewritten

```
1 void List_Init(list_t *L) {
2     L->head = NULL;
3     pthread_mutex_init(&L->lock, NULL);
4 }
5
6 void List_Insert(list_t *L, int key) {
7     // synchronization not needed
8     node_t *new = malloc(sizeof(node_t));
9     if (new == NULL) {
10         perror("malloc");
11         return;
12     }
13     new->key = key;
14
15     // just lock critical section
16     pthread_mutex_lock(&L->lock);
17     new->next = L->head;
18     L->head = new;
19     pthread_mutex_unlock(&L->lock);
20 }
21
```

Concurrent Linked List: Rewritten(Cont.)

(Cont.)

```
22     int List_Lookup(list_t *L, int key) {
23         int rv = -1;
24         pthread_mutex_lock(&L->lock);
25         node_t *curr = L->head;
26         while (curr) {
27             if (curr->key == key) {
28                 rv = 0;
29                 break;
30             }
31             curr = curr->next;
32         }
33         pthread_mutex_unlock(&L->lock);
34         return rv; // now both success and failure
35     }
```

Scaling Linked List

- Hand-over-hand locking (lock coupling) 锁耦合
 - Add **a lock per node** of the list instead of having a single lock for the entire list. 为列表的每个节点添加一个锁，而不是为整个列表添加一个锁
 - When traversing the list,
 - ▶ First grabs the next node's lock.
 - ▶ And then releases the current node's lock.
- Enable a high degree of concurrency in list operations. 在列表操作中启用高度并发
 - ▶ However, in practice, the overheads of acquiring and releasing locks for each node of a list traversal is *prohibitive*. 但在实践中，为列表遍历每个节点获取和释放锁的开销难以接受

Michael and Scott Concurrent Queues

- There are two locks.
 - One for the **head** of the queue.
 - One for the **tail**.
 - The goal of these two locks is to enable concurrency of `enqueue` and `dequeue` operations.这两个锁的目标是启用入队和出队操作的并发性

- Add a dummy node
 - Allocated in the queue initialization code 在队列初始化代码中分配
 - Enable the separation of head and tail operations 隔离头部和尾部的操作

Concurrent Queues (Cont.)

```
1     typedef struct __node_t {
2         int value;
3         struct __node_t *next;
4     } node_t;
5
6     typedef struct __queue_t {
7         node_t *head;
8         node_t *tail;
9         pthread_mutex_t headLock;
10        pthread_mutex_t tailLock;
11    } queue_t;
12
13    void Queue_Init(queue_t *q) {
14        node_t *tmp = malloc(sizeof(node_t));
15        tmp->next = NULL;
16        q->head = q->tail = tmp;
17        pthread_mutex_init(&q->headLock, NULL);
18        pthread_mutex_init(&q->tailLock, NULL);
19    }
20
(Cont.)
```

Concurrent Queues (Cont.)



(Cont.)


```
21     void Queue_Enqueue(queue_t *q, int value) {
22         node_t *tmp = malloc(sizeof(node_t));
23         assert(tmp != NULL);
24
25         tmp->value = value;
26         tmp->next = NULL;
27
28         pthread_mutex_lock(&q->tailLock);
29         q->tail->next = tmp;
30         q->tail = tmp;
31         pthread_mutex_unlock(&q->tailLock);
32     }
33     int Queue_Dequeue(queue_t *q, int *value) {
34         pthread_mutex_lock(&q->headLock);
35         node_t *tmp = q->head;
36         node_t *newHead = tmp->next;
37         if (newHead == NULL) {
38             pthread_mutex_unlock(&q->headLock);
39             return -1; // queue was empty
40         }
41         *value = newHead->value;
42         q->head = newHead;
43         pthread_mutex_unlock(&q->headLock);
44         free(tmp);
45         return 0;
46     }
```

Concurrent Hash Table

- Focus on a simple hash table
 - The hash table does not resize.
 - Built using the concurrent lists
 - It uses a **lock per hash bucket** each of which is represented by a *list*.
它对每个哈希桶使用一个锁，每个哈希桶都由一个列表表示。

Performance of Concurrent Hash Table

- From 10,000 to 50,000 concurrent updates from each of four threads.
 - iMac with four Intel 2.7GHz i5 CPUs.



The simple concurrent hash table scales magnificently.

Concurrent Hash Table

```
1 #define BUCKETS (101)
2
3 typedef struct __hash_t {
4     list_t lists[BUCKETS];
5 } hash_t;
6
7 void Hash_Init(hash_t *H) {
8     int i;
9     for (i = 0; i < BUCKETS; i++) {
10         List_Init(&H->lists[i]);
11     }
12 }
13
14 int Hash_Insert(hash_t *H, int key) {
15     int bucket = key % BUCKETS;
16     return List_Insert(&H->lists[bucket], key);
17 }
18
19 int Hash_Lookup(hash_t *H, int key) {
20     int bucket = key % BUCKETS;
21     return List_Lookup(&H->lists[bucket], key);
22 }
```

1. Concurrency Introduction
2. Locks
3. 基于Lock的并发数据结构
4. Condition Variables 条件变量
5. Semaphore 信号量
6. 常见并发问题
7. 基于事件的并发

Condition Variables 条件变量的引入

- There are many cases where a thread wishes to check whether a **condition** is true before continuing its execution. 一个线程需要检查另一个的状态，并据此决定自己是否继续执行
- Example:
 - A parent thread might wish to check whether a child thread has *completed*.
 - This is often called a `join()`.

Condition Variables (Cont.)

A Parent Waiting For Its Child

```
1      void *child(void *arg) {
2          printf("child\n");
3          // XXX how to indicate we are done?
4          return NULL;
5      }
6
7      int main(int argc, char *argv[]) {
8          printf("parent: begin\n");
9          pthread_t c;
10         Pthread_create(&c, NULL, child, NULL); // create child
11         // XXX how to wait for child?
12         printf("parent: end\n");
13         return 0;
14     }
```

What we would like to see here is:

```
parent: begin
child
parent: end
```

Parent waiting for child: Spin-based Approach

```
1     volatile int done = 0;
2
3     void *child(void *arg) {
4         printf("child\n");
5         done = 1;
6         return NULL;
7     }
8
9     int main(int argc, char *argv[]) {
10        printf("parent: begin\n");
11        pthread_t c;
12        Pthread_create(&c, NULL, child, NULL); // create child
13        while (done == 0)
14            ; // spin
15        printf("parent: end\n");
16        return 0;
17    }
```

- This is hugely inefficient as the parent spins and **wastes CPU time**.

How to wait for a condition

- Condition variable 本质上是一个队列及对该队列的操作原语
 - Waiting on the condition
 - ▶ An explicit queue that threads can put themselves on when some state of execution is not as desired. 当某些执行状态不符合要求时, Waiting让该线程将自己放入对应的显式队列中
 - Signaling on the condition
 - ▶ Some other thread, *when it changes said state*, can wake one of those waiting threads and allow them to continue. 其他一些线程在改变所述状态时, 可以唤醒其中一个等待线程并允许它们继续

Definition and Routines

□ Declare condition variable

```
pthread_cond_t c;
```

- Proper initialization is required.

□ Operation (the POSIX calls)

```
pthread_cond_wait(pthread_cond_t *c, pthread_mutex_t *m);      // wait()  
pthread_cond_signal(pthread_cond_t *c);                      // signal()
```

- The `wait()` call takes a mutex as a parameter.

- ▶ The `wait()` call release the lock and put the calling thread to sleep.
 - ▶ When the thread wakes up, it must re-acquire the lock.

Parent waiting for Child: Use a condition variable

```
1      int done = 0;
2      pthread_mutex_t m = PTHREAD_MUTEX_INITIALIZER;
3      pthread_cond_t c = PTHREAD_COND_INITIALIZER;
4
5      void thr_exit() {
6          Pthread_mutex_lock(&m);
7          done = 1;
8          Pthread_cond_signal(&c);
9          Pthread_mutex_unlock(&m);
10     }
11
12     void *child(void *arg) {
13         printf("child\n");
14         thr_exit();
15         return NULL;
16     }
17
18     void thr_join() {
19         Pthread_mutex_lock(&m);
20         while (done == 0)
21             Pthread_cond_wait(&c, &m);
22         Pthread_mutex_unlock(&m);
23     }
24 }
```

Parent waiting for Child: Use a condition variable

```
(cont.)  
25     int main(int argc, char *argv[]) {  
26         printf("parent: begin\n");  
27         pthread_t p;  
28         Pthread_create(&p, NULL, child, NULL);  
29         thr_join();  
30         printf("parent: end\n");  
31         return 0;  
32     }
```

Create the child thread and continues running itself.

Parent waiting for Child: Use a condition variable

□ Parent:

- Create the child thread and continues running itself.
- Call into `thr_join()` to wait for the child thread to complete.
 - ▶ Acquire the lock
 - ▶ Check if the child is done
 - ▶ Put itself to sleep by calling `wait()`
 - ▶ Release the lock

□ Child:

- Print the message “child”
- Call `thr_exit()` to wake the parent thread
 - ▶ Grab the lock
 - ▶ Set the state variable `done`
 - ▶ Signal the parent thus waking it.

The importance of the state variable done

```
1     void thr_exit() {
2             Pthread_mutex_lock(&m);
3             Pthread_cond_signal(&c);
4             Pthread_mutex_unlock(&m);
5     }
6
7     void thr_join() {
8             Pthread_mutex_lock(&m);
9             Pthread_cond_wait(&c, &m);
10            Pthread_mutex_unlock(&m);
11    }
```

thr_exit() and thr_join() without variable done

- Imagine the case where the *child runs immediately*.
 - ▶ The child will signal, but there is no thread asleep on the condition.
 - ▶ When the parent runs, it will call wait and be stuck.
 - ▶ **No thread will ever wake it.**

Another poor implementation

```
1     void thr_exit() {
2         done = 1;
3         Pthread_cond_signal(&c);
4     }
5
6     void thr_join() {
7         if (done == 0)
8             Pthread_cond_wait(&c);
9     }
```

- The issue here is a subtle **race condition**.
 - ▶ The parent calls `thr_join()`.
 - The parent checks the value of `done`.
 - It will see that it is 0 and try to go to sleep.
 - *Just before* it calls `wait` to go to sleep, the parent is interrupted and the child runs.
 - ▶ The child changes the state variable `done` to 1 and signals.
 - But no thread is waiting and thus no thread is woken.
 - When the parent runs again, it sleeps forever.

1. Concurrency Introduction
2. Locks
3. 基于Lock的并发数据结构
4. Condition Variables 条件变量
5. Semaphore 信号量
6. 常见并发问题
7. 基于事件的并发

Semaphore: A definition

- An object with an integer value
 - We can manipulate(操作)with two routines; `sem_wait()` and `sem_post()`.
 - Initialization

```
1 #include <semaphore.h>
2 sem_t s;
3 sem_init(&s, 0, 1); // initialize s to the value 1
```

- ▶ Declare a semaphore `s` and initialize it to the value 1
- ▶ The second argument, 0, indicates that the semaphore is shared between *threads in the same process*.
第二个参数0表示信号量在同一个进程中的线程之间共享

Semaphore: wait原语

□ sem_wait()

```
1 int sem_wait(sem_t *s) {  
2     decrement the value of semaphore s by one  
3     wait if value of semaphore s is negative  
4 }
```

- 减1，若大于0，则返回，否则挂起等待被post唤醒
- If the value of the semaphore was *one or higher* when called `sem_wait()`, **return right away.**
- It will cause the caller to suspend execution waiting for a subsequent post. **它将导致调用者暂停执行，等待后续的post操作**
- When negative, the value of the semaphore is equal to the number of waiting threads. **当为负数时，信号量的值等于等待线程的数量**

Semaphore: post原语

□ sem_post()

```
1 int sem_post(sem_t *s) {  
2     increment the value of semaphore s by one  
3     if there are one or more threads waiting, wake one  
4 }
```

- 加1，若有等待线程，则唤醒一个
- Simply **increments** the value of the semaphore.
- If there is a thread waiting to be woken, **wakes** one of them up.

Binary Semaphores (Locks)

- What should **x** be?
 - The initial value should be **1**.

```
1 sem_t m;
2 sem_init(&m, 0, X); // initialize semaphore to X; what should X be?
3
4 sem_wait(&m);
5 //critical section here 临界区
6 sem_post(&m);
```

Thread Trace: Single Thread Using A Semaphore

Value of Semaphore	Thread 0	Thread 1
1		
1	call sema_wait()	
0	sem_wait() returns	
0	(crit sect)	
0	call sem_post()	
1	sem_post() returns	

Thread Trace: Two Threads Using A Semaphore



HARBIN INSTITUTE OF TECHNOLOGY, SHENZHEN

Value	Thread 0	State	Thread 1	State
1		Running		Ready
1	call sem_wait()	Running		Ready
0	sem_wait() retruns	Running		Ready
0	(crit set: begin)	Running		Ready
0	<i>Interrupt; Switch → T1</i>	Ready		Running
0		Ready	call sem_wait()	Running
-1		Ready	decrement sem	Running
-1		Ready	(sem < 0) → sleep	sleeping
-1		Running	<i>Switch → T0</i>	sleeping
-1	(crit sect: end)	Running		sleeping
-1	call sem_post()	Running		sleeping
0	increment sem	Running		sleeping
0	wake (T1)	Running		Ready
0	sem_post() returns	Running		Ready
0	<i>Interrupt; Switch → T1</i>	Ready		Running
0		Ready	sem_wait() retruns	Running
0		Ready	(crit sect)	Running
0		Ready	call sem_post()	Running
1		Ready	sem_post() returns	Running

Semaphores As Condition Variables

```
1  sem_t s;
2
3  void *
4  child(void *arg) {
5      printf("child\n");
6      sem_post(&s); // signal here: child is done
7      return NULL;
8  }
9
10 int
11 main(int argc, char *argv[]) {
12     sem_init(&s, 0, X); // what should X be?
13     printf("parent: begin\n");
14     pthread_t c;
15     pthread_create(c, NULL, child, NULL);
16     sem_wait(&s); // wait here for child
17     printf("parent: end\n");
18     return 0;
19 }
```

A Parent Waiting For Its Child

```
parent: begin
child
parent: end
```

The execution result

- What should x be?
 - ▶ The value of semaphore should be set to is 0.

Thread Trace: Parent Waiting For Child (Case 1)

- The parent call `sem_wait()` before the child has called `sem_post()`.

Value	Parent	State	Child	State
0	Create(Child)	Running	(Child exists; is runnable)	Ready
0	call <code>sem_wait()</code>	Running		Ready
-1	decrement sem	Running		Ready
-1	$(\text{sem} < 0) \rightarrow \text{sleep}$	sleeping		Ready
-1	<i>Switch→Child</i>	sleeping	child runs	Running
-1		sleeping	call <code>sem_post()</code>	Running
0		sleeping	increment sem	Running
0		Ready	wake(Parent)	Running
0		Ready	<code>sem_post()</code> returns	Running
0		Ready	<i>Interrupt; Switch→Parent</i>	Ready
0	<code>sem_wait()</code> retruns	Running		Ready

Thread Trace: Parent Waiting For Child (Case 2)

- The child runs to completion before the parent call `sem_wait()`.

Value	Parent	State	Child	State
0	Create (Child)	Running	(Child exists; is runnable)	Ready
0	<i>Interrupt; switch→Child</i>	Ready	child runs	Running
0		Ready	call <code>sem_post()</code>	Running
1		Ready	increment sem	Running
1		Ready	wake (nobody)	Running
1		Ready	<code>sem_post()</code> returns	Running
1	parent runs	Running	<i>Interrupt; Switch→Parent</i>	Ready
1	call <code>sem_wait()</code>	Running		Ready
0	decrement sem	Running		Ready
0	$(\text{sem} < 0) \rightarrow \text{awake}$	Running		Ready
0	<code>sem_wait()</code> retruns	Running		Ready

The Producer/Consumer (Bounded-Buffer) Problem

- **Producer:** `put()` interface
 - Wait for a buffer to become *empty* in order to put data into it. 等待缓冲区变为空，以便将数据放入其中
- **Consumer:** `get()` interface
 - Wait for a buffer to become *filled* before using it. 等待缓冲区被填满后再使用它。

```
1 int buffer[MAX];
2 int fill = 0;
3 int use = 0;
4
5 void put(int value) {
6     buffer[fill] = value;      // line f1
7     fill = (fill + 1) % MAX; // line f2
8 }
9
10 int get() {
11     int tmp = buffer[use];    // line g1
12     use = (use + 1) % MAX;   // line g2
13     return tmp;
14 }
```

The Producer/Consumer (Bounded-Buffer) Problem

```
1  sem_t empty;
2  sem_t full;
3
4  void *producer(void *arg) {
5      int i;
6      for (i = 0; i < loops; i++) {
7          sem_wait(&empty);           // line P1
8          put(i);                  // line P2
9          sem_post(&full);         // line P3
10     }
11 }
12
13 void *consumer(void *arg) {
14     int i, tmp = 0;
15     while (tmp != -1) {
16         sem_wait(&full);        // line C1
17         tmp = get();            // line C2
18         sem_post(&empty);       // line C3
19         printf("%d\n", tmp);
20     }
21 }
22 ...
```

First Attempt: Adding the Full and Empty Conditions

The Producer/Consumer (Bounded-Buffer) Problem

```
21 int main(int argc, char *argv[]) {  
22     // ...  
23     sem_init(&empty, 0, MAX);           // MAX buffers are empty to begin with...  
24     sem_init(&full, 0, 0);             // ... and 0 are full  
25     // ...  
26 }
```

First Attempt: Adding the Full and Empty Conditions (Cont.)

- Imagine that `MAX` is greater than 1 .
 - ▶ If there are multiple producers, **race condition** can happen at line `f1`.
如果有多个生产者，竞态条件可能发生在line f1
 - ▶ It means that the old data there is overwritten.**这意味着旧数据被覆盖了**
- We've forgotten here is **mutual exclusion**.
 - ▶ The filling of a buffer and incrementing of the index into the buffer is a **critical section**.**缓冲区的填充和向缓冲区增加索引是临界段**

A Solution: Adding Mutual Exclusion

```
1  sem_t empty;
2  sem_t full;
3  sem_t mutex;
4
5  void *producer(void *arg) {
6      int i;
7      for (i = 0; i < loops; i++) {
8          sem_wait(&mutex);           // line p0 (NEW LINE)
9          sem_wait(&empty);         // line p1
10         put(i);                 // line p2
11         sem_post(&full);        // line p3
12         sem_post(&mutex);       // line p4 (NEW LINE)
13     }
14 }
15
16 void *consumer(void *arg) {
17     int i;
18     for (i = 0; i < loops; i++) {
19         sem_wait(&mutex);       // line c0 (NEW LINE)
20         sem_wait(&full);        // line c1
21         int tmp = get();        // line c2
22         sem_post(&empty);       // line c3
23         sem_post(&mutex);       // line c4 (NEW LINE)
24         printf("%d\n", tmp);
25     }
26 }
```

Adding Mutual Exclusion(Incorrectly)

A Solution: Adding Mutual Exclusion (Cont.)

- Imagine two threads: one producer and one consumer.
 - The consumer **acquires** the `mutex` (line c0).
 - The consumer **calls** `sem_wait()` on the full semaphore (line c1).
 - The consumer is **blocked** and **yield** the CPU.
 - ▶ The consumer still holds the mutex!
 - The producer **calls** `sem_wait()` on the binary `mutex` semaphore (line p0).
 - The producer is now **stuck** waiting too. **a classic deadlock.**

Finally, A Working Solution

```

1  sem_t empty;
2  sem_t full;
3  sem_t mutex;
4
5  void *producer(void *arg) {
6      int i;
7      for (i = 0; i < loops; i++) {
8          sem_wait(&empty);           // line p1
9          sem_wait(&mutex);         // line p1.5 (MOVED MUTEX HERE...)
10         put(i);
11         sem_post(&mutex);        // line p2
12         sem_post(&full);         // line p2.5 (... AND HERE)
13     }
14 }
15
16 void *consumer(void *arg) {
17     int i;
18     for (i = 0; i < loops; i++) {
19         sem_wait(&full);          // line c1
20         sem_wait(&mutex);         // line c1.5 (MOVED MUTEX HERE...)
21         int tmp = get();
22         sem_post(&mutex);        // line c2
23         sem_post(&empty);         // line c2.5 (... AND HERE)
24         printf("%d\n", tmp);
25     }
26 }
```

Finally, A Working Solution

(Cont.)

```
27 int main(int argc, char *argv[]) {
28     // ...
29     sem_init(&empty, 0, MAX); // MAX buffers are empty to begin with ...
30     sem_init(&full, 0, 0);    // ... and 0 are full
31     sem_init(&mutex, 0, 1);   // mutex=1 because it is a lock
32     // ...
33 }
```

Reader-Writer Locks

- Imagine a number of concurrent list operations, including **inserts** and simple **lookups**. 想象一下许多并发的列表操作，包括插入和简单的查找
 - **insert:**
 - ▶ Change the state of the list 更改列表的状态
 - ▶ A traditional critical section makes sense. 传统的临界区是有意义的。
 - **lookup:**
 - ▶ Simply *read* the data structure. 只需读取数据结构
 - ▶ As long as we can guarantee that no insert is on-going, we can allow many lookups to proceed **concurrently**. 只要能够保证没有插入正在进行，就可以允许多个查找同时进行

This special type of lock is known as a **reader-write lock**.

A Reader-Writer Locks

- Only a **single writer** can acquire the lock. 只有一个写者可以获得锁
- Once a reader has acquired a **read lock**, 一旦读者获得了读锁
 - **More readers** will be allowed to acquire the read lock too. 更多的读者也将被允许获得读锁。
 - A writer will have to wait until all readers are finished. 写者必须等到所有读者都看完。

```
1  typedef struct _rwlock_t {  
2      sem_t lock;          // binary semaphore (basic lock)  
3      sem_t writelock;    // used to allow ONE writer or MANY readers  
4      int readers;        // count of readers reading in critical section  
5  } rwlock_t;  
6  
7  void rwlock_init(rwlock_t *rw) {  
8      rw->readers = 0;  
9      sem_init(&rw->lock, 0, 1);  
10     sem_init(&rw->writelock, 0, 1);  
11  }  
12  
13 void rwlock_acquire_readlock(rwlock_t *rw) {  
14     sem_wait(&rw->lock);  
15     ...  
16 }
```

A Reader-Writer Locks (Cont.)


```
15     rw->readers++;
16     if (rw->readers == 1)
17         sem_wait(&rw->writelock); // first reader acquires writelock
18     sem_post(&rw->lock);
19 }
20
21 void rwlock_release_readlock(rwlock_t *rw) {
22     sem_wait(&rw->lock);
23     rw->readers--;
24     if (rw->readers == 0)
25         sem_post(&rw->writelock); // last reader releases writelock
26     sem_post(&rw->lock);
27 }
28
29 void rwlock_acquire_writelock(rwlock_t *rw) {
30     sem_wait(&rw->writelock);
31 }
32
33 void rwlock_release_writelock(rwlock_t *rw) {
34     sem_post(&rw->writelock);
35 }
```

A Reader-Writer Locks (Cont.)

- The reader-writer locks have fairness(公平) problem.
 - It would be relatively easy for reader to **starve writer**. 对于读者来说，饿死写者是相对容易的
 - How to prevent more readers from entering the lock once a writer is waiting? 在写者等待时，如何防止更多的读者进入锁？

The Dining Philosophers

- Assume there are five “**philosophers**” sitting around a table.
- Between each pair of philosophers is a single fork (five total). 每一对哲学家之间都有一个叉(总共五个叉)
- The philosophers each have times where they **think**, and don't need any forks, and times where they **eat**. 每个哲学家都有他们思考的时候，他们不需要叉子，也有他们吃饭的时候。
- In order to **eat**, a philosopher needs **two forks**, both the one on their *left* and the one on their *right*. 为了吃饭，哲学家需要两把叉子，一把在左边，一把在右边
- **The contention(争用) for these forks.**



The Dining Philosophers (Cont.)

□ Key challenge

- There is **no deadlock()**.
- **No** philosopher **starves** and never gets to eat. 没有一个哲学家挨饿不吃东西
- **Concurrency(并发)** is high.

```
while (1) {  
    think();  
    getforks();  
    eat();  
    putforks();  
}
```

Basic loop of each philosopher

```
// helper functions  
int left(int p) { return p; }  
  
int right(int p) {  
    return (p + 1) % 5;  
}
```

Helper functions (Downey's solutions)

- ▶ Philosopher p wishes to refer to the fork on their left \rightarrow call `left(p)`.
- ▶ Philosopher p wishes to refer to the fork on their right \rightarrow call `right(p)`.

The Dining Philosophers (Cont.)

- We need some **semaphore**, one for each fork: `sem_t forks[5]`.

```
1 void getforks() {
2     sem_wait(forks[left(p)]);
3     sem_wait(forks[right(p)]);
4 }
5
6 void putforks() {
7     sem_post(forks[left(p)]);
8     sem_post(forks[right(p)]);
9 }
```

The `getforks()` and `putforks()` Routines (Broken Solution)

- **Deadlock** occur!
 - ▶ If each philosopher happens to **grab the fork on their left** before any philosopher can grab the fork on their right. **如果每一位哲学家碰巧在任何一位哲学家抓住他们右边的叉子之前，抓住了他们左边的叉子**
 - ▶ Each will be stuck *holding one fork* and waiting for another, *forever*. **每个人都会被卡住，拿着一个叉子，永远等着另一个**

A Solution: Breaking The Dependency

- Change how forks are acquired.
 - Let's assume that philosopher 4 acquire the forks in a *different order*.

```
1 void getforks() {  
2     if (p == 4) {  
3         sem_wait(forks[right(p)]);  
4         sem_wait(forks[left(p)]);  
5     } else {  
6         sem_wait(forks[left(p)]);  
7         sem_wait(forks[right(p)]);  
8     }  
9 }
```

- ▶ There is no situation where each philosopher grabs one fork and is stuck waiting for another. **The cycle of waiting is broken.** 不存在每一位哲学家抓住一个叉子而等待另一个的情况。 等待的循环被打破。

How To Implement Semaphores



- Build our own version of semaphores called Zemaphores

```
1  typedef struct __Zem_t {
2      int value;
3      pthread_cond_t cond;
4      pthread_mutex_t lock;
5  } Zem_t;
6
7  // only one thread can call this
8  void Zem_init(Zem_t *s, int value) {
9      s->value = value;
10     Cond_init(&s->cond);
11     Mutex_init(&s->lock);
12 }
13
14 void Zem_wait(Zem_t *s) {
15     Mutex_lock(&s->lock);
16     while (s->value <= 0)
17         Cond_wait(&s->cond, &s->lock);
18     s->value--;
19     Mutex_unlock(&s->lock);
20 }
21
22 void Zem_post(Zem_t *s) {
23     Mutex_lock(&s->lock);
24     s->value++;
25     Cond_signal(&s->cond);
26     Mutex_unlock(&s->lock);
27 }
```

How To Implement Semaphores (Cont.)



- Zemaphore don't maintain the invariant that *the value of the semaphore*.
Zemaphore不保持信号量值不变
 - The value never be lower than zero. 该值不得低于0
 - This behavior is **easier** to implement and **matches** the current Linux implementation. 这种行为更易于实现，并与当前的Linux实现相匹配

例题

【例题】对于两个并发进程，设互斥信号量为mutex（初值为1），若 mutex=-1，则（ ）。

- A. 表示没有进程进入临界区
- B. 表示有一个进程进入临界区
- C. 表示有两个进程进入临界区
- D. 表示有一个进程进入临界区，另一个进程在等待进入

答案：D

P操作：进入临界区，mutex减1。

V操作：释放临界区，mutex加1。

mutex初值为1，此时mutex=-1说明已经有一个进程在临界区，另一个在等待。


例题

【例题】当 V 操作唤醒一个等待进程时，被唤醒进程变为（ ）态。

- A. 运行
- B. 阻塞
- C. 就绪
- D. 完成

答案：C

等待唤醒的进程处于阻塞态，**被唤醒后进入就绪态**。只有就绪进程能获得处理器资源，被唤醒的进程并不能直接转化为运行态。



例题

【例题】有三个进程共享同一程序段，而每次只允许两个进程进入该程序段，若用 PV 操作同步机制，则信号量 S 的取值范围是（ ）

- A. 2,1,0,-1
- B. 3,2,1,0
- C. 2,1,0,-1,-2
- D. 1,0,-1,-2

答案：A

因为每次允许两个进程进入该程序段，**信号量最大值取2**（否则三个进程可以同时进入程序段）。至多有三个进程申请，则信号量最小为-1。

例题

【例题】若一个信号量的初值为3, 经过多次PV操作后当前值为-1, 这表示等待进入临界区的进程数是()。

- A. 1
- B. 2
- C. 3
- D. 4

答案：A

信号量的初值为3, 表示可以有3个进程进入临界区。如果经过多次PV操作后, 信号量的当前值为-1, 则表示当前有4个进程尝试进入临界区, 其中3个进程已进入, 1个进程在等待。

例题

【2013统考真题】某博物馆最多可容纳500人同时参观，有一个出入口，该出入口一次仅允许一人通过，参观者的活动描述如下：

```
cobegin
    参观者进程i;
{
    ...
    进门;
    ...
    参观;
    ...
    出门;
    ...
}
coend
```

请添加必要的信号量和 P, V[或 wait(), signal()]操作，以实现上述过程中的互斥与同步。要求写出完整的过程，说明信号量的含义并赋值。

例题

解答：

出入口一次仅允许一个人通过，设置互斥信号量mutex，初值为1。博物馆最多可同时容纳500人，因此设置信号量empty，初值为500。

```
Semaphore empty = 500; // 博物馆可以容纳的最多人数
Semaphore mutex = 1; // 用于出入口资源的控制
cobegin
    参观者进程i:
    {
        ...
        P(empty); // 可容纳人数减1
        P(mutex); // 互斥使用门
        进门;
        V(mutex);
        参观;
        P(mutex); // 互斥使用门
        出门;
        V(mutex);
        V(empty); // 可容纳人数增加1
        ...
    }
coend
```

例题

【2011统考真题】某银行提供1个服务窗口和10个供顾客等待的座位。顾客到达银行时，若有空座位，则到取号机上领取一个号，等待叫号。取号机每次仅允许一位顾客使用。当营业员空闲时，通过叫号选取一位顾客，并为其服务。顾客和营业员的活动过程描述如下：

```

cobegin
{
    process 顾客i
    {
        从取号机获取一个号码;
        等待叫号;
        获取服务;
    }
}
{
    process 营业员
    {
        while(TRUE)
        {
            叫号;
            为顾客服务;
        }
    }
}
coend

```

请添加必要的信号量和 P, V[或 wait(), signal()]操作，实现上述过程中的互斥与同步。要求写出完整的过程，说明信号量的含义并赋值。

例题

解答：

互斥资源：取号机（一次只有一位顾客领号），因此设置互斥信号量mutex。

同步问题：顾客需要获得空座位等待叫号。营业员空闲时，将选取一位顾客并为其服务。是否有空座位影响等待顾客的数量，是否有顾客决定了营业员是否能开始服务，因此分别设置信号量empty和full来实现这一同步关系。

另外，顾客获得空座位后，需要等待叫号和被服务。顾客与营业员就服务何时开始又构成了一个同步关系，定义信号量service来完成这一同步过程。

```

semaphore empty = 10; // 空座位数量
semaphore full = 0; // 已占座位的数量
semaphore mutex = 1; // 互斥使用取号机
semaphore service = 0; // 等待叫号 (当前是否正在服务)

cobegin
{
    Process 顾客 i{
        P(empty); // 等空位
        P(mutex); // 申请使用取号机
        取号;
        V(mutex); // 取号结束
        V(full); // 通知营业员有新顾客
        P(service); // 等待营业员叫号
        接收服务
    }
    Process 营业员{
        while(True){
            P(full); // 没有顾客则休息
            V(empty); // 离开座位
            V(service); // 叫号
            为顾客服务;
        }
    }
}
coend
    
```

信号量问题解题步骤

1. 找出问题中所有同步与互斥的关系

- 互斥
 - 找到进程竞争的临界资源
 - 抓住“仅允许”或类似词汇
 - 博物馆出入口每次仅允许一人通过、取号机每次仅允许一位顾客使用
- 同步
 - 不同进程对资源合作处理
 - 博物馆内容纳500人、空座位10个

2. 确定信号量个数及每个信号量的初值

- 互斥
 - 用1个信号量，代表资源是否被互斥使用
 - 一般初值为0或1
- 同步
 - 用1或2个信号量，1个信号量用于判断是否资源为empty或full，2个信号量分别判断资源是否为empty和full
 - 一般初值为题目中给出的特定数值（容纳500人、空座位10个）或0

3. 用类似程序的语言描述算法

并发与同步

1. Concurrency Introduction
2. Locks
3. 基于Lock的并发数据结构
4. Condition Variables 条件变量
5. Semaphore 信号量
6. 常见并发问题
7. 基于事件的并发

Common Concurrency Problems

- More recent work focuses on studying other types of common concurrency bugs. 而不是deadlock 越来越多近期的工作专注于研究其它类型的常见并发问题，而不是死锁
- Take a brief look at some example concurrency problems found in real code bases. 简单看一下在实际代码库中发现的一些并发问题事例

What Types Of Bugs Exist?

- Focus on four major open-source applications
 - MySQL, Apache, Mozilla, OpenOffice.

Application	What it does	Non-Deadlock	Deadlock
MySQL	Database Server	14	9
Apache	Web Server	13	4
Mozilla	Web Browser	41	16
Open Office	Office Suite	6	2
Total		74	31

Bugs In Modern Applications

Non-Deadlock Bugs

- Make up a majority of concurrency bugs. 非死锁类bug占并发bug的大头!
- Two major types of non deadlock bugs:
 - Atomicity violation 违反原子性
 - Order violation 违反顺序性

Atomicity-Violation Bugs

- The desired **serializability** among multiple memory accesses is *violated*.
(违反原子性错误) 违反了多个内存访问之间所要求的序列化性
- Simple Example found in MySQL:
 - ▶ Two different threads access the field `proc_info` in the struct `thd`.

```
1  Thread1::  
2  if(thd->proc_info){  
3      ...  
4      fputs(thd->proc_info , ...);  
5      ...  
6  }  
7  
8  Thread2::  
9  thd->proc_info = NULL;
```

Atomicity-Violation Bugs (Cont.)

- **Solution:** Simply add locks around the shared-variable references.

```
1  pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;
2
3  Thread1::
4  pthread_mutex_lock(&lock);
5  if(thd->proc_info) {
6      ...
7      fputs(thd->proc_info , ...);
8      ...
9  }
10 pthread_mutex_unlock(&lock);
11
12 Thread2::
13 pthread_mutex_lock(&lock);
14 thd->proc_info = NULL;
15 pthread_mutex_unlock(&lock);
```

Order-Violation Bugs

- The desired order between two memory accesses is flipped. 两个内存访问之间所要求的顺序颠倒了
 - i.e., A should always be executed before B, but the order is not enforced during execution. 进程A应该始终在进程B之前执行，但是在执行过程中这个顺序并不是强制执行的
- Example:

```
1  Thread1::  
2      void init() {  
3          mThread = PR_CreateThread(mMain, ...);  
4      }  
5  
6  Thread2::  
7      void mMain(...) {  
8          mState = mThread->State  
9      }
```

- ▶ The code in Thread2 seems to assume that the variable mThread has already been *initialized* (and is not NULL).

Order-Violation Bugs (Cont.)

- **Solution:** Enforce ordering using **condition variables**

```
1  pthread_mutex_t mtLock = PTHREAD_MUTEX_INITIALIZER;
2  pthread_cond_t mtCond = PTHREAD_COND_INITIALIZER;
3  int mtInit = 0;
4
5  Thread 1::
6  void init() {
7      ...
8      mThread = PR_CreateThread(mMain, ...);
9
10     // signal that the thread has been created.
11     pthread_mutex_lock(&mtLock);
12     mtInit = 1;
13     pthread_cond_signal(&mtCond);
14     pthread_mutex_unlock(&mtLock);
15     ...
16 }
17
18 Thread2::
19 void mMain(...) {
20     ...
```

Order-Violation Bugs (Cont.)

```
21 // wait for the thread to be initialized ...
22 pthread_mutex_lock(&mtLock);
23 while(mtInit == 0)
24     pthread_cond_wait(&mtCond, &mtLock);
25 pthread_mutex_unlock(&mtLock);
26
27 mState = mThread->State;
28 ...
29 }
```

Deadlock Bugs

Thread 1:

lock(L1);

lock(L2);


Thread 2:

lock(L2);

lock(L1);

□ The presence of a cycle

- ▶ Thread1 is holding a lock L1 and waiting for another one, L2.
- ▶ Thread2 that holds lock L2 is waiting for L1 to be release.





死锁现象

■ 看一个实际的例子


■ 现在分析这个例子

- 竞争使用资源: 道路
- A占有道路1, 又要请求道路2, B占有...
- 形成了无限等待



死锁概念(Deadlock)

- 死锁: 多个进程 (线程) 因循环等待资源而造成无法执行的现象。



- 死锁会造成进程 (线程) 无法执行
- 死锁会造成系统资源的极大浪费(资源无法释放)

死锁特征分析—产生死锁的四个必要条件


- 多个进程因等待**资源**才造成死锁
- **资源:** 进程在完成其任务过程所需要的所有对象
 - CPU、内存、磁盘块、外设、文件、信号量 ...
- 显然有些资源不会造成死锁，而有些会
 - 只读文件是不会造成进程等待的，也就不会死锁
 - 打印机一次只能让一个进程使用，就会造成死锁

称为互斥访问资源
- 显然，**资源互斥访问是死锁的必要条件**
- **资源请求需要形成环路等待才死锁！如何描述这种等待关系？**


资源分配图

■ 资源分配图模型



- 一个进程集合 $\{P_1, P_2, \dots, P_n\}$
- 一资源类型集合 $\{R_1, R_2, \dots, R_m\}$
- 资源类型 R_i 有 W_i 个实例



- 资源请求边：有向边 $P_i \rightarrow R_j$
- 资源分配边：有向边 $R_i \rightarrow P_k$




资源分配图实例



■ 存在环路:

$1 \rightarrow A \rightarrow 2 \rightarrow B \rightarrow 3 \rightarrow C \rightarrow 4 \rightarrow D \rightarrow 1$

■ 产生死锁



■ 存在环路:

$P_1 \rightarrow R_1 \rightarrow P_3 \rightarrow R_2 \rightarrow P_1$

■ 但并不死锁，仍可继续执行

死锁的4个必要条件

■ 互斥使用(Mutual exclusion)

- 至少有一个资源互斥使用

■ 不可抢占(No preemption)


- 资源只能自愿放弃，如车开走以后

■ 请求和保持(Hold and wait)

- 进程必须占有资源，再去申请

■ 循环等待(Circular wait)

- 在资源分配图中存在一个环路



Conditional for Deadlock

- Four conditions need to hold for a deadlock to occur.

Condition	Description
Mutual Exclusion	Threads claim exclusive control of resources that they require. 进程要求对所需要的资源进行排它性控制
Hold-and-wait	Threads hold resources allocated to them while waiting for additional resources. 进程在等待其它资源时保留已分配给它们的资源
No preemption	Resources cannot be forcibly removed from threads that are holding them. 不能强制地移除进程所持有的资源
Circular wait	There exists a circular chain of threads such that each thread holds one more resources that are being requested by the next thread in the chain. 存在一种进程资源的循环等待链，链中每一个进程已获得的资源同时被链中下一个进程所请求

- If any of these four conditions are not met, **deadlock cannot occur**.

例题

【例题】某计算机系统中有 8 台打印机，由 K 个进程竞争使用，每个进程最多需要 3 台打印机。该系统可能发生死锁的 K 的最小值是 ()。

- A. 2
- B. 3
- C. 4
- D. 5

答案：C

考虑最极端情况，因为每个进程最多需要3台打印机，若每个进程已经占有了2台打印机，则只要还有多的打印机，总能满足一个进程达到3台的条件，然后顺利执行。所以将8台打印机分给 K 个进程，每个进程有2台打印机，这就是极端情况， $K = 4$ 。

死锁处理方法概述

- **死锁预防** “no smoking”，预防火灾
 - 破坏死锁的必要条件
- **死锁避免** 检测到煤气超标时，自动切断电源
 - 检测每个资源请求，如果造成死锁就拒绝
- **死锁检测+恢复** 发现火灾时，立刻拿起灭火器
 - 检测到死锁出现时，剥夺一些进程的资源
- **死锁忽略** 在太陽上可以对火灾全然不顾
 - 就好像没有出现死锁一样

死锁预防：破除死锁的必要条件之(1)(2)

■ 破坏互斥使用

- 资源的固有特性，通常无法破除，如打印机

■ 破除不可抢占

- 如果一个进程占有资源并申请另一个不能立即分配的资源，那么已分配资源就可被抢占（即持有不用即可抢占）
- 如果申请的资源得到满足，则抢占其他资源一次性分配给该进程
- 只对状态能保存和恢复的资源(如CPU，内存空间)有效，对打印机等外设不适用

实例：两个进程使用串口，都要读串口，数据不同不可恢复。

死锁预防：破除死锁的必要条件之(3)

■ 破除请求和保持

- 在进程执行前，一次性申请所有需要的资源
- 缺点1：需要预知未来，编程困难
- 缺点2：许多资源分配后很长时间后才使用，资源利用率低

死锁预防：破除死锁的必要条件之(4)

■ 破除循环等待

- 对资源类型进行排序，**资源申请必须按序进行**
- **例如：**所有的进程必须先申请磁盘驱动，再申请打印机，再....，如同日常交通中的单行道
- 缺点：如果编程时就需考虑，用户会觉得很别扭；可能需要释放某些资源(申请序号小的资源)，进程可能会无法执行
- **总之，破除死锁的必要条件会引入不合理因素，实际中很少使用。**

例题

【例题】一次性分配所有资源的方法可以预防死锁的发生，它破坏了死锁4个必要条件中的()。

- A. 互斥使用 (Mutual exclusion)
- B. 不可抢占 (No preemption)
- C. 请求和保持 (Hold-and-wait)
- D. 循环等待 (Circular wait)

答案：C

解析：一次性分配所有资源的方法是当进程需要资源时，一次性提出所有的请求，若请求的所有资源满足则分配，只要有一项不满足，就不分配任何资源。这种分配方式**不会部分地占有资源**，因此破坏了“**请求和保持**”（**进程在等待其它资源时保留已分配给它们的资源**）。

死锁避免

不死锁就成了问题的核心!

- 思想: 判断此次请求**是否造成死锁**
若会造成死锁, 则拒绝该请求

- 安全状态定义: 如果系统中的所有进程存在一个可完成的执行序列 P_1, \dots, P_n , 则称系统处于安全状态

都能执行完成当然就不死锁

- 安全序列: 上面的执行序列 P_1, \dots, P_n

如何找到这样的序列?

死锁避免之银行家算法

一个银行家：目前手里只有1亿

第A个开发商：已贷款15亿，资金紧张还需3亿。

第B个开发商：已贷款5亿，还需贷款1亿，运转良好能收回。

第C个开发商：已贷款2亿，欲贷款18亿

....

开发商B还钱，再借给A，则可以继续借给C

银行家当前可用的资金 (**Available**) ? 可以利用的资金，即可用的加上能收回的共有多少 (**work**) ? 各个开发商已贷款——已分配的资金 (**Allocation**) ?
各个开发商还需要贷款 (**need**)

死锁避免之银行家算法

■ 安全序列 P_1, \dots, P_n 应该满足的性质：

$P_i (1 \leq i \leq n)$ 需要资源 \leq 剩余资源 + 分配给 $P_j (1 \leq j < i)$ 资源

1. Banker()
2. int n,m; // 系统中进程总数n和资源种类总数m
3. int Available[m]; // 资源当前可用总量
4. int Allocation[n][m];
5. // 当前给分配给每个进程的各种资源数量
6. int Need[n][m];
7. // 当前每个进程还需分配的各种资源数量
8. int Work[m]; // 当前可分配的资源，包括可收回的
9. bool Finish[n]; // 进程是否结束

死锁避免之银行家算法

■ 安全状态判定 (思路) :

① 初始化设定:

Work = Available (动态记录当前可 (收回) 分配资源)

Finish[i]=false (设定所有进程均未完成)

② 查找这样的进程 P_i (未完成但目前剩余资源可满足其需要, 这样的进程是能够完成的) :

- a) Finish[i] == false
- b) Need[i] ≤ Work

如果没有这样的进程 P_i , 则跳转到第④步

③ (若有则) P_i 一定能完成, 并归还其占用的资源, 即:

- a) Finish[i] = true
- b) Work = Work + Allocation[i]

GOTO 第②步, 继续查找

④ 如果所有进程 P_i 都是能完成的, 即Finish[i]=ture

则系统处于安全状态, 否则系统处于不安全状态

死锁避免之银行家算法实例

■ 当前状态:

		<u>Allocation</u>	<u>Need</u>	<u>Available</u>
		A B C	A B C	A B C
	Work=[3 3 2]	P0 0 1 0	7 4 3	3 3 2
P_1	Work=[5 3 2]	P1 2 0 0	1 2 2	
P_3	Work=[7 4 3]	P2 3 0 2	6 0 0	
P_4	Work=[7 4 5]	P3 2 1 1	0 1 1	
P_0	Work=[7 5 5]	P4 0 0 2	4 3 1	
P_2	Work=[10 5 7]			

- 安全序列是 $\langle P_1, P_3, P_4, P_0, P_2 \rangle$
- 安全序列是唯一的吗?

死锁避免之银行家算法

```

1.bool Found;
2.Work = Available; Finish = false;
3.while(true){
4.    Found = false; //是否为安全序列找到一个新进程
5.    for(i=1; i<=n; i++){
6.        if(Finish[i]==false && Need[i]<=Work){
7.            Work = Work + Allocation[i];
8.            Finish[i] = true;
9.            printf("%d->",i); //输出安全序列
10.           Found = true;
11.       }
12.   } 没有安全序列或已经找到
13.   if(Found==false)break;
14. }
15. for(i=1;i<=n;i++)
16.   if(Finish[i]==false)
17.     return "deadlock";

```

$$T(n)=O(mn^2)$$

最好情形：安全状态就是P₁-P_n
最坏情形：P_n-P₁

死锁避免之资源请求算法

思想：可用的资源可以满足某个进程的资源请求，则分配，然后寻找安全序列，找到，分配成功，找不到，已分配资源收回。

```
1. extern Banker();
2. int Request[m]; /*进程Pi的资源申请*/
3. if(Request>Need[i]) return “error”;
4. if(Request>Available) sleep();
5. Available=Available-Request;
6. Allocation[i]=Allocation[i]+Request;
7. Need[i]=Need[i]-Request;
8.           /*先将资源分配给Pi*/
9. if(Banker()==“deadlock”)
10.          /*调用银行家算法判定是否会死锁*/
11. 拒绝Request; /*若算法判定deadLock则拒绝请求，资源回滚*/
```

死锁避免之资源请求实例(1)

■ P_1 申请资源(1,0,2)

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C
P_0	0 1 0	7 4 3	2 3 0
P_1	3 0 2	0 2 0	
P_2	3 0 2	6 0 0	
P_3	2 1 1	0 1 1	
P_4	0 0 2	4 3 1	

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C
P_0	0 1 0	7 4 3	3 3 2
P_1	2 0 0	1 2 2	
P_2	3 0 2	6 0 0	
P_3	2 1 1	0 1 1	
P_4	0 0 2	4 3 1	

- 序列 $\langle P_1, P_3, P_2, P_4, P_0 \rangle$ 是安全的
- 此次申请允许

死锁避免之资源请求实例(2)

■ P_0 再申请(0,2,0)

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C
P_0	0 3 0	7 2 3	2 1 0
P_1	3 0 2	0 2 0	
P_2	3 0 2	6 0 0	
P_3	2 1 1	0 1 1	
P_4	0 0 2	4 3 1	

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C
P_0	0 1 0	7 4 3	2 3 0
P_1	3 0 2	0 2 0	
P_2	3 0 2	6 0 0	
P_3	2 1 1	0 1 1	
P_4	0 0 2	4 3 1	

- 进程 P_0, P_1, P_2, P_3, P_4 一个也没法执行，死锁进程组
- 此次申请被拒绝

银行家算法讨论：

- 每个进程进入系统时必须告知所需资源的最大数量
对应用程序员要求高
- 安全序列寻找算法（安全状态判定算法）计算**时间复杂度为 $O(mn^2)$** ，过于复杂
- 若每次资源请求都要调用银行家算法，耗时过大，
系统效率降低
- 采用此算法，存在情况：当前有资源可用，尽管可能很快就会释放，由于会使整体进程处于不安全状态，而不被分配，致使**资源利用率大大降低**

死锁检测+恢复: 死锁检测

- 基本原因: 每次申请都执行 $O(mn^2)$, 效率低
- 对策: 只要可用资源足够, 则分配, **发现问题再处理**
 - 定时检测或者当发现资源利用率低时检测

```

1.bool Found;
2.int Request[n][m];
3.Work = Available; Finish = false;
4.if Allocation[i] != 0: Finish[i] = false;
5.else: Finish[i] = true;
6.while(true){
7.    Found = false; //是否为安全序列找到一个新进程
8.    for(i=1; i<=n; i++){
9.        if(Finish[i]==false && Request[i]<=Work){
10.            Work = Work + Allocation[i];
11.            Finish[i] = true;
12.            Found = true;
13.        }
14.    }
15.    if(Found==false)break;
16.}
17.for(i=1;i<=n;i++) {
18.    if(Finish[i]==false) {
19.        deadlock = deadlock + {i}; return "deadlock";
20.    }
21.}

```

//对银行家算法进行改进

对于无分配资源的进程, 不论其是否获得请求
资源, 则认为其是完成的

死锁检测+恢复: 死锁恢复

- 终止进程 选谁终止?
 - 优先级? 占用资源多的? ...

- 剥夺资源 进程需要回滚 (rollback)
 - 回滚点的选取? 如何回滚? ...

鸵鸟算法 (死锁忽略)

■ 死锁预防?

- 引入太多不合理因素...

■ 死锁避免?

- 每次申请都执行银行家算法 $O(mn^2)$, 效率太低

■ 死锁检测+恢复?

- 还要执行银行家算法 $O(mn^2)$, 且恢复并不容易

■ 鸵鸟算法: 对死锁不做任何处理.....

- 死锁出现时, 手动干预——重新启动
- 死锁出现不是确定的, 避免死锁付出的代价毫无意义
- 有趣的是大多数操作系统都用它, 如UNIX和Windows

死锁总结

- 进程竞争资源 \Rightarrow 有可能形成循环竞争 \Rightarrow 死锁
- 死锁需要处理 \Rightarrow 死锁分析 \Rightarrow 死锁的必要条件
- 死锁处理 \Rightarrow 预防、避免、检测+恢复、忽略
- 死锁预防: 破除必要条件 \Rightarrow 引入了不合理因素
- 死锁避免: 用银行家算法找安全序列 \Rightarrow 效率太低
- 死锁检测恢复: 银行家算法找死锁进程组并恢复 \Rightarrow 实现较难
- 死锁忽略: 就好像没有死锁 \Rightarrow 现在用的最多

任何思想、概念、技术的主流都会
随着时间而改变，操作系统尤为明显！

例题

【例题】在下列死锁的解决方法中，属于死锁预防策略的是()。

- A. 银行家算法
- B. 资源有序分配算法
- C. 死锁检测算法
- D. 资源分配图化简算法

答案：B

- A. 属于**死锁避免**
- B. 破坏了循环等待，破坏死锁的必要条件属于**死锁预防**
- C. 属于**死锁检测**
- D. 属于**死锁检测**



信号处理

【例题】异步信号安全的函数要么是可重入的，要么不能被信号处理程序中断，包括I/O函数（）

- A.printf
- B.sprintf
- C.write
- D.malloc

答案：

C

