



哈爾濱工業大學(深圳)

HARBIN INSTITUTE OF TECHNOLOGY, SHENZHEN

# 操作系统

# Operating Systems

刘川意 教授

[liuchuanyi@hit.edu.cn](mailto:liuchuanyi@hit.edu.cn)

哈尔滨工业大学 (深圳)

2025年10月

# Module 6: I/O与存储

## 1. I/O devices

- a. 概述 (基础)
- b. 设备交互 (轮询、中断、DMA) (基础)
- c. 文件系统 (基础)
- d. 代码实例 (进阶)

## 2. Hard Disk Drives


## 3. RAID

# I/O Devices

- I/O is **critical** to computer system to **interact with systems**.
- **I/O需要解决的重要问题：**
  - How should I/O be integrated into systems?
  - What are the general mechanisms?
  - How can we make the efficiently?

# Structure of input/output (I/O) device

CPU



CPU is attached to the main memory of the system via some kind of memory bus.

Some devices are connected to the system via a general I/O bus.

# I/O Architecture

## □ Buses

- Data paths that provided to enable information between CPU(s), RAM, and I/O devices. 不同的I/O设备（如键盘、鼠标、磁盘等）需要通过相应的接口电路与总线相连接，这些接口电路由“控制器”或“适配器”提供（后面统称为“设备控制器”）。不同的设备控制器能够支持不同的接口协议

## □ I/O bus

- Data path that connects a CPU to an I/O device.
- I/O bus is connected to I/O device by three hardware components: I/O ports, interfaces and device controllers.
- 根据接口协议的性能区别，现代计算机对I/O总线进行了分层。在上图中，图像或者其他高性能的I/O设备通过常规的I/O总线连接到系统，在许多现代系统中会是PCI或它的衍生形式。而一些相对较慢的I/O设备则通过外围总线（peripheral bus）连接到系统，比如使用SCSI、SATA或者USB等协议的I/O设备

# Module 6: I/O与存储

## 1. I/O devices

- a. 概述 (基础)
- b. 设备交互 (轮询、中断、DMA) (基础)
- c. 文件系统 (基础)
- d. 代码实例 (进阶)

## 2. Hard Disk Drives


## 3. RAID

# Device interaction

- How the OS communicates with the **device**? 主机对I/O设备进行访问的目标是I/O设备的寄存器或者内存。常见的I/O设备都只提供寄存器供主机访问，对于低速外设这样的模式是足够的，但是对于需要大量、高速数据交互的外设（如显卡、网卡），就需要主机能够直接访问外设的内存了
- Solutions: 现代计算机提供了两种方式来访问I/O设备，它们分别是PMIO和MMIO
  - PMIO: 端口映射I/O (Port-mapped I/O)。将I/O设备独立看待，并使用CPU提供的专用I/O指令访问； **I/O instructions**: a way for the OS to send data to specific device registers.
    - Ex) `in` and `out` instructions on x86
  - MMIO: 内存映射I/O (Memory-mapped I/O)。将I/O设备看作内存的一部分，不使用单独的I/O指令，而是使用内存读写指令访问； **memory-mapped I/O**
    - Device registers available as if they were memory locations.
    - The OS `load` (to read) or `store` (to write) to the device instead of main memory.

# PMIO( Port-mapped I/O )


- 端口映射I/O，又叫做被隔离的I/O (isolated I/O)，它提供了一个专门用于I/O设备“注册”的地址空间，该地址空间被称为I/O地址空间，最大寻址范围为64K



- 为了使I/O地址空间与内存地址空间隔离，**要么在CPU物理接口上增加一个I/O引脚，要么增加一条专用的I/O总线。**因此，并不是所有的平台都支持PMIO，常见的ARM平台就不支持PMIO。支持PMIO的CPU通常具有专门执行I/O操作的指令，例如在Intel-X86架构的CPU中，I/O指令是in和out，这两个指令可以读/写1、2、4个字节 (outb, outw, outl) 从内存到I/O接口上。由于I/O地址空间比较小，因此I/O设备一般只在其中“注册”自己的寄存器，之后系统可以通过PMIO对它们进行访问

# MMIO (Memory-mapped I/O)

- 在MMIO中，物理内存和I/O设备共享内存地址空间（注意，这里的内存地址空间实际指的是内存的物理地址空间）



- 当CPU访问某个虚拟内存地址时，该虚拟地址首先转换为一个物理地址，**对该物理地址的访问，会通过南北桥（现在被合并为I/O桥）的路由机制被定向到物理内存或者I/O设备上。**因此，用于访问内存的CPU指令也可用于访问I/O设备，并且在内存（的物理）地址空间上，需要给I/O设备预留一个地址区域，该地址区域不能给物理内存使用。
- MMIO是应用得最为广泛的一种I/O方式，由于内存地址空间远大于I/O地址空间，I/O设备可以在内存地址空间上暴露自己的内存或者寄存器，以供主机进行访问

# PCI设备

- PCI及其衍生的接口（如PCIE）主要服务于高速I/O设备（如显卡或网卡），使用PCI接口的设备又被称为PCI设备。与慢速I/O设备不同，**计算机既需要访问它们的寄存器，也需要访问它们的内存。**
- 每个PCI设备都有一个配置空间（实际就是设备上一组连续的寄存器），大小为256byte。配置空间中包含了6个BAR(Base Address Registers，基址寄存器)，BAR中记录了设备所需要的地址空间类型、基址以及其他属性

Memory Space BAR Layout

31 - 4	3	2 - 1	0
16-Byte Aligned Base Address	Prefetchable	Type	Always 0

I/O Space BAR Layout


31 - 2	1	0
4-Byte Aligned Base Address	Reserved	Always 1

# PCI设备 (cont.)

- 可以看到，PCI设备能够申请两类地址空间，即内存地址空间和I/O地址空间，它们用BAR的最后一一位区别开来。因此，PCI设备可以通过PMIO和MMIO将自己的I/O存储器（Registers/RAM/ROM）暴露给CPU（通常寄存器使用PMIO，而内存使用MMIO的方式暴露）。
- 配置空间中的每个BAR可以映射一个地址空间，因此每个PCI设备最多能映射6段地址空间，但实际上很多设备用不了这么多。PCI配置空间的初始值是由厂商预设在设备中的，也就是说，设备需要哪些地址空间都是其自己定的，这可能会造成不同的PCI设备所映射的地址空间冲突，因此在PCI设备枚举（也叫总线枚举，由BIOS或者OS在启动时完成）的过程中，会重新为其分配地址空间，然后写入PCI配置空间中。
- 在PCI总线之前的ISA总线是使用跳线帽来分配外设的物理地址，每插入一个新设备都要改变跳线帽以分配物理地址，这是十分麻烦且易错的，但这样的方式似乎我们更容易理解。能够分配自己总线上挂载设备的物理地址这也是PCI总线相较于I2C、SPI等低速总线一个最大的特色

# 标准外设(Canonical Device)

- Canonical Devices has two important components.
  - **Hardware interface** allows the system software to control its operation. 硬件接口本质就是I/O设备提供的各式寄存器，系统软件通过与这些寄存器进行交互，达到控制I/O设备的目的
  - **Internals** which is implementation specific. 实现硬件接口提供的功能，不同的I/O设备具有不同功能，因此它们的内部实现和包含的元器件也不尽相同
  - 使用I/O设备的目的是为了交互数据，不管是网卡、磁盘，亦或是键盘，总归要将数据进行输入输出



# Hardware interface of Canonical Device

- **status register**
    - See the current status of the device
  - **command register**
    - Tell the device to perform a certain task
  - **data register**
    - Pass data to the device, or get data from the device
- By reading and writing above **three registers**,  
the operating system can **control device behavior**.
- Typical interaction example

```
while ( STATUS == BUSY)
    ; //wait until device is not busy
write data to data register
write command to command register
    (Doing so starts the device and executes the command)
while ( STATUS == BUSY)
    ; //wait until device is done with your request
```

# Polling

- Operating system waits until the device is ready by **repeatedly** reading the status register. **Device driver**通过轮询读取设备状态
- 一般来说，主机与I/O设备要进行数据交互，会经过这样一个过程：
  - CPU通过I/O设备的硬件接口（以下简称I/O接口）获取设备状态（即状态寄存器的值），只有“就绪”状态的设备才能进行数据传输。
  - CPU通过I/O接口下达交互指令：如果是读数据，则向I/O接口的命令寄存器输入要获取的数据在I/O设备的内部位置以及读设备指令；如果是写数据，则向I/O接口的命令寄存器输入要存放的数据在I/O设备的内部位置、写设备指令，以及向数据寄存器写入数据。
  - I/O设备内部根据I/O接口中寄存器的值，开始执行数据传输工作。
  - CPU在I/O设备完成工作后，执行其他操作，完成数据传送。




# Polling Evaluation

- 标准交互流程实现起来比较简单，但是难免会有一些低效和不方便。第一个问题就是轮询过程比较低效，在等待设备是否满足某种状态时浪费大量CPU时间（下图描述的就是磁盘在执行数据传输过程中，CPU不能执行其他任务，只能等待传输完成），如果此时操作系统可以切换执行下一个就绪进程，就可以大大提高CPU的利用率。
  - Positive aspect is simple and working.
  - **However, it wastes CPU time just waiting for the device.**
    - ▶ Switching to another ready process is better utilizing the CPU.

# interrupts

- 有了中断机制，CPU向设备发出I/O请求后，就可以让对应进程进入睡眠等待，从而切换执行其他进程。当设备完成I/O请求后，它会抛出一个硬件中断，引发CPU跳转执行操作系统预先定义好的中断处理程序，中断处理程序会挂起正在执行的进程，同时唤醒等待I/O的进程并继续执行
- Put the I/O request process to sleep and context switch to another. OS调度程序让执行I/O request的进程进入sleep状态，并切换成另一个进程去执行
- When the device is finished, wake the process waiting for the I/O by interrupt. 外设通过中断机制唤醒处于sleep状态的等待进程
  - Positive aspect is allow to CPU and the disk are properly utilized.



在磁盘执行进程1的I/O过程中，CPU同时执行进程2，并且在I/O请求执行完毕后，回过头来再次执行进程1

# Polling vs interrupts

- However, “interrupts is not always the best solution”
  - If, device performs very quickly, interrupt will “slow down” the system.
  - Because **context switch is expensive (switching to another process)**

If a device is fast → **poll** is best.  
If it is slow → **interrupts** is better.

# CPU is once again over-burdened

- 在标准交互流程和引入中断流程中，数据在硬件中的移动都是通过CPU完成的，比如CPU从内存读取数据到CPU寄存器，然后将CPU寄存器的数据写入I/O设备寄存器。但是对CPU来说，它的主要功能是使用内部的算数/逻辑单元 (ALU) 执行计算，而不是做一个数据搬运工，如果CPU参与大量数据的移动，就白白浪费了宝贵的时间和算力。为了让CPU从数据移动的工作中解放出来，需要引入DMA机制
- CPU **wastes a lot of time** to copy the *a large chunk of data* from memory to the device.





Diagram of CPU utilization

# DMA (Direct Memory Access)

- DMA, 全称为direct memory access, 直接内存访问。它是I/O设备与主存之间由硬件组成的直接数据通路，用于高速I/O设备与主存之间的成组数据（即数据块）传送。实现DMA机制的硬件叫做DMA控制器。
- Copy data in memory by knowing “where the data lives in memory, how much data to copy”
- When completed, DMA raises an interrupt, I/O begins on Disk.



# DMA机制下的数据交互流程

- 引入了DMA机制之后，与I/O设备的数据交互流程变为下图





图 (a) 数据传送的三个阶段

图 (b) 第二阶段的数据传送过程

- (1) DMA预处理：在进行DMA数据传送之前要用程序做一些必要的准备工作。先由CPU执行几条IN/OUT指令，测试设备状态，向DMA控制器的设备地址寄存器中送入I/O设备地址并启动I/O设备，向主存地址寄存器中送入交换数据的主存起始地址，在数据字数寄存器中送入交换的数据个数。这些工作完成之后，CPU继续执行原来的程序。
- (2) DMA控制I/O设备与主存之间的数据交换，并且在数据交换完毕或者出错时，向CPU发出结束中断请求或出错中断请求。
- (3) CPU中断程序进行后处理，若需继续交换数据，则要对DMA控制器进行初始化；若不需要交换数据，则停止外设；若为出错，则转错误诊断及处理程序。

# DMA效果示意

- 与磁盘交互时各硬件执行进程任务的时间轴，可以看到，CPU将原本用于移动进程1的I/O数据的时间用于执行进程2，相应的，DMA代替了数据移动的工作




Diagram of CPU utilization by DMA

# Device interaction (Cont.)

- How the OS interact with **different specific interfaces**?
  - Ex) We'd like to build a file system that worked on top of SCSI disks, IDE disks, USB keychain drivers, and so on.
- Solutions: **Abstraction**
  - Abstraction encapsulate **any specifics of device interaction**. 抽象封装了设备交互的任何具体细节。

# 例题

一个典型的文本打印页面有50行，每行 80 个字符，假定一台标准的打印机每分钟能打印6页，向打印机的输出寄存器中写一个字符的时间很短，可忽略不计。若每打印一个字符都需要花费 50us 的中断处理时间(包括所有服务)，使用中断驱动 I/O 方式运行这台打印机，中断的系统开销占 CPU 的百分比为()

- A.2%                  B. 5%                  C.20%                  D.50%

答案：A

解析：

这台打印机每分钟打印  $50 \times 80 \times 6 = 24000$  个字符，即每秒打印 400 个字符。每个字符打印中断需要占用 CPU 时间 50us，所以每秒用于中断的系统开销为  $400 \times 50\text{us} = 20\text{ms}$ 。若使用中断驱动I/O，则 CPU 剩余的 980ms 可用于其他处理，中断的开销占 CPU 的 2%。

# Module 6: I/O与存储

## 1. I/O devices


- a. 概述 (基础)
- b. 设备交互 (轮询、中断、DMA) (基础)
- c. 文件系统 (基础)
- d. 代码实例 (进阶)

## 2. Hard Disk Drives

## 3. RAID

# File system Abstraction

- File system **specifics** of which disk class it is using.
  - Ex) It issues **block read** and **write** request to the generic block layer.



# Problem of File system Abstraction

- If there is a device having many special capabilities, these capabilities **will go unused** in the generic interface layer.**如果有一个设备具有许多特殊的能力，这些能力将在通用接口层中不被使用。**
  
- Over 70% of OS code is found in device drivers.
  - Any device drivers are needed because you might plug it to your system.
  - They are primary contributor to **kernel crashes**, making **more bugs**.

# A Simple IDE Disk Driver

- Four types of register
  - Control, command block, status and error
  - Memory mapped IO
  - in and out I/O instruction

# IDE Device Interface

## □ Control Register:

Address 0x3F6 = 0x80 (0000 1RE0): R=reset, E=0 means "enable interrupt"

## □ Command Block Registers:

Address 0x1F0 = Data Port

Address 0x1F1 = Error

Address 0x1F2 = Sector Count

Address 0x1F3 = LBA low byte

Address 0x1F4 = LBA mid byte

Address 0x1F5 = LBA hi byte

Address 0x1F6 = 1B1D TOP4LBA: B=LBA, D=drive

Address 0x1F7 = Command/status

## □ Status Register (Address 0x1F7):

7	6	5	4	3	2	1	0
BUSY	READY	FAULT	SEEK	DRQ	CORR	IDDEX	ERROR

## □ Error Register (Address 0x1F1): (check when Status ERROR==1)

7	6	5	4	3	2	1	0
BBK	UNC	MC	IDNF	MCR	ABRT	T0NF	AMNF
□ BBK = Bad Block				□ MCR = Media Change Requested			
□ UNC = Uncorrectable data error				□ ABRT = Command aborted			
□ MC = Media Changed				□ T0NF = Track 0 Not Found			
□ IDNF = ID mark Not Found				□ AMNF = Address Mark Not Found			

# OS跟设备交互的典型协议

- **Wait for drive to be ready.** Read Status Register (0x1F7) until drive is not busy and READY.
- **Write parameters to command registers.** Write the sector count, logical block address (LBA) of the sectors to be accessed, and drive number (master=0x00 or slave=0x10, as IDE permits just two drives) to command registers (0x1F2-0x1F6).
- **Start the I/O.** by issuing read/write to command register. Write READ—WRITE command to command register (0x1F7).
- **Data transfer (for writes):** Wait until drive status is READY and DRQ (drive request for data); write data to data port.
- **Handle interrupts.** In the simplest case, handle an interrupt for each sector transferred; more complex approaches allow batching and thus one final interrupt when the entire transfer is complete.  
**在最简单的情况下，为每一个传输的扇区处理一个中断；更复杂的方法允许分批处理，从而在整个传输完成后处理一个最终中断。**
- **Error handling.** After each operation, read the status register. If the ERROR bit is on, read the error register for details.  
**每次操作后，读取状态寄存器。如果ERROR位处于开启状态，请读取错误寄存器的详细信息。**

# Module 6: I/O与存储

## 1. I/O devices

- a. 概述 (基础)
- b. 设备交互 (轮询、中断、DMA) (基础)
- c. 文件系统 (基础)
- d. 代码实例 (进阶)

## 2. Hard Disk Drives

## 3. RAID



# Wait for drive to be ready

```
1. static int ide_wait_ready()
2. /* ensure the drive is ready before issuing a request to it */
3. {
4.     while (((int r = inb(0x1f7)) & IDE_BSY) ||
5.             !(r & IDE_DRDY))
6.         ; // loop until drive isn't busy
7. }
```

# Write parameters to command registers

```
1. static void ide_start_request(struct buf *b)
2. /* send a request (and perhaps data, in the case of a write) to the disk, in
   and out x86 instructions are called to read and write device registers */
3. {
4.     ide_wait_ready();
5.     outb(0x3f6, 0); // generate interrupt
6.     outb(0x1f2, 1); // how many sectors?
7.     outb(0x1f3, b->sector & 0xff); // LBA goes here ...
8.     outb(0x1f4, (b->sector >> 8) & 0xff); // ... and here
9.     outb(0x1f5, (b->sector >> 16) & 0xff); // ... and here!
10.    outb(0x1f6, 0xe0 |((b->dev&1)<<4)|((b->sector>>24)&0x0f));
11.    if(b->flags & B_DIRTY)
12.    {
13.        outb(0x1f7, IDE_CMD_WRITE); // this is a WRITE
14.        outsl(0x1f0, b->data, 512/4); // transfer data too!
15.    }
16.    else
17.    {
18.        outb(0x1f7, IDE_CMD_READ); // this is a READ (no data)
19.    }
20. }
```



# IO interface

```
1.void ide_rw(struct buf *b)
2. // queues a request (if there are others pending)
3. // or issues it directly to the disk (via ide_start_request())
4.{  
5.     acquire(&ide_lock);  
6.     for (struct buf **pp = &ide_queue; *pp; pp=&(*pp)->qnext) ; // walk queue  
7.     *pp = b; // add request to end  
8.     if (ide_queue == b) // if q is empty  
9.         ide_start_request(b); // send req to disk  
10.    while ((b->flags & (B_VALID|B_DIRTY)) != B_VALID)  
11.        sleep(b, &ide_lock); // wait for completion  
12.    release(&ide_lock);  
13.}
```



# Handle interrupts

```
1.void ide_intr()
2./* invoked when an interrupt takes place; it reads data from the device (if
the request is a read, not a write), wakes the process waiting for the I/O to
complete, and (if there are more requests in the I/O queue), launches the next
I/O via ide_start_request() */
3.{  
4.    struct buf *b;  
5.    acquire(&ide_lock);  
6.    if (!(b->flags & B_DIRTY) && ide_wait_ready(1) >= 0)  
7.        insl(0x1f0, b->data, 512/4); // if READ: get data  
8.    b->flags |= B_VALID;  
9.    b->flags &= ~B_DIRTY;  
10.   wakeup(b); // wake waiting process  
11.   if ((ide_queue = b->qnext) != 0) // start next request  
12.       ide_start_request(ide_queue); // (if one exists)
13.   release(&ide_lock);
14. }
```



# 总结一些I/O系统要完成的工作!



**write(buf, 10);**

OS需要提供系统调用  
接口

**DMA.addr = buf;**  
**DMA.count = 10;**

.....

**sleep\_on(Disk);**

查一下手册就可以找  
到该写什么命令?该向  
哪里写?

让出CPU?

需要写  
中断处  
理程序!

**do\_write\_end()//中断处理**

{  
    }

**wakeup(Disk);**

■ **总的感觉:** 很简单  
处理流程是很简单, 复  
杂的是一些**细节问题**,  
如滚屏



# I/O设备管理总结

- 如何实现交互? ⇒ 首先需要了解I/O的工作原理
- 从用户如何I/O开始 ⇒ 用户发送一个命令(**read**)
- 系统调用**read** ⇒ 被展开成**给一些寄存器发送命令**的代码
- 发送完命令以后... ⇒ **CPU轮询**, **CPU干其它事情并等中断**
- 中断方案最常见 ⇒ 相比其他设备, **CPU太快了**
- 实现独享设备的共享 ⇒ 假脱机系统 (**SPOOLING**)



# 例题

系统将数据从磁盘读到内存的过程包括以下操作：

- ①DMA控制器发出中断请求
- ②初始化DMA控制器并启动磁盘
- ③从磁盘传输一块数据到内存缓冲区
- ④执行“DMA结束”中断服务程序

正确的执行顺序是( )

- A. ③→①→②→④
- B. ②→③→①→④
- C. ②→①→③→④
- D. ①→②→④→③

答案 B

解析：在开始DMA传输时，主机向内存写入DMA命令块，向DMA控制器写入该命令块的地址，启动 I/O 设备。然后，CPU继续其他工作，DMA控制器则继续直接操作内存总线，将地址放到总线上开始传输。整个传输完成后，DMA控制器中断CPU，即正确执行顺序为：2, 3, 1, 4。

# Module 6: I/O与存储

## 1. I/O devices

## 2. Hard Disk Drives

- a. 磁盘的基本原理 (基础)
- b. 磁盘调度算法(FCFS、SSTF、SCAN、C-SCAN、C-LOOK) (基础)
- c. 寻址方式 (基础)

## 3. RAID

---

# Tape is Dead Disk is Tape Flash is Disk RAM Locality is King


Jim Gray

Microsoft


# Hard Disk Drives

- 硬盘: **the main form of persistent data storage** in computer systems for decades.
  - The drive consists of a large number of **sectors** (512-byte blocks).
  - **Address Space :**
    - ▶ We can view the disk with  $n$  sectors as an array of sectors; 0 to  $n-1$ .

# 认识一下磁盘



# 认识一下磁盘



机械臂杆



盘片高速旋转产生气流非常强，足以使磁头托起，并与盘面保持一个微小的距离。


现在的水平已经达到  $0.005\mu\text{m} \sim 0.01\mu\text{m}$ ，这只是人类头发直径的千分之一。

# Interface

- The only guarantee is that a single 512-byte write is **atomic**.
- Multi-sector operations are possible.
  - Many file systems will read or write 4KB at a time.
  - **Torn write:**
    - ▶ If an untimely power loss occurs, only a portion of a larger write may complete.
- Accessing blocks in a **contiguous**(连片的) **chunk** is the fastest access mode.
  - A sequential read or write
  - Much faster than any more random access pattern 大读大写特性


# Basic Geometry

- **Platter**(盘面) (Aluminum(铝) coated with a thin magnetic layer)
  - A circular hard surface
  - Data is stored persistently by inducing magnetic changes to it.
  - Each platter has 2 sides, each of which is called a **surface**. Track (12 sectors)
- **Spindle**(主轴)
  - Spindle is connected to a motor that spins the platters around.
  - The rate of rotations is measured in **RPM** (Rotations Per Minute).
    - Typical modern values : 7,200 RPM to 15,000 RPM.
    - E.g., 10000 RPM : A single rotation takes about 6 ms.
- **Track**(轨道)
  - Concentric circles(同心圆) of sectors
  - Data is encoded on each surface in a track.
  - A single surface contains many thousands and thousands of tracks.



A Disk with Just A Single


# A Simple Disk Drive



A Single Track Plus A Head

- **Disk head(磁盘头)** (One head per surface of the drive)
  - The process of *reading* and *writing* is accomplished by the **disk head**.
  - Attached to a single disk arm, which moves across the surface.


# Single-track Latency: The Rotational Delay



A Single Track Plus A Head

- **Rotational delay(旋转延迟):** Time for the desired sector to rotate
  - Ex) Full rotational delay is  $R$  and we start at sector 6
    - ▶ Read sector 0: Rotational delay =  $\frac{R}{2}$
    - ▶ Read sector 5: Rotational delay =  $R-1$  (**worst case**)

# Multiple Tracks: Seek Time



**Three Tracks Plus A Head (Right: With Seek)**  
(e.g., read to sector 11)

- **Seek:** Move the disk arm to the correct track
  - **Seek time:** Time to move head to the track contain the desired sector.
  - One of the most costly disk operations.

# Phases of Seek

- Acceleration → Coasting → Deceleration → Settling
  - **Acceleration:** The disk arm gets moving.
  - **Coasting:** The arm is moving at full speed.
  - **Deceleration:** The arm slows down.
  - **Settling:** The head is *carefully positioned* over the correct track.
    - ▶ The settling time is often quite significant, e.g., 0.5 to 2ms.

# Transfer

---

- The final phase of I/O
  - Data is either *read from* or *written* to the surface.
- Complete I/O time:
  - **Seek**
  - Waiting for the **rotational delay**
  - **Transfer**

## 例题

设一个磁盘的平均寻道时间为12ms，传输速率是200MB/s，控制器开销是0.2ms，转速为每分钟5400转。求读写一个512KB大小数据块的平均磁盘访问时间？

答案：

$$\text{平均旋转延时} = 0.5/5400 \text{ 转/分} = 0.0056 \text{ 秒} = 5.6 \text{ ms} \quad (\text{平均转半圈})$$

$$\begin{aligned}\text{平均磁盘访问时间} &= \text{平均寻道时间} + \text{平均旋转延时} + \text{传输时间} + \text{控制器延时} \\ &= 12 \text{ ms} + 5.6 \text{ ms} + 512 \text{ KB}/200 \text{ MB/s} + 0.2 \text{ ms} = (12 + 5.6 + 2.5 + 0.2) \text{ ms} = 20.3 \text{ ms}\end{aligned}$$

# 例题

在磁盘中读取数据的下列时间中，影响最大的是（ ）。

- A. 处理时间
- B. 延迟时间
- C. 传送时间
- D. 寻找时间/寻道时间

答案：D

解析： 磁盘寻道过程是机械运动，需要移动磁头，时间较长，因此选择D。

# Module 6: I/O与存储

## 1. I/O devices

## 2. Hard Disk Drives

- a. 磁盘的基本原理 (基础)
- b. 磁盘调度算法(FCFS、SSTF、SCAN、C-SCAN、C-LOOK) (基础)
- c. 寻址方式 (基础)

## 3. RAID

# I/O过程是解开许多磁盘问题的钥匙

## ■ 磁盘调度：

前两项可以忽略!

$$\text{磁盘访问延迟} = \text{队列时间} + \text{控制器时间} + \\ \text{寻道时间} + \text{旋转时间} + \text{传输时间}$$

12 ms to 8 ms

8 ms to 4 ms

约0.25ms

■ 多个磁盘访问请求出现在请求队列怎么办? 调度

■ 调度的目标是什么? 调度时主要考察什么?

目标当然是平均  
访问延迟小!

寻道时间是主要  
矛盾!

■ 磁盘调度: 输入多个磁道请求, 给出服务顺序!

# FCFS磁盘调度


$$130+146+85+108+110+59+2=640$$

■ 最直观、最公平的调度：

■ 一个实例：磁头开始磁道位置 = 53，

请求队列 = 98, 183, 37, 122, 14, 124, 65, 67

FCFS: 磁头共移  
动640磁道！





# SSTF磁盘调度

## ■ Shortest-seek-time First最短寻道时间优先:

■ 继续该实例: 磁头开始位置=53;

请求队列=98, 183, 37, 122, 14, 124, 65, 67



SSTF: 磁头共移动  
236(14+53+169)  
磁道, 要少很多!

如果在处理183之  
前又来一些中间磁  
道的请求, 则...





# SCAN磁盘调度(扫描/电梯算法)

## ■ SSTF+中途不回折：每个请求都有处理机会

■ 继续该实例：磁头开始位置=53；

请求队列=98, 183, 37, 122, 14, 124, 65, 67



SCAN: 磁头共移动  
 $53 + 183 = 236$ 磁道，  
和SSTF一样！

这些请求的等待时  
间较长，只因所在  
方向不够幸运！

根据其特征，  
SCAN也被称为  
电梯算法！



■ SCAN导致延迟不均

# C-SCAN磁盘调度

- SCAN+直接移到另一端：两端请求都能很快处理


- 继续该实例：磁头开始位置=53；

- 请求队列=98, 183, 37, 122, 14, 124, 65, 67



CSCAN中的Circular  
是环的意思!

CSCAN: 磁头共移动  
53+199+134磁道!  
其中199会较快!





14→0(183→199)  
没有必要

# C-LOOK磁盘调度

## ■ CSCAN+看一看：前面没有请求就回移

- 继续该实例：磁头开始位置=53；

请求队列=98, 183, 37, 122, 14, 124, 65, 67



- LOOK和C-LOOK是比较合理的缺省算法



操作系统中所有的算法都要因地制宜!



# C-LOOK磁盘调度

继续该实例: 磁头开始位置=53;

请求队列=98, 183, 37, 122, 14, 124, 65, 67



- 1) 磁道请求队列的形式
- 2) 新磁道请求如何入队列

$C[i+1] < X < c[i]$  或者  
 $X > C[i+1] > c[i]$

# 例题

【2018统考真题】系统总是访问磁盘的某个磁道而不响应对其他磁道的访问请求，这种现象称为磁道黏着。下列磁盘调度算法中，不会导致磁道黏着的是（ ）。

- A. 先来先服务(FCFS)
- B. 最短寻道时间优先(SSTF)
- C. 扫描算法(SCAN)
- D. 循环扫描算法(CSCAN)

答案：A

解析：当系统总是持续出现某个磁道的访问请求时，均持续满足SSTF、SCAN、CSCAN的访问条件，会一直服务该访问请求。而FCFS按照请求次序进行调度，比较公平，因此选A。

# Module 6: I/O与存储

## 1. I/O devices

## 2. Hard Disk Drives

- a. 磁盘的基本原理 (基础)
- b. 磁盘调度算法(FCFS、SSTF、SCAN、C-SCAN、C-LOOK) (基础)
- c. 寻址方式 (基础)

## 3. RAID

# 磁盘编址

如何管理磁盘，  
首先对磁盘的扇区进行编号！


- 出厂的磁盘需要低级格式化(物理格式化):  
将连续的磁性记录材料分成物理扇区
- 扇区 = 头 + 数据区 + 尾
- 头、尾中包含只有磁盘控制器能识别的扇区号码和纠错等信息

什么是磁盘的逻辑格式化?  
第12章 文件系统!



# I/O过程是解开许多磁盘问题的钥匙

- 磁盘寻址：对于内存，我们往往更关心存放内容的地址

- 实际上就是扇区怎么编址?
- 显然这个地址是(盘面 (H) + 磁道 (C) + 扇区 (S))
- 寻道和旋转费时多  $\Rightarrow$  花最少时间访问最多扇区的方案: 磁臂不动、磁盘旋转一周, 访问磁头遇到的所有扇区。
- 让这些扇区的编址邻近:  
因为局部性!
- 扇区编址(1): CHS(Cylinder/Head/Sector)
- 扇区编址(2): 扇区编号(Logical Block Addressing LBA)



# 扇区编号—现代磁盘的常见寻址方式



- 扇区编号，按照 $(C, H, S)$ 将扇区形成一维扇区数组，数组索引就是扇区编号

某扇区 $(c, h, s)$ 编号  $A = c * H * S + h * S + s$  扇区总数  $= C * H * S$

已知  $A$ , 则  $s = A \% S$ ;  $h = [A / S] \% H$ ;  $c = [A / (H * S)]$

# 扇区编号—现代磁盘的常见寻址方式

- chs(Cylinder/Head/Sector)模式
- 以前, 硬盘的容量还非常小, 采用与软盘类似的结构生产硬盘.
- 也就是硬盘盘片的每一条磁道都具有相同的扇区数
- 由此产生了所谓的3D参数 (Disk Geometry).:
- 磁柱面数(Cylinders), 头数(Heads), 扇区数(Sectors per track), 以及相应的寻址方式.

# 扇区编号—现代磁盘的常见寻址方式

- chs(Cylinder/Head/Sector)模式
- 磁头数(Heads) 表示硬盘总共有几个磁头,也就是有几面盘片, 最大为 256 (用 8 个二进制位存储);
- 柱面数(Cylinders) 表示硬盘每一面盘片上有几条磁道, 最大为 1024(用 10 个二进制位存储);
- 扇区数(Sectors per track) 表示每一条磁道上有几个扇区, 最大为 63 (用 6 个二进制位存储).
- 每个扇区一般是 512个字节;
- 所以磁盘最大容量为:
- $256 * 1024 * 63 * 512 / 1048576 = 8064 \text{ MB}$

# 磁盘速度与内存速度的差异

- 1) 磁盘往往不是严格按需读取，而是每次都会预读，即使只需要一个字节，磁盘也会从这个位置开始，顺序向后读取**一定扇区长度**的数据放入内存。
- 2) 这样做的理论依据是计算机科学中著名的**局部性原理**：当一个数据被用到时，其附近的数据也通常会马上被使用。

# 例题

下列关于驱动程序的叙述中，不正确的是()。

- A. 驱动程序与 IO 控制方式无关
- B. 初始化设备是由驱动程序控制完成的
- C. 进程在执行驱动程序时可能进入阻塞态
- D. 读/写设备的操作是由驱动程序控制完成的

答案：A

解析：

厂家在设计一个设备时，通常会为该设备编写驱动程序，主机需要先安装驱动程序，才能使用设备。当一个设备被连接到主机时，驱动程序负责初始化设备(如将设备控制器中的寄存器初始化)

# Module 6: I/O与存储

1. I/O devices
2. Hard Disk Drives
3. RAID
  - a. 概述 (基础)
  - b. RAID-0 (基础)
  - c. RAID-1 (基础)
  - d. RAID-4 (基础)
  - e. RAID-5 (基础)

# RAID (Redundant Array of Inexpensive Disks)

- Use multiple disks in concert to build a faster, bigger, and more reliable disk system.
  - RAID just looks like a big disk to the host system.
- Advantage
  - Performance & Capacity: Using multiple disks in parallel
  - Reliability: RAID can tolerate the loss of a disk.

RAIDs provide these advantages transparently  
to systems that use them.

# RAID Interface

- When a RAID receives I/O request,
  1. The RAID **calculates** which disk to access.
  2. The RAID **issue** one or more **physical I/Os** to do so.
- RAID example: A mirrored RAID system
  - Keep two copies of each block (each one on a separate disk)
  - Perform two physical I/Os for every one logical I/O it is issued.

# RAID Internals

- A microcontroller
  - Run firmware to direct the operation of the RAID
- Volatile memory (such as DRAM)
  - Buffer data blocks
- Non-volatile memory
  - Buffer writes safely
- Specialized logic to perform parity calculation, 以固件的形态

# Fault Model

- RAIDs are designed to **detect** and **recover** from certain kinds of disk faults.
- **Fail-stop** fault model
  - A disk can be in one of two states: *Working* or *Failed*.
    - ▶ Working: all blocks can be read or written.
    - ▶ Failed: the disk is permanently lost.
  - RAID controller can immediately observe when a disk has failed.

# How to evaluate a RAID

- **Capacity**

- How much useful capacity is available to systems?

- **Reliability**

- How many disk faults can the given design tolerate?

- **Performance**

# Module 6: I/O与存储

1. I/O devices
2. Hard Disk Drives
3. RAID
  - a. 概述 (基础)
  - b. RAID-0 (基础)
  - c. RAID-1 (基础)
  - d. RAID-4 (基础)
  - e. RAID-5 (基础)

# RAID Level 0: Striping

- RAID Level 0 is the simplest form as **striping** blocks.
  - **Spread the blocks** across the disks in a round-robin fashion.
  - No redundancy
  - Excellent performance and capacity

Disk 0	Disk 1	Disk 2	Disk 3	
0	1	2	3	→ Stripe (The blocks in the same row)
4	5	6	7	
8	9	10	11	
12	13	14	15	

RAID-0: Simple Striping  
(Assume here a 4-disk array)

# RAID Level 0 (Cont.)

- Example) RAID-0 with a bigger chunk size
  - Chunk size : 2 blocks (8 KB)
  - A Stripe: 4 chunks (32 KB)

Disk 0	Disk 1	Disk 2	Disk 3	
0	2	4	6	chunk size: 2blocks
1	3	5	7	
8	10	12	14	
9	11	13	15	

**Striping with a Bigger Chunk Size**

# Chunk Sizes

- Chunk size mostly affects performance of the array
  - **Small chunk size**
    - ▶ Increasing the parallelism
    - ▶ Increasing positioning time to access blocks
  - **Big chunk size**
    - ▶ Reducing intra-file parallelism(文件内并行)
    - ▶ Reducing positioning time(定位时间)

Determining the “best” chunk size is hard to do.  
Most arrays use larger chunk sizes (e.g., 64 KB)

# RAID Level 0 Analysis

$N$  : the number of disks

- **Capacity** → RAID-0 is perfect.
  - Striping delivers  $N$  disks worth of useful capacity.
- **Performance** of striping → RAID-0 is excellent.
  - All disks are utilized often in parallel.
- **Reliability** → RAID-0 is bad.
  - Any disk failure will lead to data loss.

# Evaluating RAID Performance

- Consider two performance metrics
  - Single request latency
  - Steady-state throughput
- Workload
  - **Sequential:** access 1MB of data (block (B) ~ block (B + 1MB))
  - **Random:** access 4KB at random logical address
- A disk can transfer data at
  - $S$  MB/s under a sequential workload
  - $R$  MB/s under a random workload

# Evaluating RAID Performance Example

- sequential ( $S$ ) vs random ( $R$ )
  - **Sequential** : transfer 10 MB on average as continuous data.
  - **Random** : transfer 10 KB on average.
  - Average seek time: 7 ms
  - Average rotational delay: 3 ms
  - Transfer rate of disk: 50 MB/s
- Results:
  - $S = \frac{\text{Amount of Data}}{\text{Time to access}} = \frac{10 \text{ MB}}{210 \text{ ms}} = 47.62 \text{ MB /s}$
  - $R = \frac{\text{Amount of Data}}{\text{Time to access}} = \frac{10 \text{ KB}}{10.195 \text{ ms}} = 0.981 \text{ MB /s}$

# Evaluating RAID-0 Performance

$N$  : the number of disks

- Single request latency
  - Identical to that of a single disk.
- Steady-state throughput
  - **Sequential** workload :  $N \cdot S$  MB/s
  - **Random** workload :  $N \cdot S$  MB /s

# Module 6: I/O与存储

1. I/O devices
2. Hard Disk Drives
3. RAID
  - a. 概述 (基础)
  - b. RAID-0 (基础)
  - c. RAID-1 (基础)
  - d. RAID-4 (基础)
  - e. RAID-5 (基础)

# RAID Level 1 : Mirroring

- RAID Level 1 tolerates **disk failures**.
  - Copy more than one of **each block** in the system.
  - Copy block places on a separate disk.

Disk 0	Disk 1	Disk 2	Disk 3
0	0	1	1
2	2	3	3
4	4	5	5
6	6	7	7

Simple RAID-1: Mirroring (Keep two physical copies)

- ▶ RAID-10 (RAID 1+0) : mirrored pairs and then stripe
- ▶ RAID-01 (RAID 0+1) : contain two large striping arrays, and then mirrors

# RAID-1 Analysis

$N$  : the number of disks

- **Capacity:** RAID-1 is Expensive
  - The useful capacity of RAID-1 is  $N/2$ .
- **Reliability:** RAID-1 does well.
  - It can tolerate the failure of any one disk (up to  $N/2$  failures depending on which disk fail).

# Performance of RAID-1

- Two physical writes to complete
  - It suffers the worst-case seek and rotational delay of the two request.
  - Steady-state throughput
    - ▶ **Sequential Write** :  $\frac{N}{2} \cdot S$  MB/s
      - Each logical write must result in two physical writes.
    - ▶ **Sequential Read** :  $\frac{N}{2} \cdot S$  MB/s
      - Each disk will only deliver half its peak bandwidth.
    - ▶ **Random Write** :  $\frac{N}{2} \cdot R$  MB/s
      - Each logical write must turn into two physical writes.
    - ▶ **Random Read** :  $N \cdot R$  MB/s
      - Distribute the reads across all the disks.

# Module 6: I/O与存储

1. I/O devices
2. Hard Disk Drives
3. RAID
  - a. 概述 (基础)
  - b. RAID-0 (基础)
  - c. RAID-1 (基础)
  - d. RAID-4 (基础)
  - e. RAID-5 (基础)

# RAID Level 4 : Saving Space With Parity

- Add a single parity block

- A Parity block(奇偶校验块) stores the *redundant information* for that stripe of blocks.

\* P: Parity

Disk 0	Disk 1	Disk 2	Disk 3	Disk 4
0	0	1	1	P0
2	2	3	3	P1
4	4	5	5	P2
6	6	7	7	P3

Five-disk RAID-4 system layout

# RAID Level 4 (Cont.)

- Compute parity : the XOR of all of bits

C0	C1	C2	C3	P
0	0	1	1	XOR(0,0,1,1)=0
0	1	0	0	XOR(0,1,0,0)=1

- Recover from parity

- Imagine the bit of the C2 in the first row is lost.
  1. Reading the other values in that row : 0, 0, 1
  2. The parity bit is 0 → even number of 1's in the row
  3. What the missing data must be: a 1.

# RAID-4 Analysis

$N$  : the number of disks

## □ Capacity

- The useful capacity is  $(N - 1)$ .

## □ Reliability

- RAID-4 tolerates 1 disk failure and no more.

# RAID-4 Analysis (Cont.)

## □ Performance

### □ Steady-state throughput

- ▶ Sequential read:  $(N - 1) \cdot S$  MB/s
- ▶ Sequential write:  $(N - 1) \cdot S$  MB/s

Disk 0	Disk 1	Disk 2	Disk 3	Disk 4
0	1	2	3	P0
4	5	6	7	P1
8	9	10	11	P2
12	13	14	15	P3

### Full-stripe(全条带) Writes In RAID-4

- ▶ Random read:  $(N - 1) \cdot R$  MB/s

# Random write performance for RAID-4

- Overwrite a block + update the parity
- **Method 1: additive parity**
  - Read in all of the other data blocks in the stripe
  - XOR those blocks with the new block (1)
  - **Problem:** the performance scales with the number of disks
- **Method 2: subtractive parity**

C0	C1	C2	C3	P
0	0	1	1	XOR(0,0,1,1)=0

$$P(\text{new}) = (C2(\text{old}) \text{ } XOR \text{ } C2(\text{new})) \text{ } XOR \text{ } P(\text{old})$$

- Update C2(old) → C2(new)
  1. Read in the old data at C2 ( $C2(\text{old})=1$ ) and the old parity ( $P(\text{old})=0$ )
  2. Calculate  $P(\text{new})$ :
    - If  $C2(\text{new})==C2(\text{old}) \rightarrow P(\text{new})==P(\text{old})$
    - If  $C2(\text{new})!=C2(\text{old}) \rightarrow$  Flip the old parity bit

# Small-write problem

- The parity disk can be a **bottleneck**(瓶颈).
- Example: update blocks 4 and 13 (marked with \*)

Disk 0	Disk 1	Disk 2	Disk 3	Disk 4
0	1	2	3	P0
*4	5	6	7	+P1
8	9	10	11	P2
12	*13	14	15	+P3

Writes To 4, 13 And Respective Parity Blocks.

- ▶ Disk 0 and Disk 1 can be accessed in parallel.
- ▶ Disk 4 prevents any parallelism.

RAID-4 throughput(吞吐量) under random small writes is  $(\frac{R}{2})$  MB/s (**terrible**)

# A I/O latency in RAID-4

- A single read
  - Equivalent to the latency of a single disk request.
- A single write
  - Two reads and then two writes
    - Data block + Parity block
    - The reads and writes can happen in parallel.
  - Total latency *is about twice* that of a single disk.

# Module 6: I/O与存储

1. I/O devices
2. Hard Disk Drives
3. RAID
  - a. 概述 (基础)
  - b. RAID-0 (基础)
  - c. RAID-1 (基础)
  - d. RAID-4 (基础)
  - e. RAID-5 (基础)

# RAID Level 5: Rotating Parity

- RAID-5 is solution of small write problem.
  - Rotate the parity blocks across drives.
  - Remove the parity-disk bottleneck for RAID-4

Disk 0	Disk 1	Disk 2	Disk 3	Disk 4
0	1	2	3	P0
5	6	7	P1	4
10	11	P2	8	9
15	P3	12	13	14
P4	16	17	18	19

RAID-5 With Rotated Parity

# RAID-5 Analysis

$N$  : the number of disks

## □ Capacity

- The useful capacity for a RAID group is  $(N - 1)$ .

## □ Reliability

- RAID-5 tolerates 1 disk failure and no more.

# RAID-5 Analysis (Cont.)

$N$  : the number of disks

## □ Performance

- Sequential read and write }
- A single read and write request      Same as RAID-4
- Random read : a little better than RAID-4
  - ▶ RAID-5 can utilize all of the disks.
- Random write :  $\frac{N}{4} \cdot R$  MB/s
  - ▶ The factor of four loss is cost of using parity-based RAID.

# RAID Comparison: A Summary

$N$  : the number of disks

$D$  : the time that a request to a single disk take

S : 顺序读写的集成带宽

R : 随机读写的集成带宽

	RAID-0	RAID-1	RAID-4	RAID-5
<b>Capacity</b>	$N$	$N/1$	$N-1$	$N-1$
<b>Reliability</b>	0	$\frac{1}{2}$ (for sure) $\frac{N}{2}$ (if lucky)	1	1
<b>Throughput</b>				
Sequential Read	$N \cdot S$	$(N/2) \cdot S$	$(N-1) \cdot S$	$(N-1) \cdot S$
Sequential Write	$N \cdot S$	$(N/2) \cdot S$	$(N-1) \cdot S$	$(N-1) \cdot S$
Random Read	$N \cdot R$	$N \cdot R$	$(N-1) \cdot R$	$N \cdot R$
Random Write	$N \cdot R$	$(N/2) \cdot R$	$\frac{1}{2}R$	$\frac{N}{4}R$
<b>Latency</b>				
Read	$D$	$D$	$D$	$D$
Write	$D$	$D$	$2D$	$2D$

# RAID Comparison: A Summary

- **Performance** and do not care about reliability → RAID-0 (Striping)
- Random I/O performance and **Reliability** → RAID-1 (Mirroring)
- **Capacity** and **Reliability** → RAID-5
- **Sequential I/O** and Maximize **Capacity** → RAID-5

# 謝 謝！